

A Tale of Three Deep Learning Frameworks: TensorFlow, Keras, and PyTorch

Brooke Wenig
Jules S. Damji

Spark + AI Summit, London 4 October 2018



About Us ...

Jules S. Damji

Apache Spark Developer & Community Advocate @Databricks

Program Chair Spark + AI Summit

Software engineering @ Sun Microsystems, Netscape, @Home, VeriSign, Scalix, Centrify, LoudCloud/Opware, ProQuest

<https://www.linkedin.com/in/dmatrix>

[@2twitme](#)

Brooke Wenig

Machine Learning Practice Lead @ Databricks

Data Science @ Splunk & MyFitnessPal

MS Machine Learning (UCLA)

Fluent in Chinese

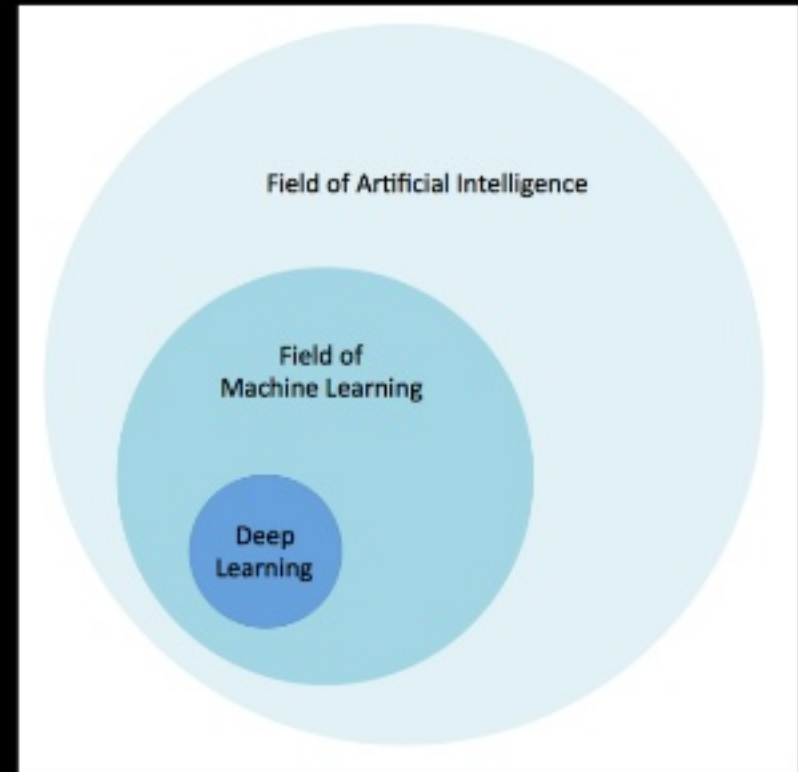

<https://www.linkedin.com/in/brookewenig/>

Agenda for Today's Talk

- What's Deep Learning and Why?
- Short Survey of 3 DL Frameworks
 - TensorFlow
 - Keras
 - PyTorch
- Training Options
 - Single Node
 - Distributed
- Q&A

What is Deep Learning?

“Composing representations of data in a hierarchical manner”



Why Deep Learning?

MENU ▾

nature

International journal of science

Review Article | Published: 27 May 2015

Deep learning

Yann LeCun , Yoshua Bengio & Geoffrey Hinton

Nature **521**, 436–444 (28 May 2015) | [Download Citation](#) 

Abstract

Deep learning allows computational models that are composed of multiple processing layers to learn representations of data with multiple levels of abstraction. These methods have dramatically improved the state-of-the-art in speech recognition, visual object recognition, object detection and many other domains such as drug discovery and genomics. Deep learning discovers intricate structure in large data sets by using the backpropagation algorithm to indicate how a machine should change its internal parameters that are used to compute the representation in each layer from the representation in the previous layer. Deep convolutional nets have brought about breakthroughs in processing images, video, speech and audio, whereas recurrent nets have shone light on sequential data such as text and speech.

Applications


Machine Learning is Everywhere?




AlphaGo



Recommendation systems



Drug discovery




Character recognition



Hedge fund stock predictions



Voice assistants



Assisted driving


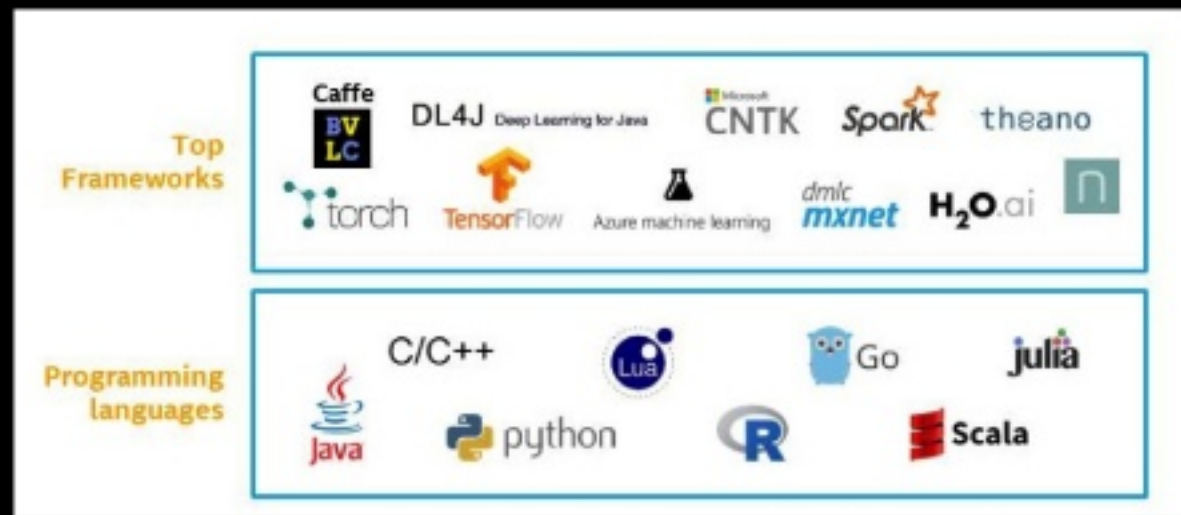


Face detection/recognition



Cancer diagnosis

Zoo of DL Frameworks: Which One?



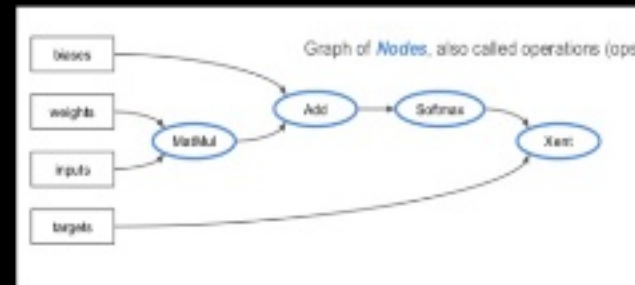


Survey of Three Deep Learning Frameworks

What's TensorFlow?



- Open source from Google, 2015
 - [Current v1.12 API](#)
 - 2.0 Coming Soon... :)
 - Declarative Toolkit
- Fast: Backend C/C++
- Data flow graphs
 - Nodes are functions/operators
 - Edges are input or data (tensors)
 - Lazy execution
 - Eager execution (1.7)



TensorFlow Key API Concepts



- Constants
- Variables
- Placeholders
- Operations
- Sessions
- Tensors
- Graphs

```
x = tf.constants(42, name='x')
```

```
w = tf.Variable(1.34, name='w')
```

```
input = tf.Placeholder("float")
```

```
c = tf.add(x, w); m = tf.matmul(a, b) ...
```

```
with tf.Session([URI]) as sess:
```

```
1, [1, 2], [[2, 3], [4, 5]] ...
```

```
g = tf.Graph("my_graph")
```

```
with g.as_default():
```

```
c = tf.add(x, w)
```

```
m = tf.matmul(a, b)
```

TensorFlow Code



```
import tensorflow as tf
```

```
a = tf.placeholder(tf.float32, shape=(2,1))
```

```
b = tf.placeholder(tf.float32, shape=(1,2))
```



Create TF placeholder types, a & b
Define their input shape as tensors

```
c = tf.matmul(a, b)
```



TF matmul matrix operation

```
sess = tf.Session()
```

```
print(sess.run(c, {a: [[1],[2]], b: [[3,4]]}))
```



Create a TF Session
Run the session, with input parameters
for place holders 'a' & 'b'

'c' as an operation won't to run until
sess.run() Lazily evaluated.

```
[[3. 4.] [6. 8.]]
```



TF session output

TensorFlow Code: MNIST



```
import tensorflow as tf
```

```
from tensorflow.examples.tutorials.mnist import input_data  
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

← Use tf input_data modules for MNIST

```
# Create the model
```

```
x = tf.placeholder(tf.float32, [None, 784])
```

← TF placeholders & variables

```
W = tf.Variable(tf.zeros([784, 10]))
```

```
b = tf.Variable(tf.zeros([10]))
```

```
y = tf.matmul(x, W) + b
```

← Define our model

```
y_ = tf.placeholder(tf.int64, [None])
```

TF variable for predicted value y'

```
...
```

```
# Define loss and optimizer
```

```
cross_entropy = tf.losses.sparse_softmax_cross_entropy(labels=y_, logits=y)  
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

← Define our loss function: cross_entropy
Use Gradient Descent Optimizer

```
#create session, train, and evaluate
```

```
sess = tf.InteractiveSession()
```

```
tf.global_variables_initializer().run()
```

```
# Train
```

```
for _ in range(1000):
```

```
    batch_xs, batch_ys = mnist.train.next_batch(100)
```

```
    sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
```

← Train or evaluate the model

```
# Test trained model
```

```
correct_prediction = tf.equal(tf.argmax(y, 1), y_)
```

```
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

```
...
```

```
print(sess.run(accuracy, feed_dict={
```

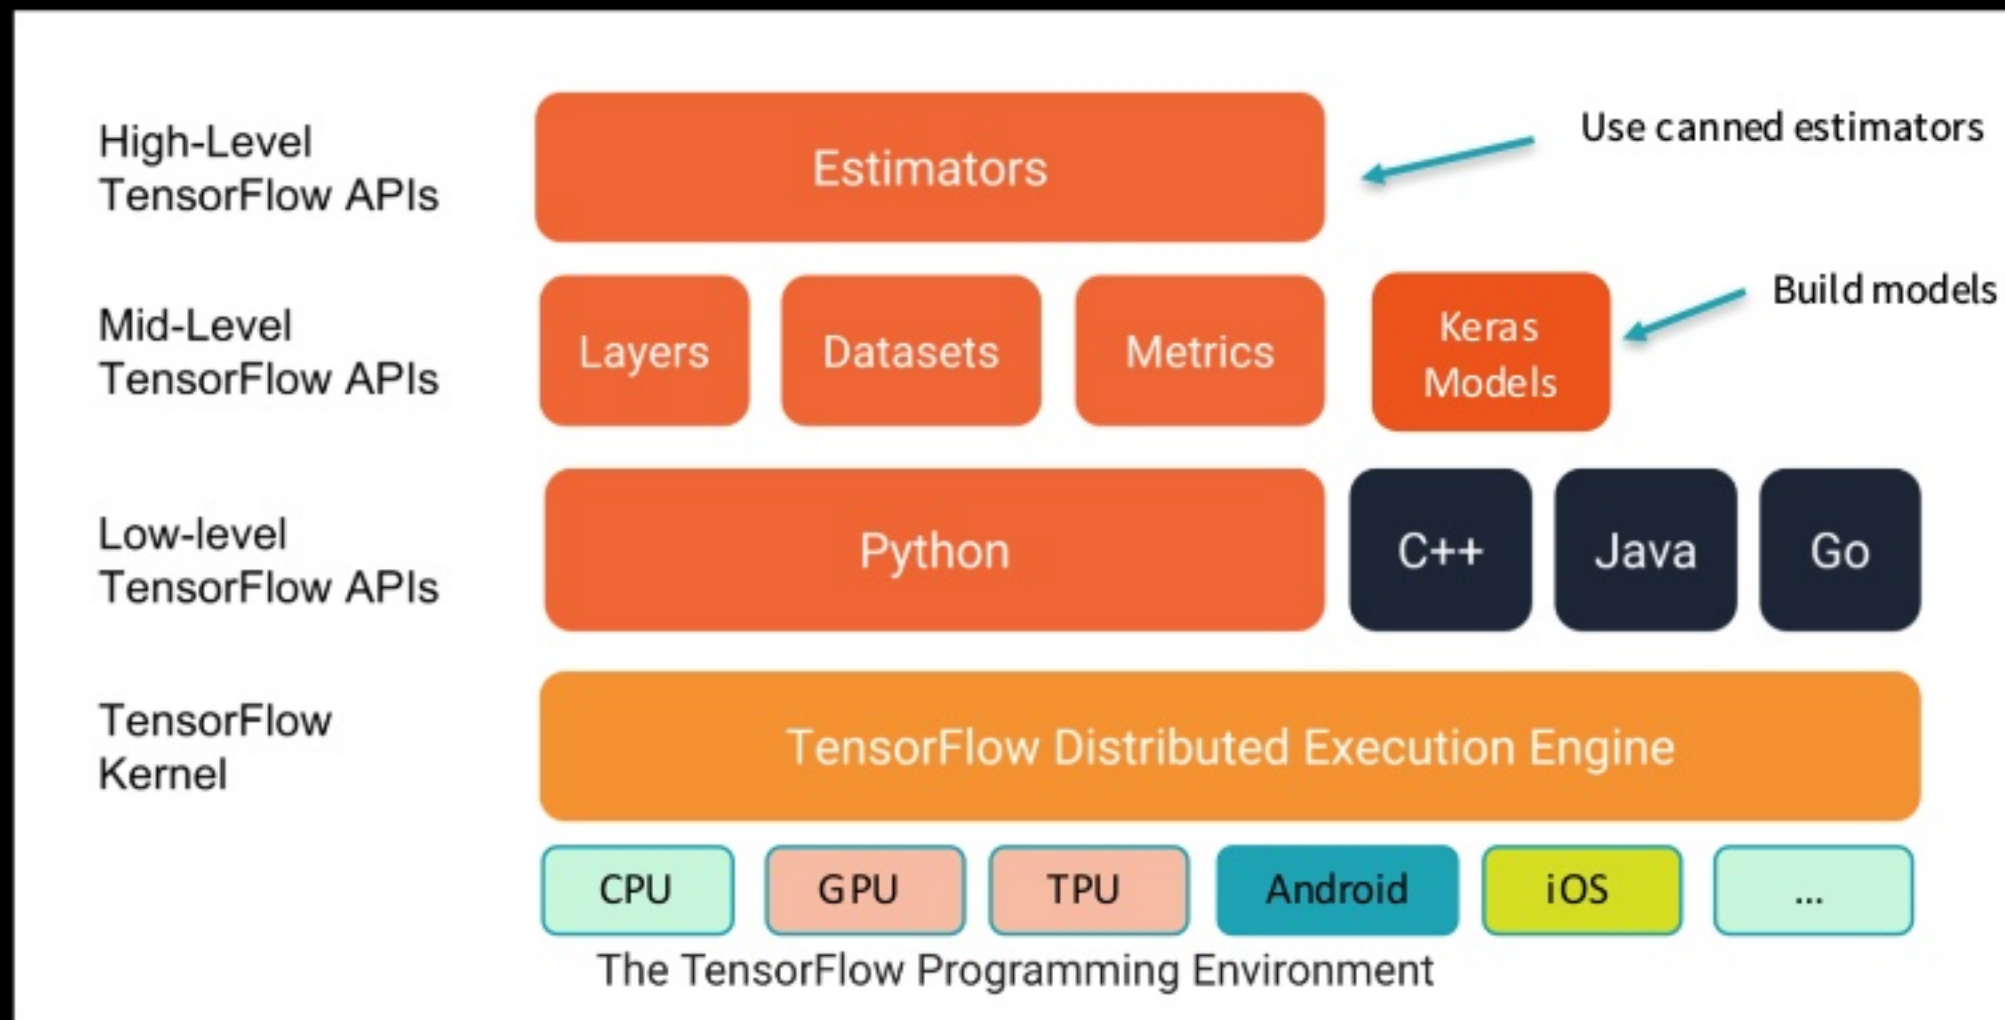
```
    x: mnist.test.images,
```

```
    y_: mnist.test.labels
```

```
}))
```

https://github.com/tensorflow/tensorflow/blob/master/tensorflow/examples/tutorials/mnist/mnist_softmax.py

TensorFlow Programming Stack



Why TensorFlow: Community



tensorflow / tensorflow

Watch

8,179

Star

105,292

Fork

65,530

Code

Issues 1,535

Pull requests 271

Projects 0

Insights

Computation using data flow graphs for scalable machine learning <https://tensorflow.org>

tensorflow

machine-learning

python

deep-learning

deep-neural-networks

neural-network

ml

distributed

36,584 commits

23 branches

62 releases

1,549 contributors

Apache-2.0

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download

- 105K+ stars!
- 11+M downloads
- Popular open-source code
- TensorFlow Hub & Blog
 - Code Examples & Tutorials!
 - Learn + share from others

Get Started with TensorFlow

TensorFlow is an open source machine learning library for research and production. TensorFlow offers APIs for beginners and experts to create for desktop, mobile, web, and cloud. See the sections below to get started.

Learn and use ML

The HighLevel Keras API provides building blocks to create and train deep learning models. Start with the Keras Getting Started notebook examples, then read the TensorFlow Getting Guide.

1. Basic classification
2. Text classification
3. Regression
4. Generating and understanding
5. State and feed

Research and experimentation

TensorFlow provides an imperative, define-by-run interface for advanced operations. For faster research, prototyping, and testing loops with auto-differentiation. Start with these notebooks, then read the TensorFlow Getting Guide.

1. Image classification basics
2. Automatic differentiation and gradient descent
3. Customizing layers
4. Custom layers
5. Custom training workflow
6. Example: Neural machine translation with attention

ML at product on scale

Distributions can train large models on multiple machines in a production environment. TensorFlow provides a collection of pre-made C++ code to implement common ML algorithms. See the TensorFlow Getting Guide.

1. Build a basic model with Estimators
2. Write and deploy training with Estimator v2
3. Customized model
4. How to build a deep neural network with TF Hub
5. Build a distributed neural network using Estimators

Introducing Swift For TensorFlow

Posted by the Swift for TensorFlow team at Google

Introducing TensorFlow.js: Machine Learning in Javascript

Posted by Josh Gordon and Sara Robinson, Developer Advocates

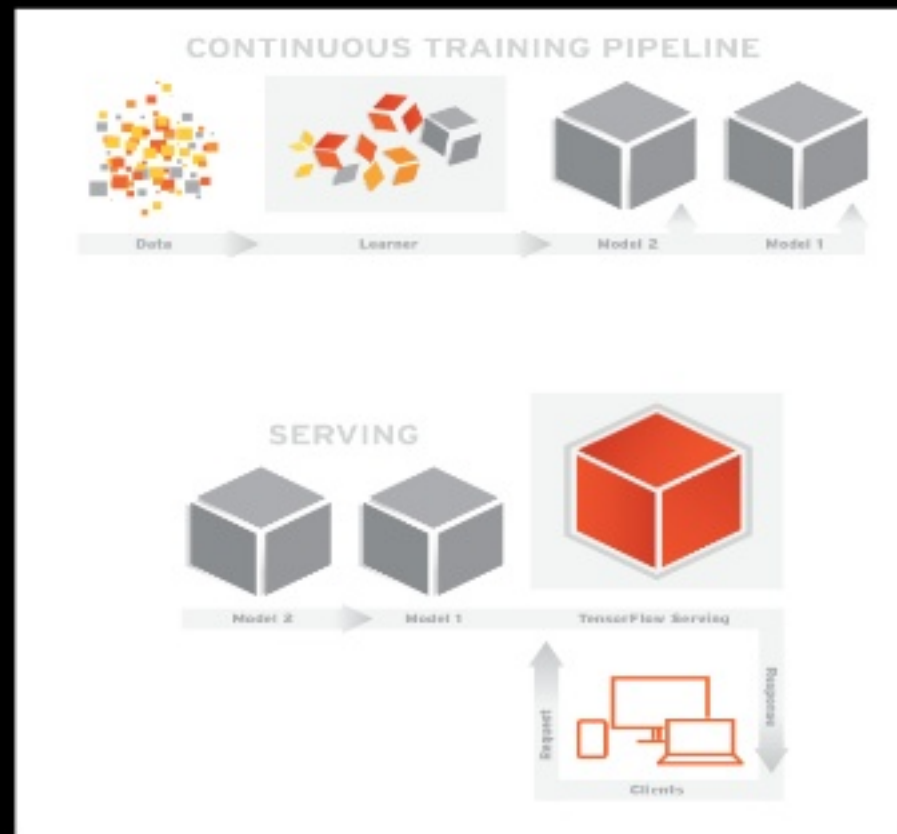
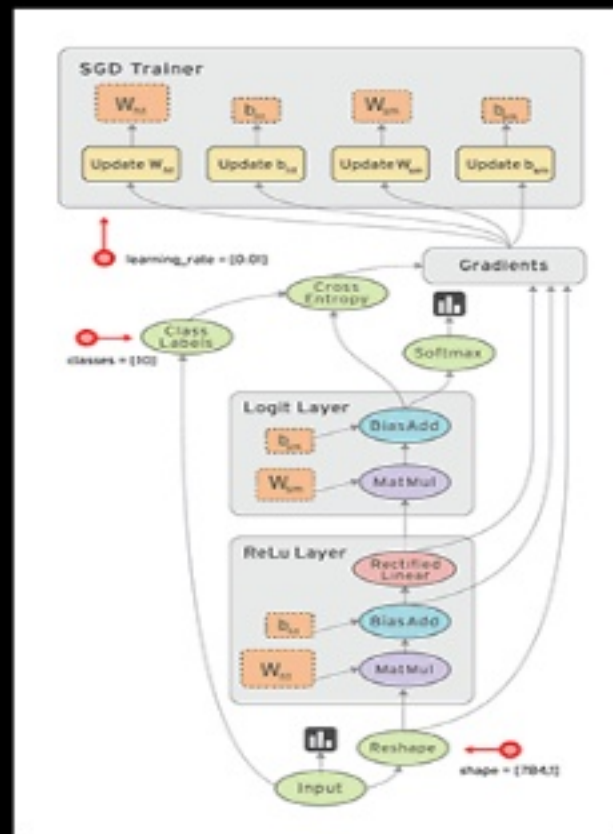
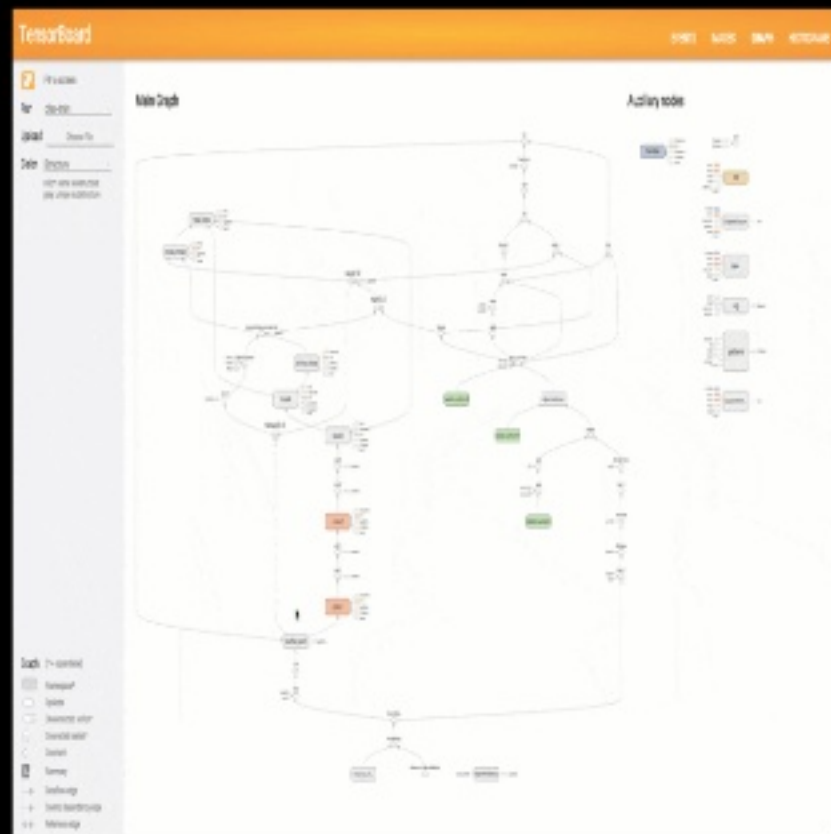
Introducing TensorFlow Probability

Posted by Josh Dillon, Software Engineer, Mike Shaw, Product Manager, and Dustin Tran, Research Scientist — on behalf of the TensorFlow...

OSCON TensorFlow Community Day: Call for participation

I'm happy to announce the call for participation for TensorFlow Community Day at the O'Reilly Open Source Convention (OSCON), on July 17.

Why TensorFlow: Tools



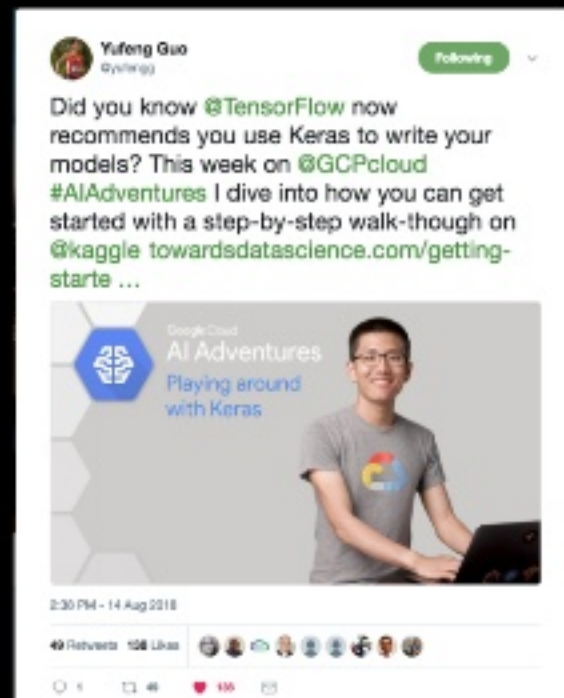
Visualize Tensors flow

Deploy + Serve Models
TFX

TensorFlow: We Get it ... So What?



- Steep learning curve, *but powerful!!*
- Low-level APIs, *but offers control!!*
- Expert in Machine Learning, *just learn!!*
- Yet, high-level Estimators help, *you bet!!*
- Yeah, TensorFlow 2.0, *ease-of-use, eager execution!*
- Better, Keras integration helps, *indeed!!*

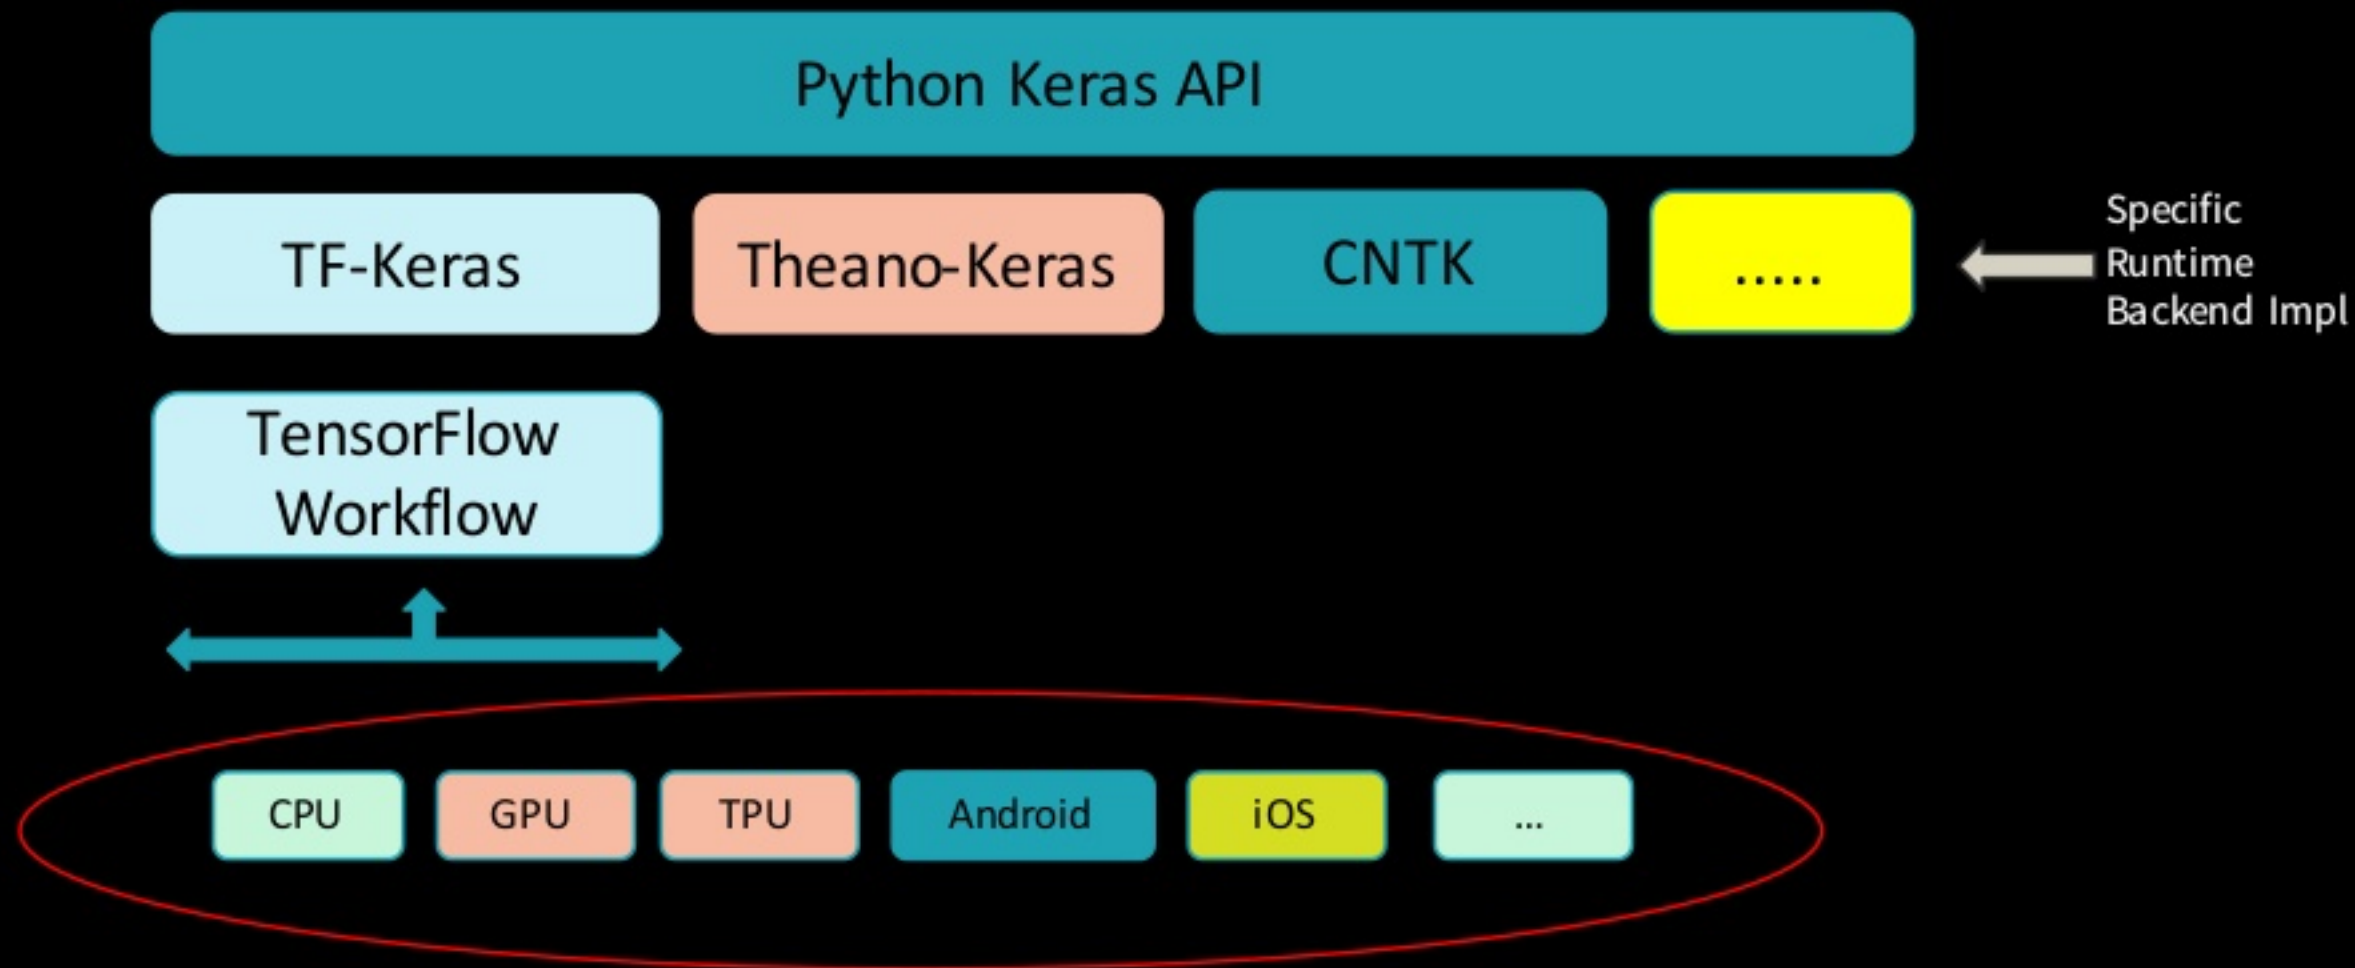


What's Keras?



- Open source Python Library APIs for Deep Learning
 - Current v2.2.2 APIs François Chollet (Google)
- APIs : with TensorFlow, CNTK and Theano Backends
- *Easy to Use High-Level Declarative APIs!*
 - Build layers
 - Great for Neural Network Applications
 - CNN
 - RNN & LSTM
- *Fast Experimentation, Modular & Extensible!*

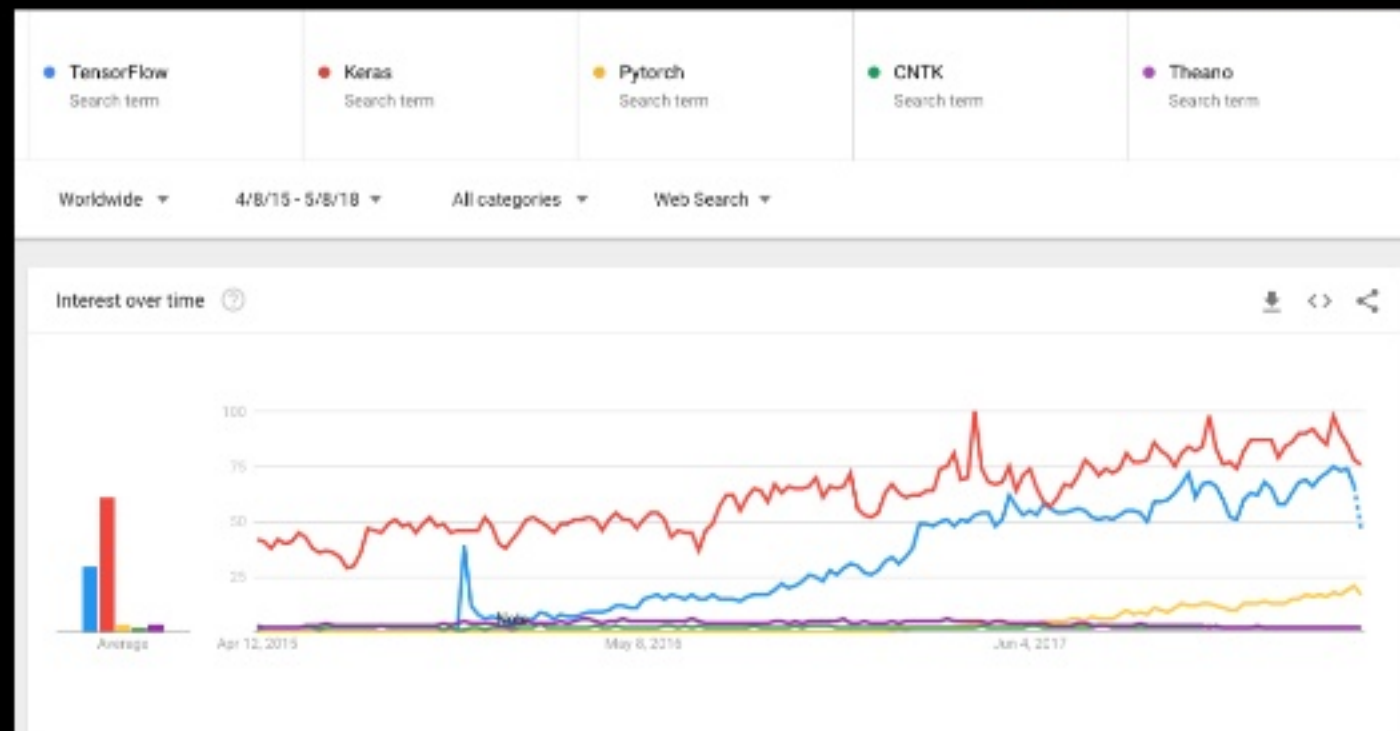

Keras Programming Stack



Why Keras?

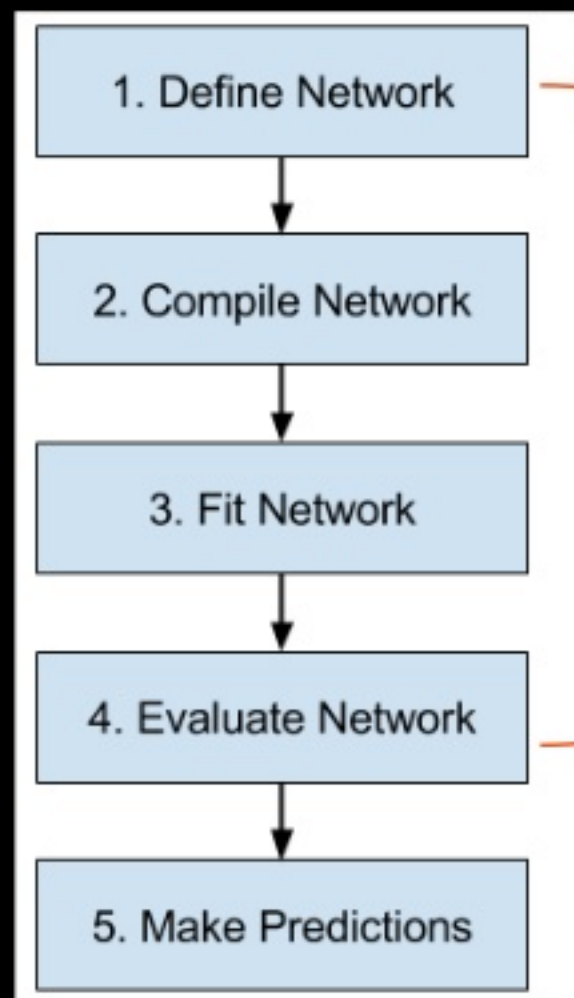


- *Focuses on Developer Experience*
- Popular & Broader Community
- Supports multiple backends
- Modularity
 - Sequential
 - Functional



```
model = Sequential()  
model.add(Dense(32, input_dim=784))  
model.add(Activation('relu'))  
model.add(Dense, 32, activation='softmax')  
...
```

Keras Code: *MNIST*



```
from keras import models
from keras import layers
```

```
mnist = tf.keras.datasets.mnist
(train_images, train_labels), (test_images, test_labels) =
    prepare_data(mnist.load_data())
```

```
network = models.Sequential()
network.add(layers.Dense(512, activation='relu',
                        input_shape=(28 * 28,)))
network.add(layers.Dense(10, activation='softmax'))
```

```
network.compile(optimizer='rmsprop',
                loss='categorical_crossentropy',
                metrics=['accuracy'])
```

```
network.fit(train_images, train_labels, epochs=5,
            batch_size=128)
```

```
results = network.evaluate(test_images, test_labels)
```

```
predictions = network.predict(new_images)
```

← Set up code & use dataset

← Define Network

← Compile Network

← Fit Network

← Evaluate Network

← Make Predictions

Another a Simple Network Model: Which code is easier to read?

Python: The Language of Deep Learning?

```
with tf.variable_scope('conv1') as scope:
    kernel = _variable_with_weight_decay('weights',
                                         shape=[5, 5, 3, 64],
                                         stddev=5e-2,
                                         wd=None)
    conv = tf.nn.conv2d(images, kernel, [1, 1, 1, 1], padding='SAME')
    biases = _variable_on_cpu('biases', [64], tf.constant_initializer(0.0))
    pre_activation = tf.nn.bias_add(conv, biases)
    conv1 = tf.nn.relu(pre_activation, name=scope.name)
    _activation_summary(conv1)

# pool1
pool1 = tf.nn.max_pool(conv1, ksize=[1, 3, 3, 1], strides=[1, 2, 2, 1],
                       padding='SAME', name='pool1')

# norm1
norm1 = tf.nn.lrn(pool1, 4, bias=1.0, alpha=0.001 / 9.0, beta=0.75,
                  name='norm1')

# conv2
with tf.variable_scope('conv2') as scope:
    kernel = _variable_with_weight_decay('weights',
                                         shape=[5, 5, 64, 64],
                                         stddev=5e-2,
                                         wd=None)
    conv = tf.nn.conv2d(norm1, kernel, [1, 1, 1, 1], padding='SAME')
    biases = _variable_on_cpu('biases', [64], tf.constant_initializer(0.1))
    pre_activation = tf.nn.bias_add(conv, biases)
    conv2 = tf.nn.relu(pre_activation, name=scope.name)
    _activation_summary(conv2)

# norm2
```

TensorFlow

```
model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same',
                 input_shape=x_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3, 3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))
```

Keras

```
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

PyTorch

What is PyTorch?



- Open source from Facebook, 2017
 - v1.0 dev release
- Primarily a Python Package
- Tensor Computations
 - Torch.tensor -> CPU, GPU/CUDA
- *Dynamic* NN: Tape-based Autograd
- Graph Based *Dynamic* Computations
- Imperative Toolkit

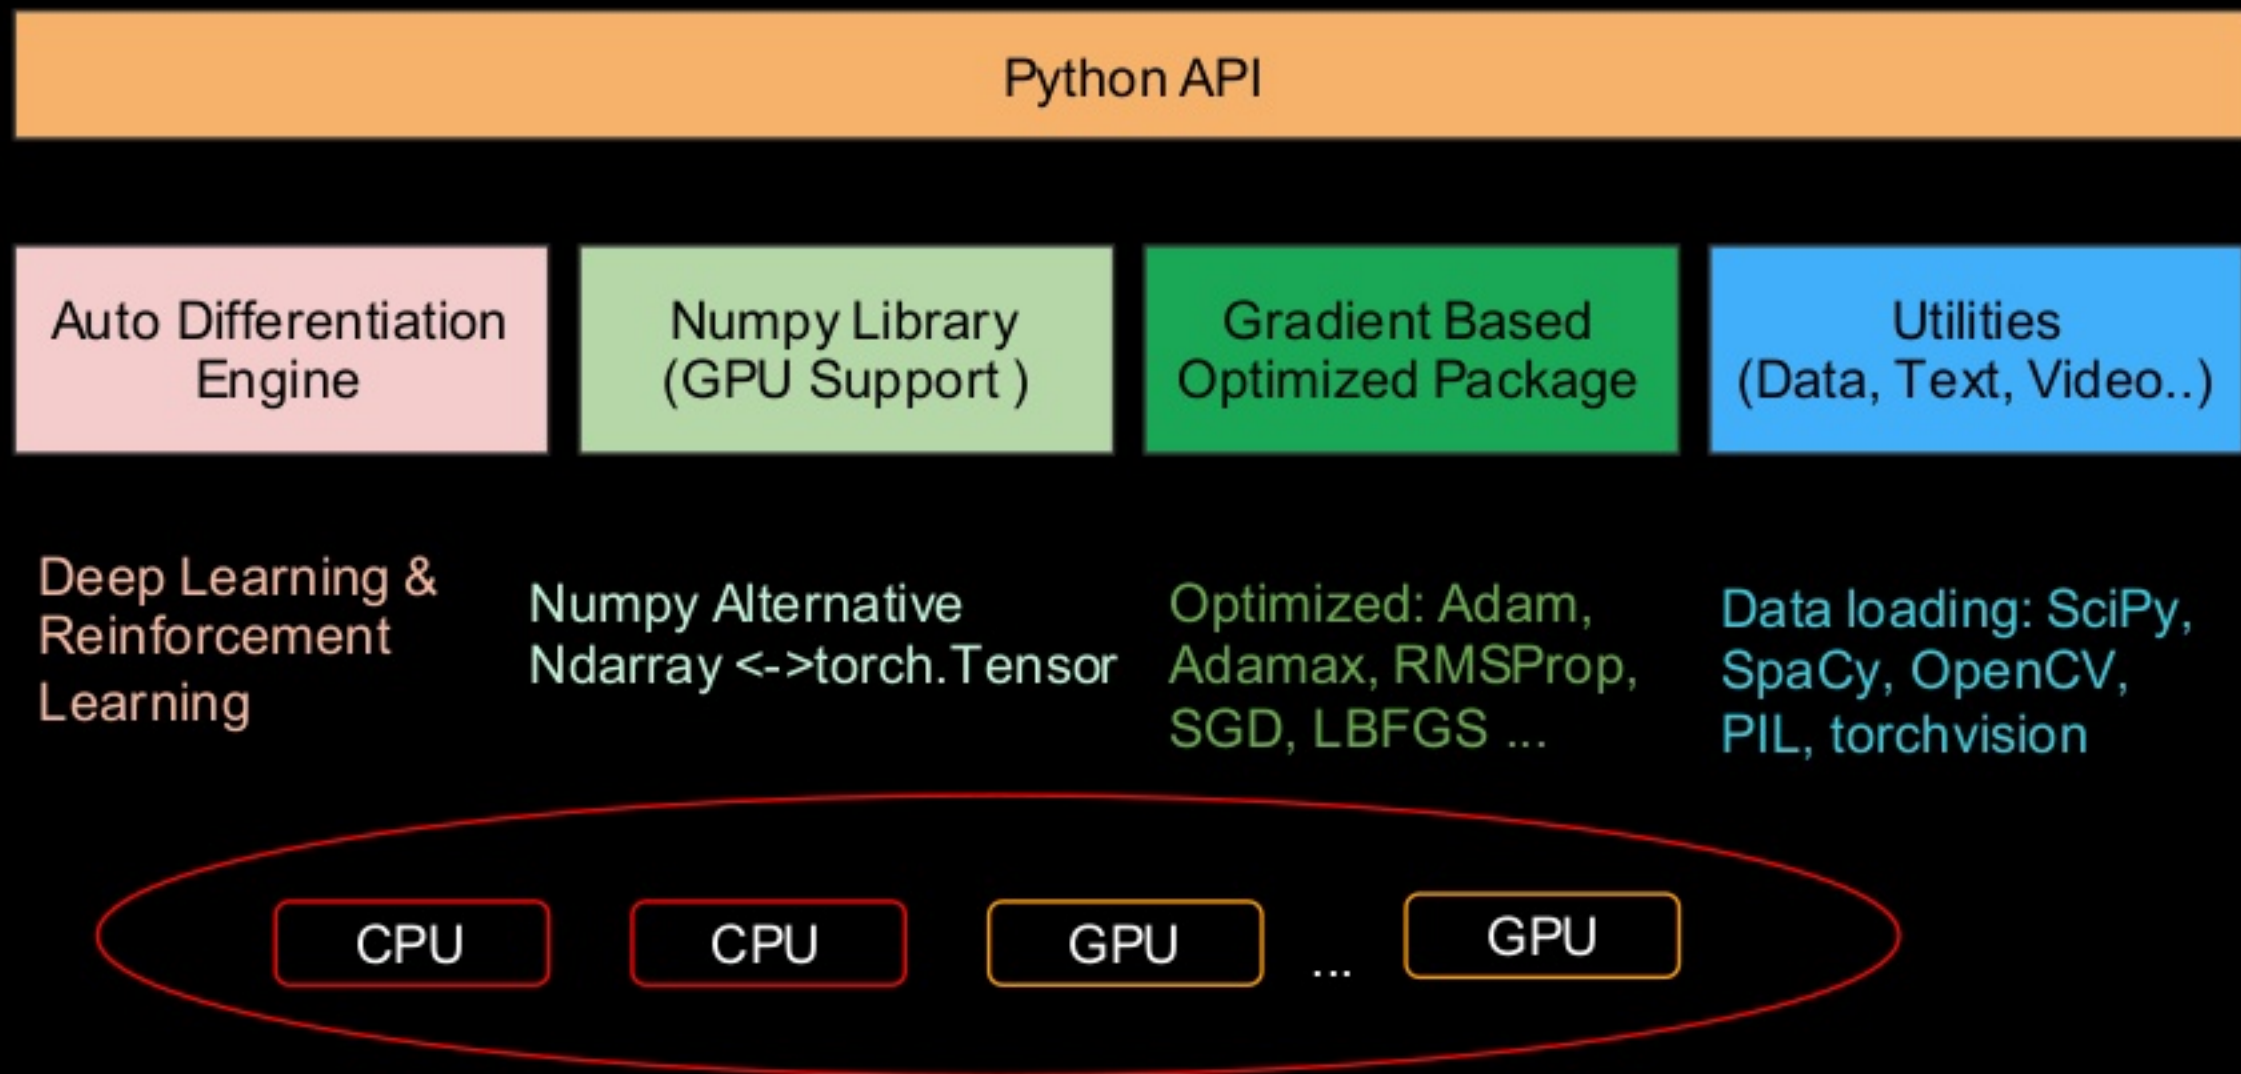
A graph is created on the fly



```
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
W_h = torch.randn(20, 20)
W_x = torch.randn(20, 10)
```



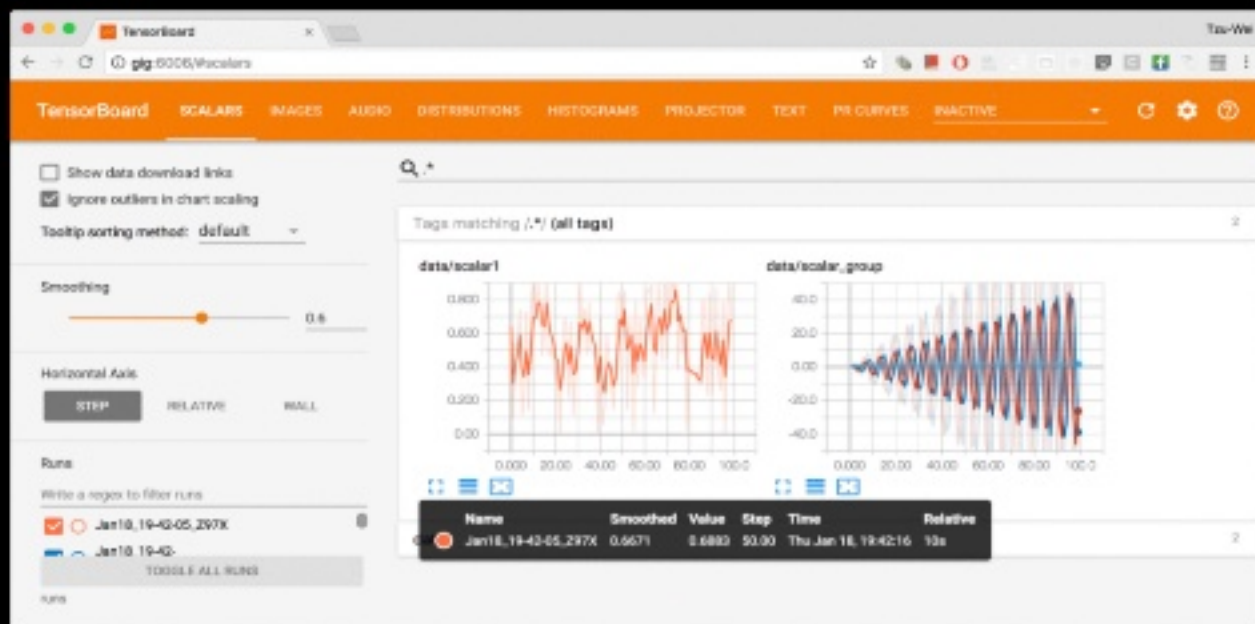
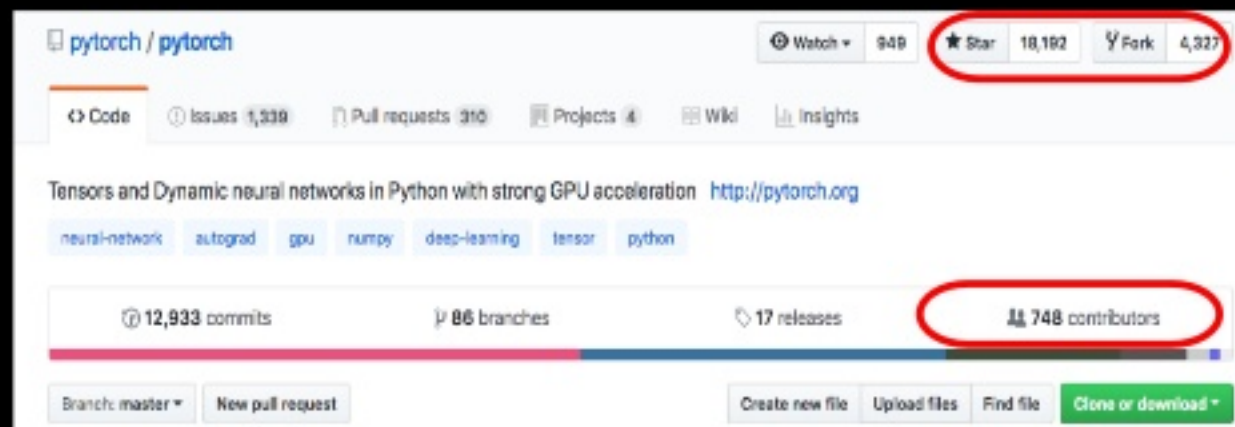
PyTorch Programming Stack



Why PyTorch?



- Imperative Experience
 - Rapid Prototyping for Research
 - Easy Debugging & TBX
- Quick Ramp-up time
- Decent Docs & Community
 - 275K Downloads
 - 1900 Community Repos
 - 13+K Blogs Posts
- *Pythonic!*



PyTorch Key API Concepts



- Variables & Autograd
- Torch Tensors
- Operations

```
from torch import Variable
x = Variable(torch.Tensor([2]),requires_grad=True)
y = 5*x**4 + 3*x**3 + 7*x**2 + 9*x - 5
y.backward() #compute gradient and backpropagate
x.grad
```

```
#different kinds of Torch Tensors
x = torch.rand(5, 3);
y = torch.rand(5, 3)
t = torch.tensor([5.5, 3])
n = torch.tensor(np.array([[1, 2, 3], [4, 5, 6]]))
```


#operations or element-wise operations

```
a = (x + y)
m = (x * y)
if torch.cuda.is_available():
    x = x.cuda()
    y = y.cuda()
torch.add(x, y, out=result)
print(a, m, out)
```

PyTorch Rhythm ...



1. Design a model using PyTorch Autograd Variables
2. Construct a loss function and optimizer with PyTorch APIs
3. Train your model: forward, backward, and update steps



PyTorch Rhythm : Linear Regression



Logistic regression

```
x_data = Variable(torch.tensor([[1.0], [2.0], [3.0], [4.0]]))
y_data = Variable(torch.tensor([[0.], [0.], [1.], [1.])))


class Model(torch.nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.linear = torch.nn.Linear(1, 1) # One in and one out

    def forward(self, x):
        y_pred = F.sigmoid(self.linear(x))
        return y_pred

# our model
model = Model()
```

1

Design your model using class



```
criterion = torch.nn.BCELoss(size_average=True)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

2

Construct loss and optimizer
(select from PyTorch API)

Training loop

```
for epoch in range(1000):
    # Forward pass: Compute predicted y by passing x to the model
    y_pred = model(x_data)

    # Compute and print Loss
    loss = criterion(y_pred, y_data)
    print(epoch, loss.data[0])

    # Zero gradients, perform a backward pass, and update the weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

# After training
hour_var = Variable(torch.Tensor([[1.0]]))
print("predict 1 hour ", 1.0, model(hour_var).data[0][0] > 0.5)
hour_var = Variable(torch.Tensor([[7.0]]))
print("predict 7 hours ", 7.0, model(hour_var).data[0][0] > 0.5)
```

3

Training cycle
(forward, backward, update)



Training Options

Options

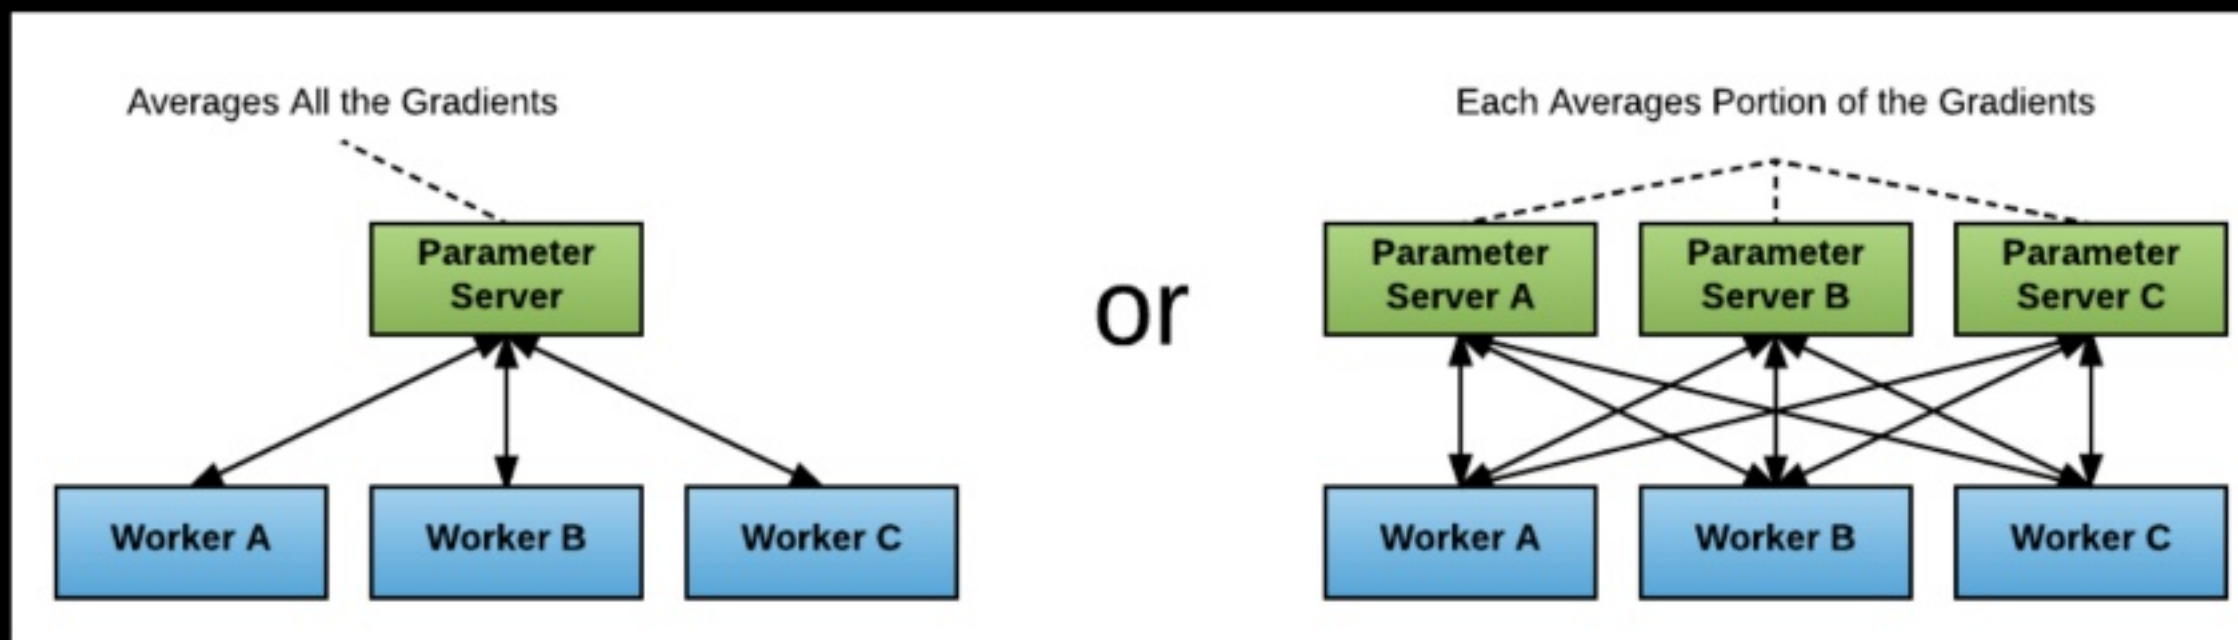
- 1) Train on single node
- 2) Train on single node, distributed inference
- 3) Distributed training

Horovod

- Created by Alexander Sergeev of Uber, [open-sourced](#) in 2017
- Simplifies distributed neural network training
- Supports TensorFlow, Keras, and PyTorch




Classical Parameter Server



All-Reduce

Horovod Technique



Patarasuk, P., & Yuan, X. (2009). Bandwidth optimal all-reduce algorithms for clusters of workstations. *Journal of Parallel and Distributed Computing*, 69(2), 117-124. doi:10.1016/j.jpdc.2008.09.002

UBER

data science | data engineering

Minimal Code Change

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))


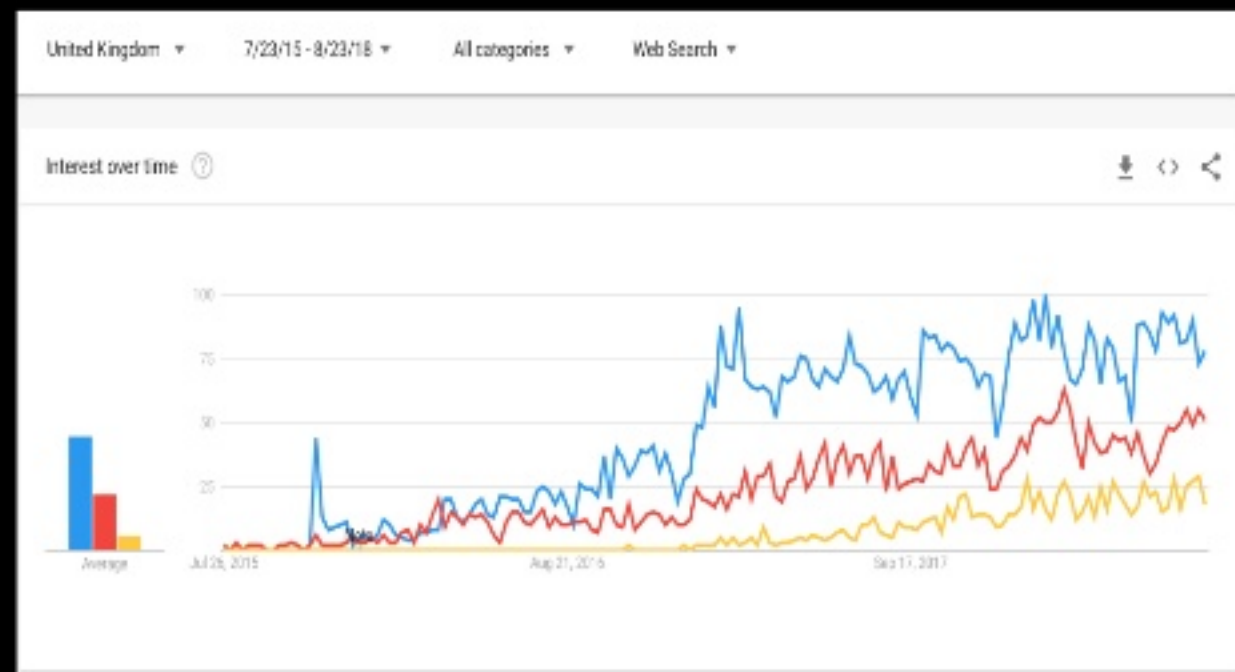

# Horovod: adjust learning rate based on number of GPUs.
opt = keras.optimizers.Adadelta(1.0 * hvd.size())

# Horovod: add Horovod Distributed Optimizer.
opt = hvd.DistributedOptimizer(opt)

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=opt,
              metrics=['accuracy'])
```


TensorFlow, Keras, or
PyTorch?

Takeaways: Gaining Momentum...



Keras
TensorFlow
PyTorch

Takeaways: When to Use TF, Keras or PyTorch

- Low-level APIs & Control
- Model Serving
- Supports multiple languages

TensorFlow

- High-level APIs
- Multiple Backends
- Love Python
- Rapid Experimentation

Keras

- *Pythonic!*
- Imperative Programming
- Rapid Experimentation

PyTorch

Databricks Runtime for Machine Learning

Ready to use clusters with built-in ML Frameworks

including TensorFlow, Keras, Horovod, and more




Horovod Estimator

for simplified distributed training on TensorFlow with Horovod using Apache Spark on Databricks

GPU support

on AWS (P2/P3) and Azure (NC/NC-v3) instances now supported!




Resources & Books

Blog posts Talk, & webinars (<http://databricks.com/blog>)

- [GPU acceleration in Databricks](#)
- [Deep Learning and Apache Spark](#)
- [fast.ai](#)
- [TensorFlow Tutorials](#)
- [TensorFlow Dev Summit](#)
- [Keras/TensorFlow Tutorials](#)
- [PyTorch Docs & Tutorials](#)
 - [Talk-1 from Soumith Chintala](#)
 - [Talk-2 from Soumith Chintala](#)
- [MLflow.org](#)


Docs for Deep Learning on Databricks (<http://docs.databricks.com>)

- [Databricks Runtime ML](#)
- [HorovodEstimator](#)



fast.ai

Making neural nets
uncool again



Thank You!
Questions?

brooke@databricks.com
[@2twitme](mailto:jules@databricks.com)