

How to Avoid Drowning in Logs

Streaming 180 Billion Events/Day and
Batching 150 TB/Hour

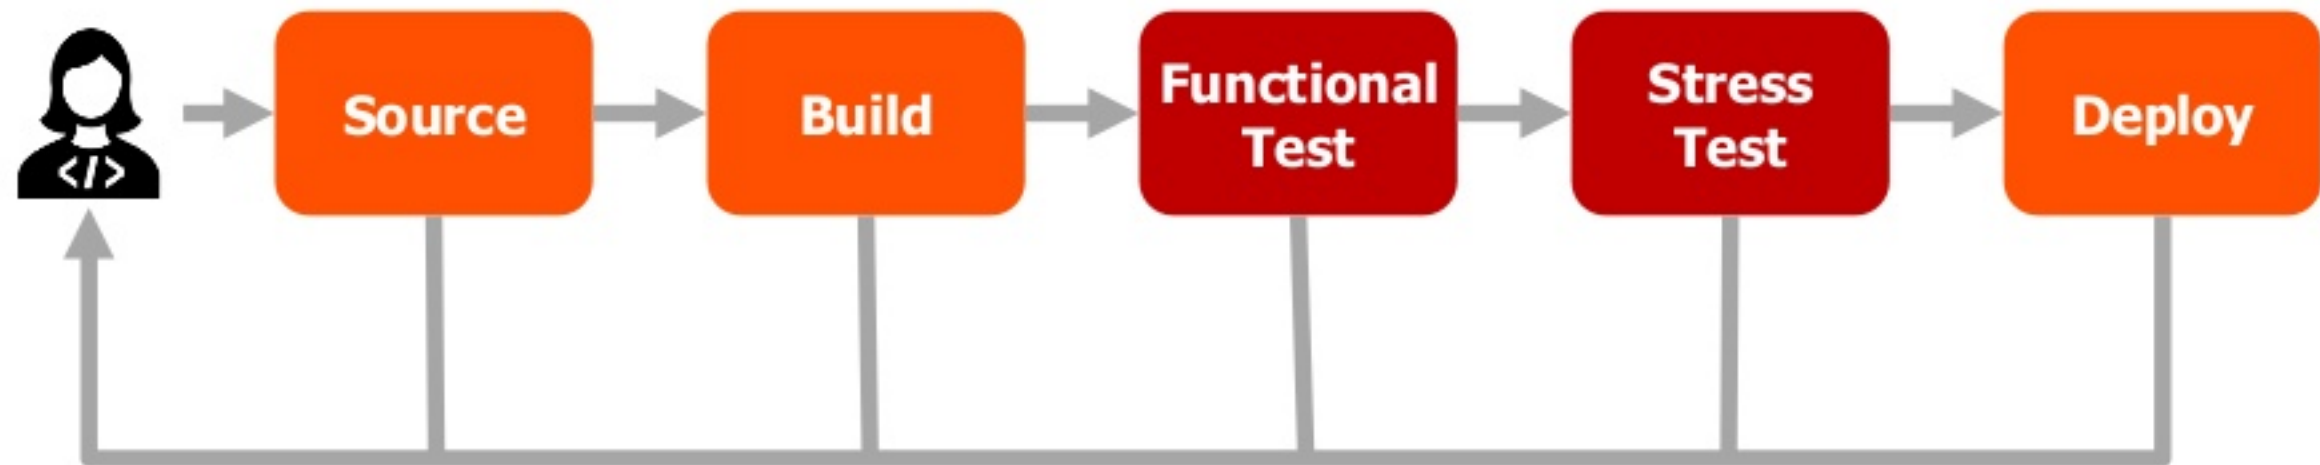
Joshua Robinson

Founding Engineer, FlashBlade


Log Analytics Pipeline in Numbers

- ✓ **2M** events / second
- ✓ **5** seconds SLA
- ✓ **0.5 - 1 PB** of data / day

Continuous Integration & Continuous Deployment



CI/CD works!




Scale Problems



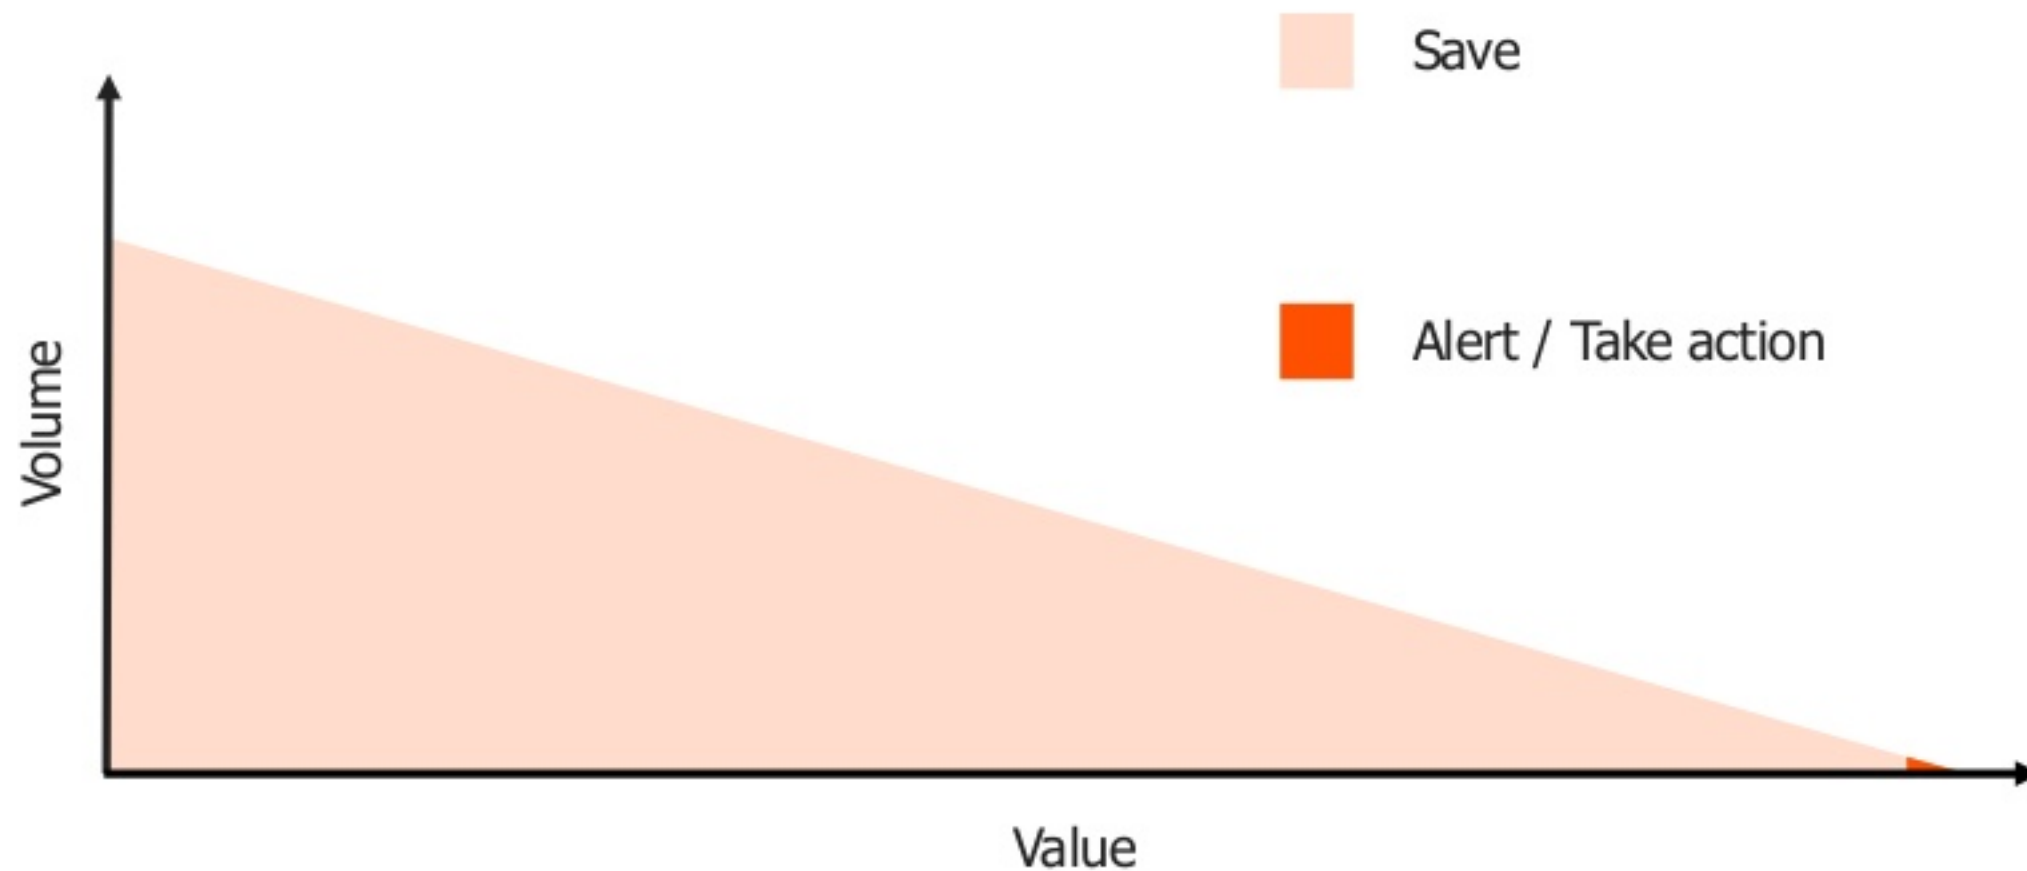
Log Analysis Dream

1. Automate triaging of failures
2. Extract performance metrics
3. Save our logs for future use
4. Do all of this in a scalable system
5. Real-time results!


Log Analysis



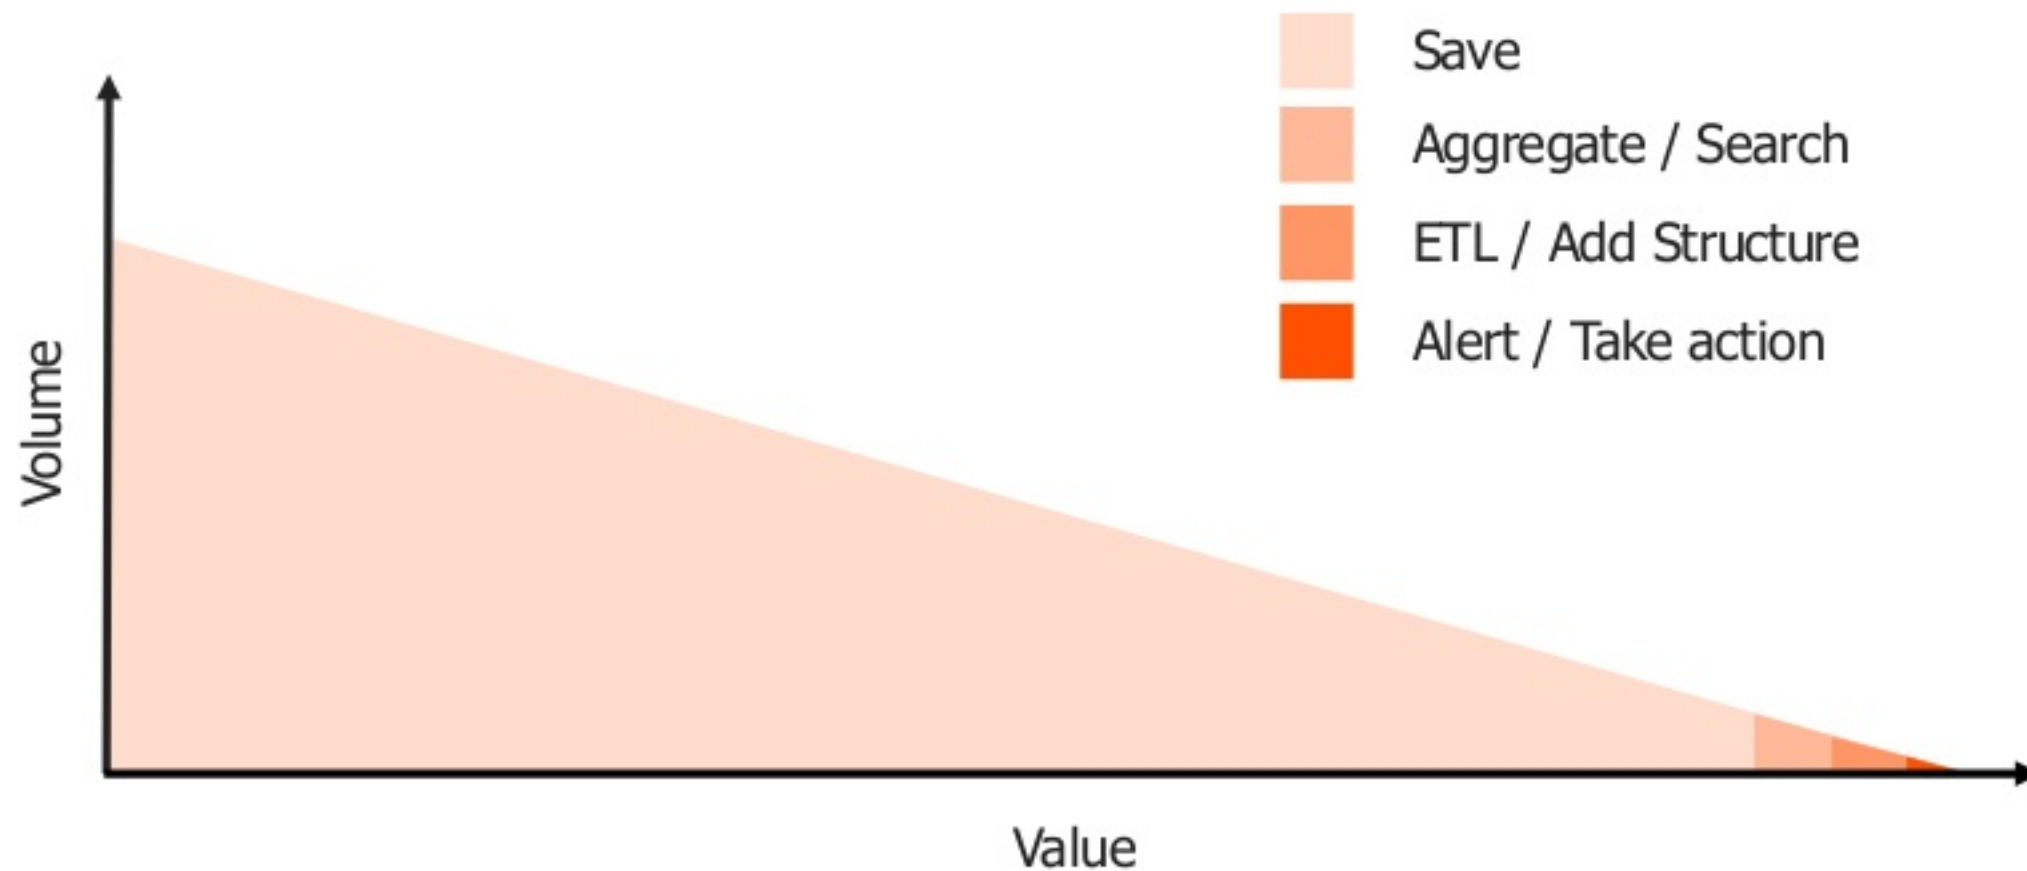
Log Analysis v1




Log Analysis v2



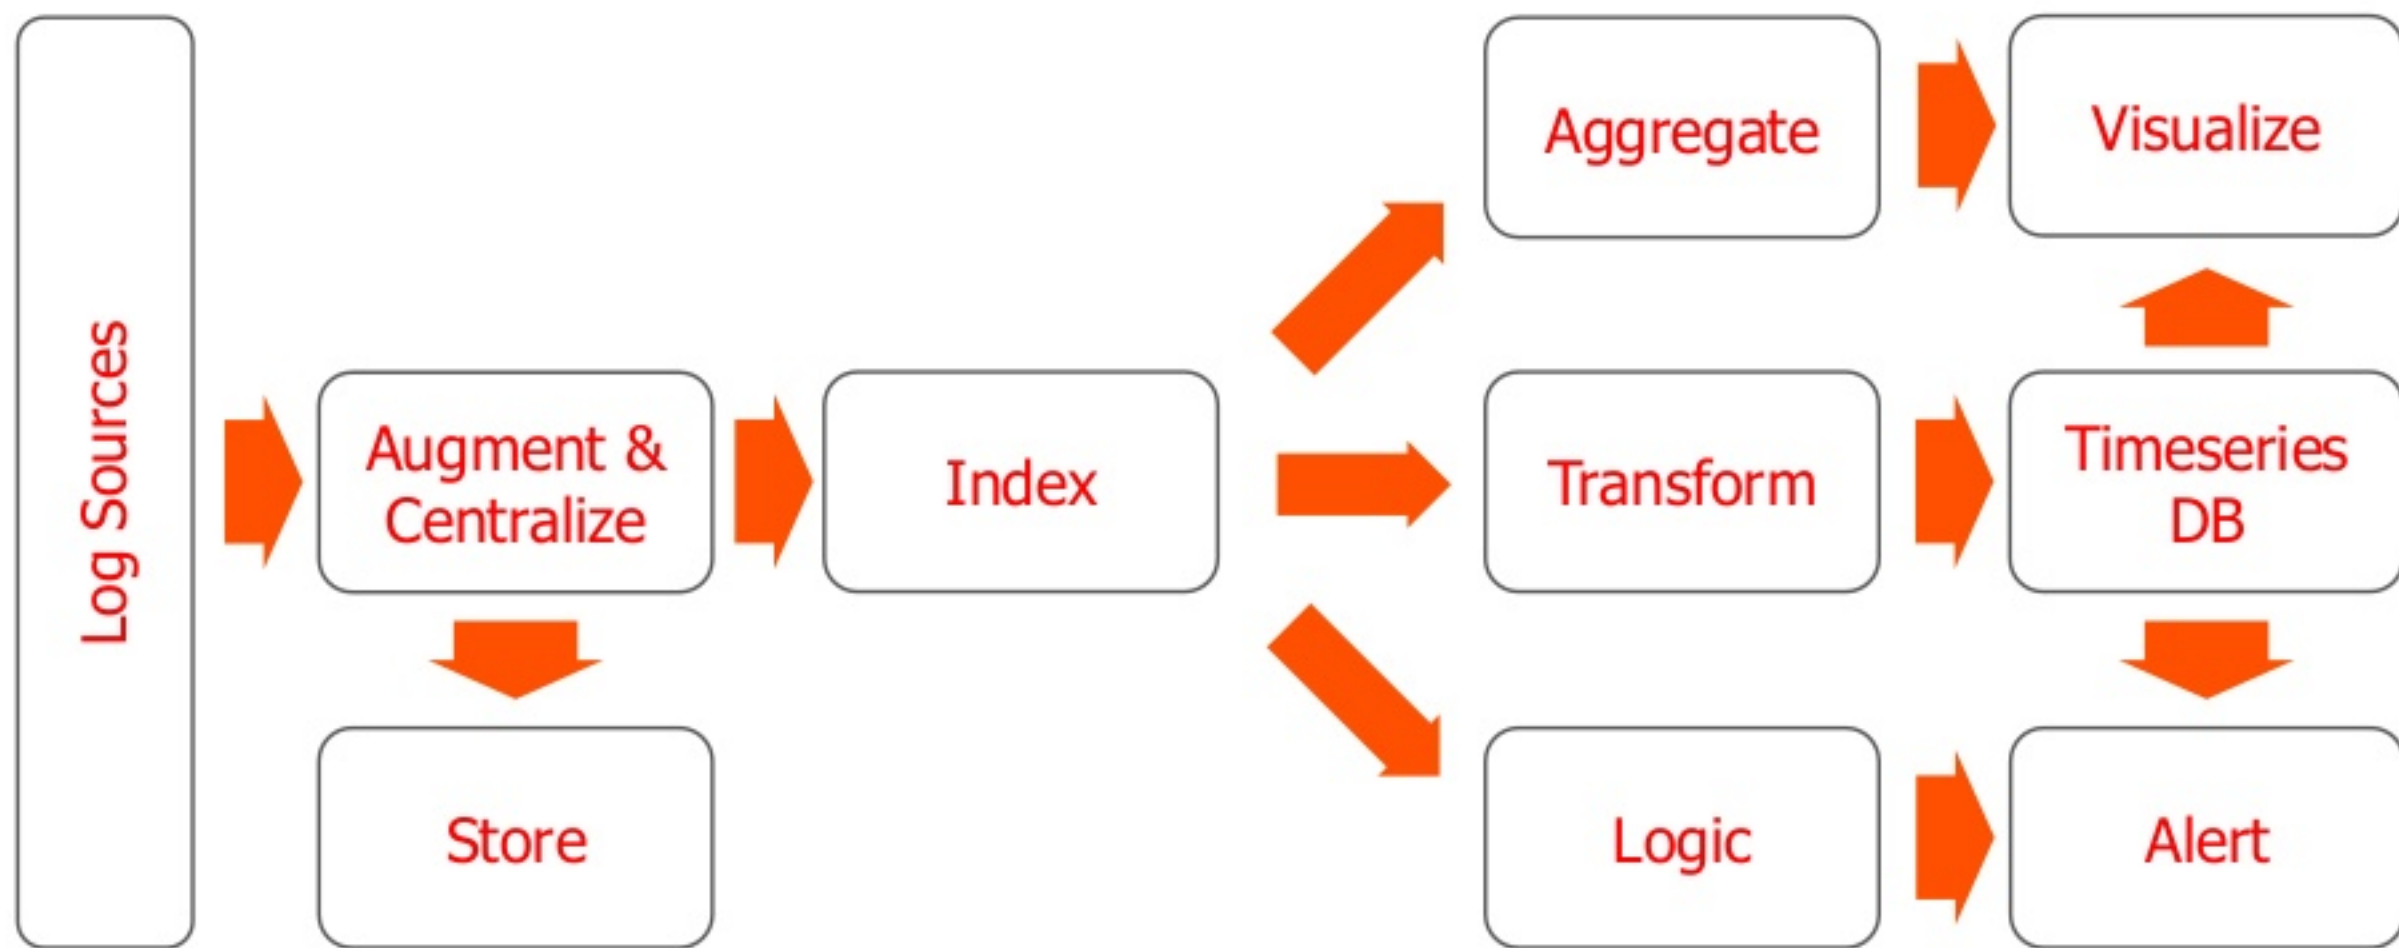
Log Analysis v3



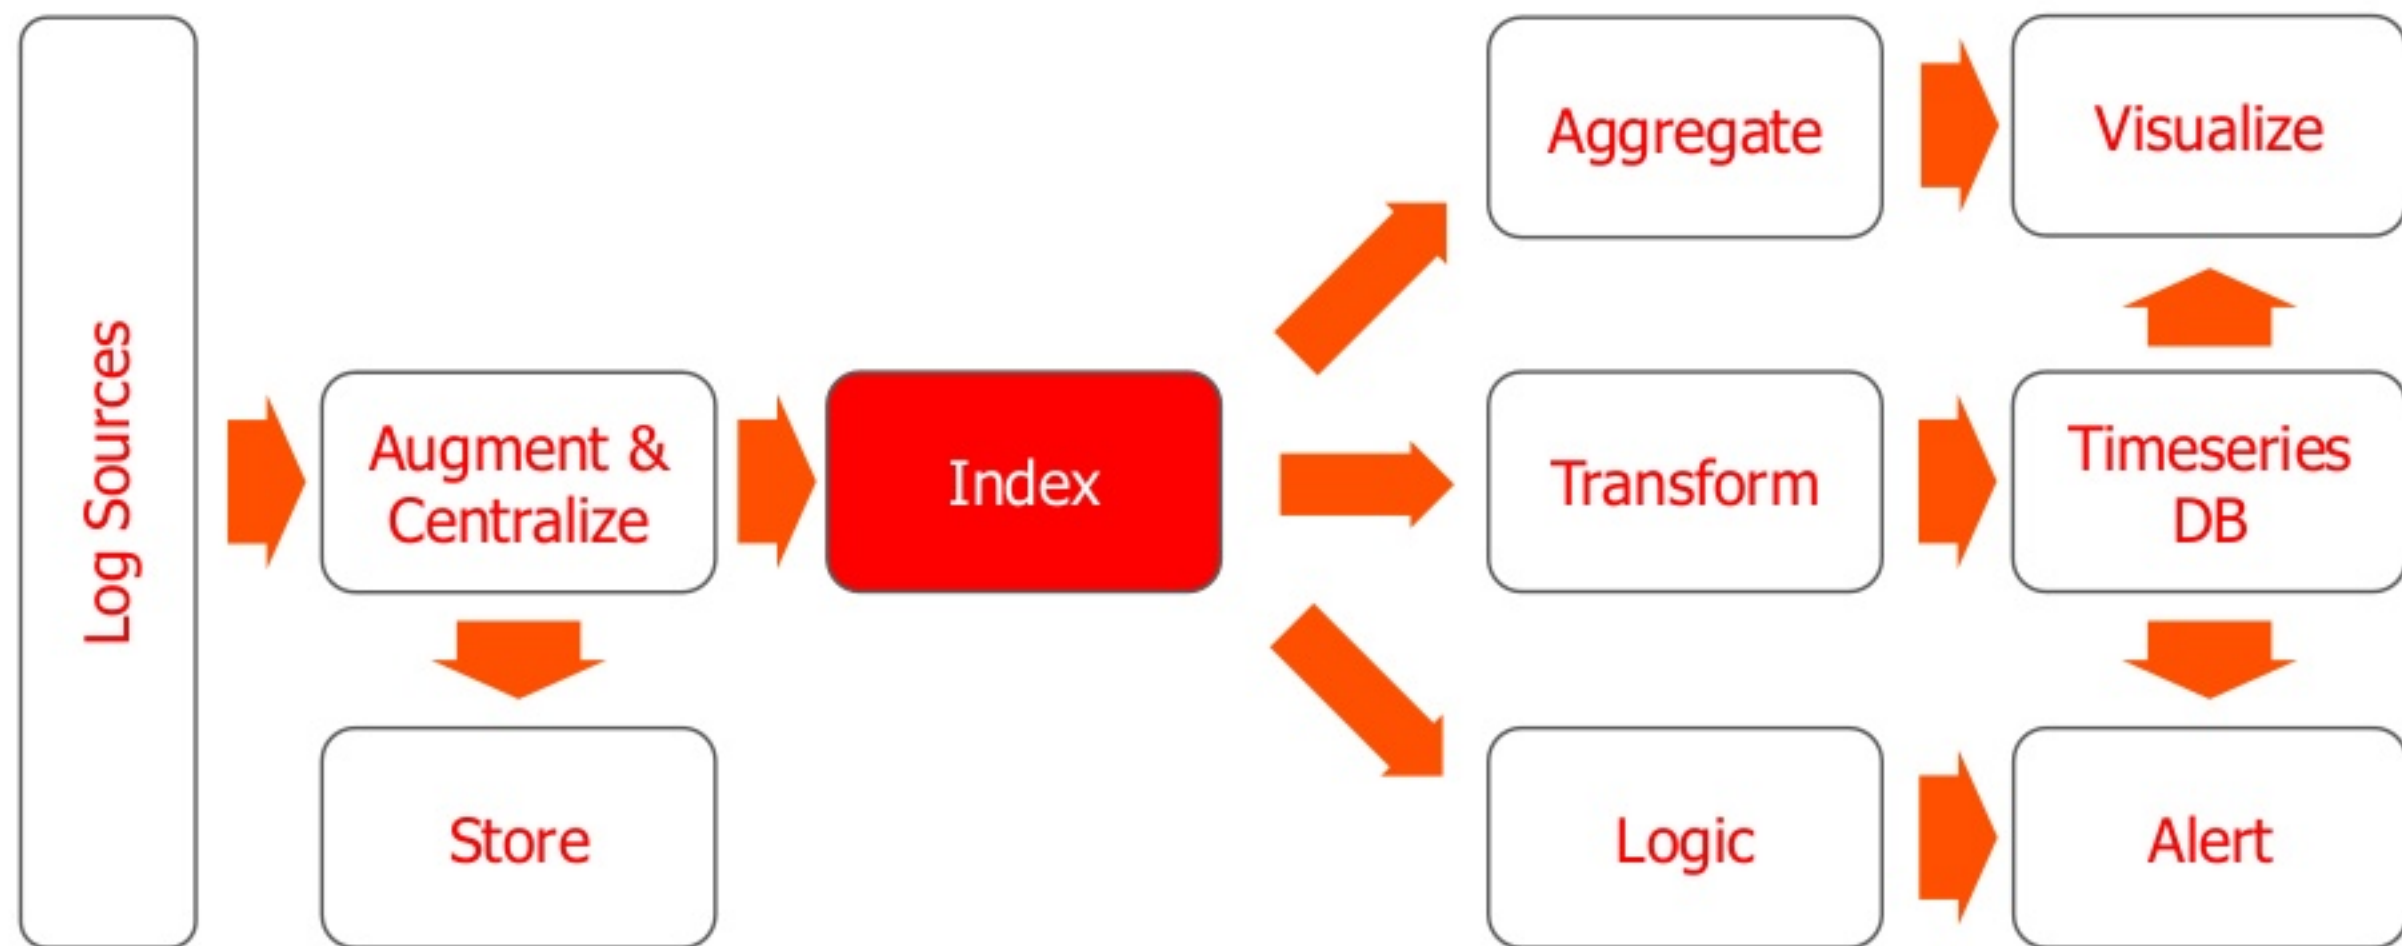
Log Analysis v10



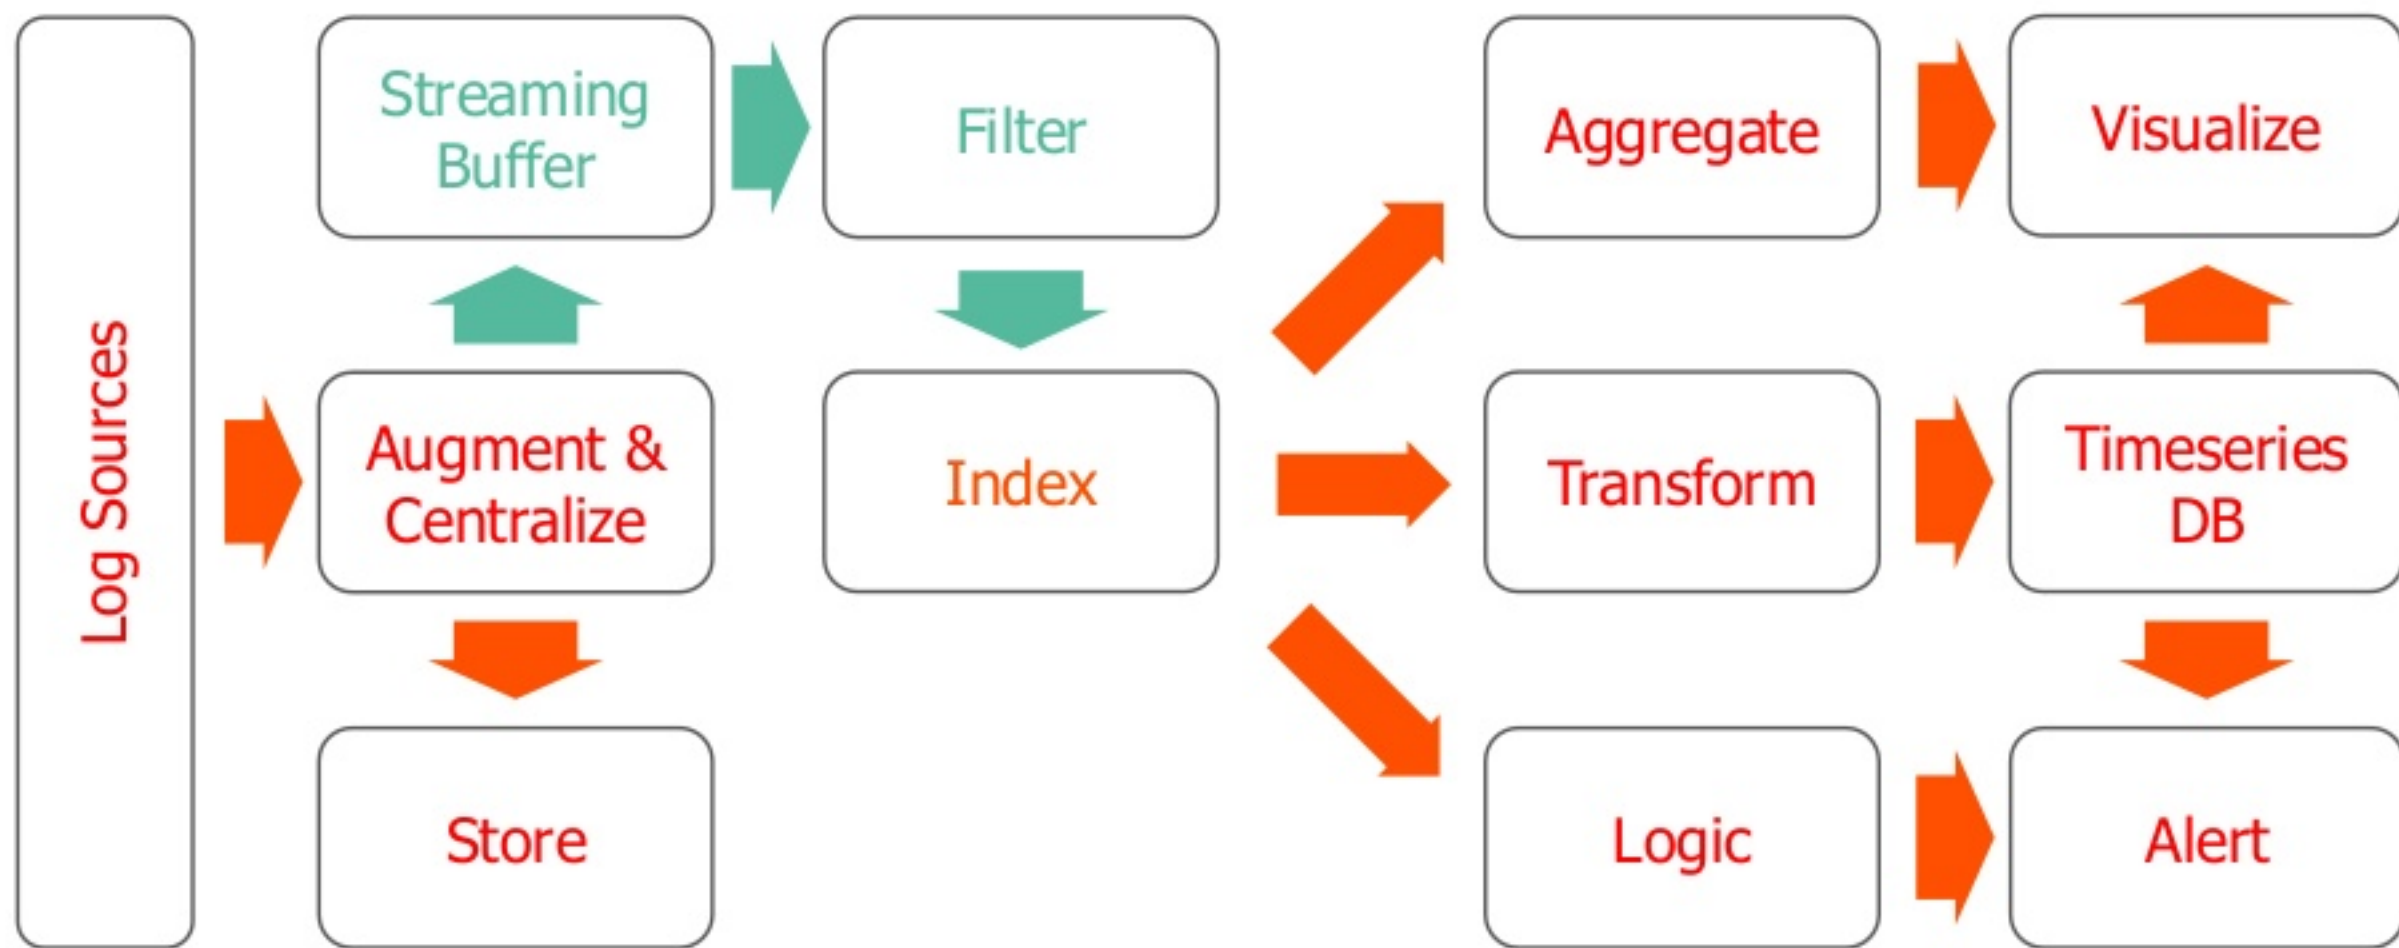
Log Analysis Pipeline



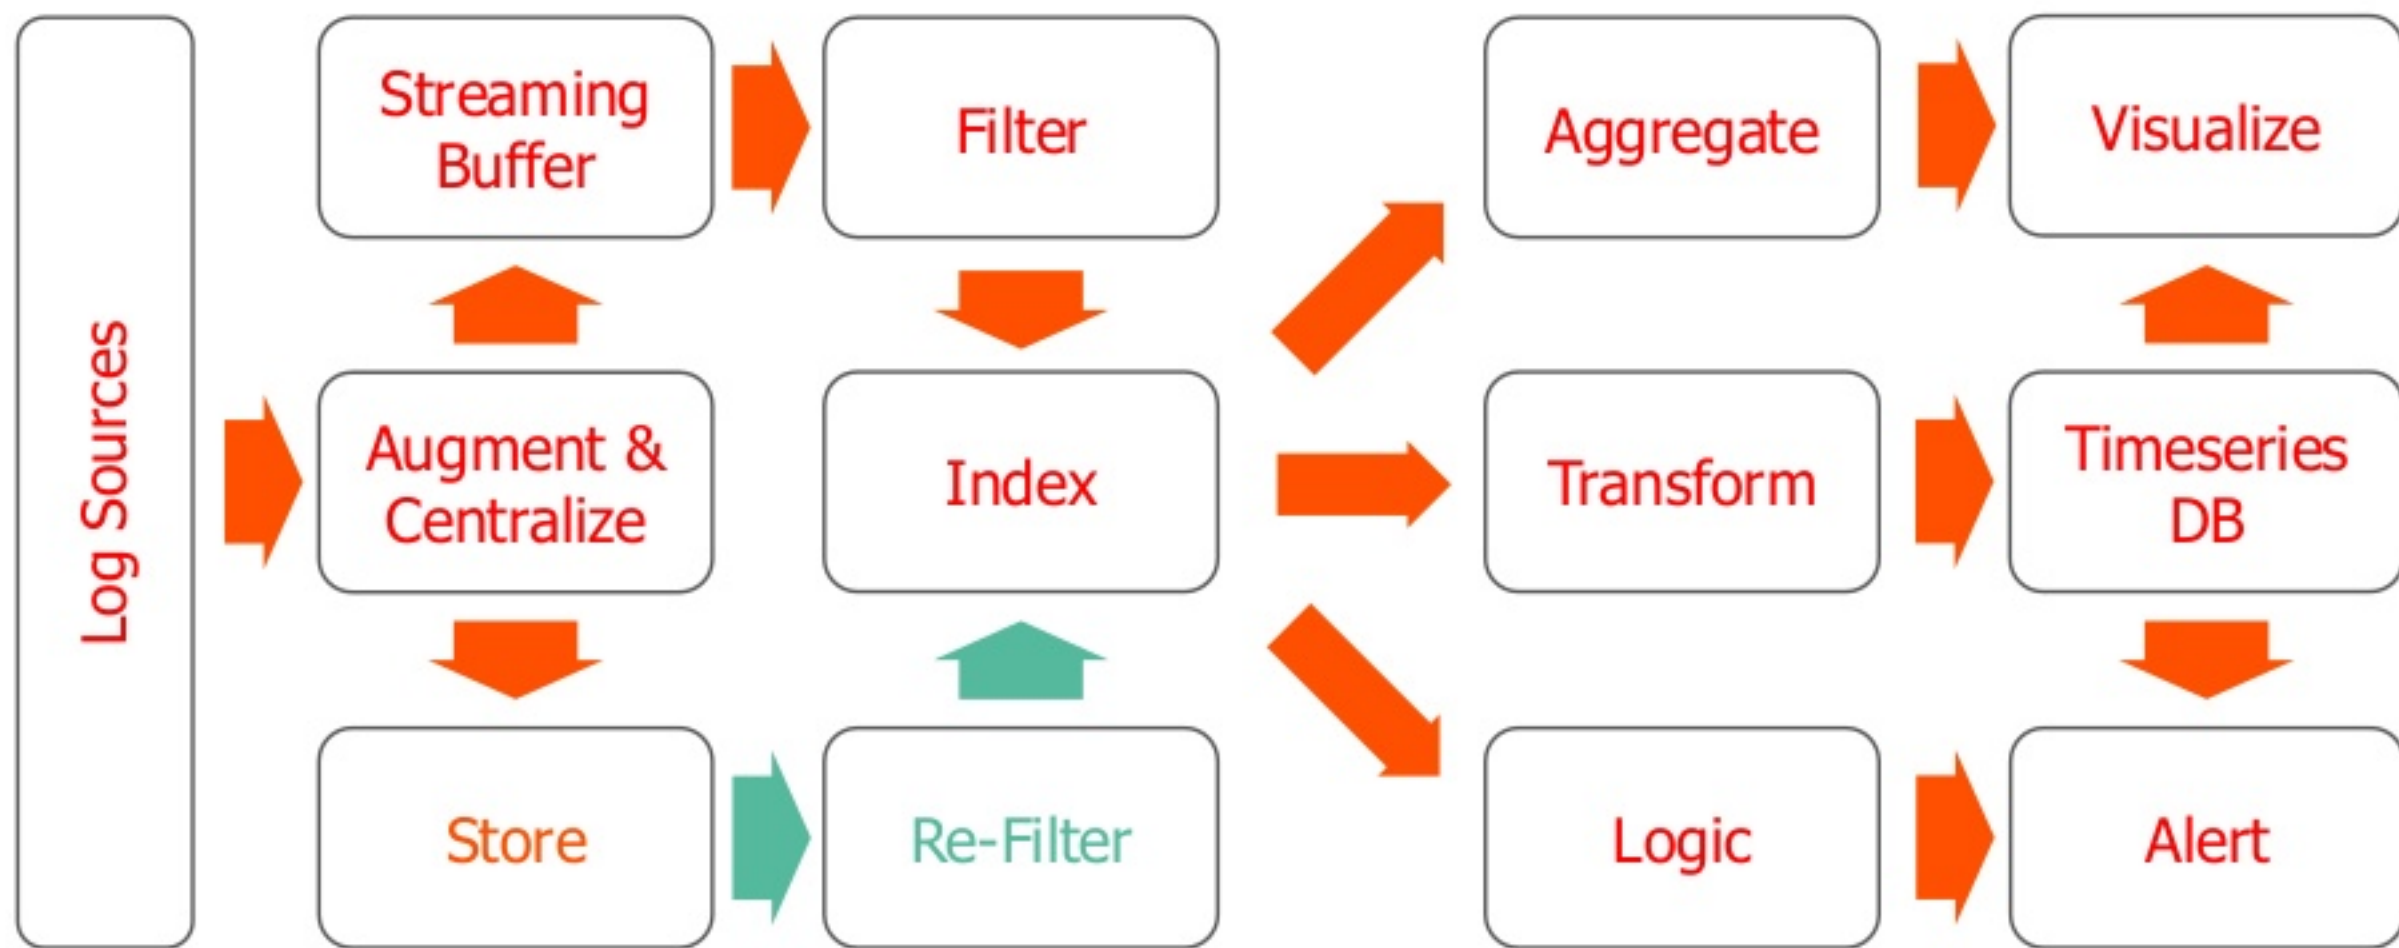
Log Analysis Pipeline



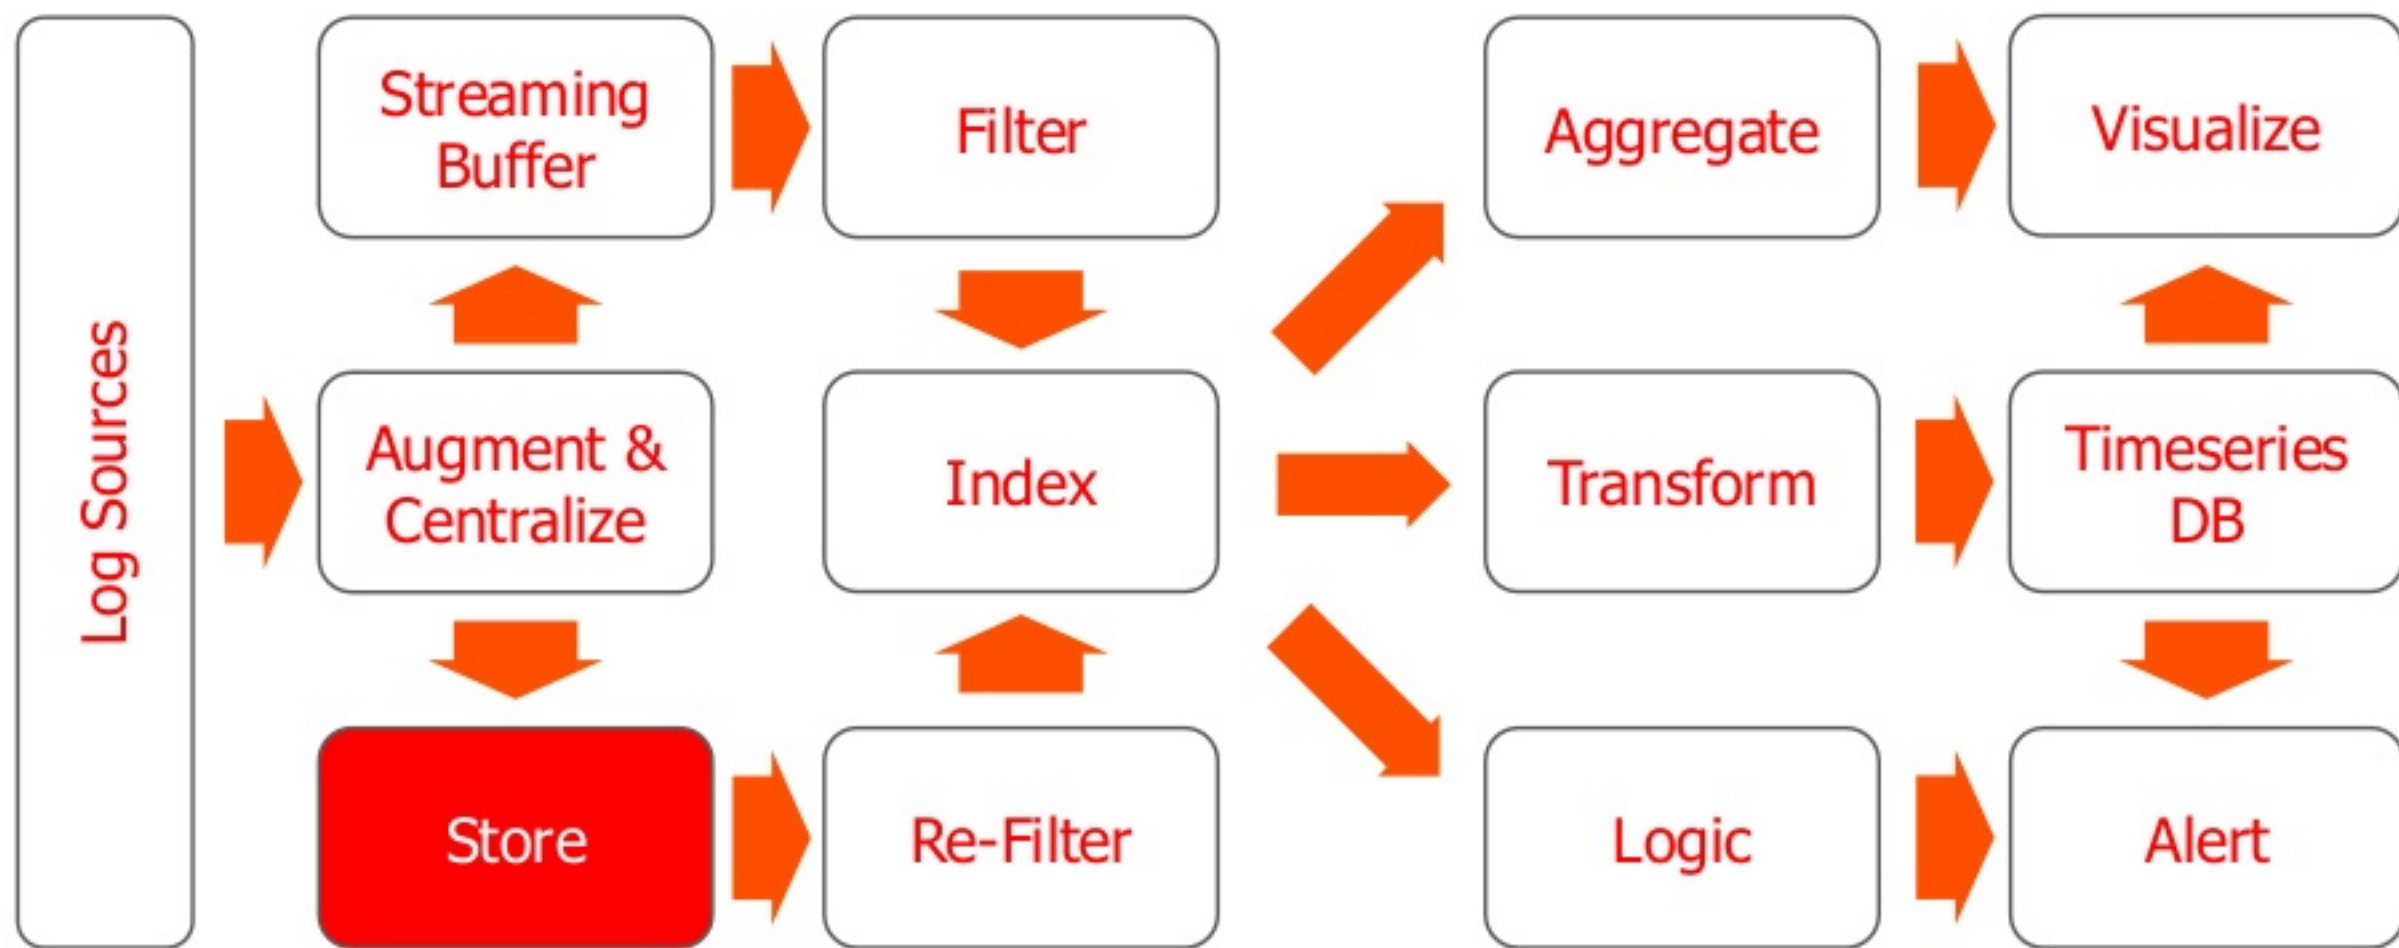
Log Analysis Pipeline



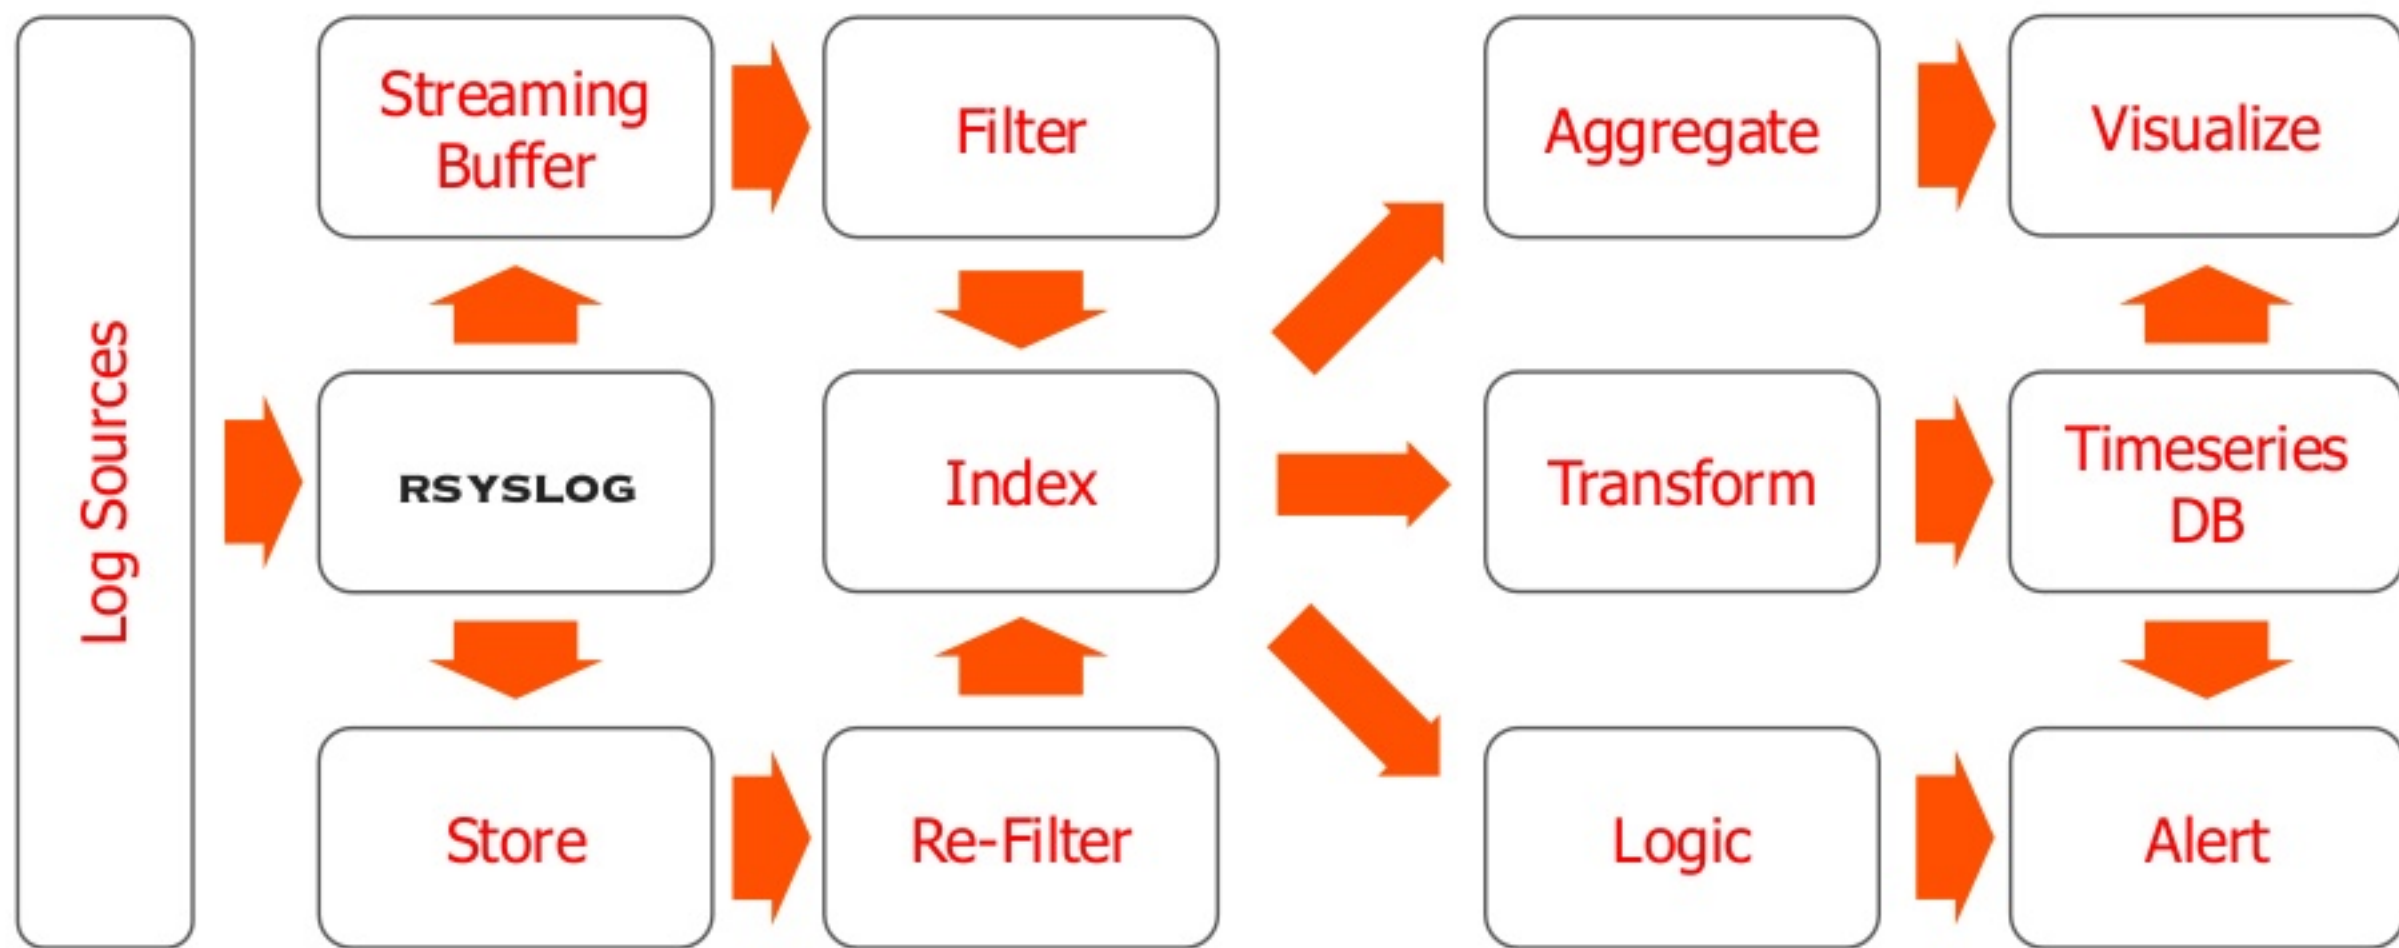
Log Analysis Pipeline



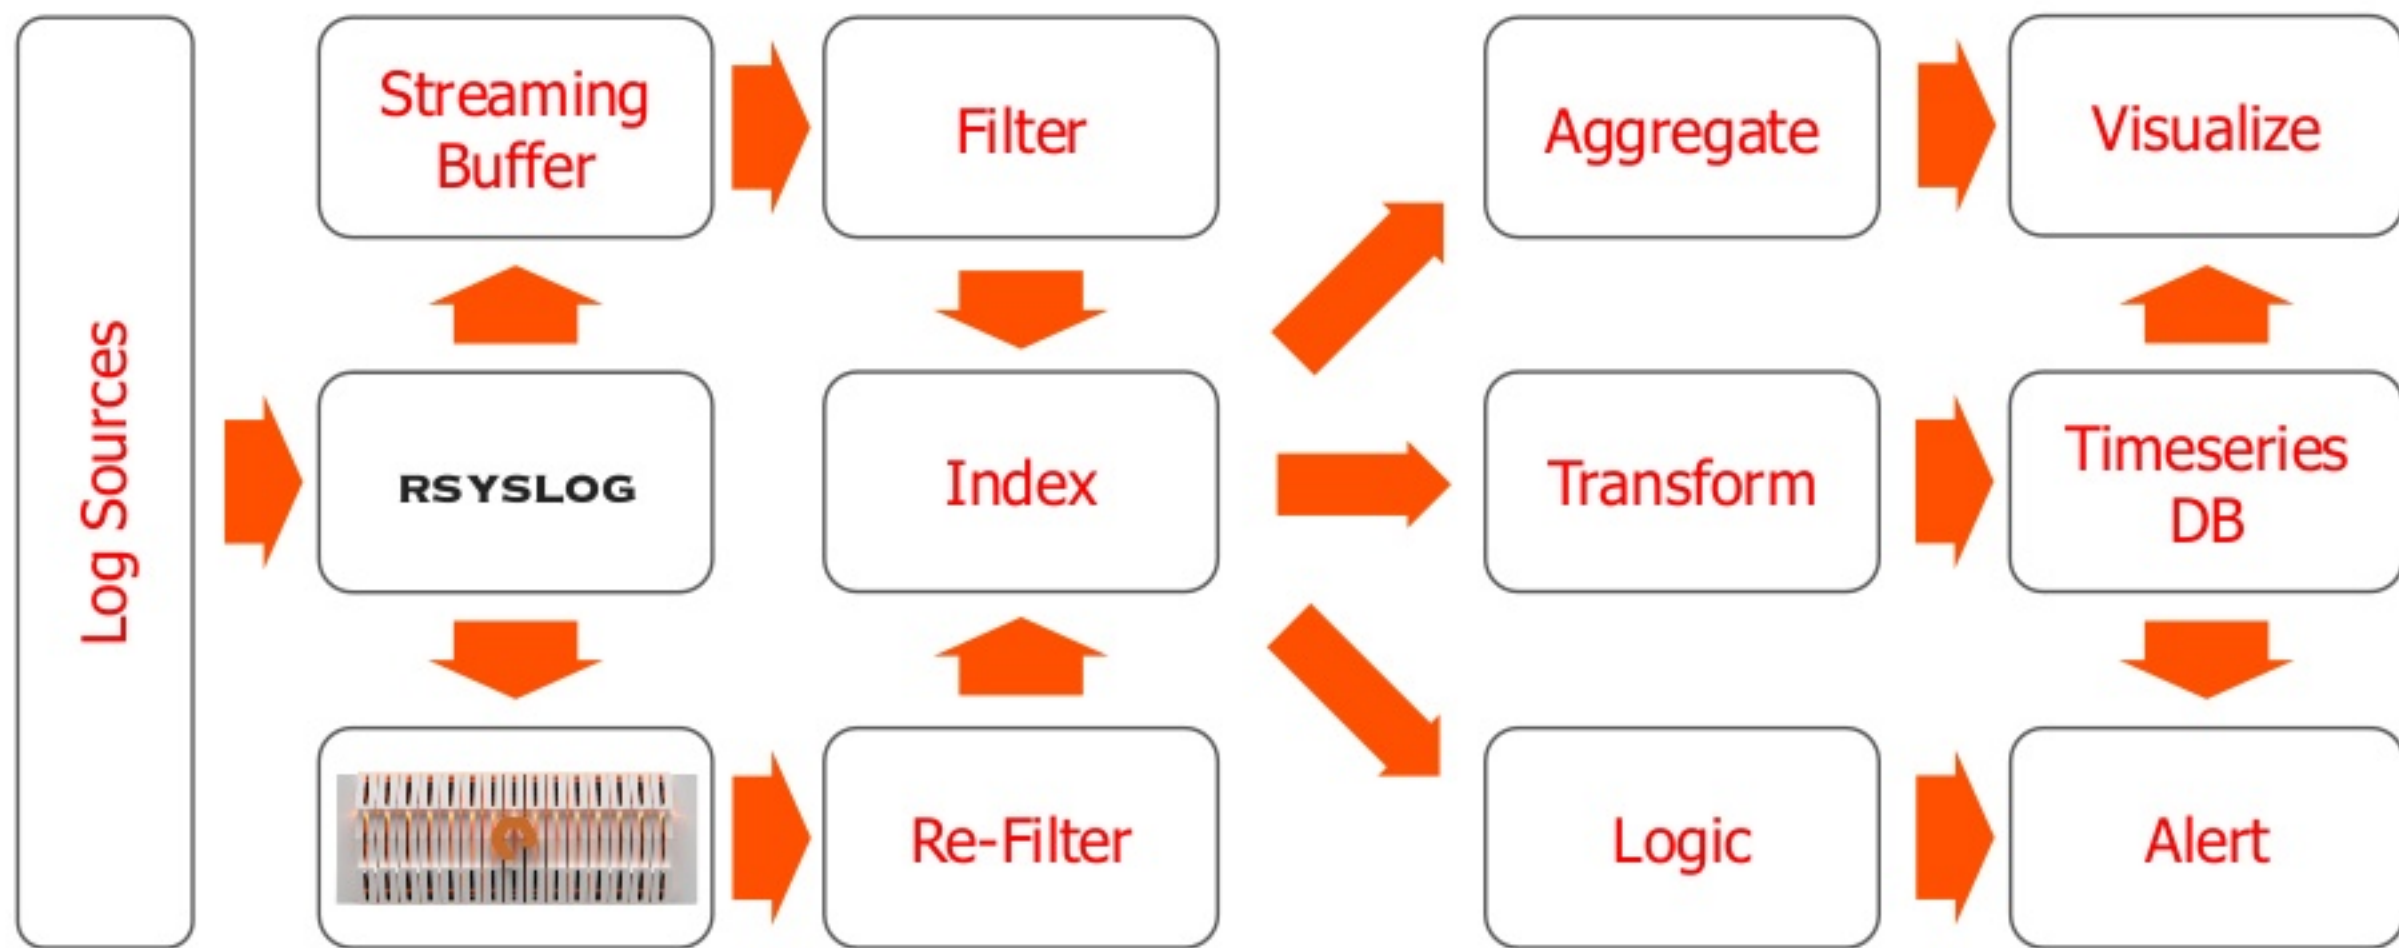
Log Analysis Pipeline



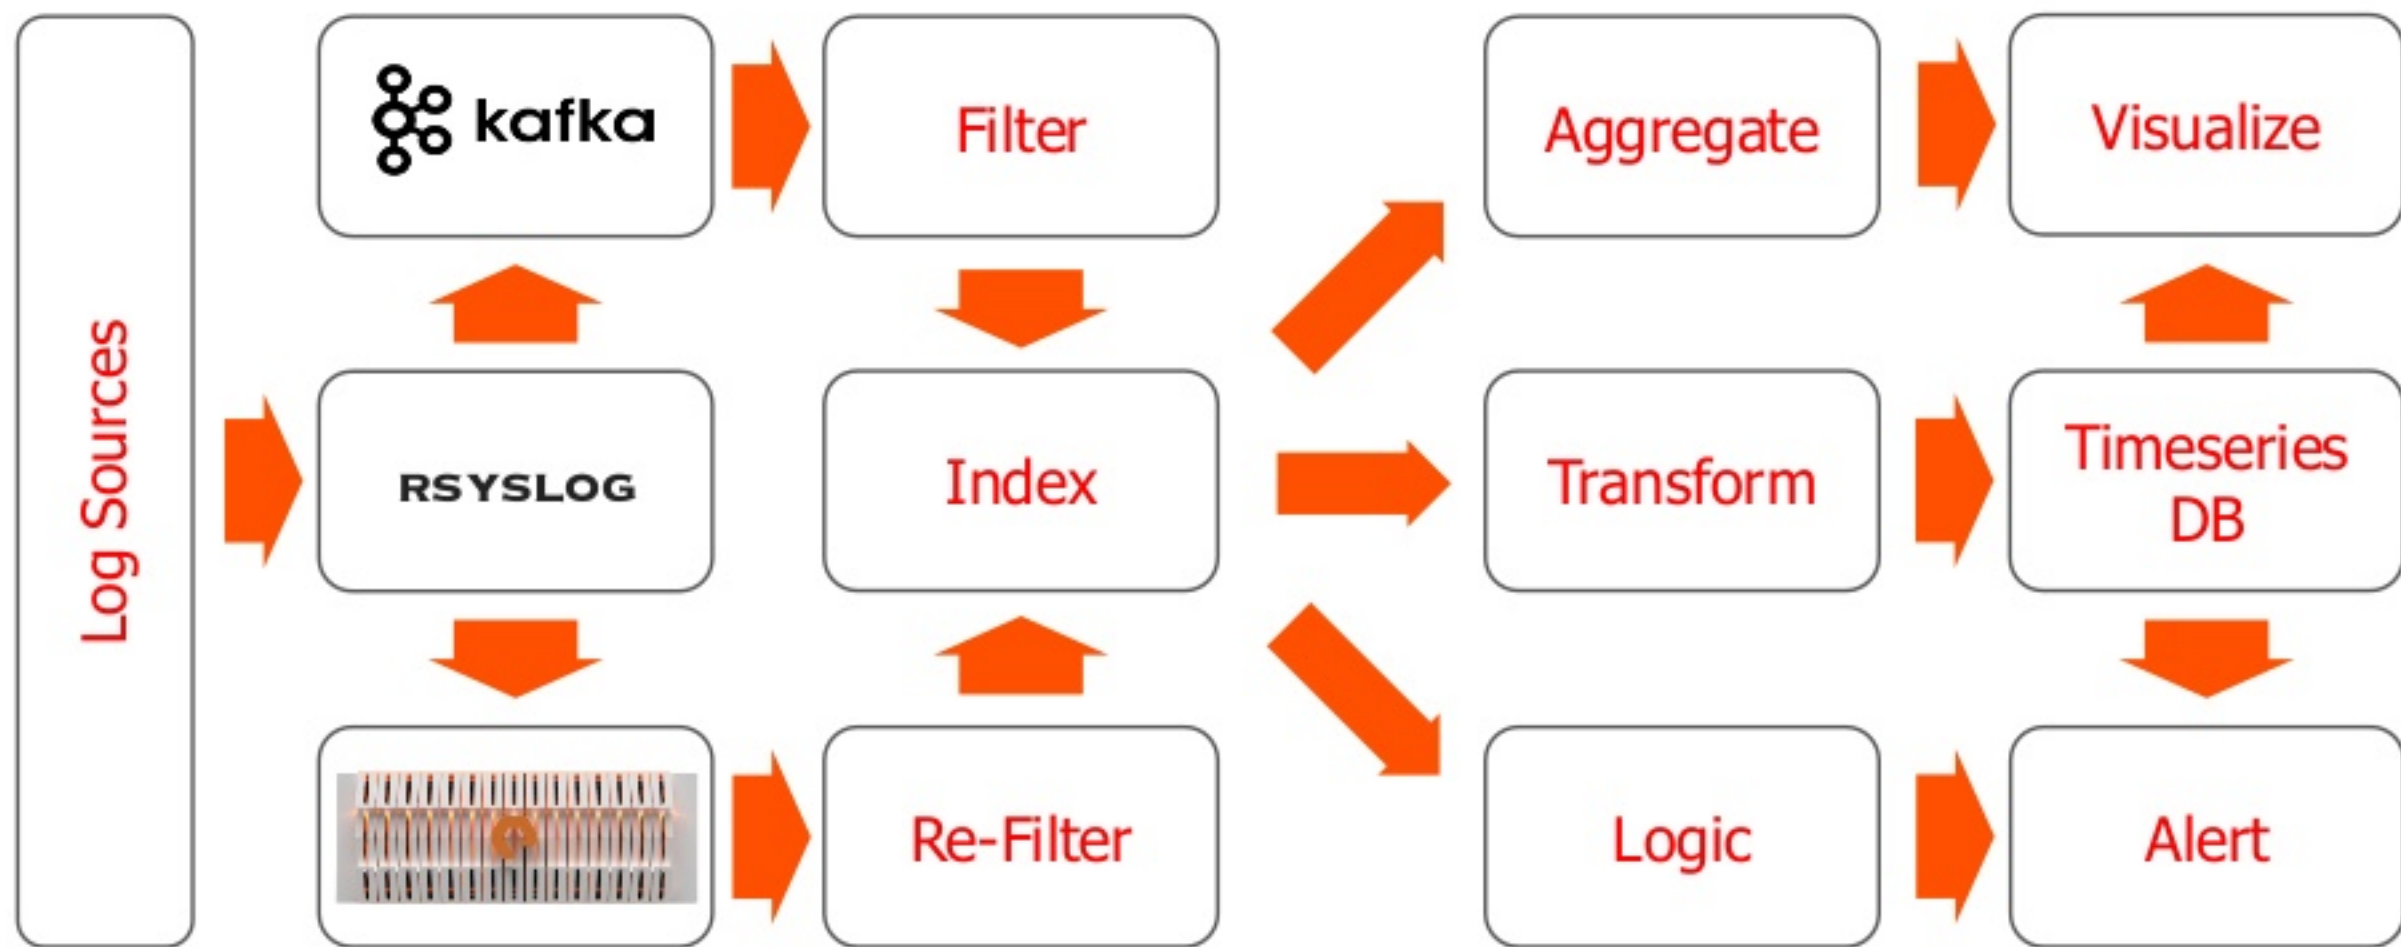
Log Analysis Pipeline



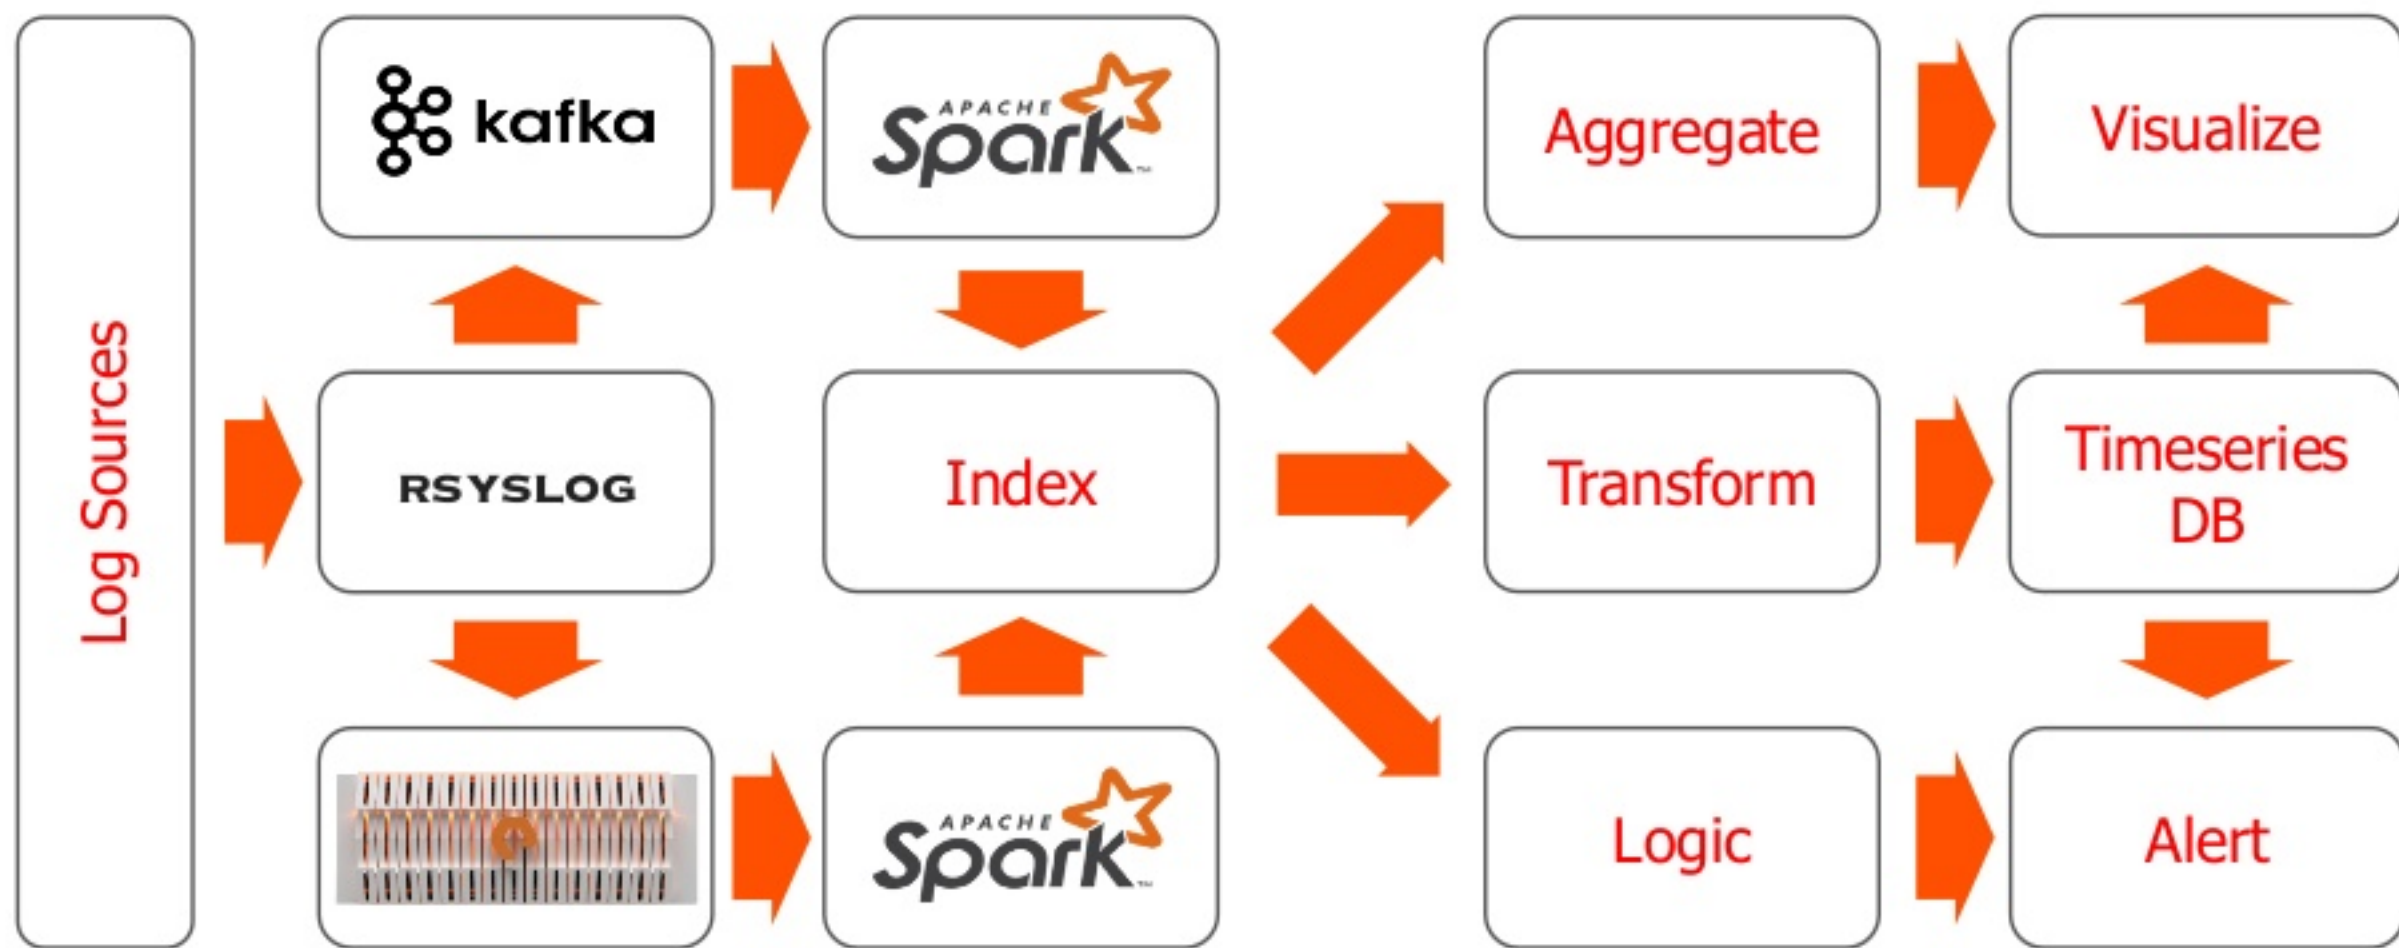
Log Analysis Pipeline



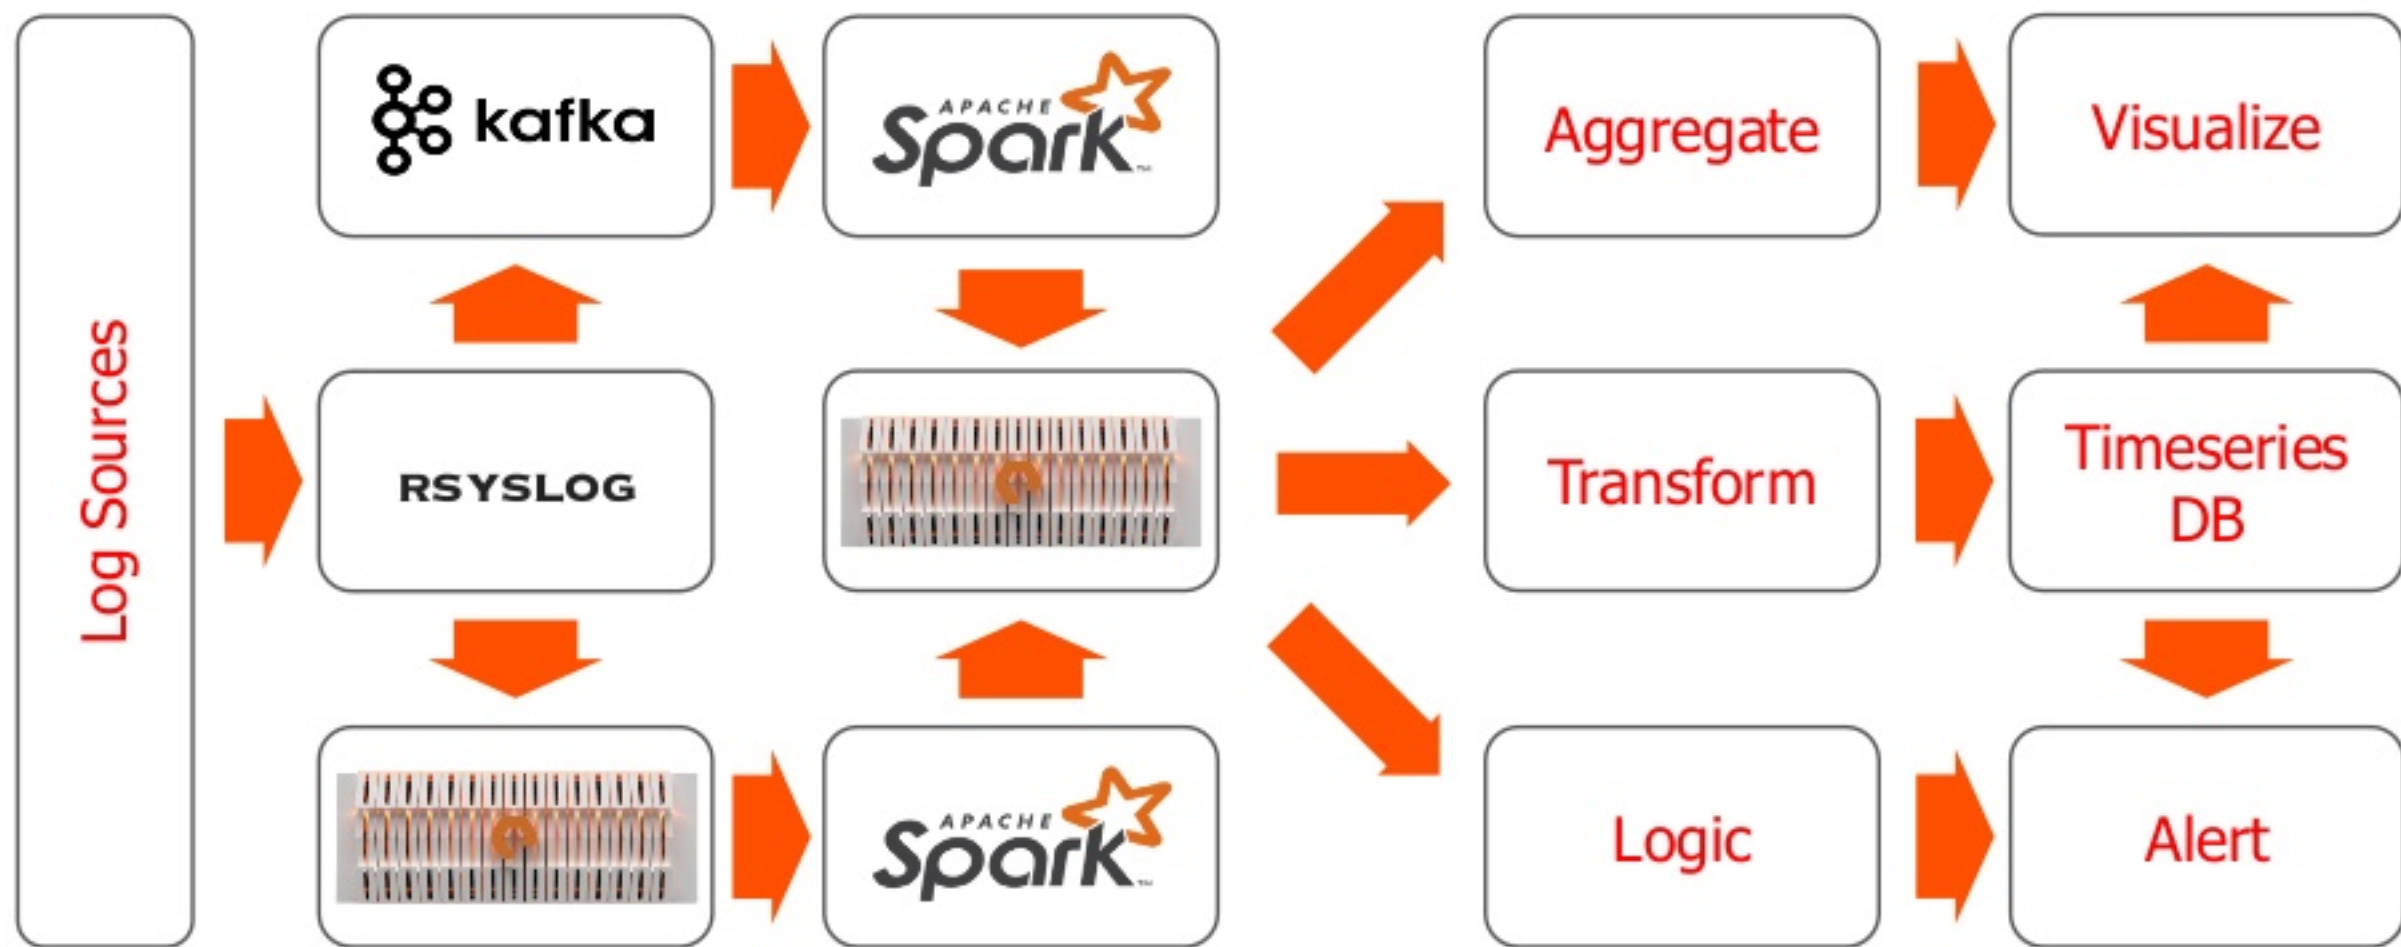
Log Analysis Pipeline



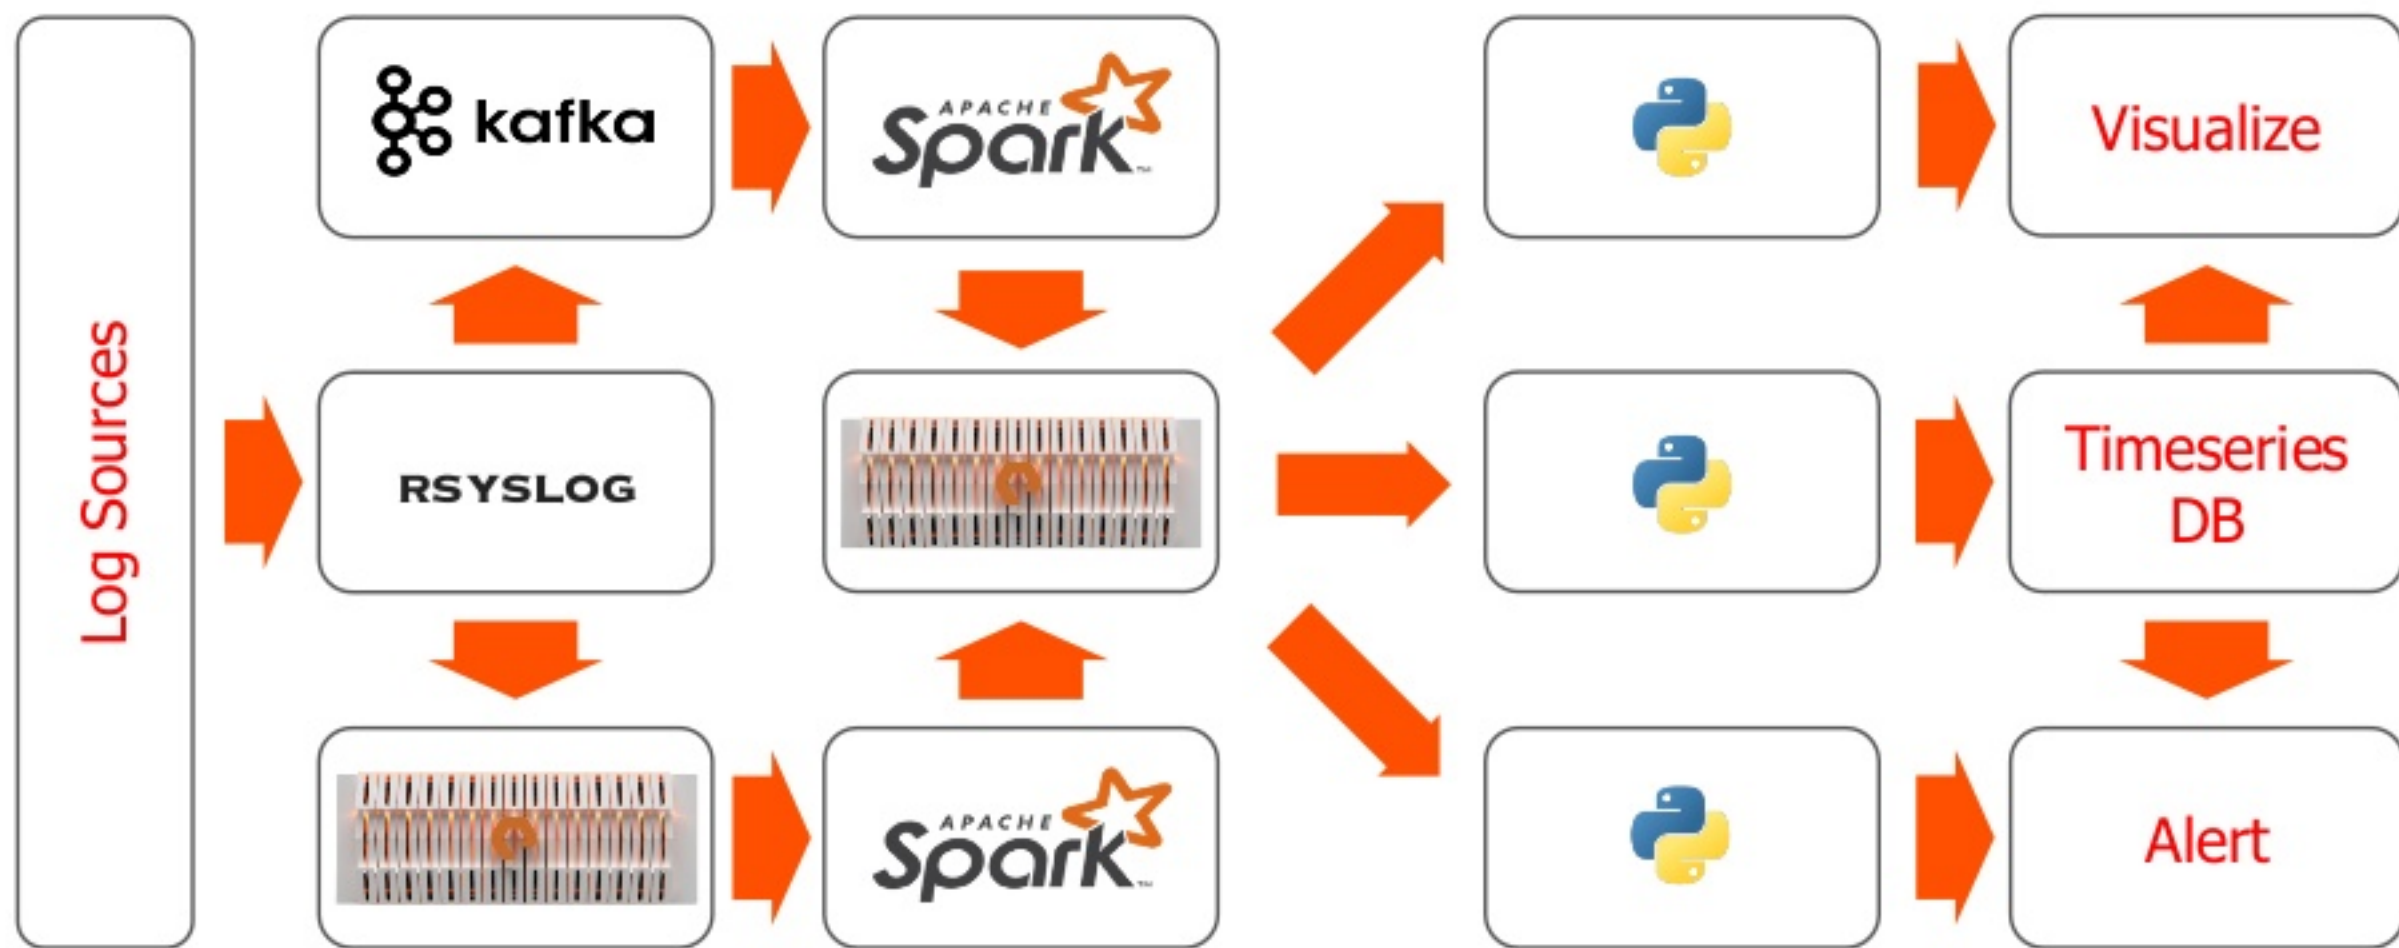
Log Analysis Pipeline



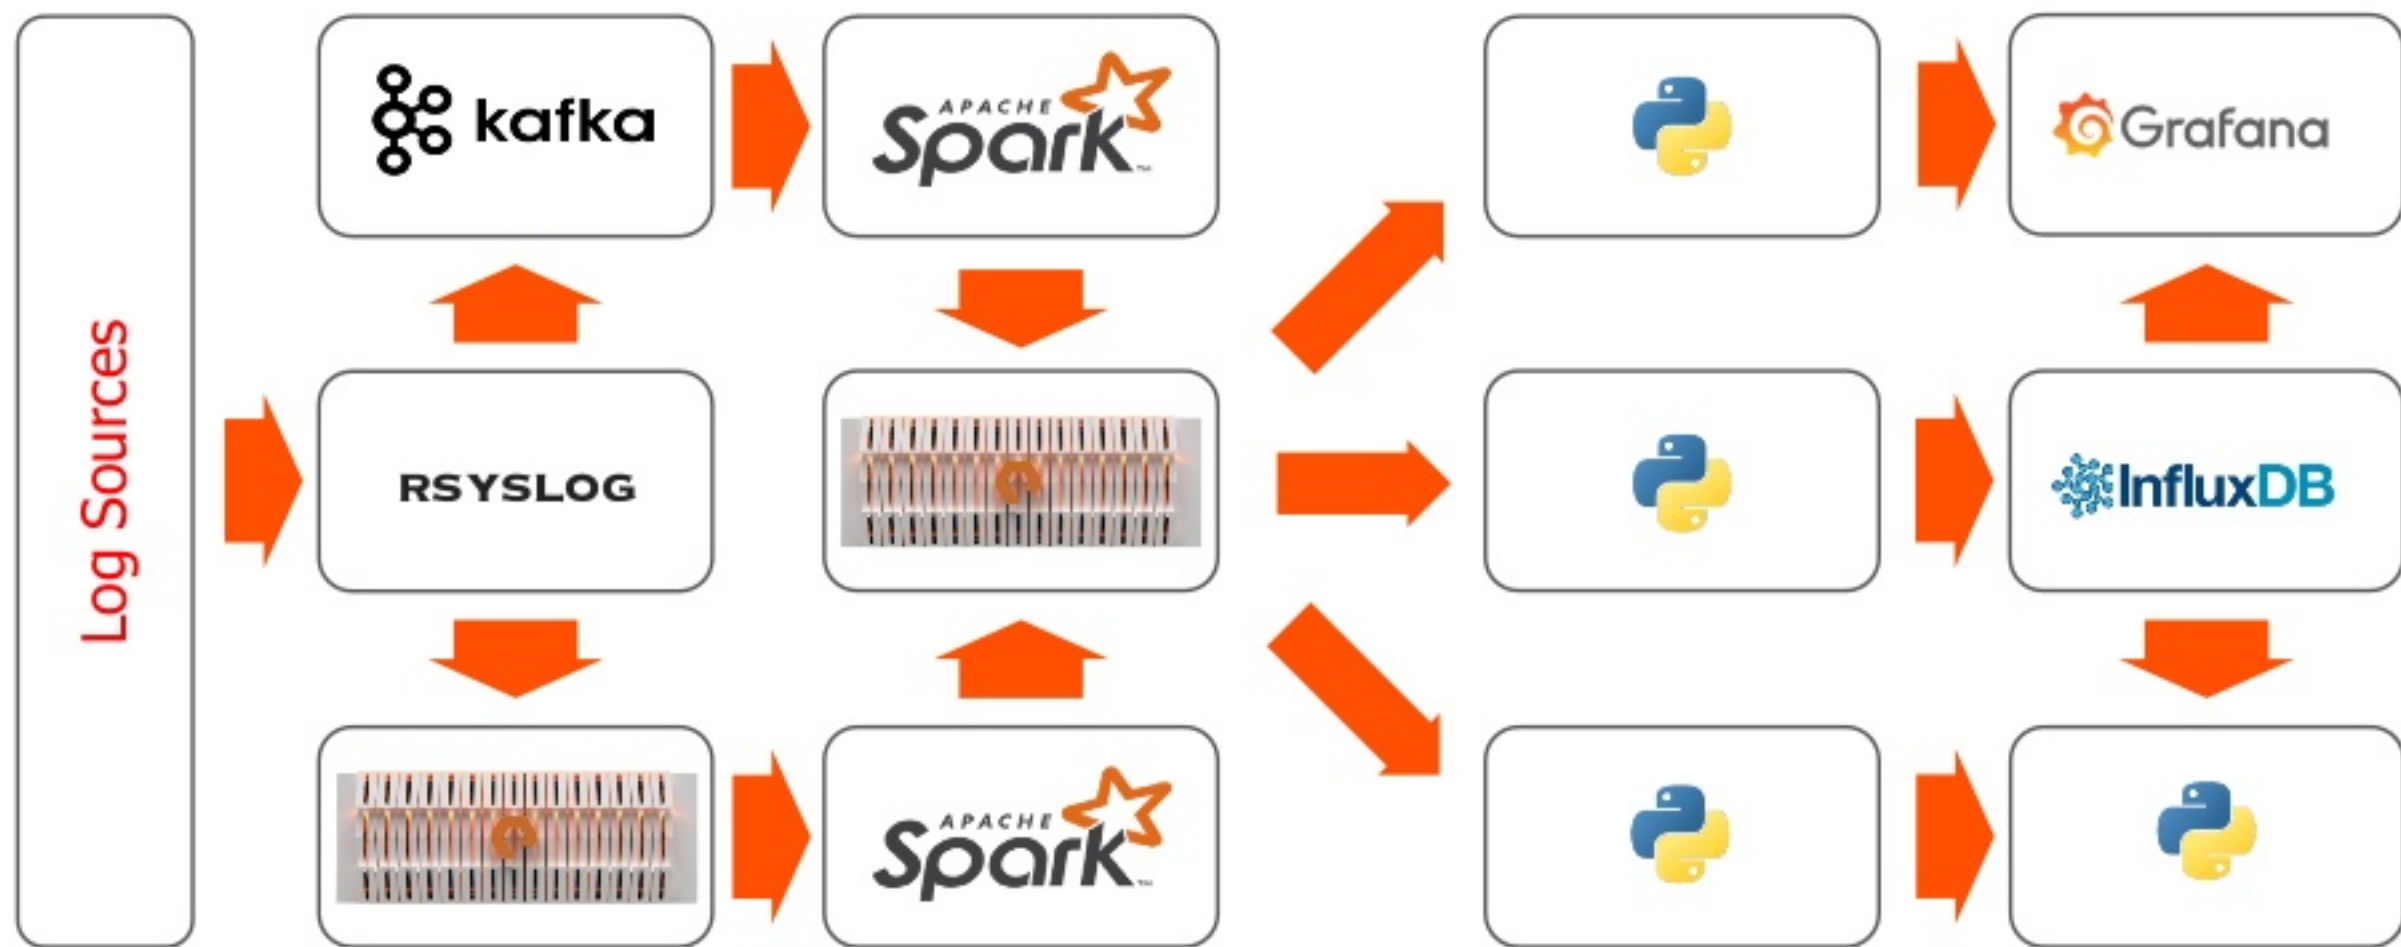
Log Analysis Pipeline



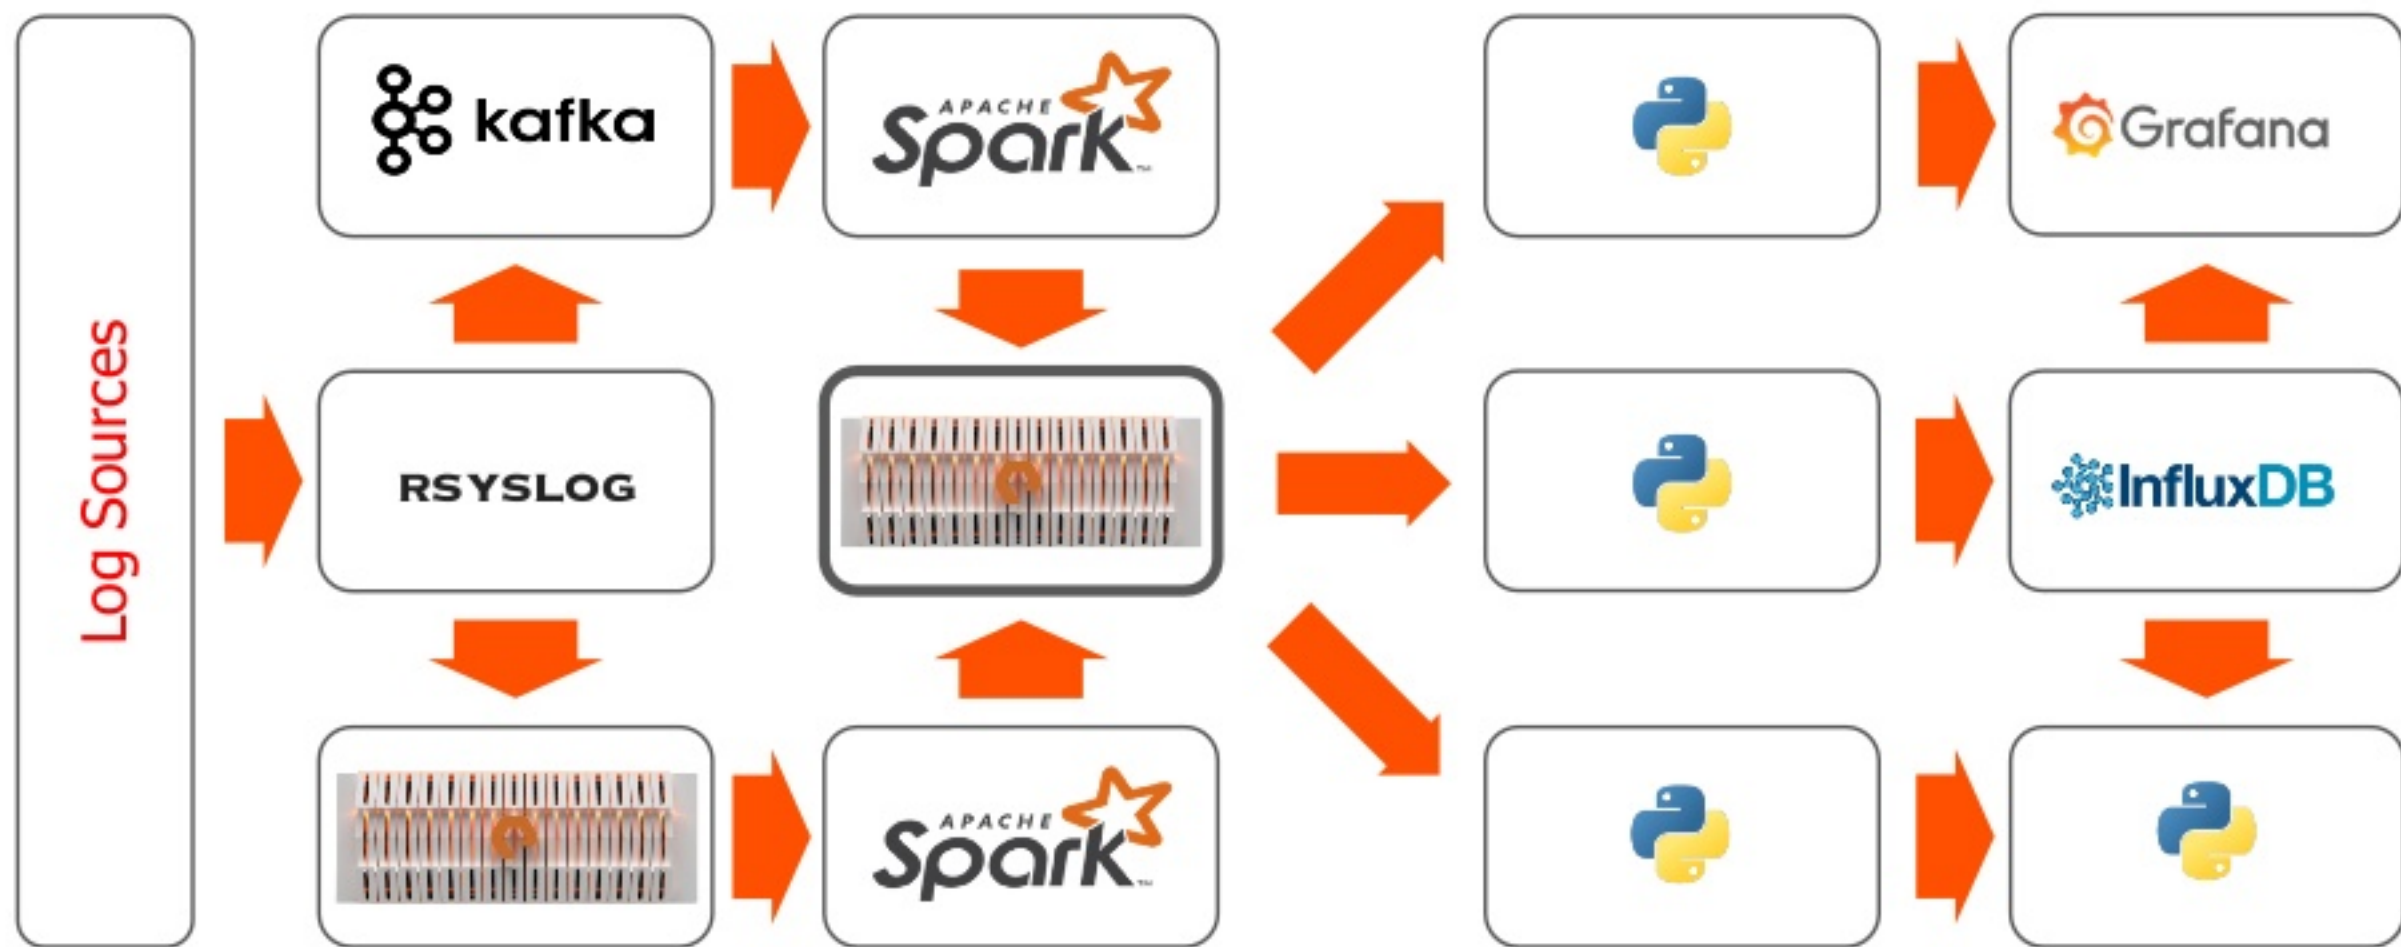
Log Analysis Pipeline



Log Analysis Pipeline



Log Analysis Pipeline



Indexing

Use filesystem directory structure to encode metadata

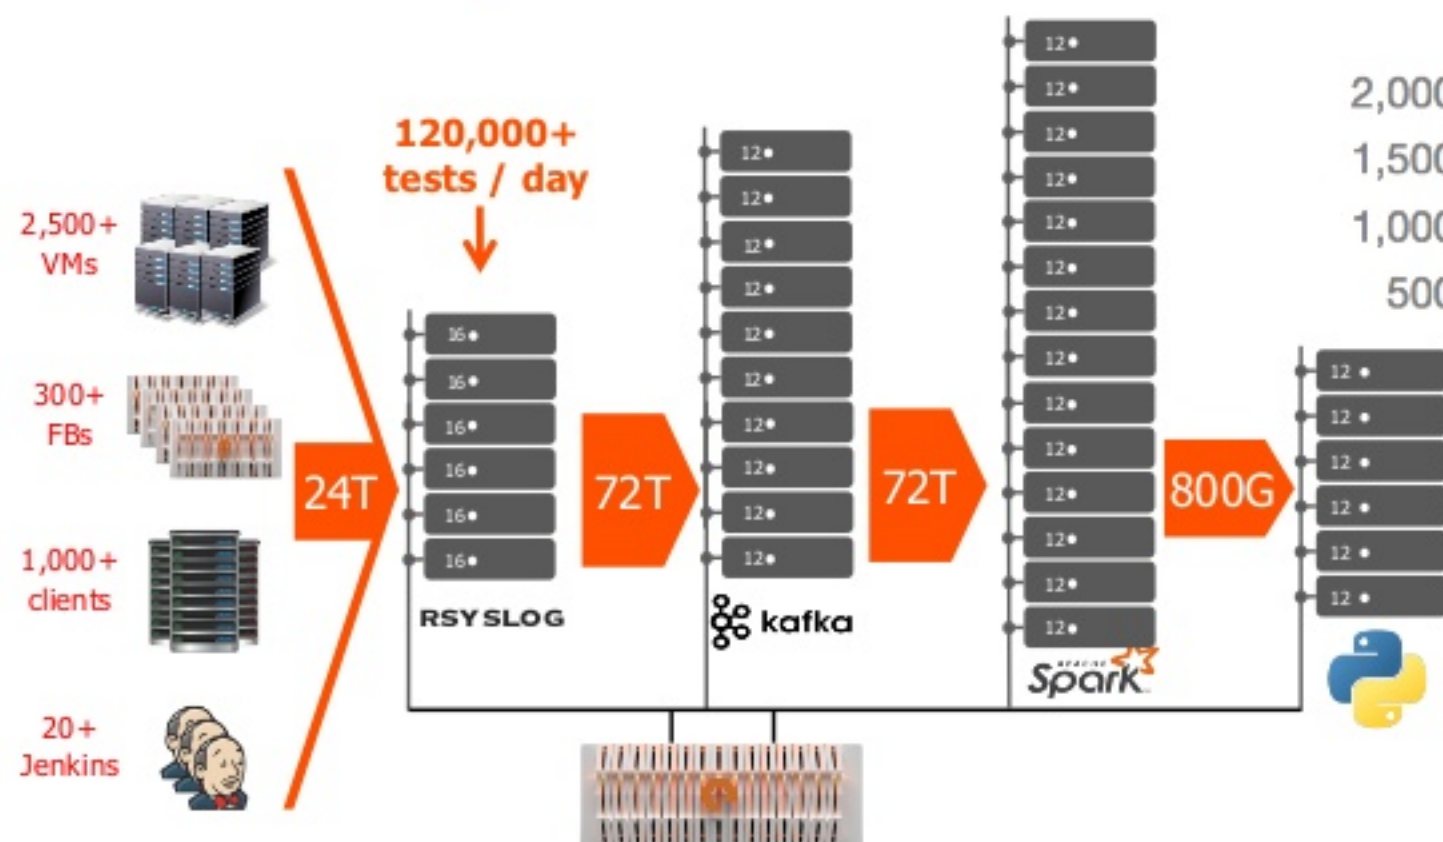
- **Raw data:** <host>/<year>/<month>/<day>/<flat files>
 - Producer: Rsyslog
 - Consumer: Spark batch (re-filter or custom lookbacks)
- **Indexed data:** <pattern>/<year>/<month>/<day>/<hour>/<host>/<flat files>
 - Producer: Spark streaming (filter)
 - Consumer: Python services (e.g. ETL, alert, searchability)

Querying

Find and load data


- FlashBlade NFS protocol. < 1ms latency
- **Listing**
 - "ls -alR" is still SLOW
 - NFS client in kernel sequentially discovers filesystem structure.
 - Solution: Skip the kernel. Use libnfs to create our own parallelized discovery. 1000x faster for 1M files
- **Reading**
 - Buffering: Create input pipeline to optimize for throughput and hide latency away

Full Pipeline



Timelines (Last 1000 batches, 0 active, 1000 co

2,033,456.10 records/sec at 12:16:20



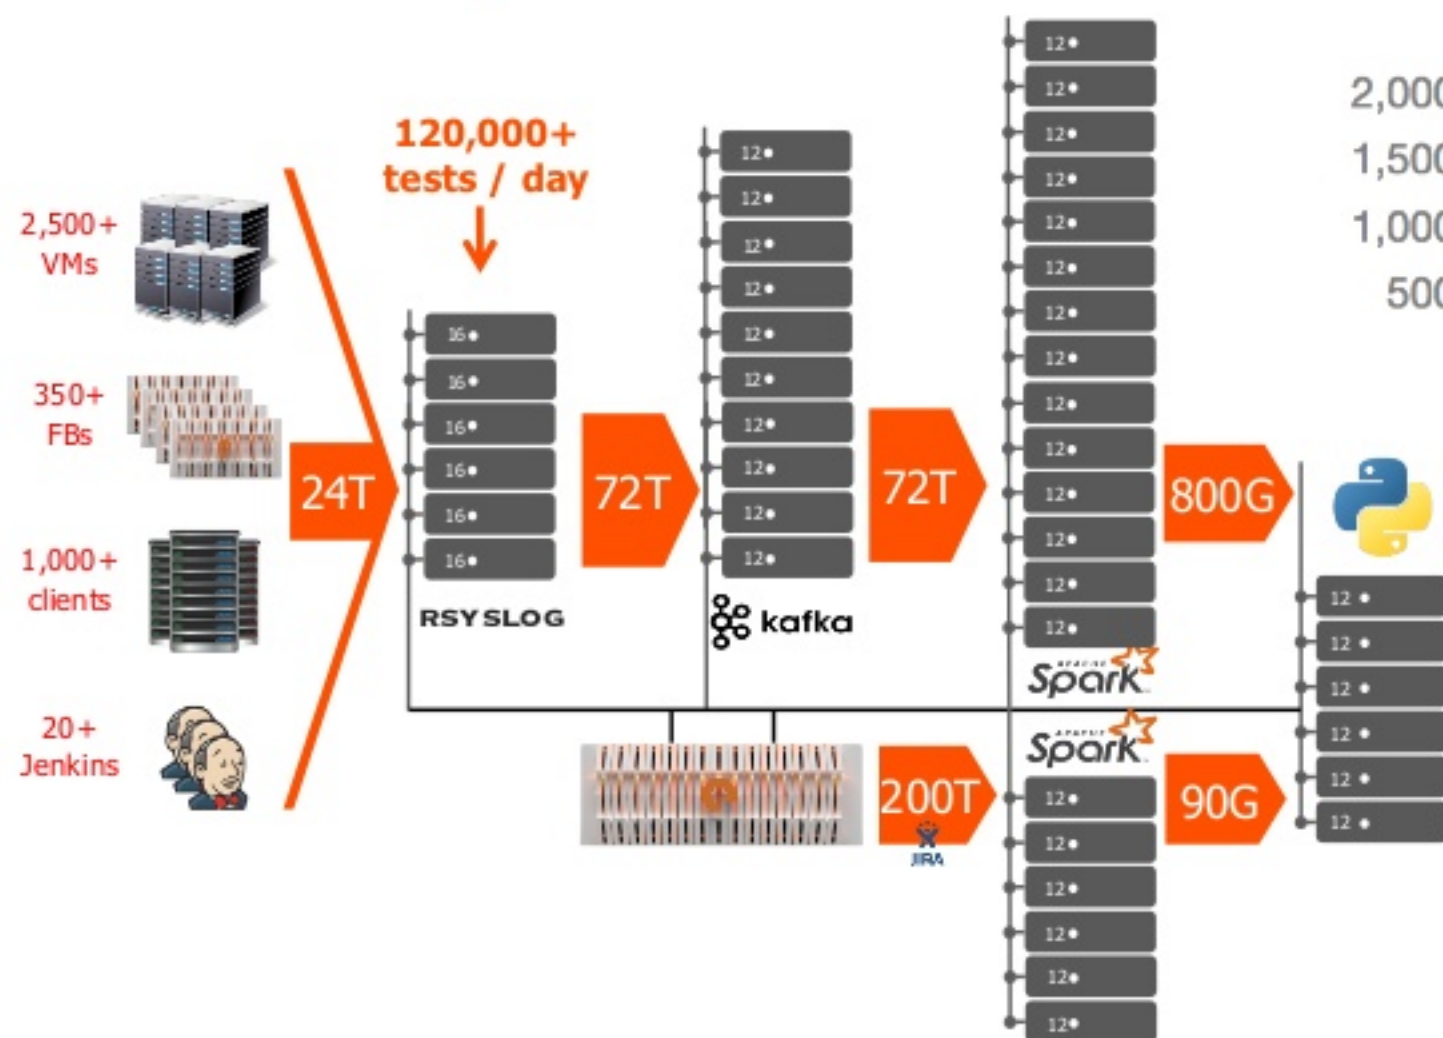
JIRA



Grafana

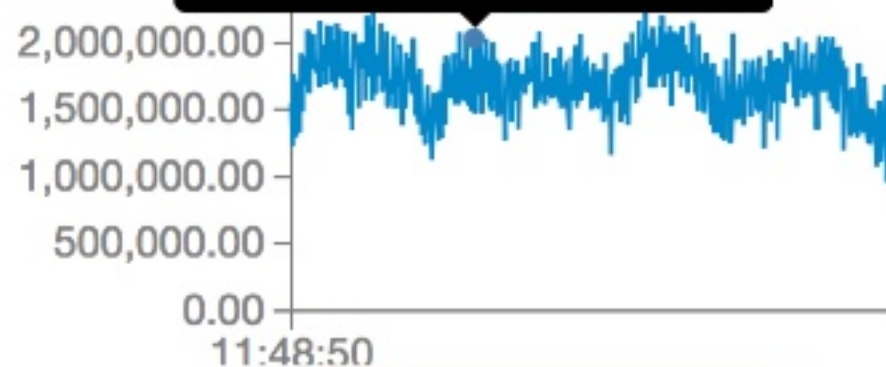
- ✓ Duplicate bug
- ✓ Infrastructure failure
- ✓ Performance regression

Full Pipeline



Timelines (Last 1000 batches, 0 active, 1000 co

2,033,456.10 records/sec at 12:16:20



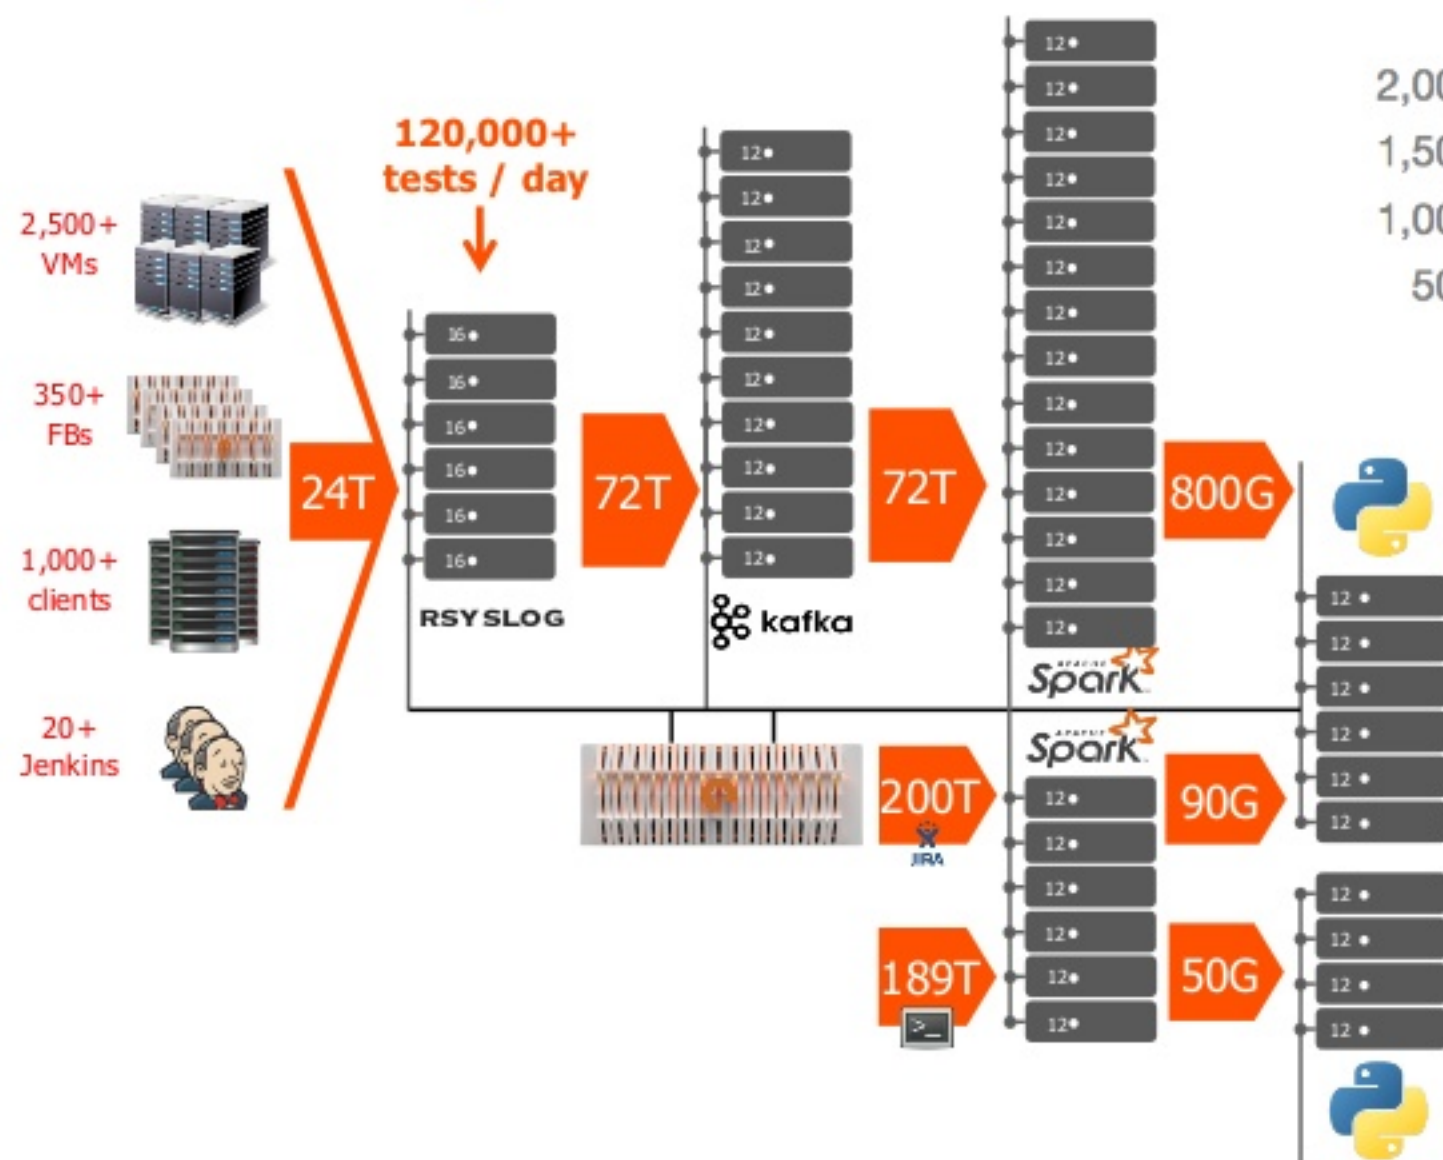
JIRA




Grafana

- ✓ Duplicate bug
- ✓ Infrastructure failure
- ✓ Performance regression

Full Pipeline



Timelines (Last 1000 batches, 0 active, 1000 c)



Takeaways

- ✓ **Index only what you need, store the rest**
(in a storage layer that scales in throughput and to billions of files/objects)
- ✓ **Optimize for throughput and not latency**
- ✓ **Disaggregation of compute and storage for scalability of subsystems**



QUESTIONS?