

# Efficient Spark Analytics on Encrypted Data

Gidon Gershinsky, IBM Research - Haifa  
[gidon@il.ibm.com](mailto:gidon@il.ibm.com)

#SAISDev14

# Overview

- What problem we are solving?
- Parquet modular encryption
- Spark encryption with Parquet
  - connected car scenario: demo screenshots
  - performance implications of encryption
  - getting from prototype to real thing

# What Problem Are We Solving?

- Protect sensitive data-at-rest
  - data confidentiality: encryption
  - data integrity
  - in any storage - untrusted, cloud or private, file system, object store, archives
- Preserve performance of Spark analytics
  - advanced data filtering (projection, predicate) with encrypted data
- Leverage encryption for fine-grained access control
  - per-column encryption keys
  - key-based access in any storage: private -> cloud -> archive



Read only the data  
you need!

a	b	c
a1	b1	c1
a2	b2	c2
a3	b3	c3
a4	b4	c4
a5	b5	c5



# Background

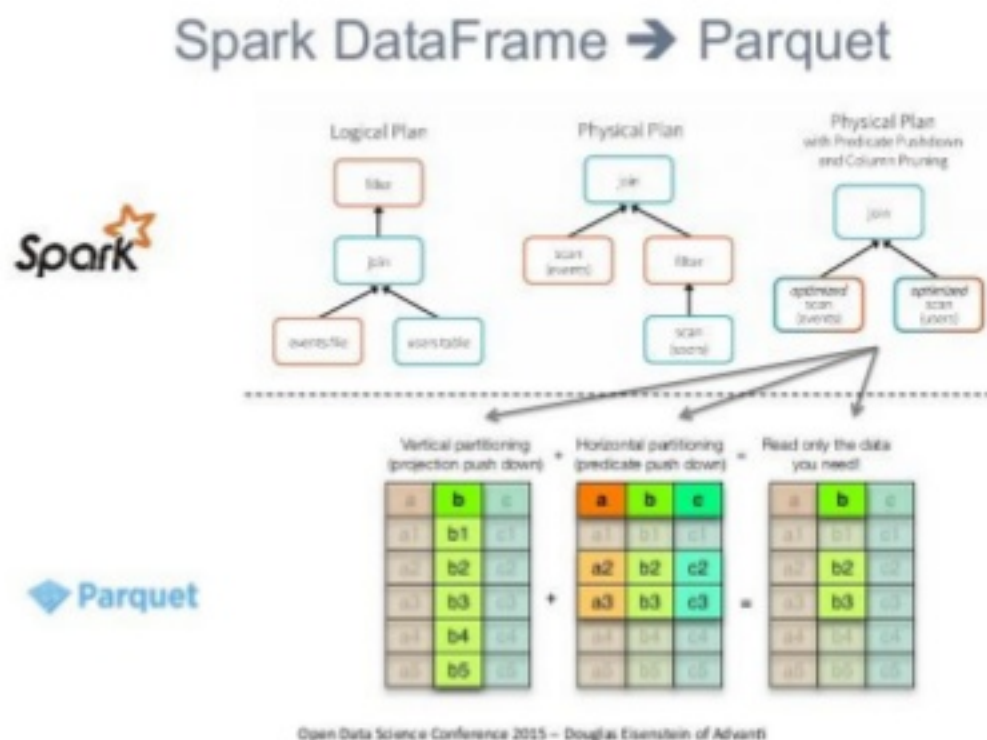


- EU Horizon2020 research project
- EU partners: Adaptant, IT Innovation, OCC, Thales, UDE
  - usage-based car insurance, social services
- Collaboration with UC Berkeley RISELab
  - Opaque technology
- Secure cloud analytics (Spark and H/W Enclaves)
  - how to protect “data-in-use”?
  - how to protect “data-at-rest” – **without losing analytics efficiency?**



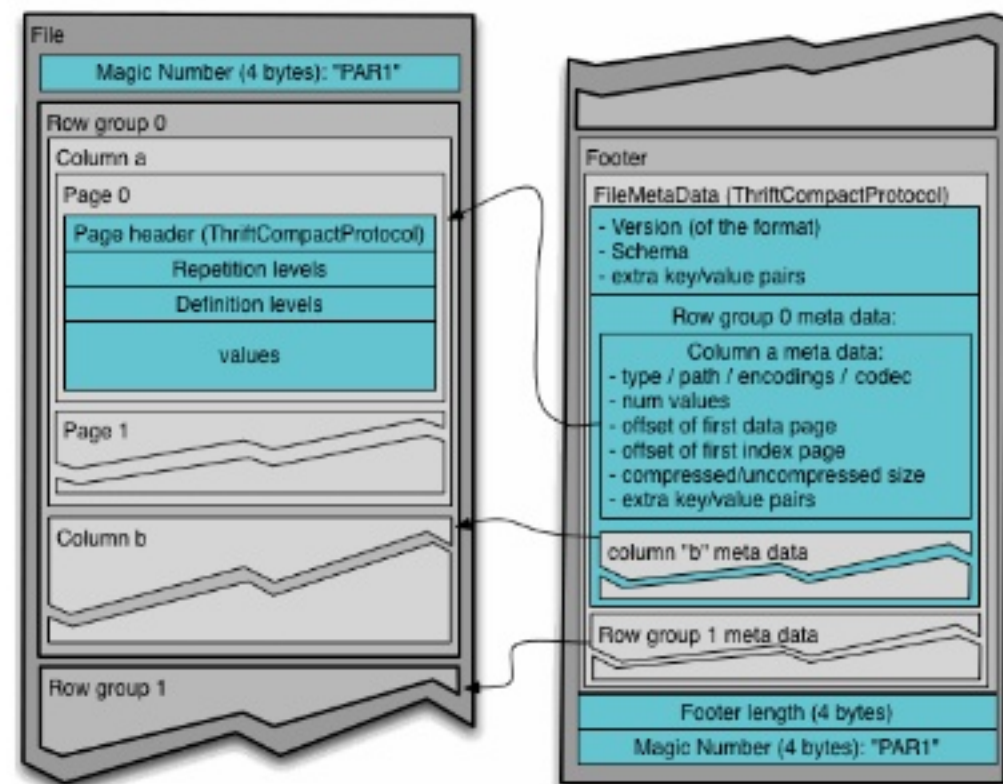
# Apache Parquet

- Integral part of Apache Spark
- Encoding, compression
- Advanced data filtering
  - *columnar projection*: skip columns
  - *predicate push down*: skip files, or row groups, or data pages
- Performance benefits
  - less data to fetch from storage: I/O, latency
  - less data to process: CPU, latency
- How to protect sensitive Parquet data?
  - in any storage - keeping performance, supporting column access control, data tamper-proofing etc.



# Parquet Modular Encryption

- Apache Parquet community work
- Full encryption: all data and metadata modules
  - metadata: indexes (*min/max*), **schema**, size and number of secret values, encryption key metadata, etc ..
  - encrypted and **tamper-proof**
- Enables columnar projection and predicate pushdown
- Storage server / admin never sees encryption keys or cleartext data
  - “client-side” encryption (e.g., Spark-side)



# Parquet Modular Encryption

- Works in any storage
  - supporting range-read
- Multiple encryption algorithms
  - different security and performance requirements
- Data integrity verification
  - AES GCM: “authenticated encryption”
  - data not tampered with
  - file not replaced with wrong version
- Column access control
  - encryption with column-specific keys

```
929 union EncryptionAlgorithm {
930     1: AesGcmV1 AES_GCM_V1
931     2: AesGcmCtrV1 AES_GCM_CTR_V1
932 }
933
934 struct FileCryptoMetaData {
935     1: required EncryptionAlgorithm encryption_algorithm
936
937     /** Parquet footer can be encrypted, or left as plaintext */
938     2: required bool encrypted_footer
939
940     /** Retrieval metadata of key used for encryption of footer,
941      * and (possibly) columns */
942     3: optional binary footer_key_metadata
943
944     /** Offset of Parquet footer (encrypted, or plaintext) */
945     4: required i64 footer_offset
946 }
```

# AES Modes: GCM and CTR

- AES: symmetric encryption standard supported in
  - Java and other languages
  - CPUs (x86 and others): hardware acceleration of AES
- GCM (Galois Counter Mode)
  - encryption
  - integrity verification
    - basic mode: make sure data not modified
    - AAD mode: “additional authentication data”, e.g. file name with a version/timestamp
      - prevent replacing with wrong (old) version
- CTR (Counter Mode)
  - encryption only, no integrity verification
  - useful when AES hardware acceleration is not available (e.g. Java 8)



# Status & Roadmap

- **PARQUET-1178**

- full encryption functionality layer, with API access
- format definition PR merged in Apache Parquet feature branch
- Java and C++ implementation
  - one PR merged, others being reviewed

- **Next: tools on top**

- **PARQUET-1373: key management tools**

- number of different mechanisms
    - options for storing key material

- **PARQUET-1376: data obfuscation layer**

- per-cell masking and aggregated anonymization
    - reader alerts and choice of obfuscated data

- **PARQUET-1396: schema service interface**

- "no-API" interface to encryption
    - all parameters fetched via loaded class

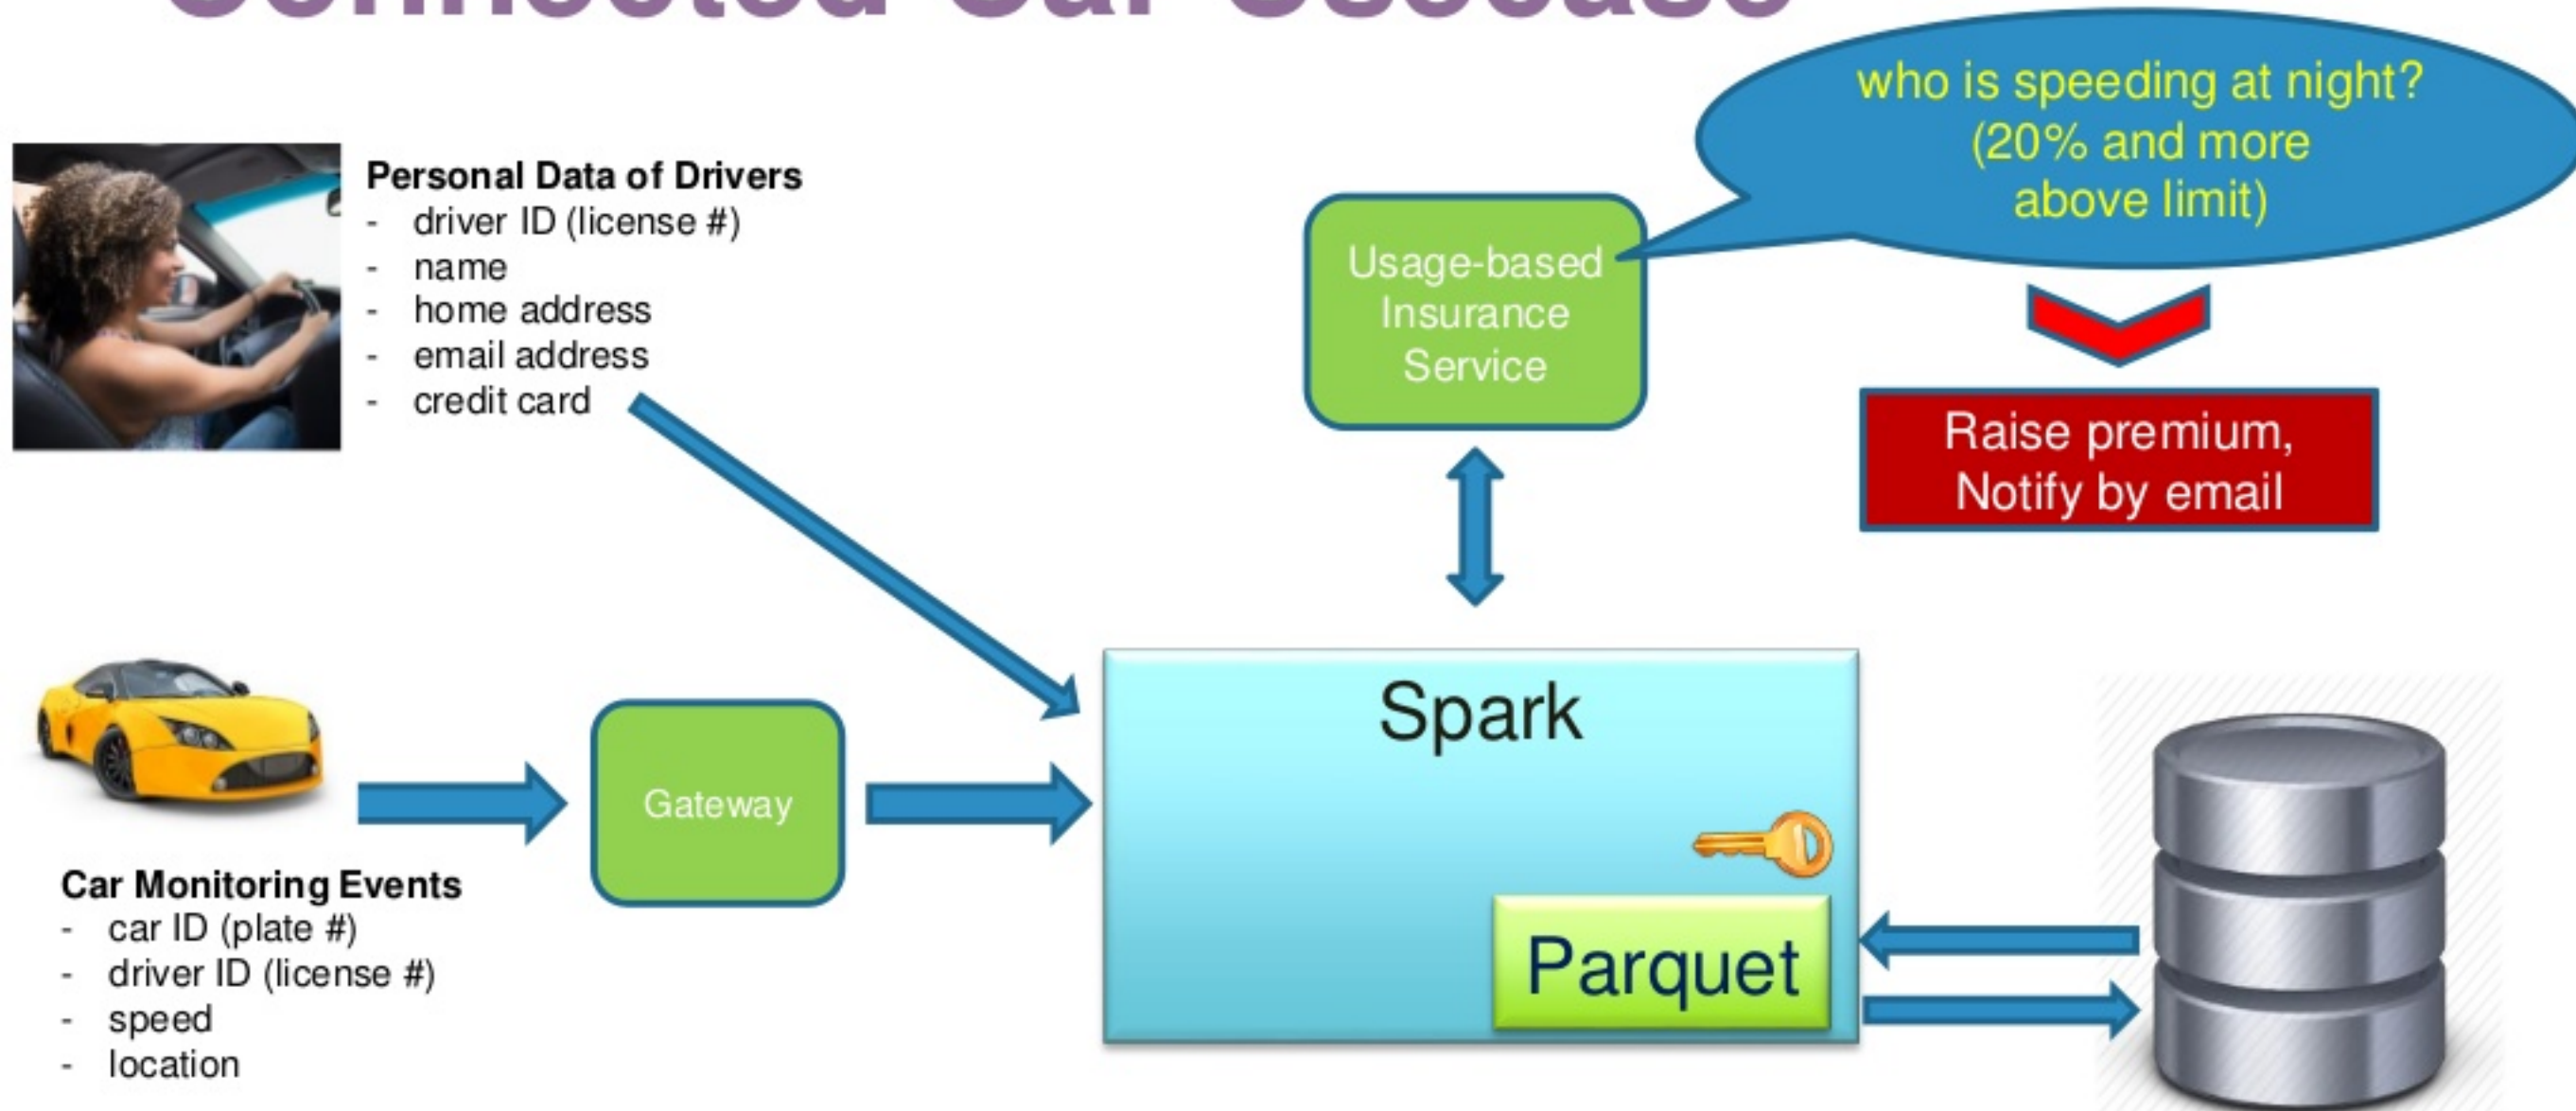
- **Demo feature: Property interface**

- "no-API" interface to encryption
    - most parameters passed via (Hadoop) config properties
    - key server access via loaded class

# Spark Integration Prototype

- Standard Spark 2.3.0 is built with Parquet 1.8.2
- Encryption prototype: standard Spark 2.3.0 with Parquet 1.8.2-**E**
  - Spark-side encryption and decryption
    - storage (admin) never sees data key or plain data
  - Column access control (key-per-column)
  - Cryptographic verification of data integrity
    - replaced with wrong version, tampered with
  - Filtering of encrypted data
    - column projection, predicate push-down

# Connected Car Usecase





# Car Event Schema

Secret data, to  
be encrypted

```
scala> val events_01_09_2018 = spark.read.schema(  
  |   StructType(Seq(  
  |     StructField("CarID", IntegerType),  
  |     StructField("DriverID", IntegerType),  
  |     StructField("Timestamp", TimestampType),  
  |     StructField("CarIDHash", StringType),  
  |     StructField("DriverIDHash", StringType),  
  |     StructField("Latitude", DoubleType),  
  |     StructField("Longitude", DoubleType),  
  |     StructField("Speed", IntegerType),  
  |     StructField("SpeedLimit", IntegerType),  
  |     StructField("TransmissionGearPosition", IntegerType),  
  |     StructField("AcceleratorPedalPosition", IntegerType),  
  |     StructField("BrakePedalStatus", BooleanType),  
  |     StructField("Odometer", LongType),  
  |     StructField("FuelLevel", IntegerType),  
  |     StructField("EngineSpeed", IntegerType),  
  |     StructField("HeadlampStatus", BooleanType),  
  |     StructField("HighBeamStatus", BooleanType),  
  |     StructField("DoorStatus", IntegerType),  
  |     StructField("WindshieldWiperStatus", BooleanType),  
  |     StructField("CarServiceMessage", StringType)  
  |   )).csv("/home/gidon/Summit/Demo/CarEvents_01_09_2018.csv")
```

Unencrypted  
columns



# Driver Table Schema

Secret data, to  
be encrypted

```
scala> val peopleTable = spark.read.schema(  
  |   StructType(Seq(  
  |     StructField("PersonID", IntegerType),  
  |     StructField("PersonIDHash", StringType),  
  |     StructField("FullName", StringType),  
  |     StructField("Gender", StringType),  
  |     StructField("DateOfBirth", DateType),  
  |     StructField("PhoneNumber", StringType),  
  |     StructField("Email", StringType),  
  |     StructField("StreetAddress", StringType),  
  |     StructField("City", StringType),  
  |     StructField("State", StringType),  
  |     StructField("CreditCard", StringType),  
  |     StructField("MaskedCreditCard", StringType),  
  |   )).csv("/home/gidon/Summit/Demo/People_09_2018.csv")  
peopleTable: org.apache.spark.sql.DataFrame = [PersonID: int,  
more fields]
```

Masked data

# Writing Encrypted Parquet Files

```
scala> sc.hadoopConfiguration.set("encryption.key.list",
|   "k0:iKwfmI5rDf7HwVBcqeNE6w==," +
|   "k1:LjxH/aXxMduX6IQcwQgOlw==," +
|   "k2:rnZHCxhUhr79Y6zvQnxSEQ==," +
|   "k3:6b2G9UsRmCAsoGsd3IMQrA==," +
|   "k4:mORupskWLHafuvDJbXrCEw==," +
|   "k5:NG6Hi85MW04sqMlXJHt5lg==," +
|   "k6:AAECAwQFBgcICQoLDA0ODw==")

scala>

scala> events_01_09_2018.write.
|   option("encryption.footer.key", "k0").
|   option("encryption.column.keys",
|         "CarID:k1, DriverID:k2, Longitude:k3, Latitude:k3, Speed:k4").
|   option("encryption.writer.aad", "E_01_09_18.parquet.encrypted").
|   parquet("/home/gidon/Summit/Demo/CarEvents_Sept2018/E_01_09_18.parquet.encrypted")
```

Data Keys  
(base64)

# Writing Encrypted Parquet Files

```
scala> peopleTable.write.  
|   option("encryption.footer.key" , "k5").  
|   option("encryption.column.keys" ,  
|       "PersonID:k2, FullName:k2, DateOfBirth:k2, PhoneNumber:k2, Email:k2, StreetAddress:k2, " +  
|       "CreditCard:k6").  
|   option("encryption.writer.aad" , "People_09_2018.parquet.encrypted").  
|   parquet("/home/gidon/Summit/Demo/People_09_2018.parquet.encrypted")
```



# Reading Encrypted Parquet Files

```
scala> val carEventsEncrypted = spark.read.parquet("/home/gidon/Summit/Demo/CarEvents_Sept2018/*")
18/09/06 14:10:51 ERROR Executor: Exception in task 0.0 in stage 0.0 (TID 0)
Caused by: java.lang.RuntimeException: Trying to read file with encrypted footer. No keys available
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:658)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:582)
    at org.apache.parquet.hadoop.ParquetFileReader.readFooter(ParquetFileReader.java:553)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetFileFormat$$anon$1.readParquetFooter(ParquetFileFormat.scala:100)
```

table schema and other metadata are protected



```
scala> sc.hadoopConfiguration.set("encryption.key.list" ,
  |   "k0:iKwfmI5rDf7HwVBcqeNE6w==," +
  |   "k1:LjxH/aXxMduX6IQcwQgOlw==," +
  |   "k2:rnZHCxhUhr79Y6zvQnxSEQ==," +
  |   "k4:mORupskWLHafuvDJbXrCEw==," +
  |   "k5:NG6Hi85MW04sqMlXJHt5lg==" )
```

location key: not eligible

credit card key: not eligible

```
scala>
```

```
scala> val carEventsEncrypted = spark.read.parquet("/home/gidon/Summit/Demo/CarEvents_Sept2018/*")
carEventsEncrypted: org.apache.spark.sql.DataFrame = [CarID: int, DriverID: int ... 18 more fields]
```



# Hidden Columns

```
scala> carEventsEncrypted.createOrReplaceTempView("encryptedEvents")

scala>

scala> spark.sql("SELECT DriverID, Speed FROM encryptedEvents").show(2)
+-----+-----+
| DriverID|Speed|
+-----+-----+
| 237857628| 90|
| 237462937| 107|
+-----+-----+
only showing top 2 rows

scala>

scala> spark.sql("SELECT DriverID, Speed, Latitude, Longitude FROM encryptedEvent
18/09/06 09:02:59 ERROR Executor: Exception in task 0.0 in stage 4.0 (TID 4)
org.apache.parquet.crypto.HiddenColumnException: [Latitude]
```



# Query Execution

```
scala> val speedingAtNight = spark.sql(
  | "SELECT DriverID, MAX(ROUND(Speed/SpeedLimit - 1.0, 2)) AS AboveLimit, MAX(Speed) AS MaxSpeed " +
  | "FROM encryptedEvents WHERE Speed > 1.2 * SpeedLimit AND HOUR(Timestamp) < 5 " +
  | "GROUP BY DriverID ORDER BY MAX(Speed/SpeedLimit) DESC"
  | )
speedingAtNight: org.apache.spark.sql.DataFrame = [DriverID: int, AboveLimit: double ... 1 more field]

scala> speedingAtNight.createOrReplaceTempView("speedingAtNight")

scala> spark.sql(
  | "SELECT FullName, AboveLimit, MaxSpeed, Email, MaskedCreditCard FROM (speedingAtNight " +
  | "INNER JOIN encryptedPeopleTable ON speedingAtNight.DriverID=encryptedPeopleTable.PersonID)"
  | ).show
+-----+-----+-----+-----+-----+
| FullName|AboveLimit|MaxSpeed|Email|MaskedCreditCard|
+-----+-----+-----+-----+-----+
| Jack Wilson|0.8|162|jackw@example.com|XX-8727|
| Anthony Grossel|0.39|125|antonyg@example.com|XX-5613|
| Kayla Woodcock|0.39|125|kaylaw@example.com|XX-2338|
| Amy Trefl|0.37|123|amyt@example.com|XX-7263|
| Jonathan Ruppl|0.34|121|jonathanr@example.com|XX-9283|
| Eva Muirden|0.28|115|evam@example.com|XX-8127|
+-----+-----+-----+-----+-----+
```

who is  
speeding  
at night?

# Tampering with Data



```
$ cp People_08_2018.csv People_09_2018.csv
```



```
scala> peopleTableCSVSept18.createOrReplaceTempView("unprotectedPeopleTableSept18")
scala> spark.sql(
  | "SELECT FullName, AboveLimit, MaxSpeed, Email, MaskedCreditCard FROM (speedingAtNight " +
  | "INNER JOIN unprotectedPeopleTableSept18 ON speedingAtNight.DriverID=unprotectedPeopleTableSept18.PersonID)"
  | ).show
+-----+-----+-----+-----+-----+-----+
| FullName|AboveLimit|MaxSpeed|Email|MaskedCreditCard|
+-----+-----+-----+-----+-----+
|Anthony Grosse|0.39|125|antonyg@example.com|XX-5613|
|Kayla Woodcock|0.39|125|kaylaw@example.com|XX-2338|
|Amy Trefl|0.37|123|amyt@example.com|XX-7263|
|Jonathan Rupp|0.34|121|jonathanr@example...|XX-9283|
|Eva Muirden|0.28|115|evam@example.com|XX-8127|
+-----+-----+-----+-----+-----+

```

missed Jack  
Wilson, 80%  
above speed  
limit !!



# Parquet Data Integrity Protection



```
$ rm -rf People_09_2018.parquet.encrypted/  
$ cp -r People_08_2018.parquet.encrypted/ People_09_2018.parquet.encrypted/
```



```
scala> val peopleTableEncrypted =  
  | spark.read.parquet("/home/gidon/Summit/Demo/People_09_2018.parquet.encrypted")  
18/09/06 10:36:59 WARN TaskSetManager: Lost task 0.0 in stage 112.0 (TID 3963, localhost, executor dri  
ver): java.io.IOException: Failed to decrypt  
  at org.apache.parquet.crypto.AesDecryptor.decrypt(AesDecryptor.java:134)
```

```
Caused by: javax.crypto.AEADBadTagException: Tag mismatch!  
  at com.sun.crypto.provider.GaloisCounterMode.decryptFinal(GaloisCounterMode
```





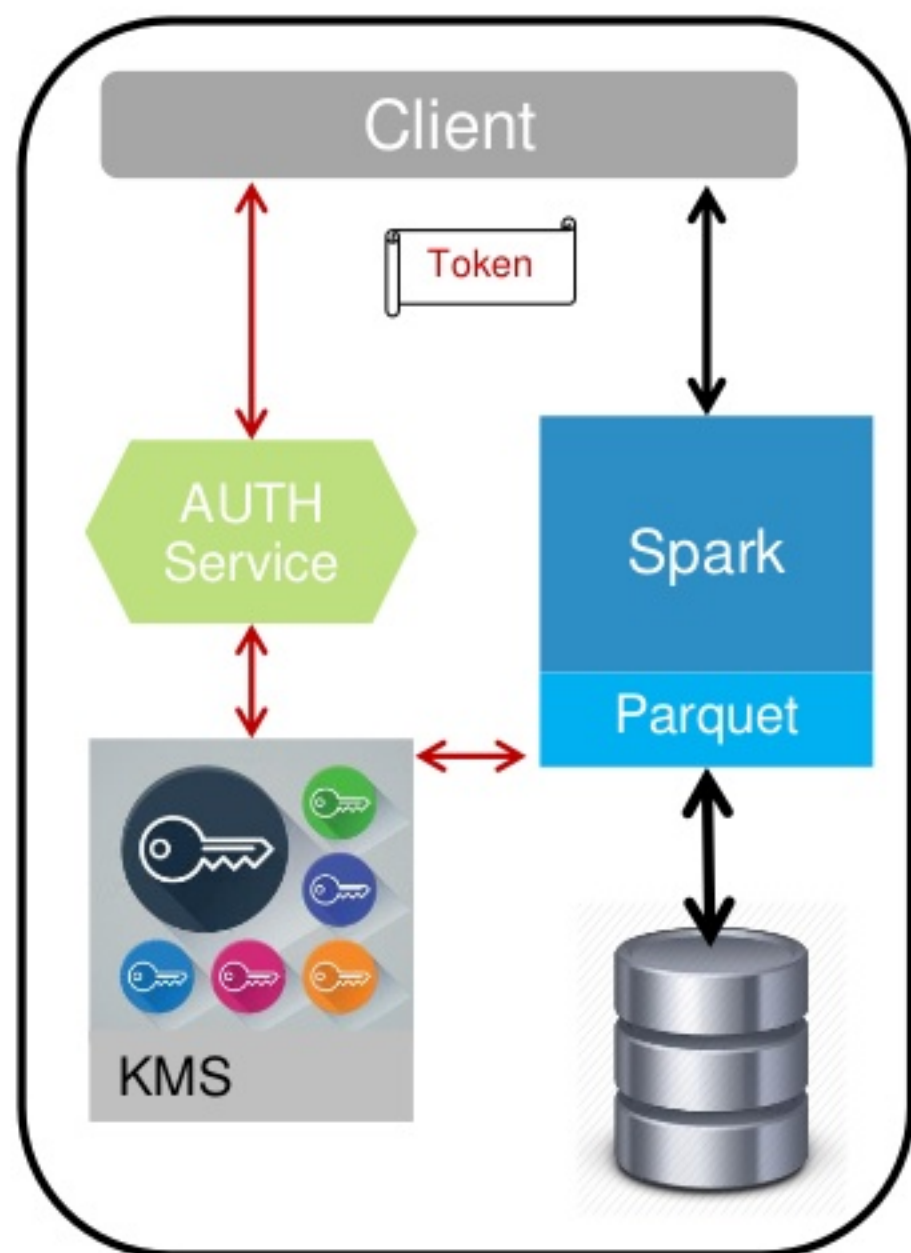
# Key Management

- PARQUET-1373
  - randomly generated Data keys
  - wrapped in KMS with Master keys
  - storage options for wrapped keys
    - in file itself: easier to find
    - in separate location: easier to re-key
      - compromised Master key
  - client: specify Master key ID per column

```
scala> sc.hadoopConfiguration.set("encryption.user.token", token)
scala> sc.hadoopConfiguration.set("encryption.kms.client.class", "com.ibm.demo.KeyProtectClient")
scala> sc.hadoopConfiguration.set("ibmcloud.keyprotect.instance", "2164109a-6779-4024-b0a4-91e59a0f67e7")
scala> val k1 = "d1ae3f02-4b7d-4a93-abba-e44e02933734"
k1: String = d1ae3f02-4b7d-4a93-abba-e44e02933734
scala> val k2 = "cead1331-2a78-4988-a7c2-57c481f58d5c"
k2: String = cead1331-2a78-4988-a7c2-57c481f58d5c
scala> val k3 = "8eaceidd-d0e8-43df-9baa-793eab8a8e08"
k3: String = 8eaceidd-d0e8-43df-9baa-793eab8a8e08
scala> peopleTable.write
  | option("encryption.footer.key", k1).
  | option("encryption.column.keys",
  |   "PersonID:" + k2 + ", FullName:" + k2 + ", DateOfBirth:" + k2 + ", PhoneNumber:" + k2 +
  |   ", Email:" + k2 + ", StreetAddress:" + k2 + ", CreditCard:" + k2 + ").
  | option("encryption.writer.out", "people_vv_2018_kv.parquet:encrypted").
  | parquet("/home/rgidon/Summit/Demo/People_vv_2018_KP.parquet:encrypted")
```

Master Key IDs

random Data Keys generated in Parquet, wrapped by Master Keys, written in file or elsewhere



# IBM Spark-based Services

- IBM Analytics Engine
  - on-demand Spark (and Hadoop) clusters in IBM cloud
- Watson Studio Spark Environments
  - cloud tools for data scientists and application developers
  - dedicated Spark cluster per Notebook
- DB2 Event Store
  - rapidly ingest, store and analyze streamed data in event-driven applications
  - analyze with a combination of Spark and C++ engines, store with C++ Parquet
- SQL Query Service
  - serverless scale-out SQL execution on TB of data
  - Spark SQL engine, running in IBM Cloud



# Column Access Control with Encryption @Uber

- Access control to sensitive columns only
- PARQUET-1178 provides fundamentals to encrypt columns
- PARQUET-1396 provides schema and key retrieval interface
- Currently under development - stay tuned!

Uber

# Encryption Performance



Not a Benchmark

- Ubuntu 16.04 box
  - 8-core Intel Core i7 6820HQ CPU / 2.70GHz
  - 32GB memory
- Raw Java AES-GCM speed
  - no Spark or Parquet, just encrypting buffers in loop
    - 8 threads, 1MB buffers
  - **Java 8: ~ 250 MB/sec**
  - **Java 9 and later: ~ 2-5 GB/sec**
    - AES-NI in Hotspot (hardware acceleration)
    - Slow warmup in decryption: Bug report JDK-8201633
      - can be bypassed with a workaround
      - *if you know someone who knows someone in Java Hotspot team...*



# Spark Encryption Performance

- Local mode, 1 executor: 8 cores, 32GB memory
- Storage: Local RamDisk
- Write / encryption test
  - load carEvents.csv files (10GB)
  - cache in Spark memory!
  - write as parquet{.encrypted}: measure time (of second write)
- Query / decryption test
  - 'read' carEvents.parquet{.encrypted} files (2.7GB)
  - run "speeding at night" query: measure time

# Spark Encryption Performance

Test	Write (sec)	Query (sec)	Notes
no encryption	26	2	query on 4 columns: input ~12% of data
encryption (GCM)	28.5	2.5	GCM on data and metadata
encryption (GCM_CTR)	26.8	2.2	CTR on data, GCM on metadata

- Typical usecases - only 5-10% of columns are sensitive (encrypted)
- Java 8... If needed, use GCM\_CTR instead of GCM
- **Next Java version in Spark will support AES hardware acceleration**
- Encryption is not a bottleneck
  - app workload, data I/O, encoding, compression

To Be Benchmarked!

# From Prototype to Real Thing

- Spark encryption prototype
  - standard Spark code, built with modified Parquet
  - encryption is triggered via Hadoop config
    - transparent to Spark
- That's it? Just upgrade to new Parquet version? *Not really ..*
  - partitions!
    - exception on encrypted columns? obfuscate or encrypt sensitive partition values?
  - Spark support for other things above file format
- TBD!

Parquet encryption: Community work. On track. Feedback welcome!

Spark encryption: Prototype, with gaps. Open challenge.



# Backup

# Existing Solutions

	Flat file encryption	Storage server encryption with range read	Storage client encryption with range read
Preserves analytics performance	X	V	V
Works in any storage system	V	X	X
Hides data and keys from storage system	V	X	V
Supports data integrity verification	V	V	X
Enables column access control (column keys)	X	X	X

- Bottom line: Need data protection mechanism built in the file format

# Obfuscated Data

Transportation  
Safety  
Agency



no keys for  
personal data

anti-speeding ad campaign:  
budget per State?

```
scala> val speedingAtNight = spark.sql(
  | "SELECT DriverIDHash MAX(ROUND(Speed/SpeedLimit, 2)) AS AboveLimit " +
  | "FROM encryptedEvents WHERE Speed > 1.2 * SpeedLimit AND HOUR(Timestamp) < 5 " +
  | "GROUP BY DriverIDHash ORDER BY MAX(Speed/SpeedLimit) DESC"
  | )
speedingAtNight: org.apache.spark.sql.DataFrame = [DriverIDHash: string, AboveLimit: double]

scala> speedingAtNight.createOrReplaceTempView("speedingAtNight")

scala> spark.sql(
  | "SELECT State, AVG(AboveLimit) FROM (speedingAtNight " +
  | "INNER JOIN encryptedPeopleTable ON speedingAtNight.DriverIDHash=encryptedPeopleTable.PersonIDHash " +
  | "GROUP BY State ORDER BY AVG(AboveLimit) DESC"
  | ).show

+-----+-----+
|State|avg(AboveLimit)|
+-----+-----+
| NJ |          1.595 |
| PA |          1.365 |
| NY |          1.28 |
+-----+-----+
```

where the worst  
drivers live?