



Large-scale, Near real-time pipelines at Uber

Nishith Agarwal, Vinoth Chandar

Session hashtag: #SAISEco10

The Uber logo is a black square with the word "UBER" in white, bold, sans-serif capital letters.A woman with curly hair, wearing a white shirt and blue overalls, is walking across a city street. She is carrying a brown bag. In the background, there are trees, buildings, and a white van. The Uber logo is overlaid on the bottom left of this image.

Speaker Intros

Who are we?



Nishith Agarwal

Snr Engineer on the Hadoop Platform team,
working on Hudi & hadoop ingestion at large.



Vinoth Chandar

Staff engineer on Infrastructure org, working
on two core technologies : Hudi/BigData
storage, Network protocols for the edge

Uber's scale

Smarter cities of the future

On a snowy Paris evening in 2008, Travis Kalanick and Garrett Camp had trouble hailing a cab. So they came up with a simple idea—press a button, get a ride.

700+


Cities

70+

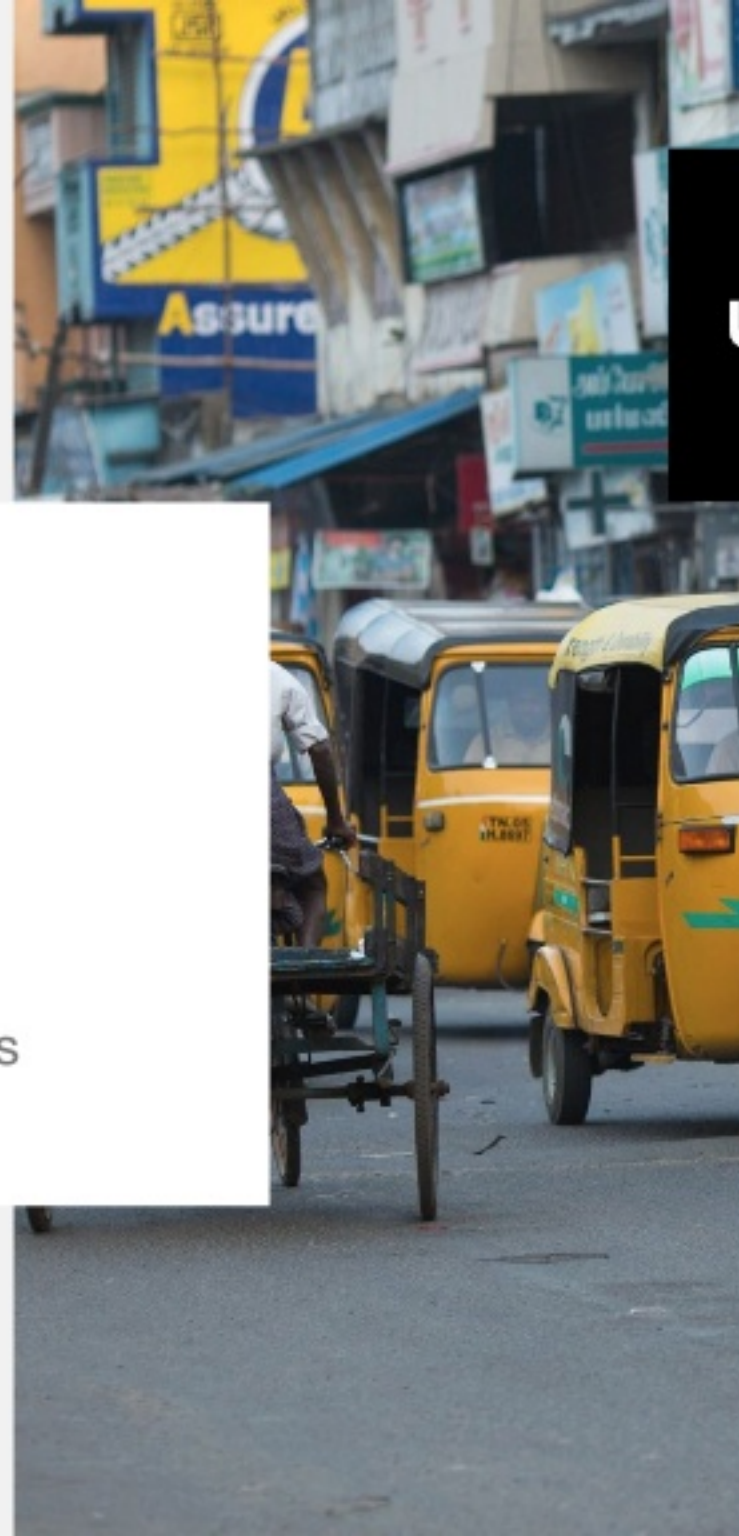
Countries

2M+

Driver partners



UBER



Agenda

- What's Incremental Processing?
- Hudi - Overview
- Recipe : Database Ingestion
- Recipe : Deduping Logs
- Recipe : Incremental ETL
- OSS Community

A photograph of a woman with curly hair, wearing a white shirt and blue overalls, walking across a city street. She is carrying a brown bag. In the background, there are trees, buildings, and a white van. A black rectangular box with the word "UBER" in white capital letters is overlaid on the bottom left of the image.

UBER

Spark Pipelines

Different Shapes & Sizes

Read, Transform, Write

- Row level (Projections/Filters)
- Aggregations (group-by, sums, averages)
- Windowing/Joins

Batch Model

- Read : All data for day/hour(s)
- Transform : Recompute all results
- Write : Replace entire day/hour(s) results

Streaming Model

- Read : New data since last compute
- Transform : Compute result deltas using state store
- Write : Update new results in sink

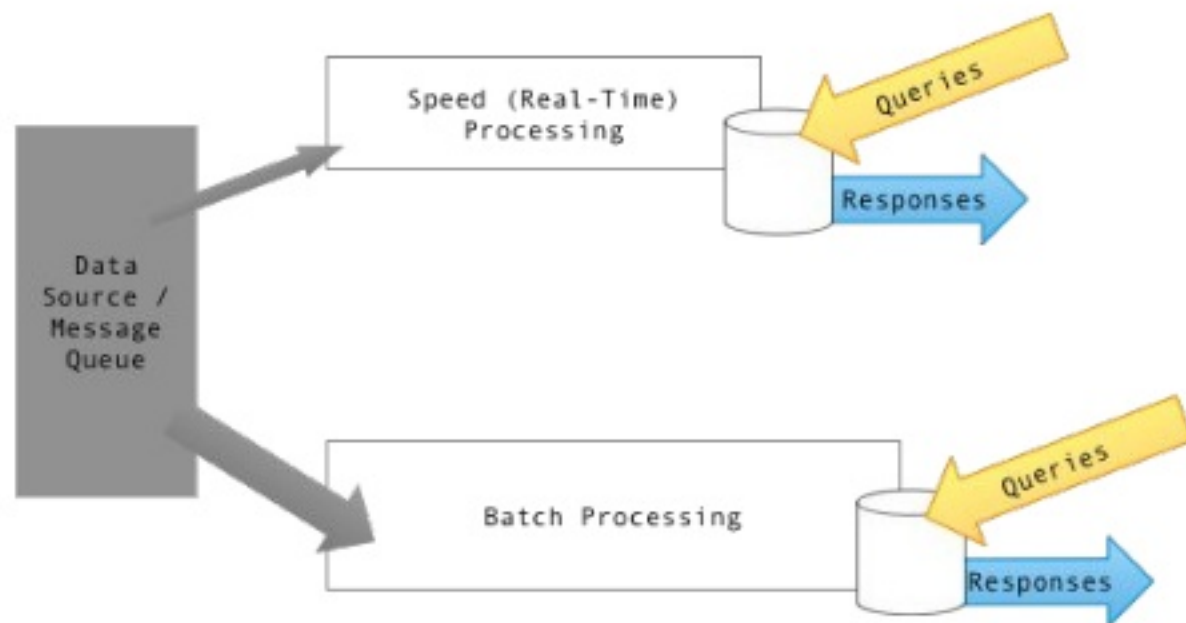


Fig: Typical lambda architecture employing both pipelines

Spark Pipelines

Key Levers

Latency

How fresh are the computed results?

Completeness

How accurate are the computed results?

Scale/Cost

For given latency/completeness, how much data can be read/written at reasonable cost?

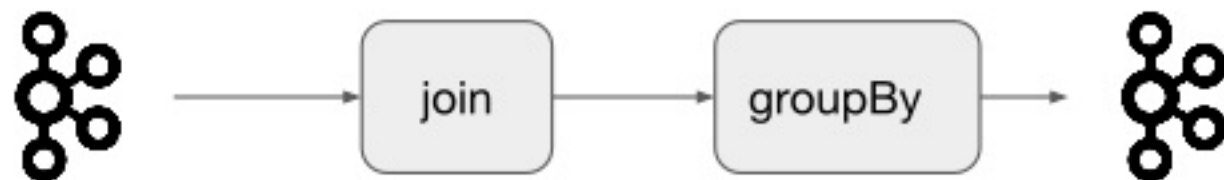


Fig: Multi way join/aggregation, at low latency, high cost and moderate/low completeness

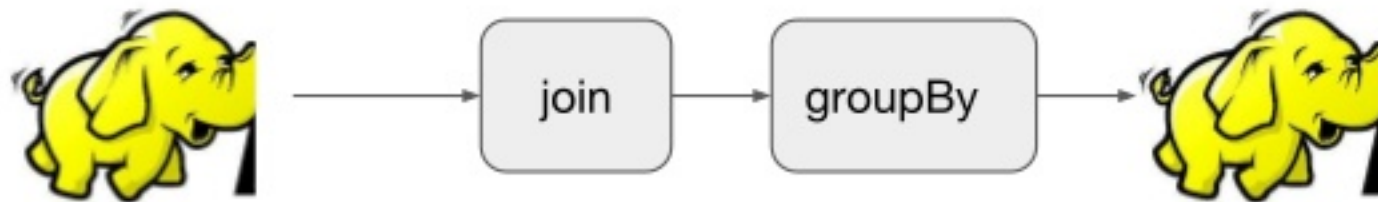


Fig: Multi way join/aggregation, at high latency, low cost and high completeness

How about large data@low latency ?

Terabytes to Petabytes

Streaming Model

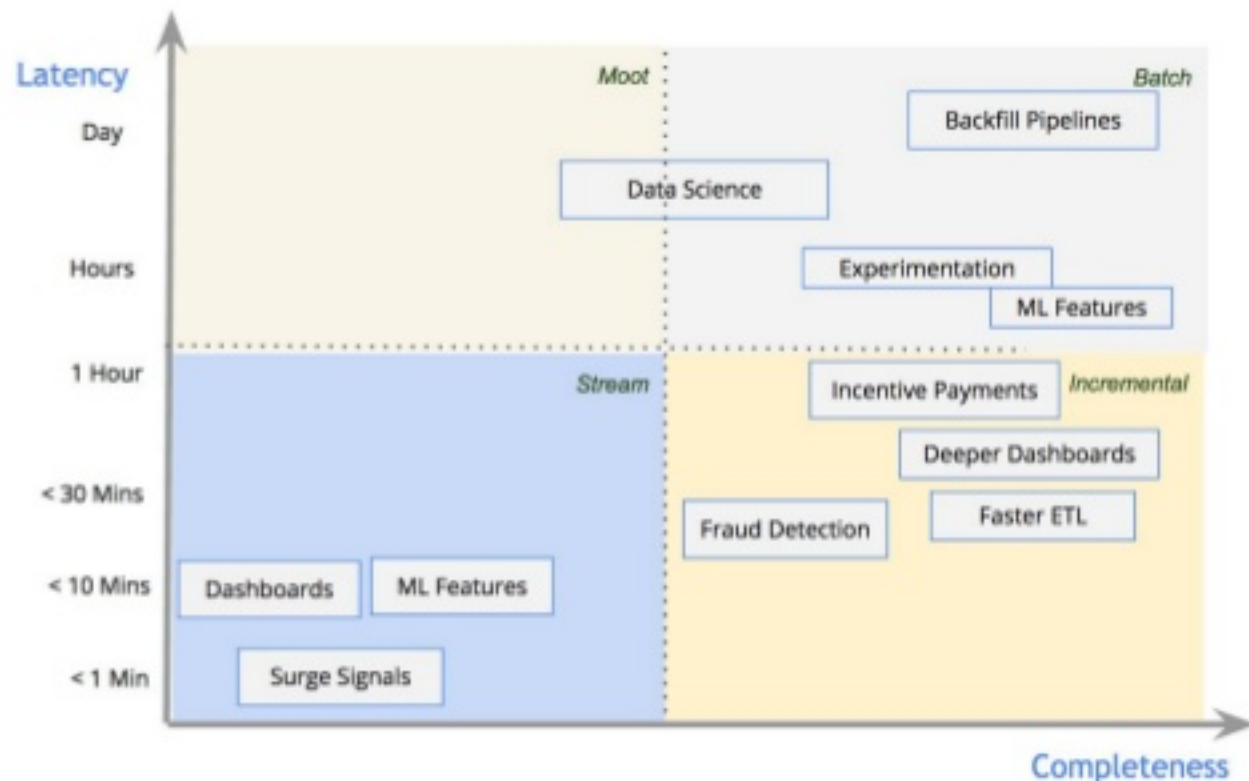
- Great latency
- Smaller datasets (TBs)

Batch Model

- Great efficiency
- Large datasets

It's needn't be a dichotomy!

- Some don't fit in either!
- Support entire spectrum of use-cases



Incremental Model

Streaming Style Processing On Batch Data

Near real-time results

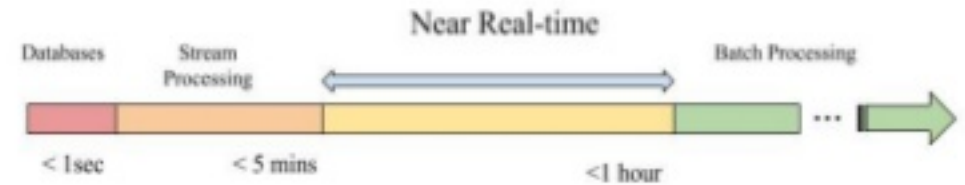
- Mini batch jobs, every few minutes

Upsert (Primitive #1)

- Modify processed results
- Like state stores in stream processing

Incremental Pull (Primitive #2)

- Log stream of changes, avoids costly scans
- Faster flow of data to next stage in dataflow



Incremental Model

Trade-offs

Latency

Streaming < Incremental < Batch

Completeness

Batch* < Incremental < Streaming

Scale

Streaming < Incremental < Batch

Cost

Batch < Incremental < Streaming

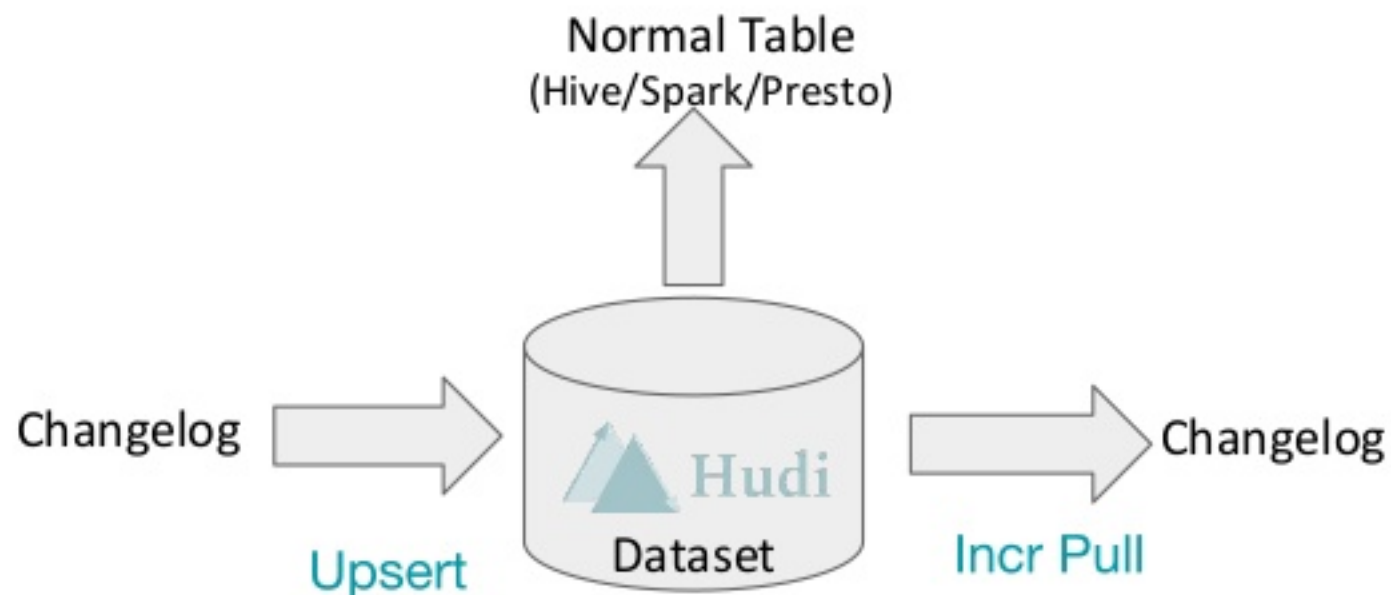
HUDI : Hadoop Upserts and Incrementals

Turn batch jobs to incremental model

- Improve latency by incorporating only deltas
- Scale more by avoiding recomputation, lowering cost

Spark Library for

- Mutations to datasets
- Changelogs from datasets
- Manage files efficiently
- Provide snapshot isolation



Open Source

- <https://github.com/uber/hudi>
- <https://eng.uber.com/Hudi>

HUDI

@Uber

10s PB

Total Storage

1000s

Ingest Pipelines

100s TB

Ingested/Day

100s

ETL Pipelines

10000+

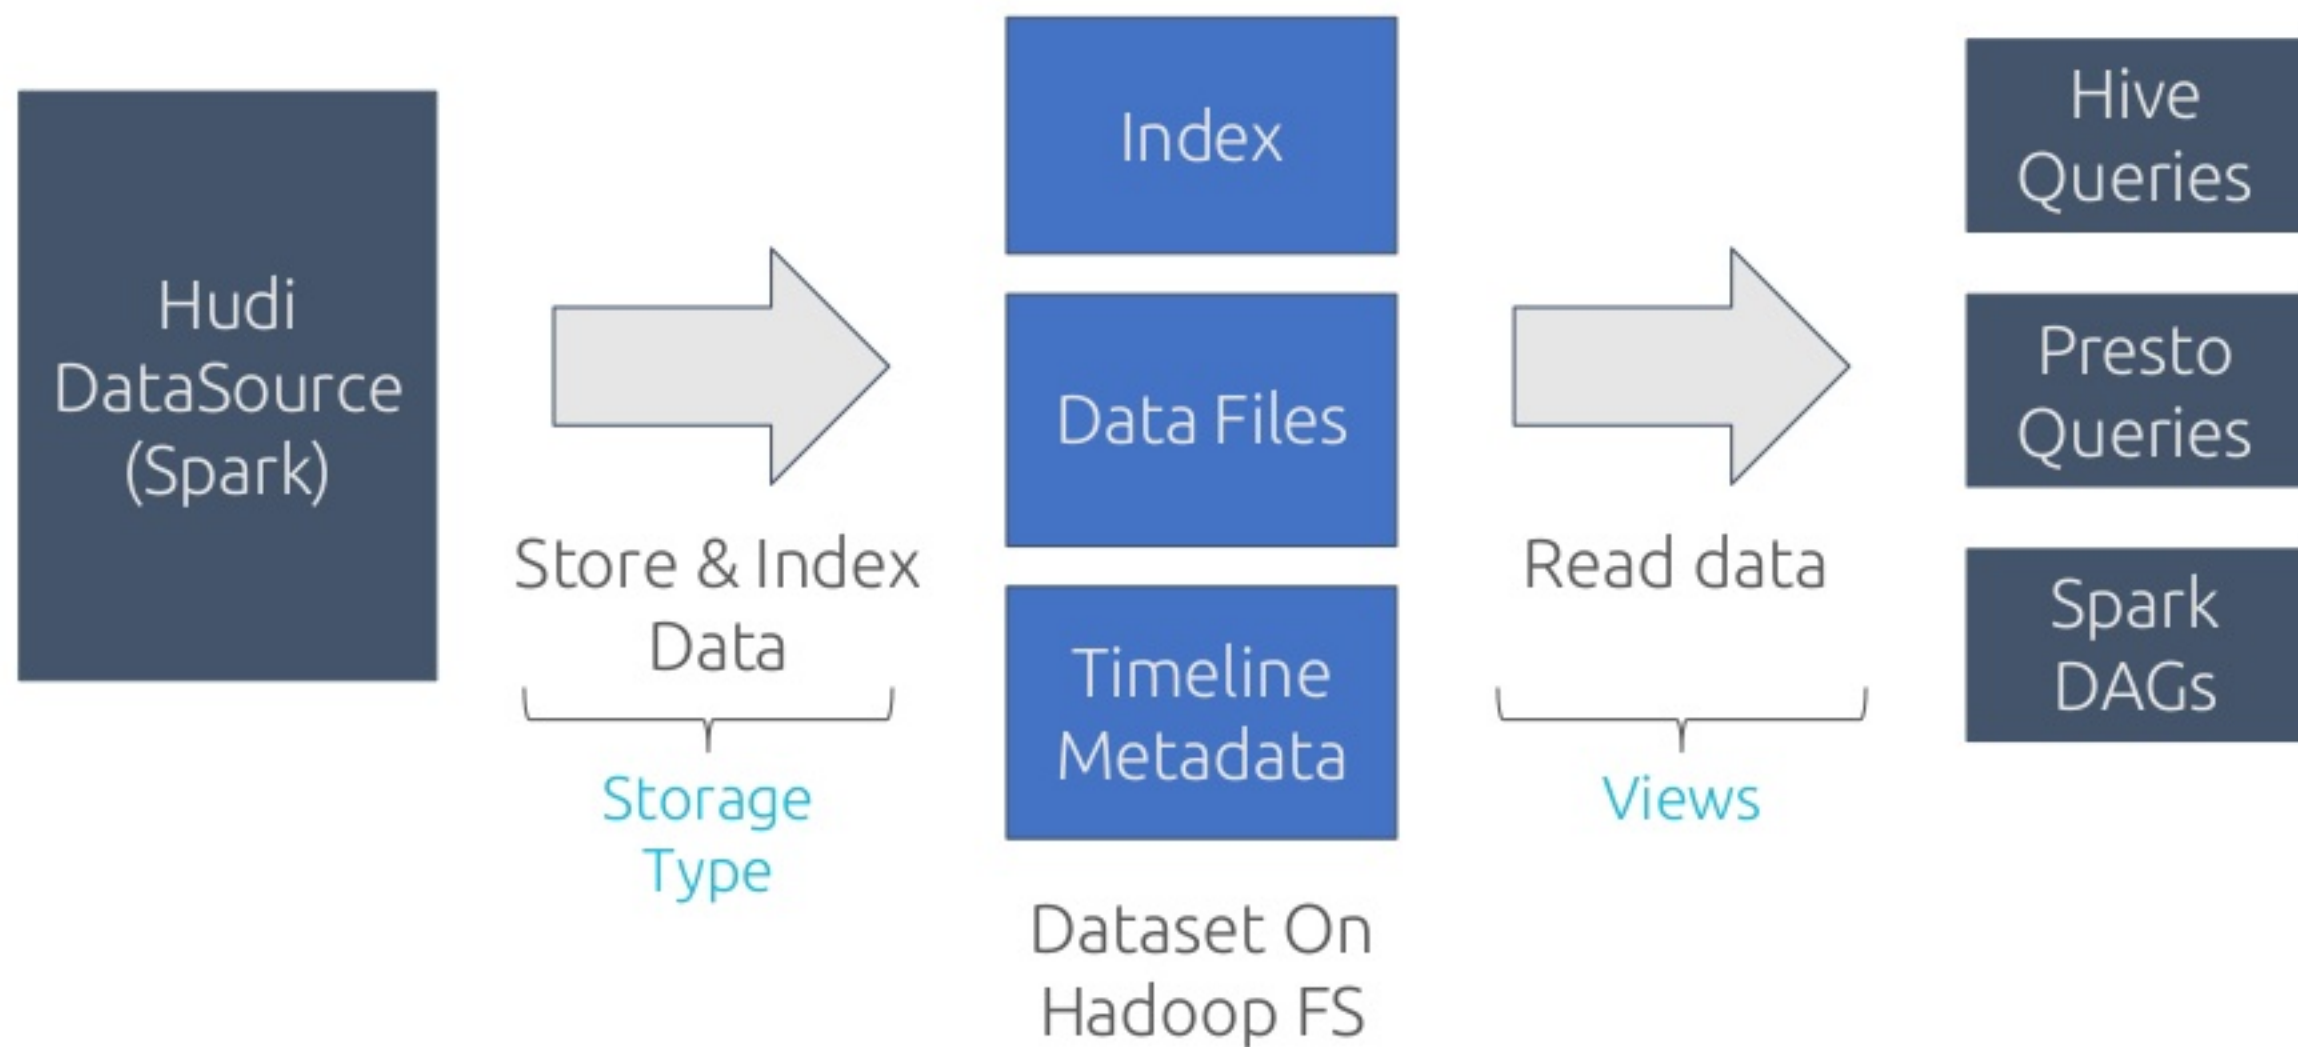
Yarn Containers

4-30 mins

Latency Range

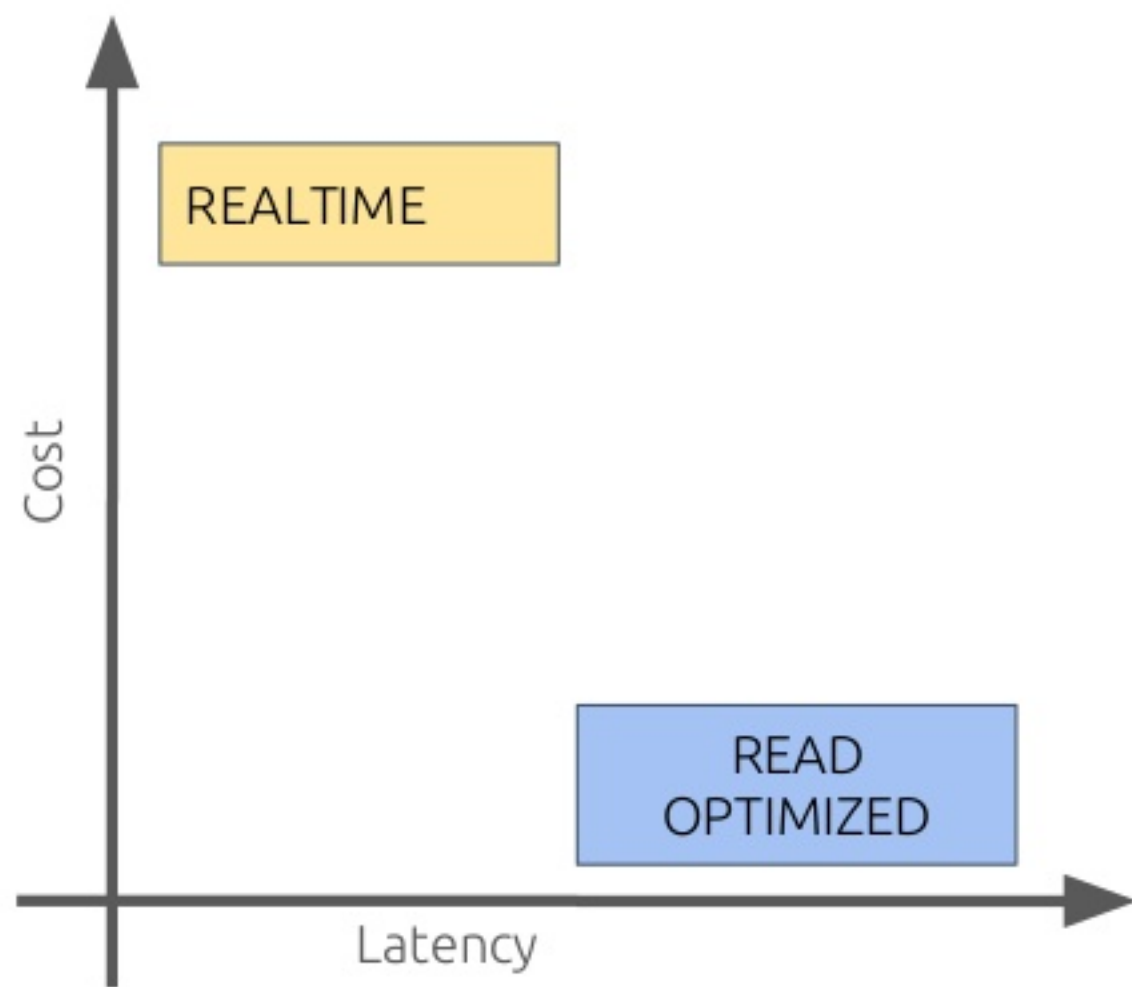
HUDI

Components



Hudi Views

Different options to view the data



3 Logical views Of Dataset

Read Optimized View

- Raw Parquet Query Performance
- Targets existing Hive tables

Real Time View

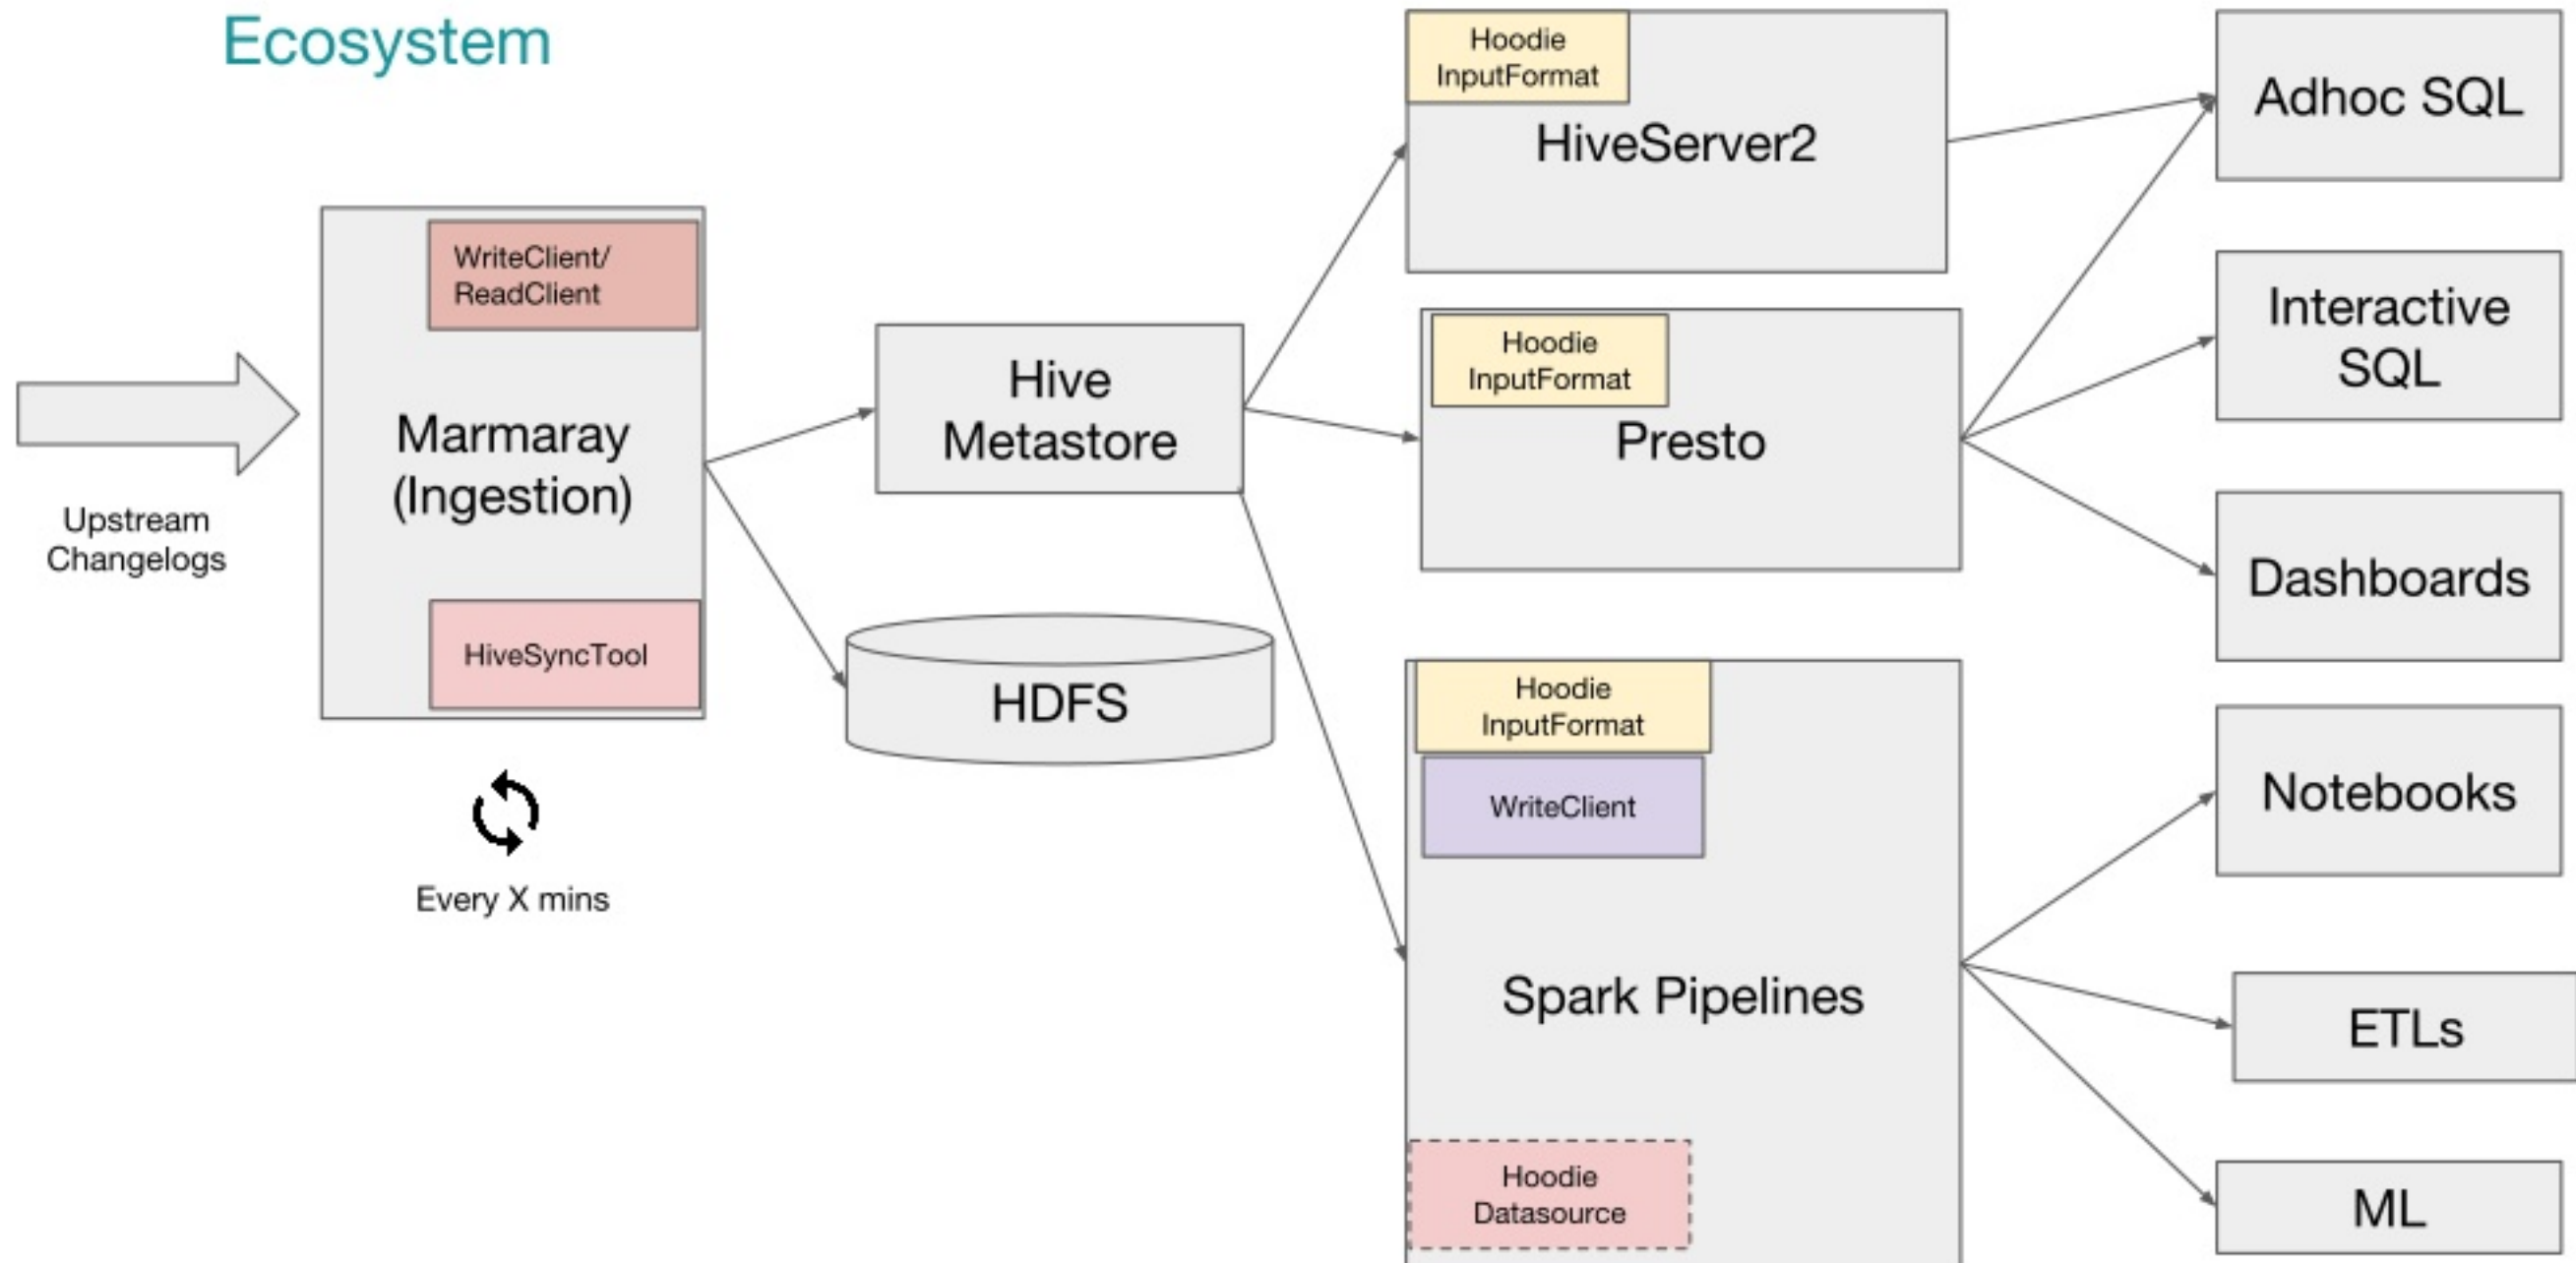
- Hybrid of row & columnar data
- Brings near-real time tables

Incremental View

- Stream of changes to a dataset
- Enables Incremental Pull

HUDI

Ecosystem



HUDI

Tools

HiveSyncTool

- Register Hive tables to access Hudi views
- Registers a RO & RT table by default
- Handles schema evolution

DeltaStreamer

- Ingest tool supporting DFS & Kafka sources
- Json and Avro data formats supported

Spark Datasource

- Custom incremental pipelines
- Jobs needing upserts into sink

CLI

- CLI for inspecting timeline
- Repair tools
- 1-time data imports

HiveIncrementalPuller

- Incrementally consume data from Hive via HQL

SnapshotCopier

- Consistently Copy/Backup Hudi datasets

Recipe: Database Ingestion

Problem

OLTP Databases

- MySQL
- NoSQL

High-value data

- User information
- Trip, transaction logs
- Heavily joined

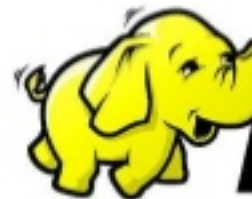
Replicate CRUD operations

- Hadoop/Hive tables need to be kept upto date

Inserts, updates, deletes



Replicate



userID	int
country	string
last_mod	long
...	...

Recipe: Database Ingestion

Typical Approach

Bulk Load

- Every few hours, copy all data
- Inefficient, expensive

Incremental Extract + Bulk Merge

- Tail table or redo logs to obtain changes C
- Merge C into existing table T
- Needs full table scan of T
- Typically done every few hours!

What If we want fresher data?

- Do analytics on a read-only copy!
- But, scales poorly for analytical scans!

Recipe: Database Ingestion

Sqoop + Datasource

```
// Command to extract incrementals using sqoop
bin/sqoop import \
  -Dmapreduce.job.user.classpath.first=true \
  --connect jdbc:mysql://localhost/users \
  --username root \
  --password ***** \
  --table users \
  --as-avrodatafile \
  --target-dir \
s3:///tmp/sqoop/import-1/users
```

Extract new changes to users table in MySQL, as
avro data files on DFS

```
// Spark Datasource
Import com.uber.hoodie.DataSourceWriteOptions._
// Use Spark datasource to read avro
Dataset<Row> inputDataset
spark.read.avro('s3:///tmp/sqoop/import-1/users/*');

// save it as a Hoodie dataset
inputDataset.write.format("com.uber.hoodie")
  .option(HoodieWriteConfig.TABLE_NAME, "hoodie.users")
  .option(RECORDKEY_FIELD_OPT_KEY(), "userID")
  .option(PARTITIONPATH_FIELD_OPT_KEY(), "country")
  .option(PRECOMBINE_FIELD_OPT_KEY(), "last_mod")
  .option(OPERATION_OPT_KEY(), UPSERT_OPERATION_OPT_VAL())
  .mode(SaveMode.Append);
  .save("/path/on/dfs")
```

Use your fav datasource to read extracted data
and directly “upsert” the users table on DFS/Hive

Completes in few mins!!

Recipe: Database Ingestion

Sqoop + DeltaStreamer

```
// DeltaStreamer command to ingest sqoop incrementals
```

```
spark-submit \  
  --class com.uber.hoodie.utilities.deltastreamer.HoodieDeltaStreamer \  
  /path/to/hoodie-utilities-*-SNAPSHOT.jar \  
  --props s3://path/to/dfs-source.properties \  
  --schemaprovider-class com.uber.hoodie.utilities.schema.FilebasedSchemaProvider \  
  --source-class com.uber.hoodie.utilities.sources.AvroDFSSource \  
  --source-ordering-field last_mod \  
  --target-base-path s3:///path/on/dfs \  
  --target-table uber.employees \  
  --op UPSERT
```

```
// dfs-source-properties
```

```
include=base.properties
```

```
# Key generator props
```

```
hoodie.datasource.write.recordkey.field=_userID
```

```
hoodie.datasource.write.partitionpath.field=country
```

```
# Schema provider props
```

```
hoodie.deltastreamer.filebased.schemaprovider.source.schema.file=s3:///path/to/users.avsc
```

```
# DFS Source
```

```
hoodie.deltastreamer.source.dfs.root=s3:///tmp/sqoop
```

Recipe: Deduping Logs

Problem

High-scale timeseries data

- Several billions/day
- Few millions/sec
- Heavily aggregated

Cause of duplicates

- Client retries/failures/network errors
- Atleast-once data pipes

Overcounting problems

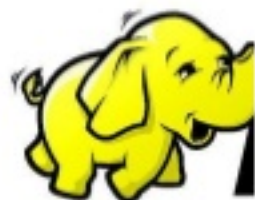
- More impressions => more \$
- Low fidelity data

Produce impression events



Impressions

Replicate
w/o
duplicates



Impressions

id	string
datestr	string
time	long
...	...

Recipe: Deduping Logs

Typical Approach

Pass it on to clients

- Very inefficient
- Often inconsistent

Key-Value Stores

- Non trivial scaling issues
- Expensive $O(\text{num_events})$ storage/lookup

Periodic data cleansing

- Duplicates have leaked to consumers anyway!
- Expensive

Recipe: Deduping Logs

Filtering Duplicates

```
// Deltastreamer command to ingest kafka events, dedupe, ingest
spark-submit --class com.uber.hoodie.utilities.deltastreamer.HoodieDeltaStreamer \
  /path/to/hoodie-utilities-*-SNAPSHOT.jar \
  --props s3://path/to/kafka-source.properties \
  --schemaprovider-class com.uber.hoodie.utilities.schema.SchemaRegistryProvider \
  --source-class com.uber.hoodie.utilities.sources.AvroKafkaSource \
  --source-ordering-field time \
  --target-base-path s3:///hoodie-deltastreamer/impressions --target-table uber.impressions \
  --op BULK_INSERT
  --filter-dupes
```

```
// kafka-source-properties
include=base.properties
# Key fields, for kafka example
hoodie.datasource.write.recordkey.field=id
hoodie.datasource.write.partitionpath.field=datestr
# schema provider configs
hoodie.deltastreamer.schemaprovider.registry.url=http://localhost:8081/subjects/impressions-value/versions/latest
# Kafka Source
hoodie.deltastreamer.source.kafka.topic=impressions
#Kafka props
metadata.broker.list=localhost:9092
auto.offset.reset=smallest
schema.registry.url=http://localhost:8081
```

Recipe: Incremental ETL

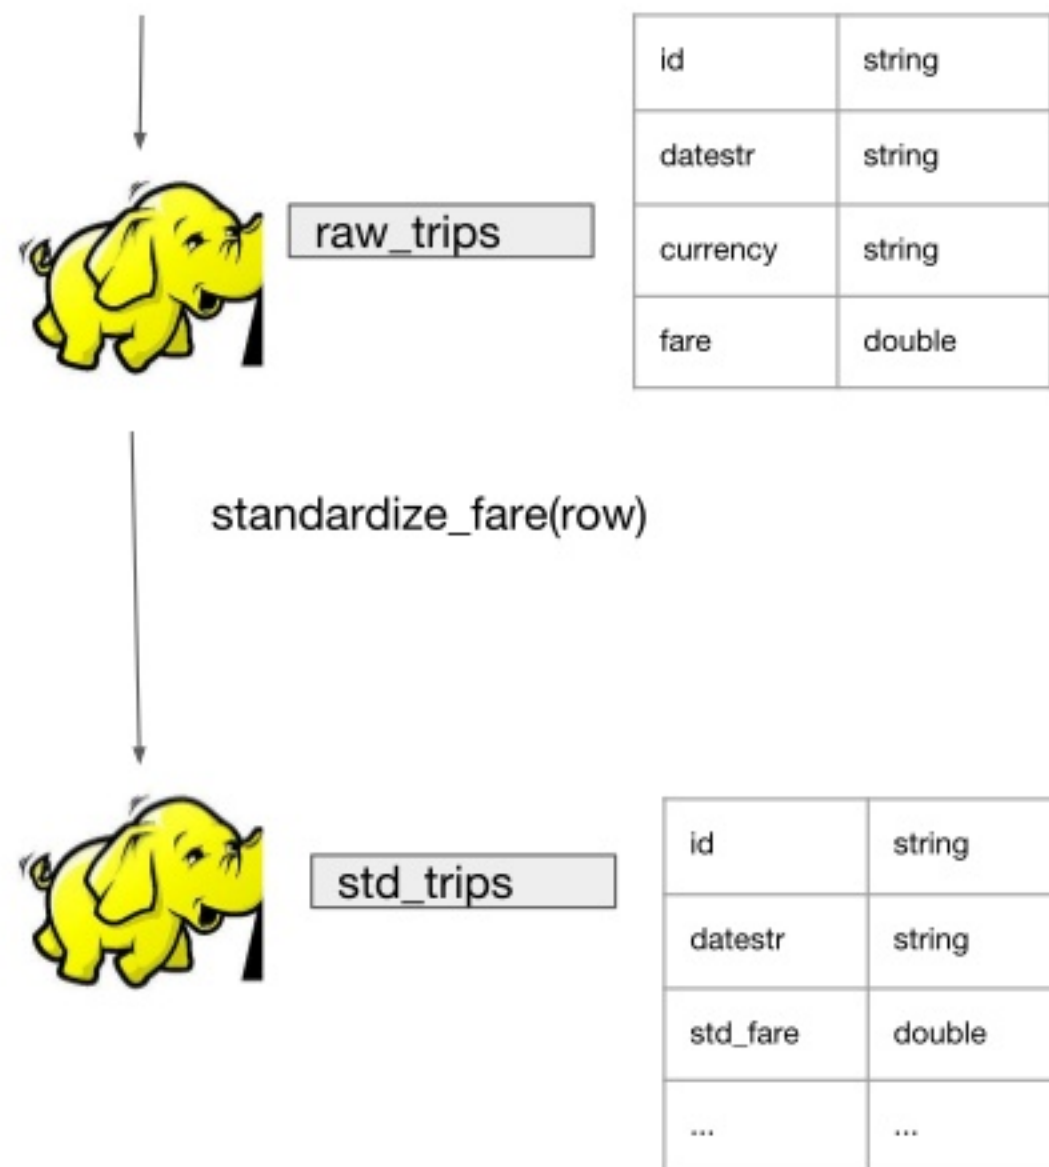
Problem

Multi stage ETL DAGS

- Very common in batch analytics
- Large amount of data
- Big Spark use-case!

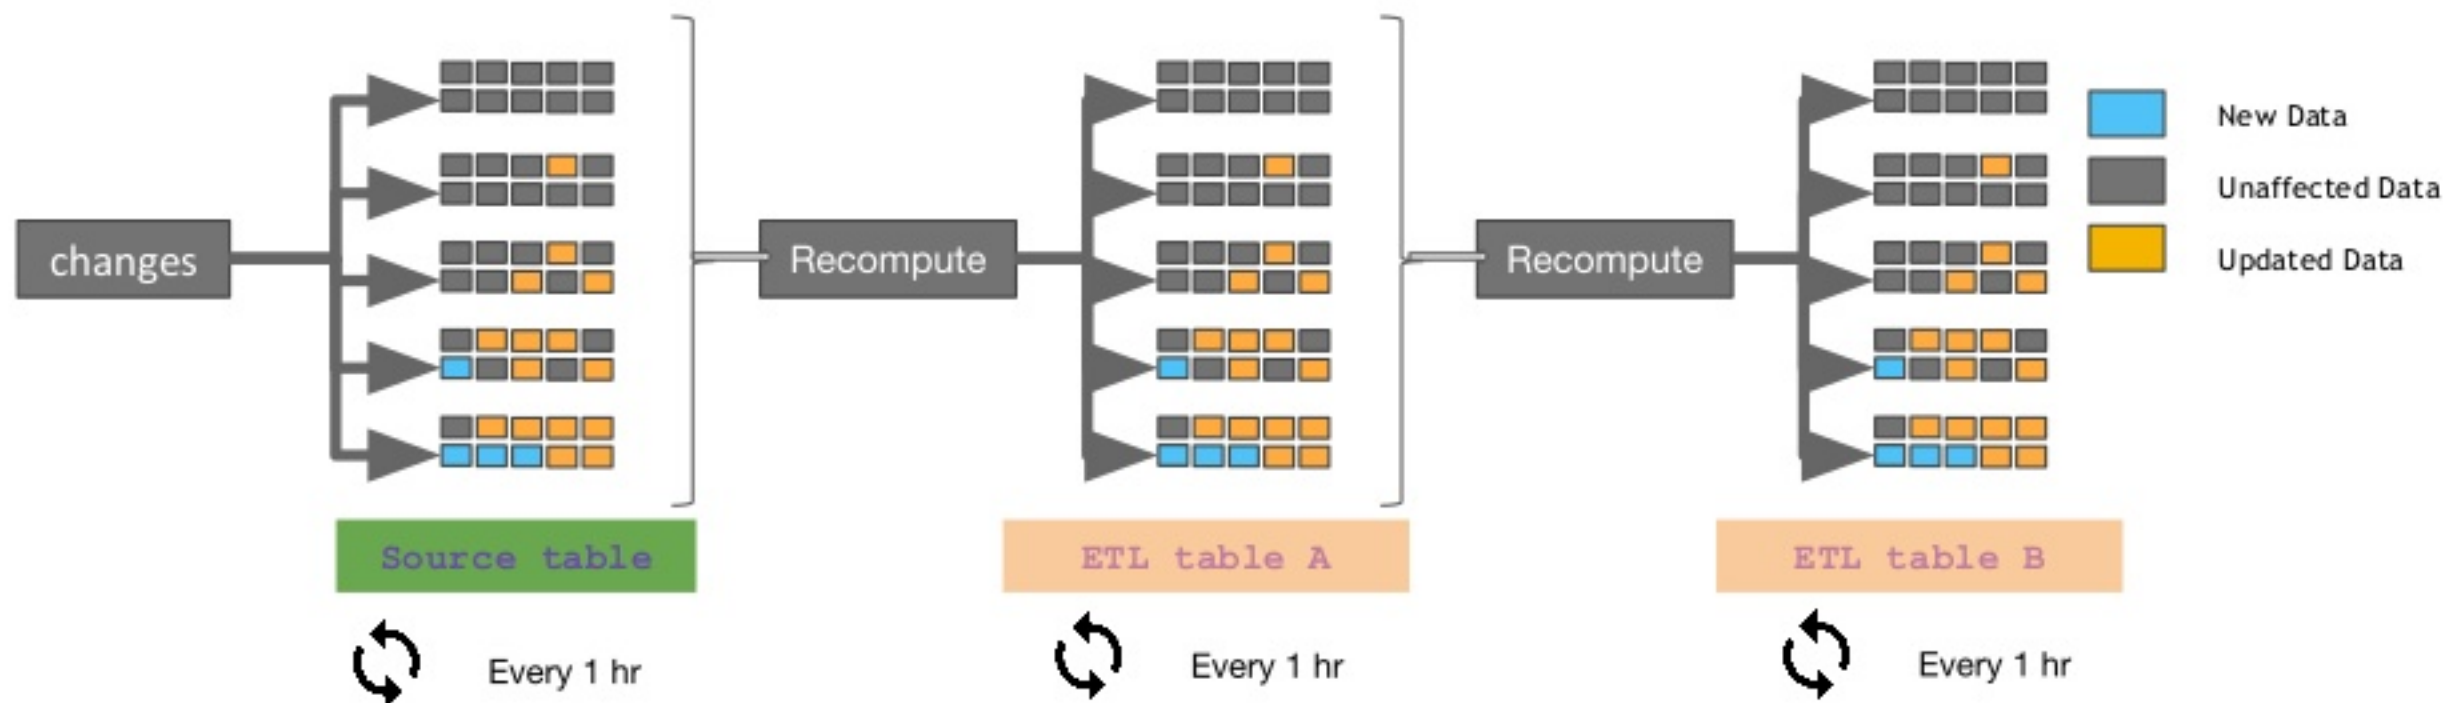
Derived/ETL tables

- Keep afresh with new/changed source data
- Star schema/warehousing



Recipe: Incremental ETL

Typical Approach



Recompute in every stage

- Heuristics => limit scan range
- Slowwww

Minimum 3 hr latency!

- Waste of resourcing in recomputing

Recipe: Incremental ETL

Incremental Spark Pipelines

Incrementally pull

- Avoid recomputes!
- Finishes in few mins!

Transform + upsert

- Avoid rewriting all data

```
// Spark Datasource
Import com.uber.hoodie.{DataSourceWriteOptions, DataSourceReadOptions}._

// Use Spark datasource to read avro
Dataset<Row> hoodieIncViewDF = spark.read().format("com.uber.hoodie")
    .option(VIEW_TYPE_OPT_KEY(), VIEW_TYPE_INCREMENTAL_OPT_VAL())
    .option(DataSourceReadOptions.BEGIN_INSTANTTIME_OPT_KEY(), commitInstantFor8AM)
    .load("s3://tables/raw_trips");

Dataset<Row> stdDF = standardize_fare(hoodieIncViewDF)

// save it as a Hoodie dataset
inputDataset.write.format("com.uber.hoodie")
    .option(HoodieWriteConfig.TABLE_NAME, "hoodie.std_trips")
    .option(RECORDKEY_FIELD_OPT_KEY(), "id")
    .option(PARTITIONPATH_FIELD_OPT_KEY(), "datestr")
    .option(PRECOMBINE_FIELD_OPT_KEY(), "time")
    .option(OPERATION_OPT_KEY(), UPSERT_OPERATION_OPT_VAL())
    .mode(SaveMode.Append);
    .save("/path/on/dfs")
```

Open Source Community

We Love Contributions!

Github

- GitHub -> <https://github.com/uber/hudi>
- Quickstart -> <https://uber.github.io/hudi/quickstart.html>
- Documentation -> <https://uber.github.io/hudi/>
- Slack channel -> <https://hoodielib.slack.com/>

Contributions

- Small & large, code & design & feature asks
- 10+ organizations incl Uber, Shopify, DoubleVerify, Vungle, Udemy..

Roadmap

- Global index to ease migrations
- Ease of managing compactions
- RTView performance

WE ARE HIRING!

Reach out to us

hadoop-platform-jobs@uber.com

Thank you. Questions?

nagarwal@uber.com

vinoth@uber.com

Proprietary and confidential © 2016 Uber Technologies, Inc. All rights reserved. No part of this document may be reproduced or utilized in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval systems, without permission in writing from Uber. This document is intended only for the use of the individual or entity to whom it is addressed and contains information that is privileged, confidential or otherwise exempt from disclosure under applicable law. All recipients of this document are notified that the information contained herein includes proprietary and confidential information of Uber, and recipient may not make use of, disseminate, or in any way disclose this document or any of the enclosed information to any person other than employees of addressee to the extent necessary for consultations with authorized personnel of Uber.



UBER