

Machine Learning for AdTech in Action

Cyrille Dubarry
Han Ju

#SAISML11

AdTech introduction

Search and Performance


kitchen knife

All Images Shopping News Maps More Settings Tools


About 290,000,000 results (0.62 seconds)

See kitchen knife


Sponsored ⓘ




ZWILLING Bloc de couteaux,...
€209.00
Zwilling Staub
Free shipping
By Google




Set de 5 couteaux wasa...
€196.70
FourniResto.com
★★★★★ (10)
By Google



Couteau de cuisine Nagiri...
€81.60
Nippon Co
By Kelkoo



Couteau santoku damas shun...
€150.40
FourniResto.com
★★★★★ (124)
By Google



Couteau santoku wasabi black 1...
€47.68
FourniResto.com
★★★★★ (107)
By Google

»

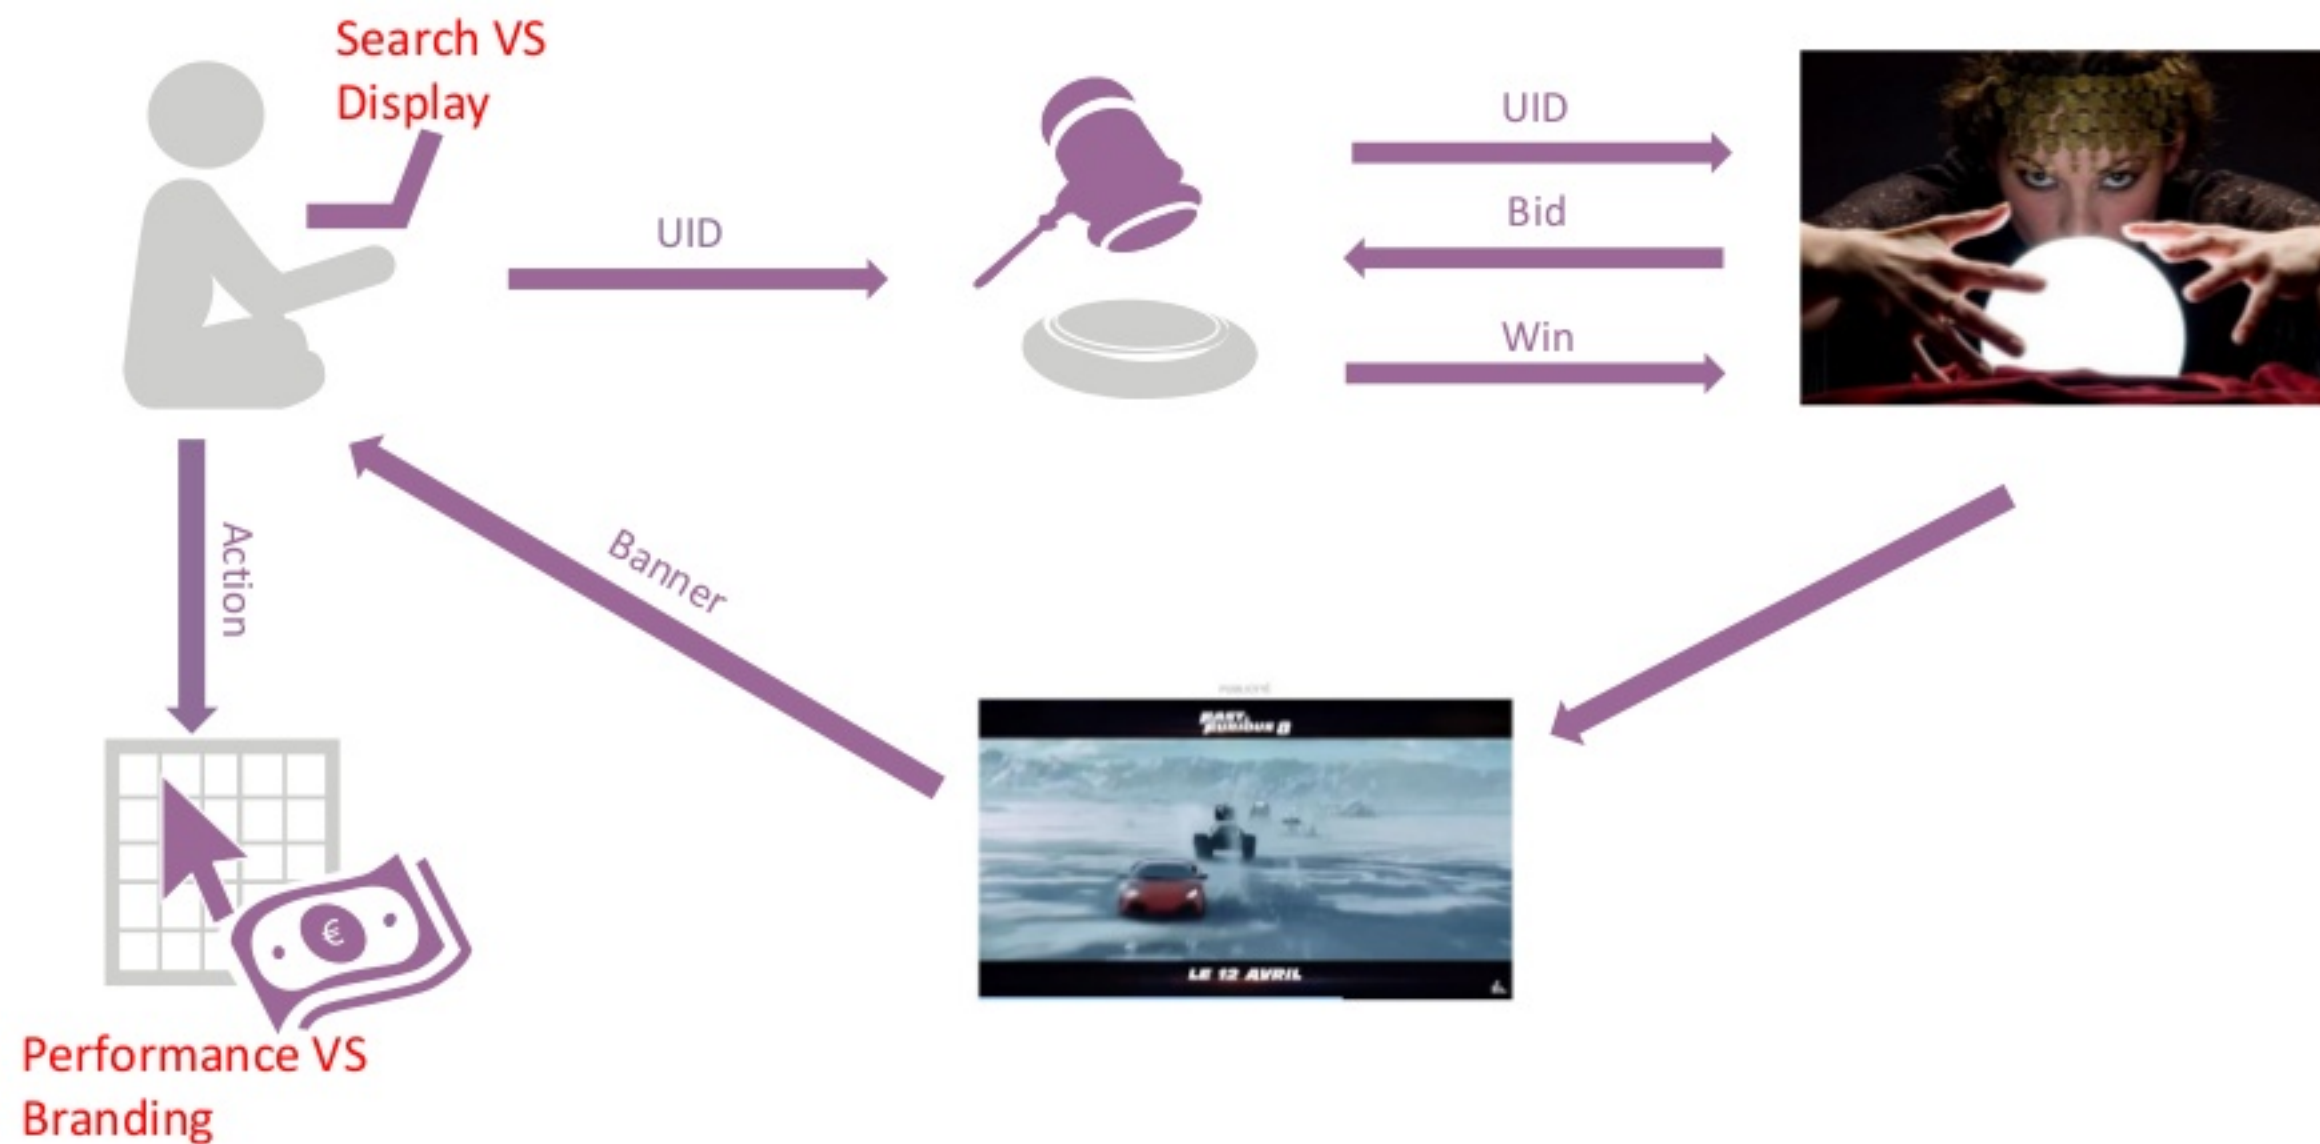
Video/Display and Branding

Cette première phase des négociations fait l'objet d'un « *rapport conjoint* » de quinze pages qui revient sur les éléments concernant les trois dossiers jugés prioritaires par Bruxelles et Londres : le règlement financier de la séparation, les droits des citoyens expatriés et la gestion de la frontière entre la république d'Irlande et la province britannique d'Irlande du Nord.




Le dossier est désormais entre les mains du Conseil européen, l'instance qui regroupe les dirigeants des Etats membres : ces derniers devront valider l'accord lors d'un sommet à Bruxelles, les 14 et 15 décembre.

Workflow of a user browsing






Real Time Bidding

RTB workflow



What does the winner pay?

2nd price auction



True Value for Video Branding

Assumption: the bidder is paid a fixed amount CPV if and only if the video ad is watched entirely.

True Value?

$$\mathbb{E}[Revenue] = \mathbb{E}[CPV \times 1_{watched}] = CPV \times \mathbb{P}\{Watched\}$$

Estimating the True Value

Which constraints?

- **Features:** several hundred thousands modalities
- **Training frequency:** several times a day
- **Number of predictions:** a million per second

Logistic regression VS DNN



Validation metric:
$$-LLH = -\sum_{i=1}^N O_i \ln(p_i) + (1 - O_i) \ln(1 - p_i)$$

Prediction model	Score = -LLH	Training time	Prediction complexity
Logistic regression	0.1293	70 min	40
Logistic regression with crosses	0.1272	142 min	86
Deep learning (11 layers x 1062 neurons) 1 epoch	0.1257	26h	$40 \times 1062 + 10 \times 1062^2 + 1062$ $= 11\,321\,982$
Deep learning (11 layers x 1062 neurons) 3 epochs	0.1250	78h	11 321 982


<https://cloud.google.com/blog/big-data/2017/02/using-google-cloud-machine-learning-to-predict-clicks-at-scale>

Reminder: Logistic Regression

$$\mathbb{P}\{Watched|x\} = \sigma(< w.x >) = \frac{1}{1 + e^{-<w.x>}}$$

x: features including historical and contextual information

w: model parameters

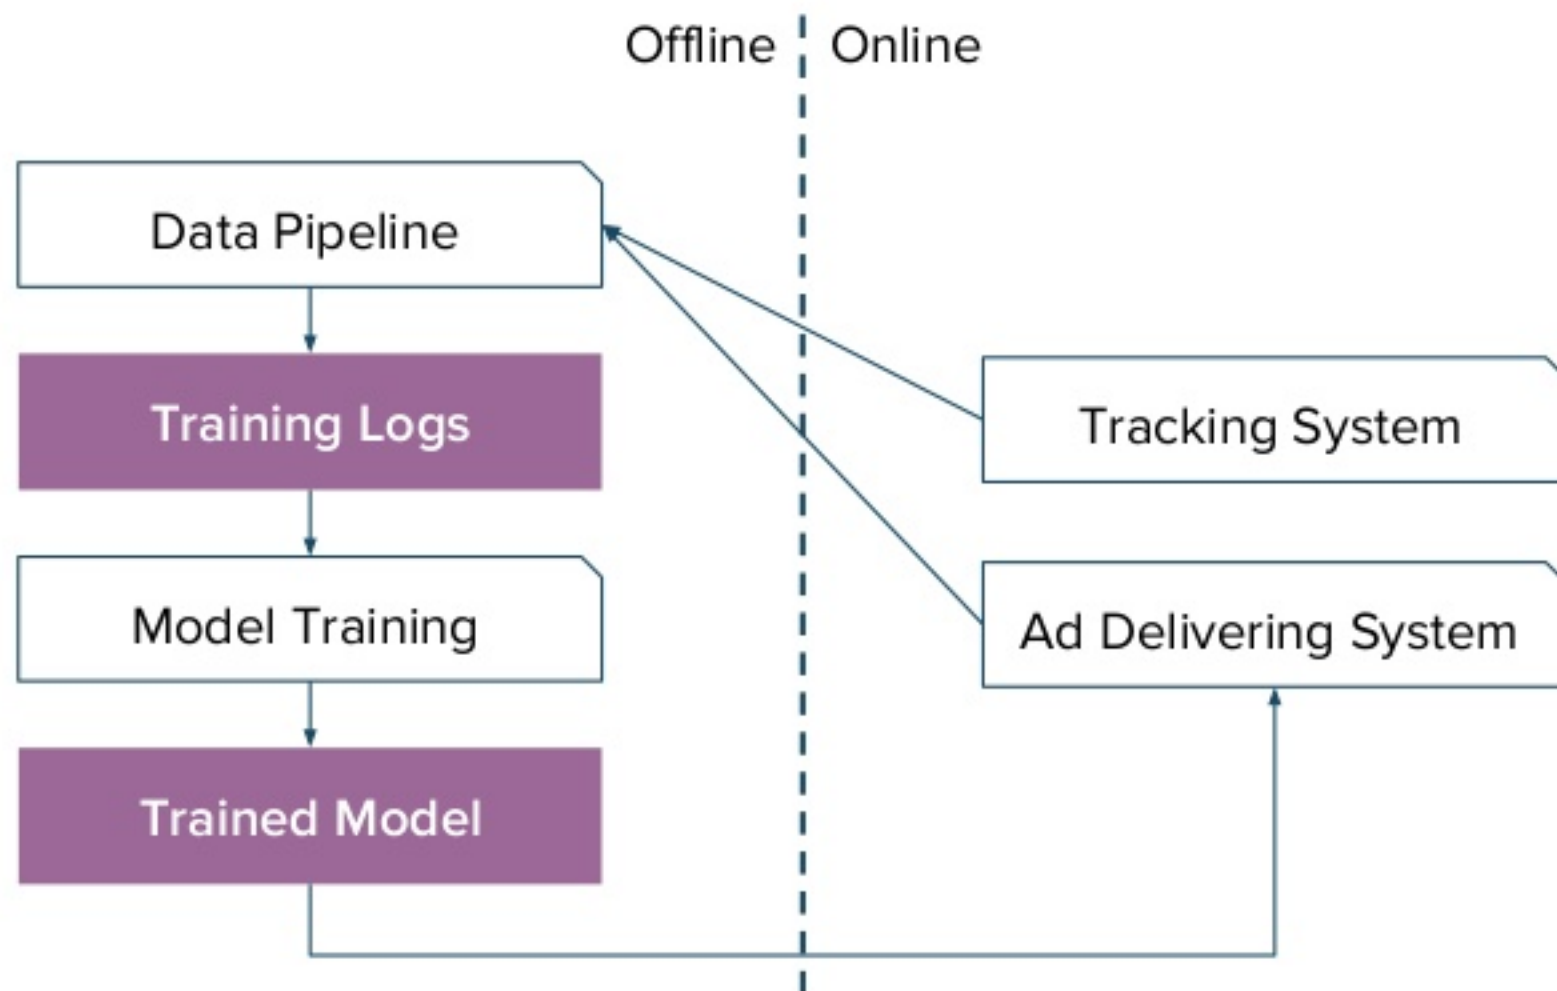


Machine Learning Stack at Teads

ML stack at Teads

- Motivation
- Model training
- Model serving
- Offline experiments

Classic setup



Challenges

Real-time prediction



150 ms response time in RTB business



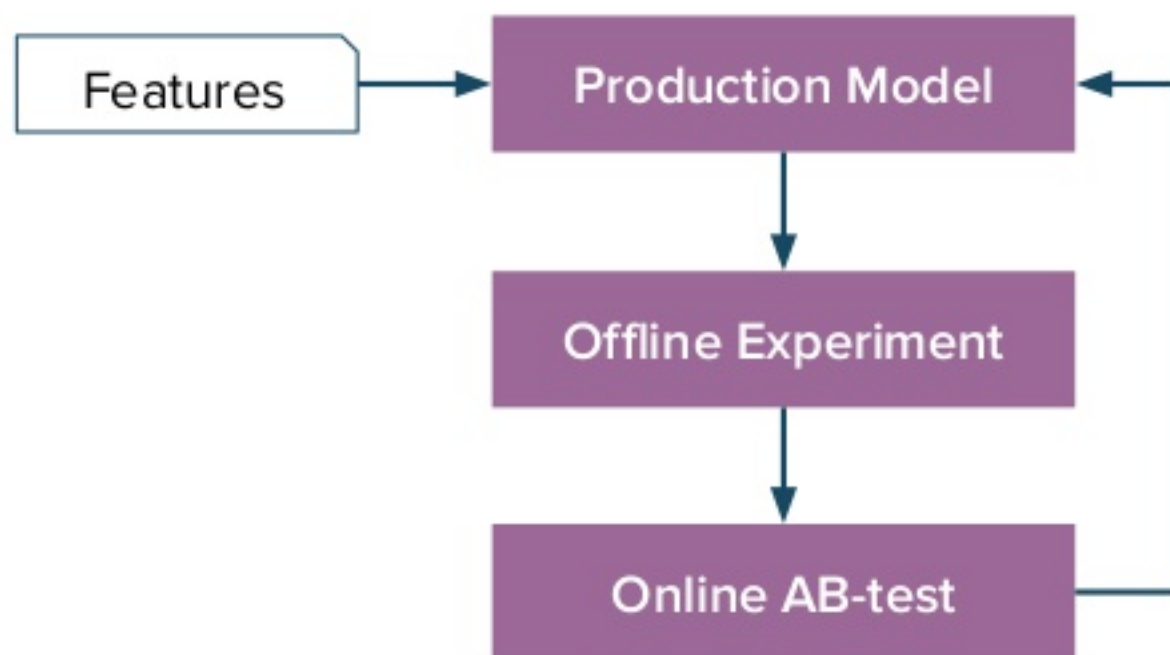
40,000 requests / second
1,000,000 model predictions / second



Model serving outside of Spark

Challenges

Flexible and reliable model improvement cycle



Why build our own stack?

No available framework/library back in 2015 that satisfied all our needs

MLlib limitations

- GLM vector size limitations
- Model serving with spark dependency

Decided to do it ourselves

- Scala and Spark

Model Training

Model training



Different use cases in different teams

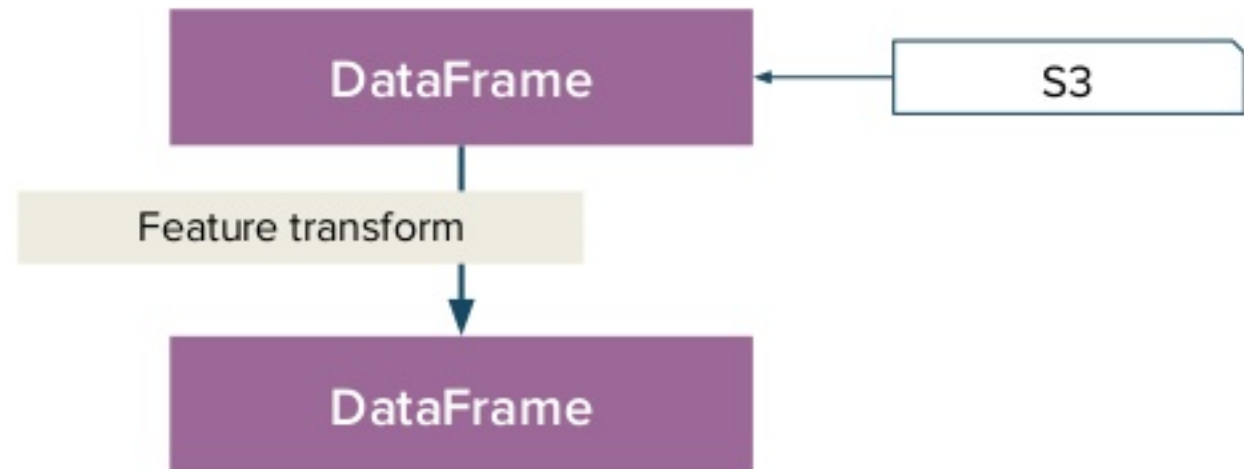
- Not the same features
- Not the same model settings



Model training steps should be configurable by using a **configuration object**

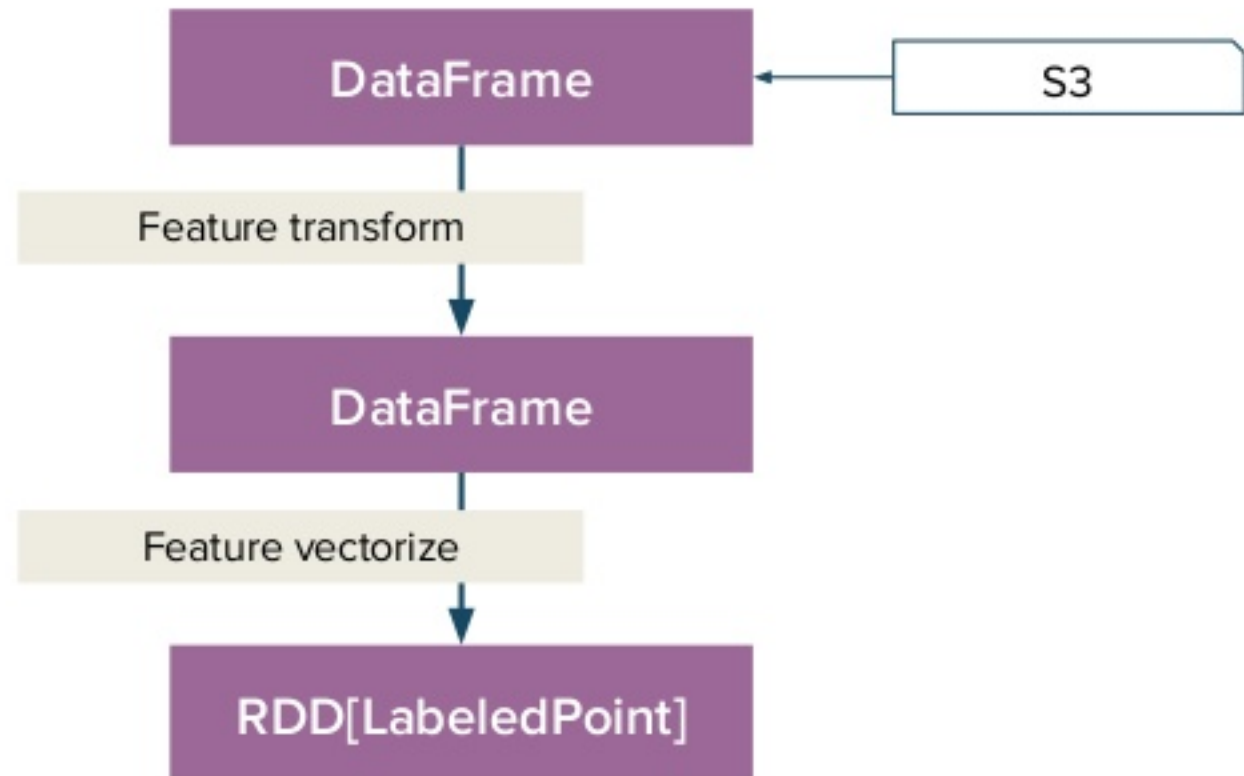
Feature transformation

```
"transform" : [  
  {  
    "type" : "Bucketize",  
    "input" : "ad_duration",  
    "output" : "ad_duration_bucket",  
    "thresholds" : [  
      10.0,  
      15.0,  
      20.0,  
      25.0,  
      30.0,  
      45.0,  
      75.0,  
      120.0  
    ],  
  },  
  ...  
],  
...
```



Feature vectorization

```
...,
"vectorize" : [
  {
    "type" : "ConstantVectorize",
    "value" : 1.0
  },
  {
    "type" : "HashVectorize",
    "inputs" : [
      "ad_duration_bucket",
      ...
    ],
    "inputTypes" : ...,
  ],
  "size" : 2145483000
},
],
"label" : {
  "input" : "event_billable_dbl"
},
...
```

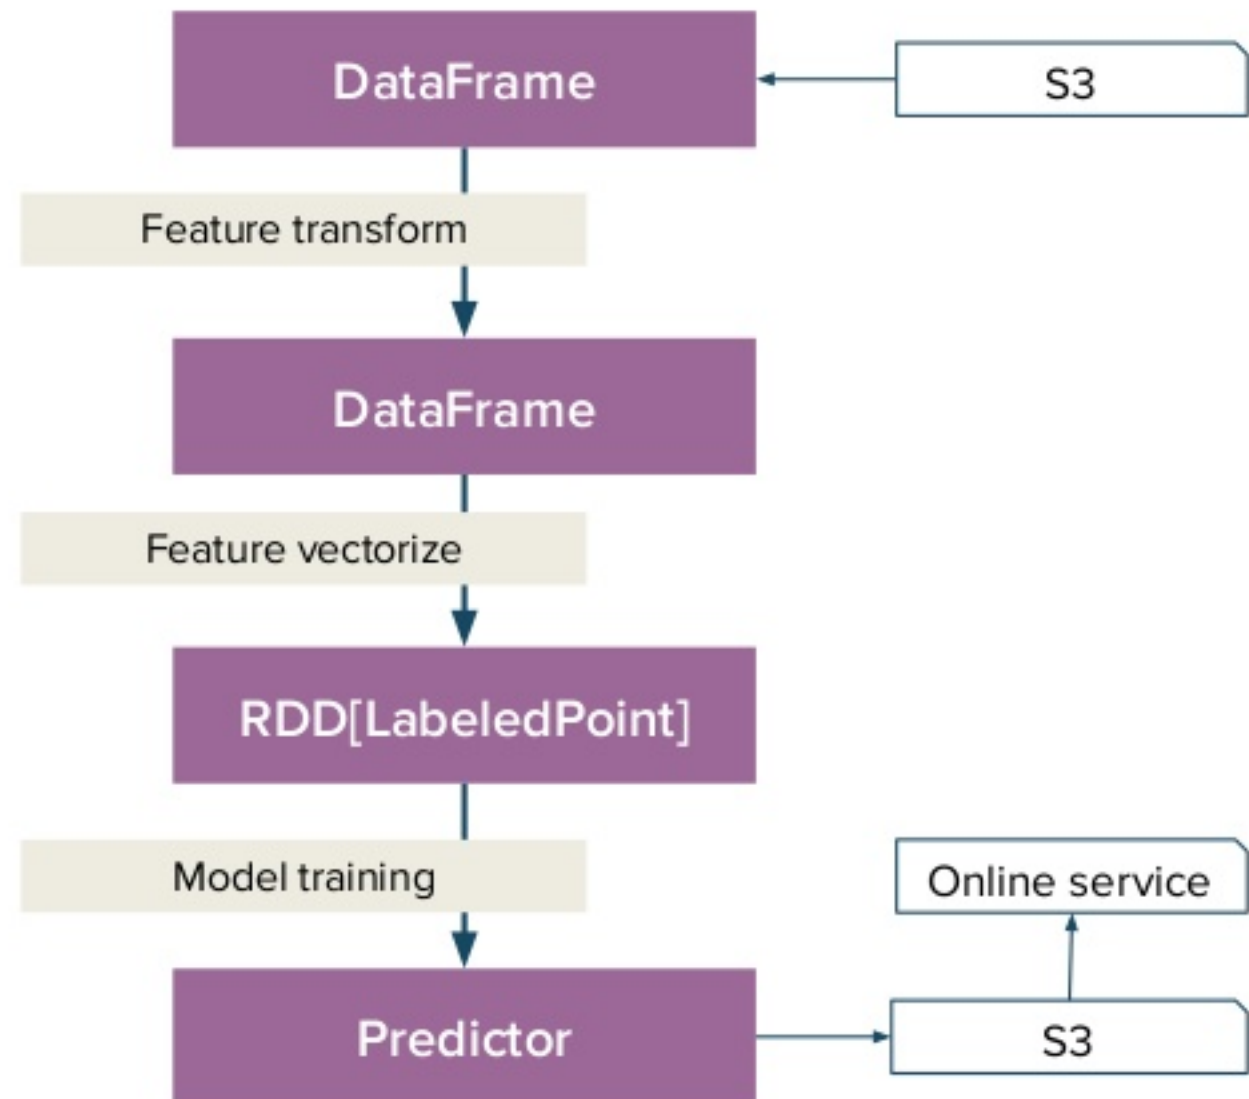


Model training

- Initial model
- Model seeding

```
...  
"initialModel" : {  
  "type" : "Logistic",  
  "intercept" : 0.0,  
  "slopes" : {  
    "type" : "Sparse",  
    "indices" : [  
    ],  
    "values" : [  
    ],  
    "length" : 2145483000  
  }  
},  
...  

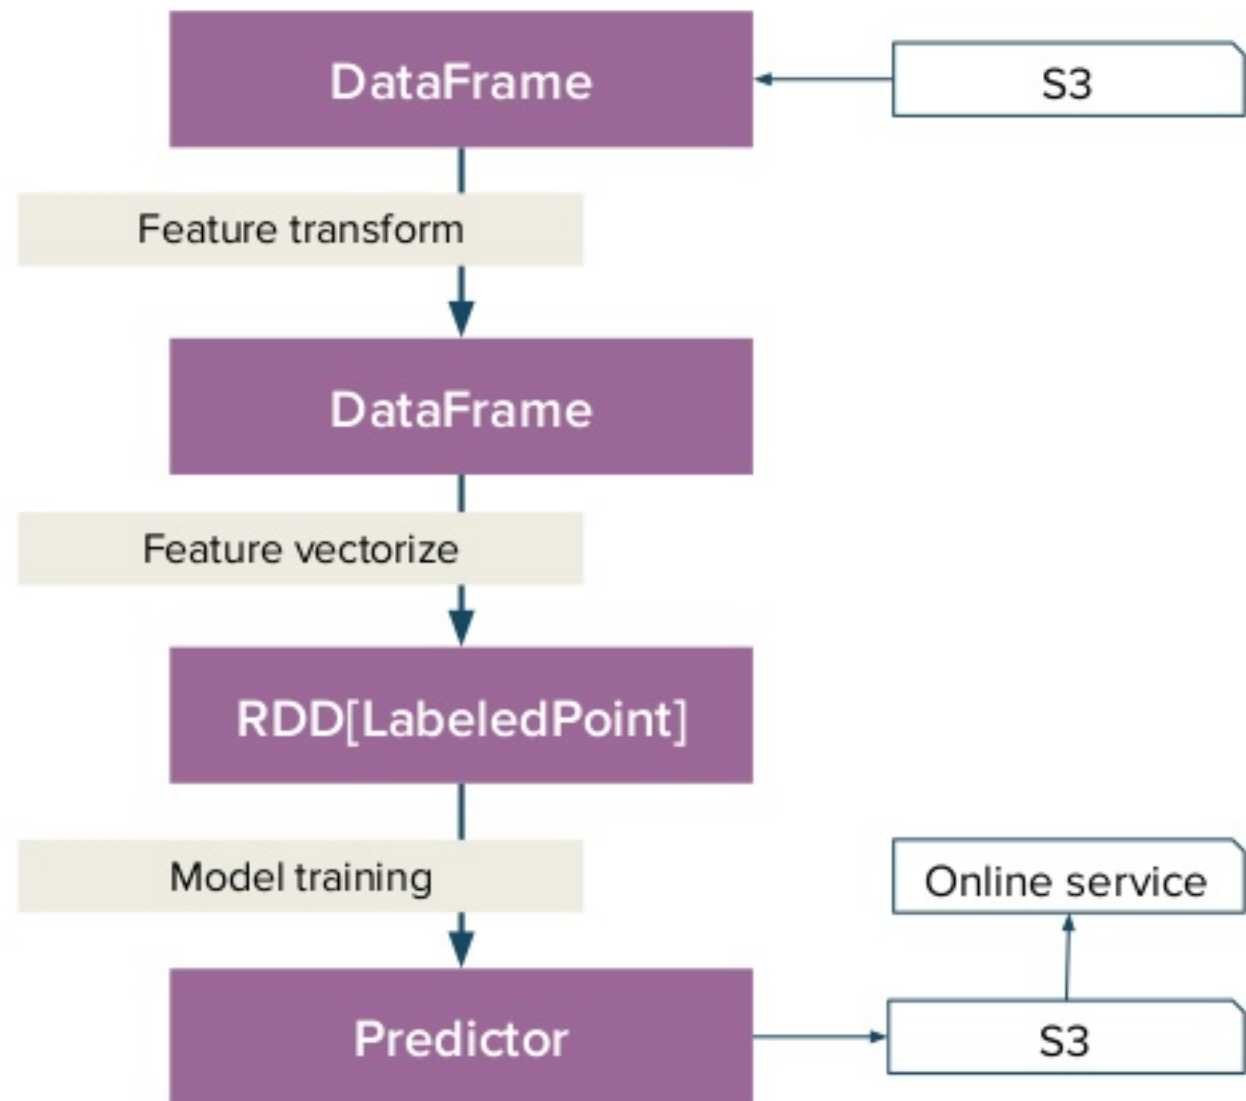
```



Model training

- Other settings

```
...,
"loss" : {
  "type" : "LogisticLoss"
},
"regularization" : {
  "type" : "L2",
  "alpha" : 29.0,
  ...
},
"optimization" : {
  "type" : "LBFGS",
  "maxIterations" : 150,
  "tolerance" : 1.0E-4
},
...
```



Model training

(DataFrame, TrainingConfig) => Predictor

Model training deployment

Model **name tag**

- Unique identifier of a production model

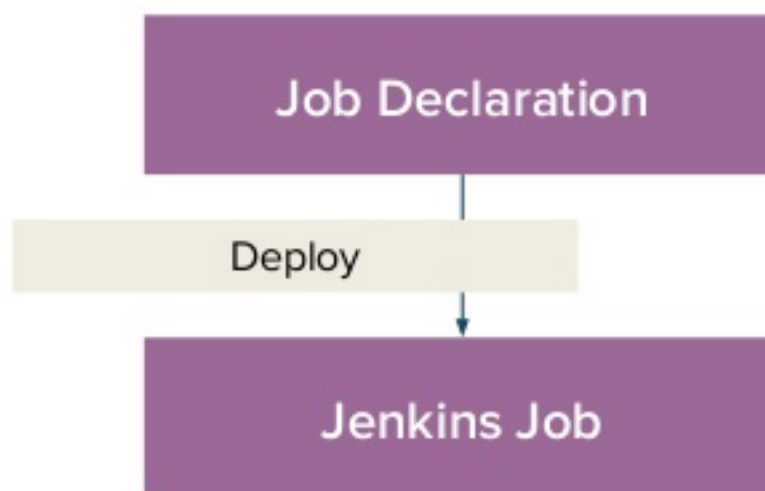
Example:

prod-team1-vtr-prediction-2018-07-15

```
def getTrainingConfig(tag: String): TrainingConfig  
def getLatest(tag: String): Predictor
```

Model training deployment

```
TrainingJob(  
  modelTag = "prod-be-vtr-prediction-2018-07-15",  
  jobLevel = Level.Day,  
  defaultClusterOptions = EmrOptions(coreInstanceCount  
= 200),  
  timeout = Some(24.hours),  
  group = Group.beProd,  
  module = Module.libPrediction("clustering"),  
  logPath = "s3a://...",  
  logSize = 1,  
  defaultConf = Seq(  
    "spark.sql.shuffle.partitions" -> "20000"  
  ),  
  mainClass = Some("...")  
)
```



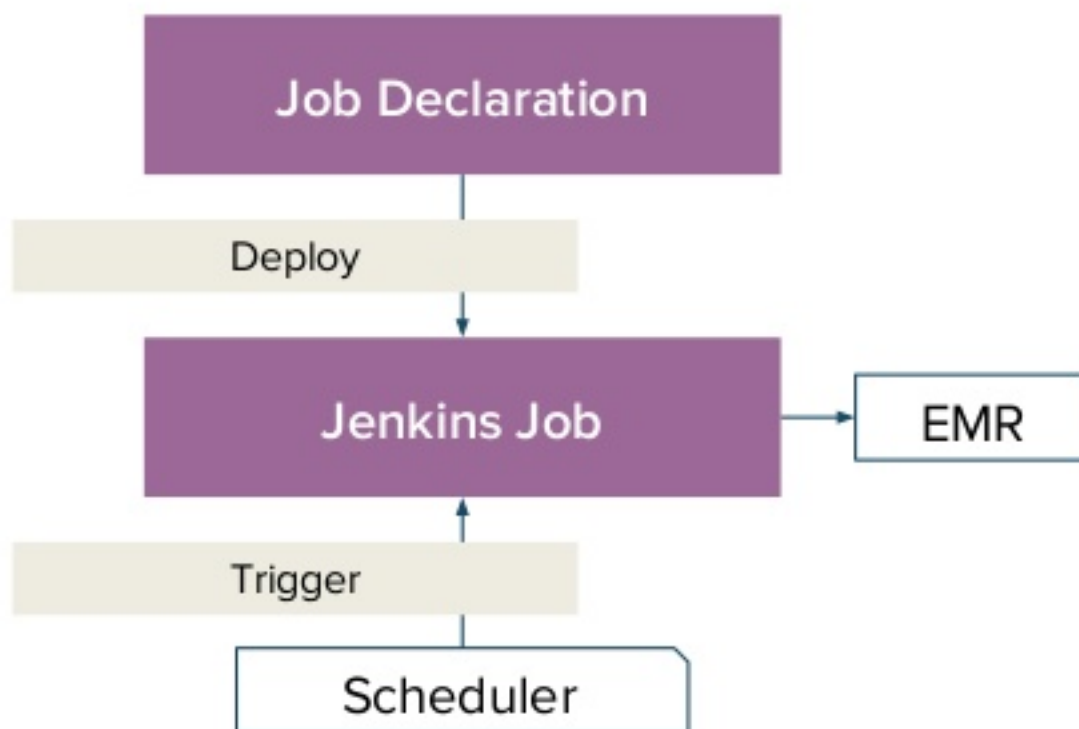
Model training deployment


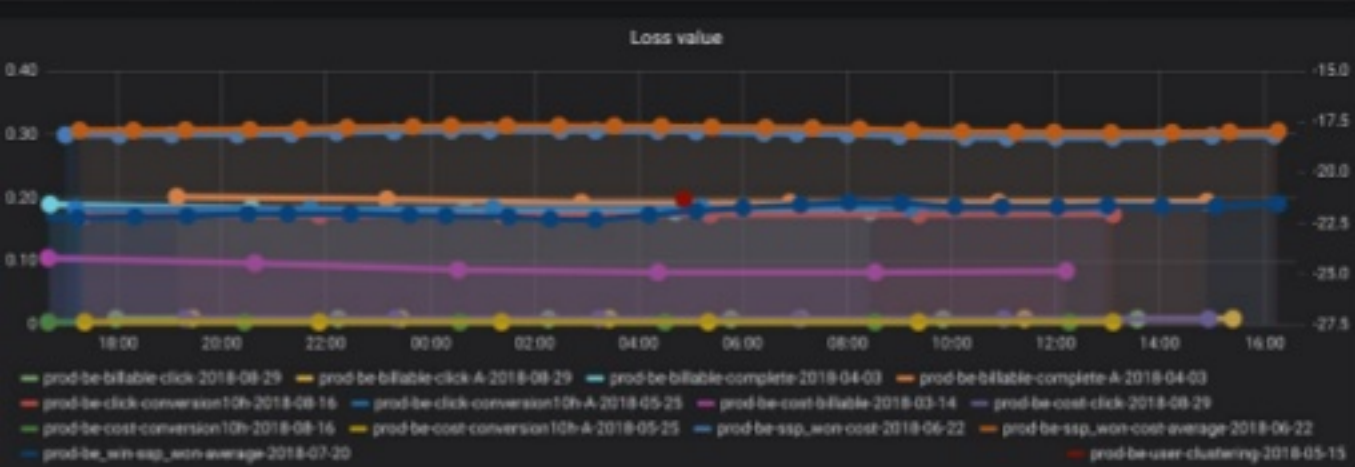
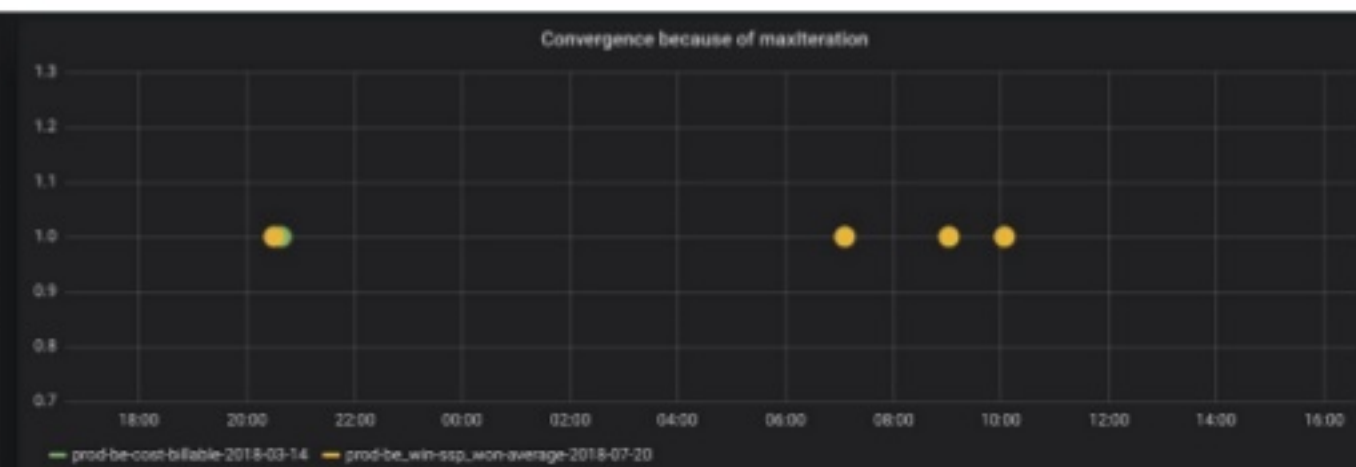
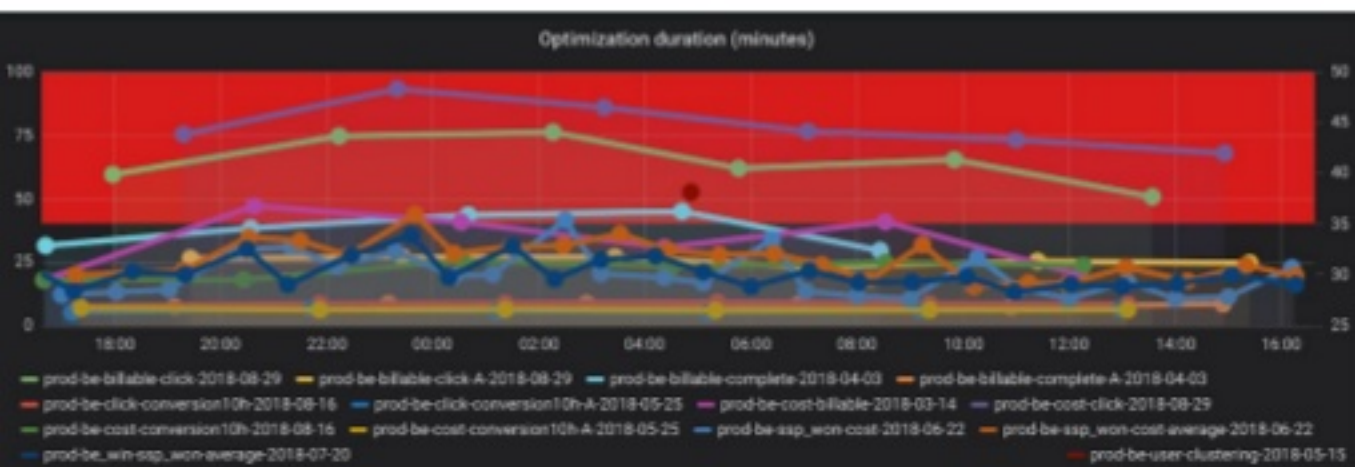
Project training-be-user-clustering-2018-05-15

This build requires parameters:

LAUNCHER_VERSION	0.3.7-2432-g9a05276e
MEMORY	128m
hour	
version	
conf	
log_level	
driver_cores	
driver_memory	
executor_count	
executor_cores	
executor_memory	
cluster_id	
cluster_name	
cluster_tag	

Build





Model Serving

Model serving



Impl as a scala library

- Dependency of online system codebase

Fast and lightweight

- Prefer native code than a prediction service
- Avoid http calls
- No Spark dependency

Predictor

Configuration of model scoring

Components

- Feature transform
- Feature vectorization
- Trained model

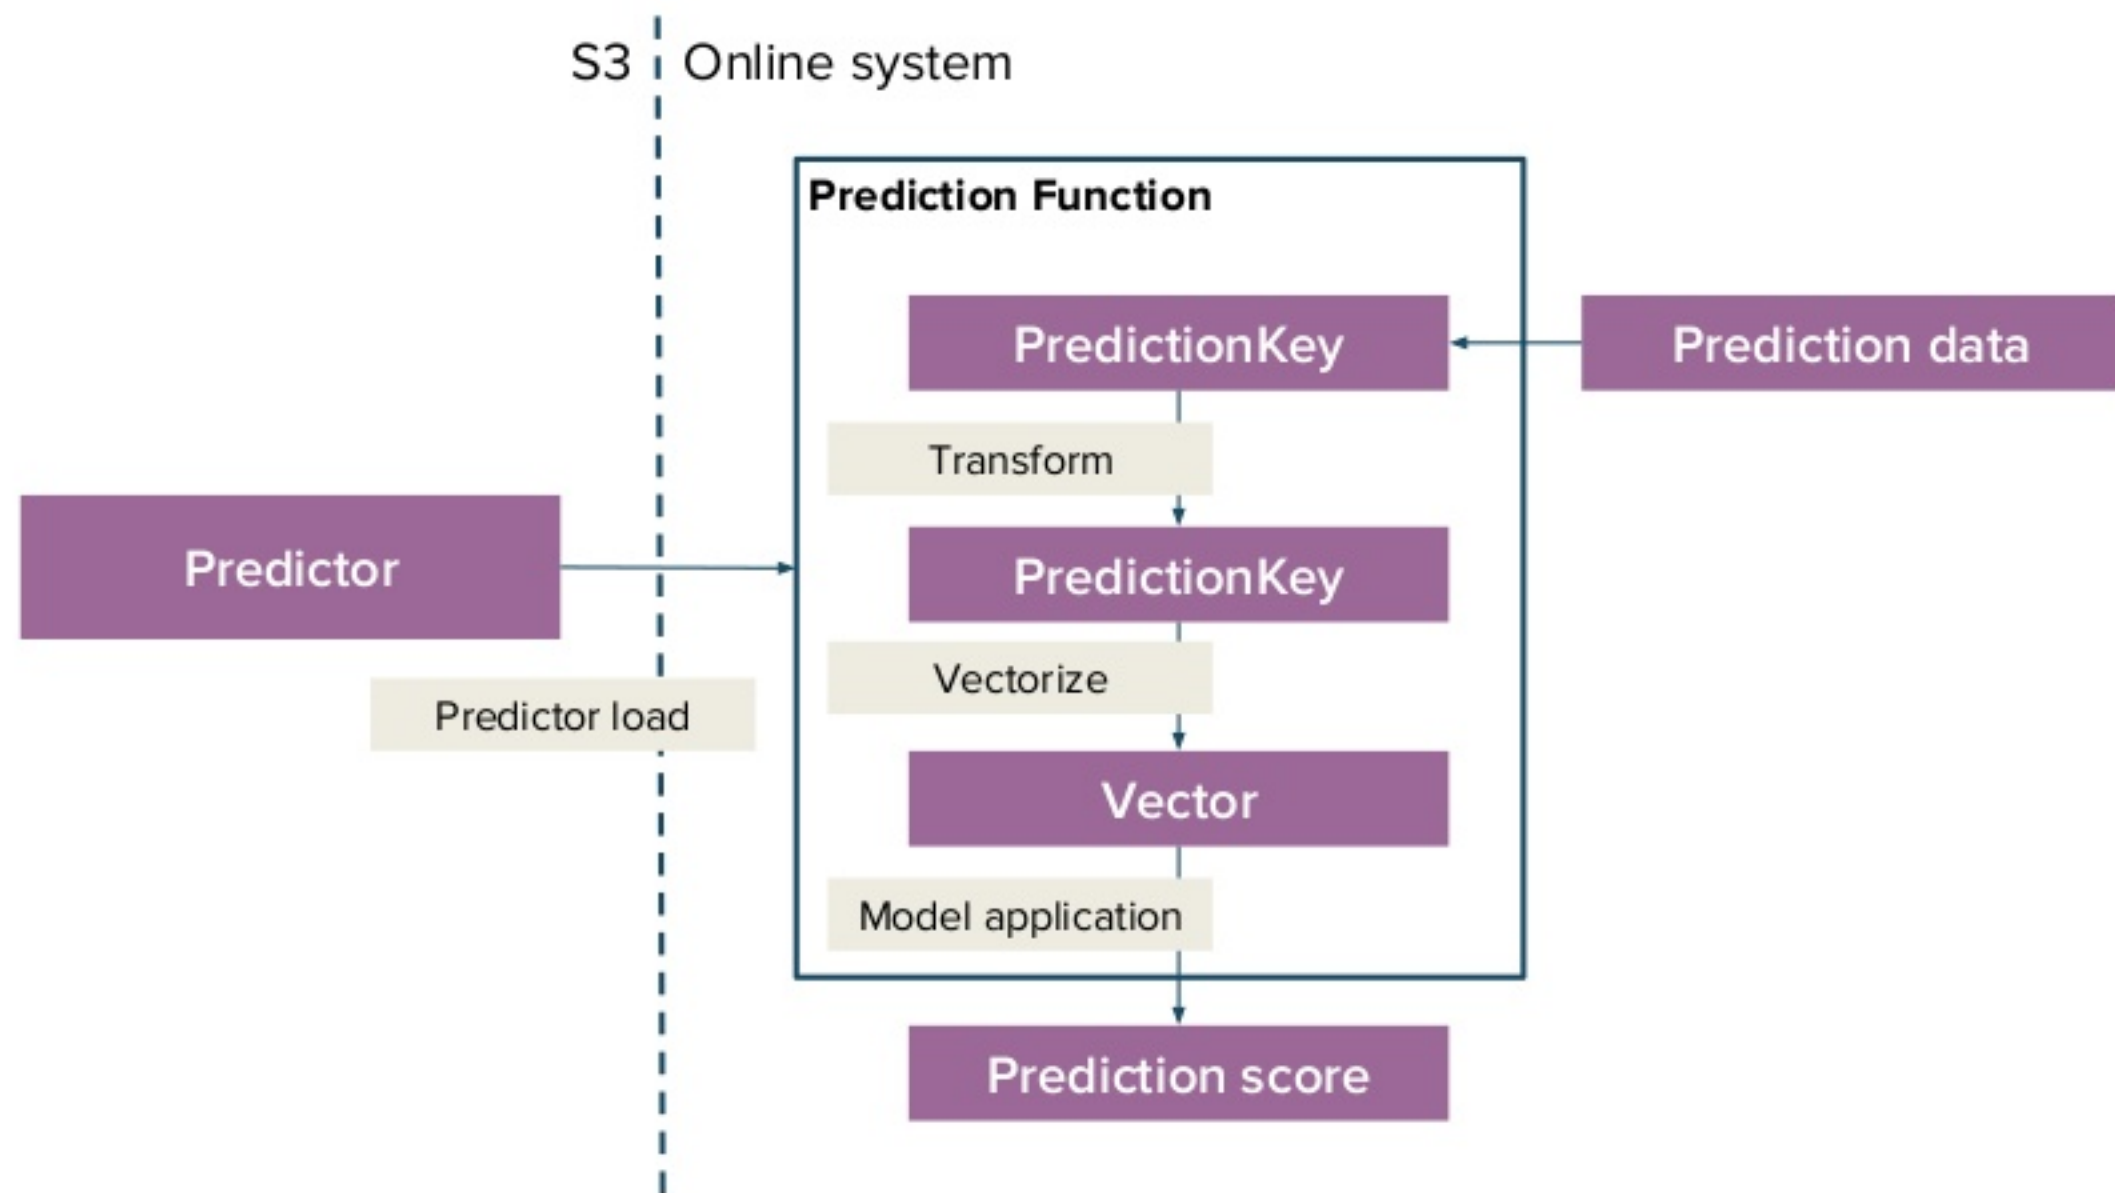
Same feature generation as in training

- Vectors are the same online and offline

```
{  
  "transform" : [...],  
  "vectorize" : [...],  
  "model" : {  
    "type" : "Logistic",  
    "intercept" : -0.59580942,  
    "slopes" : {  
      "type" : "Sparse",  
      "indices" : [  
        5385, 6271, ...  
      ],  
      "values" : [  
        -0.002536684,  
        0.026973965,  
        ...  
      ],  
      "length" : 2145483000  
    }  
  }  
}
```

Online feature transform

```
case class PredictionKey private (values: Vector[Any], schema: Type) {  
  def add[T: Convert](name: String, t: T): PredictionKey = {...}  
  
  def value(idx: Int): Any = {...}  
  
  def get[T: Convert](name: String): T = {...}  
  
  def get[T: Convert](idx: Int): T = {...}  
}
```



Average model prediction scoring takes **55 microseconds** in our production system
(prediction function application time)



Offline Experiments

Offline experiment

Internal tool to evaluate ML models on historical data

Simplify experiment
process

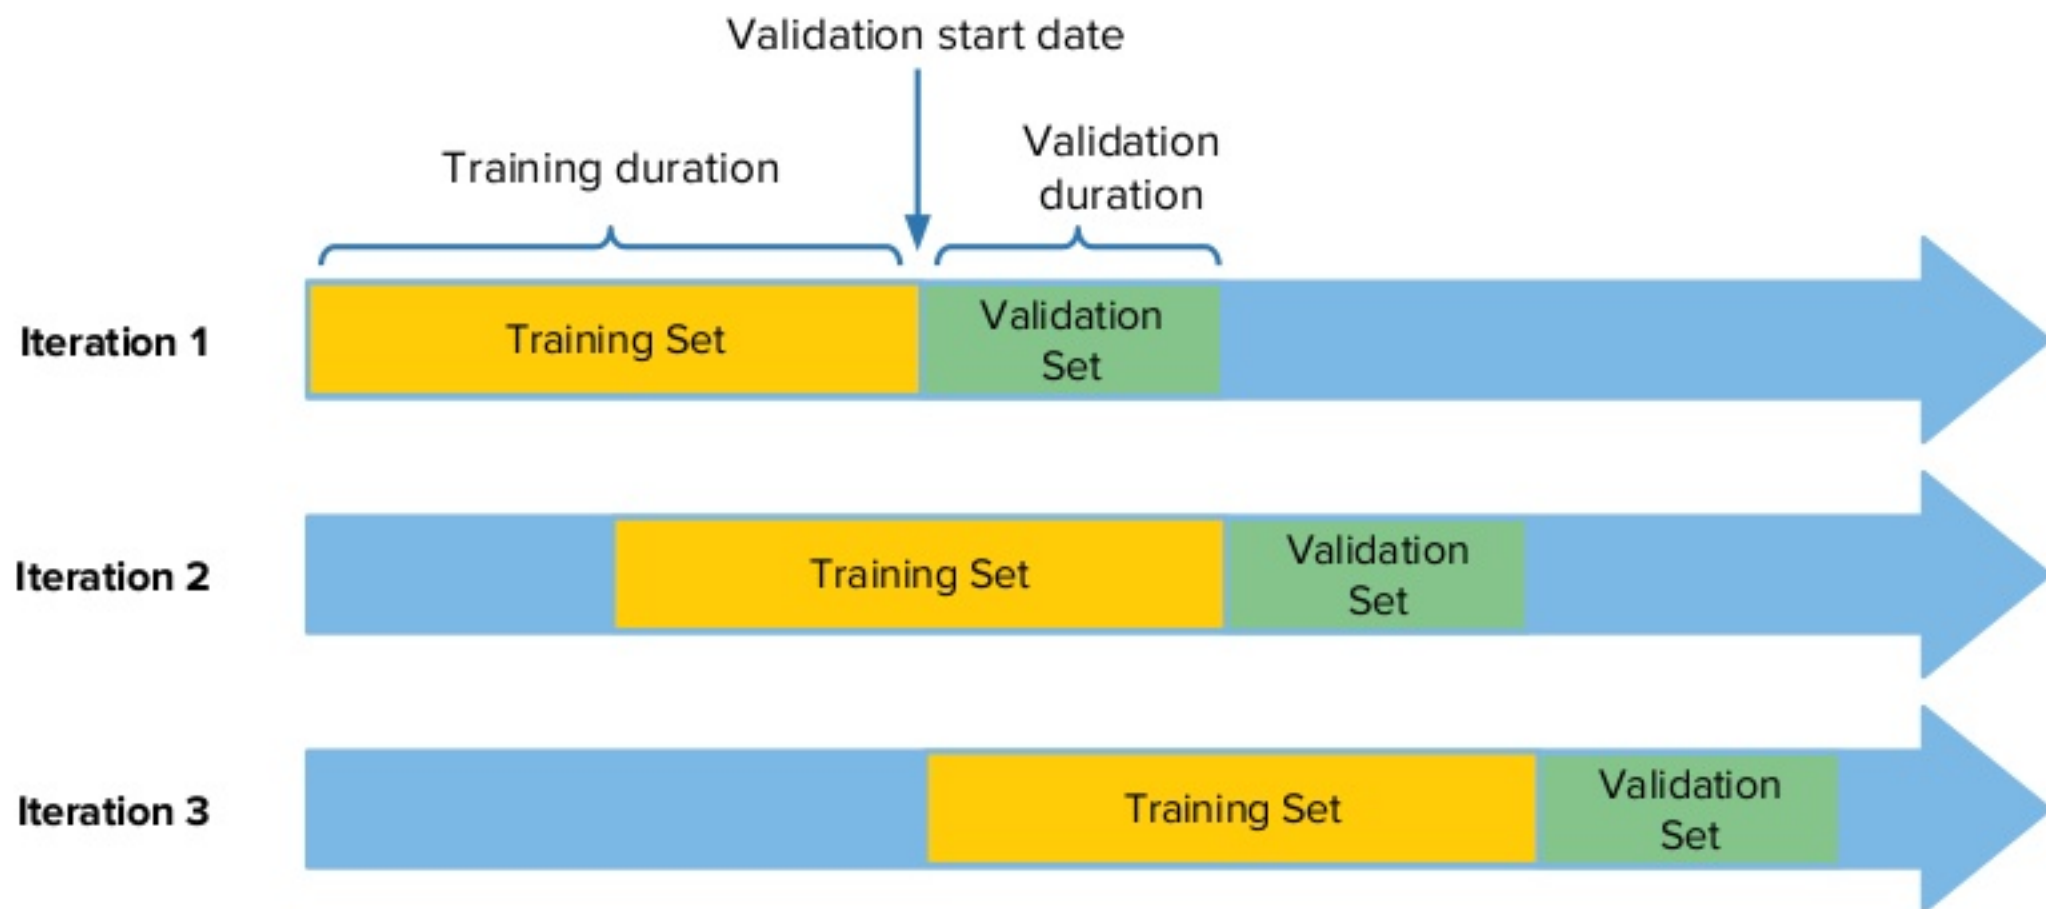
Formalize the
testing protocol



Keep track of the
experiment results

Wide range of
metrics with
confidence interval

Offline experiment



Experiment creation

```
In [8]: datak.experiments.create(  
    estimatorId = Id("b9c1cdbdb5fa2a680ba3e85aa8359a12c236689a"),  
    name = "offline experiment test example",  
    owner = "datakinator-team",  
    trainingDuration = 4,  
    trainingLogId = 43,  
    validationDuration = 4,  
    validationLogId = 43,  
    nbIter = 1,  
    validationStart = "2018-05-01 00:00:00",  
    offset = 0,  
    validationOnlyFeatures = "constantOne",  
    metric_weights = "constantOne",  
    projections = "constantOne,business_model"  
)
```

```
Out[8]: res7: Long = 3316L
```

Notebook as main interface of experiments

- **jupyter** notebook
- **jupyter-scala** kernel

Weighted metrics computation

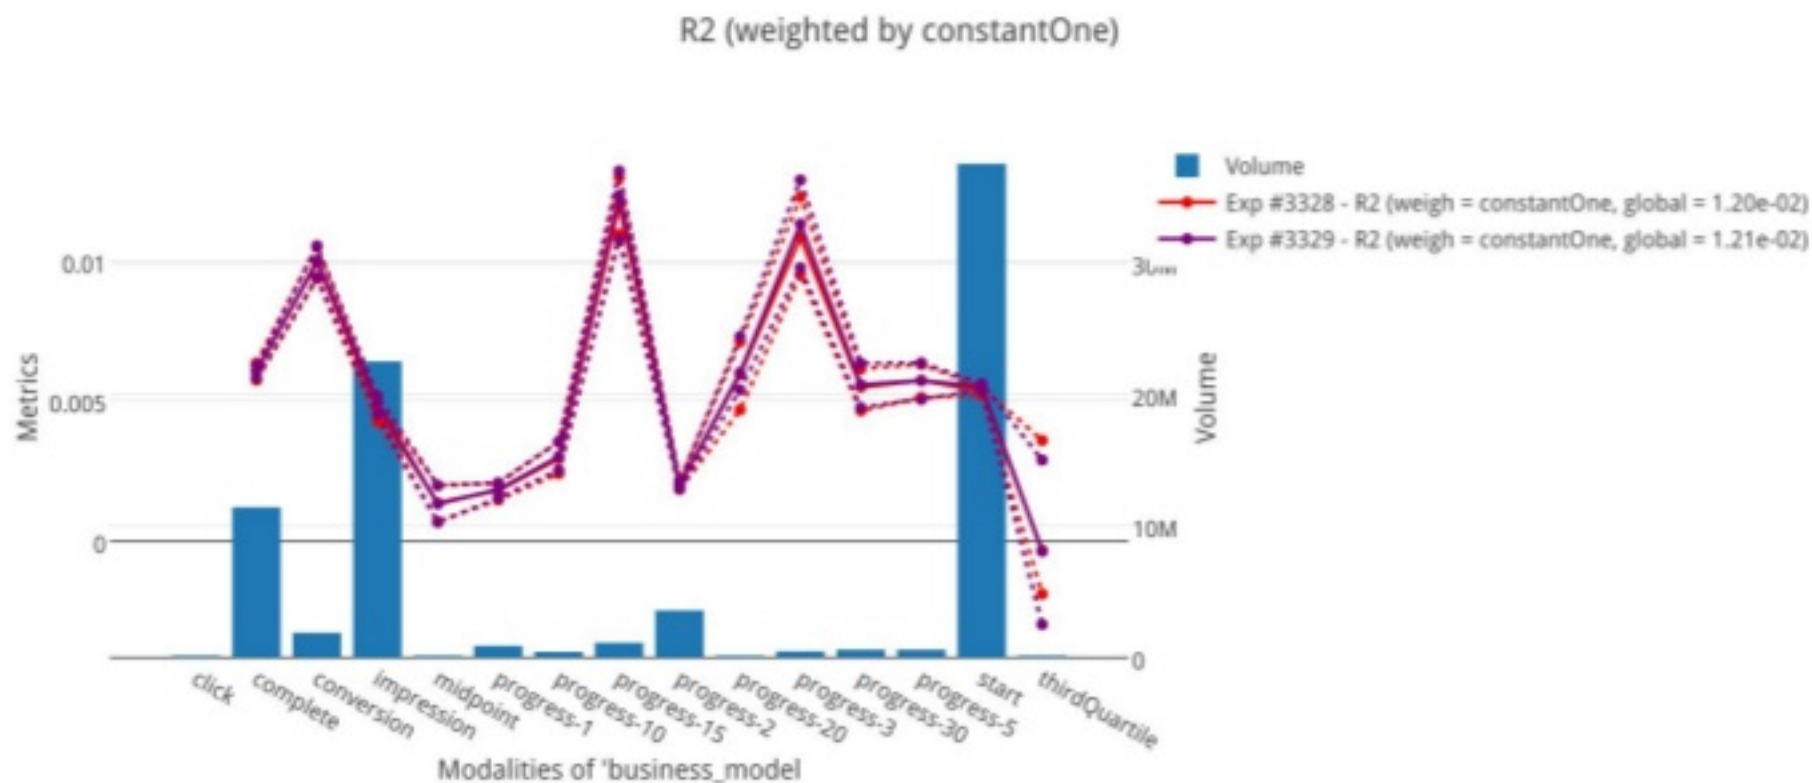
Result projections

- Aggregate the result metrics by categories (day, business model, country, etc.)

Jupyter-scala kernel: <https://github.com/jupyter-scala/jupyter-scala/>

Result visualization

```
viz.plot(  
  subSetIds = allIds,  
  metrics = Seq(RegressionMetric.R2),  
  weights = Seq('constantOne'),  
  projection = 'business_model'  
)
```



Offline experiment

3000

experiments launched
by different teams
during last 18 months

480000

hours of logs analysed,
all experiments
combined

Thanks!

Follow us!

Our engineering blog on medium:

<https://medium.com/teads-engineering>