

# Pitfalls of Apache Spark at Scale

**Cesar Delgado** @hpcfamer

**DB Tsai** @dbtsai

Spark + AI Summit Europe, London, Oct 4, 2018

© 2018 Apple Inc. All rights reserved. Redistribution or public display not permitted without written permission from Apple.

## Apple Siri Open Source Team

- We're Spark, Hadoop, HBase PMCs / Committers / Contributors
- We're the advocate for open source
- Pushing our internal changes back to the upstreams
- Working with the communities to review pull requests, develop new features and bug fixes

# Apple Siri

The world's largest virtual assistant service powering every iPhone, iPad, Mac, Apple TV, Apple Watch, and HomePod

## Apple Siri Data

- Machine learning is used to personalize your experience throughout your day
- We believe privacy is a fundamental human right

## Apple Siri Scale

- Large amounts of requests, Data Centers all over the world
- Hadoop / Yarn Cluster has thousands of nodes
- HDFS has hundred of PB
- 100's TB of raw event data per day

## Siri Data Pipeline

- Downstream data consumers were doing the same expensive query with expensive joins
- Different teams had their own repos, and built their jars - hard to track data lineages
- Raw client request data is tricky to process as it involves deep understanding of business logic

## Unified Pipeline

- Single repo for Spark application across the Siri
- Shared business logic code to avoid any discrepancy
- Raw data is cleaned, joined, and transformed into one standardized data model for data consumers to query on

## Technical Details about Strongly Typed Data

- Schema of data is checked in as case class, and CI ensures schema changes won't break the jobs
- Deeply nested relational model data with 5 top level columns
- The total fields are around 2k
- Stored in Parquet format partitioned by UTC day
- Data consumers query on the subset of data



## Review of APIs of Spark

**DataFrame:** Relational untyped APIs introduced in Spark 1.3. From Spark 2.0,

**type** `DataFrame` = `Dataset[Row]`

**Dataset:** Support all the untyped APIs in DataFrame  
+ typed functional APIs

## Review of DataFrame

```
val df: DataFrame = Seq(  
  (1, "iPhone", 5),  
  (2, "MacBook", 4),  
  (3, "iPhone", 3),  
  (4, "iPad", 2),  
  (5, "appleWatch", 1)  
).toDF("userId", "device", "counts")  
  
df.printSchema()  
"""  
  | root  
  |-- userId: integer (nullable = false)  
  |-- device: string (nullable = true)  
  |-- counts: integer (nullable = false)  
  |  
  """  
df.stripMargin
```

## Review of DataFrame

```
df.withColumn("counts", $"counts" + 1)
  .filter($"device" === "iPhone").show()
"""
  | +-----+-----+-----+
  | |userId|device|counts|
  | +-----+-----+-----+
  | |      1|iPhone|      6|
  | |      3|iPhone|      4|
  | +-----+-----+-----+
"""
.stripMargin
// $"counts" == df("counts")
// via implicit conversion
```

## Execution Plan - Dataframe Untyped APIs

```
df.withColumn("counts", $"counts" + 1).filter($"device" === "iPhone").explain(true)
=====
== Parsed Logical Plan ==
'Filter ('device = iPhone)
+- Project [userId#3, device#4, (counts#5 + 1) AS counts#10]
   +- Relation[userId#3,device#4,counts#5] parquet

== Physical Plan ==
*Project [userId#3, device#4, (counts#5 + 1) AS counts#10]
+- *Filter (isnotnull(device#4) && (device#4 = iPhone))
   +- *FileScan parquet [userId#3,device#4,counts#5]
      Batched: true, Format: Parquet,
      PartitionFilters: [],
      PushedFilters: [IsNotNull(device), EqualTo(device,iPhone)],
      ReadSchema: struct<userId:int,device:string,counts:int>

=====
```

## Review of Dataset

```
case class ErrorEvent(userId: Long, device: String, counts: Long)

val ds = df.as[ErrorEvent]

ds.map(row => ErrorEvent(row.userId, row.device, row.counts + 1))
  .filter(row => row.device == "iPhone").show()
"""
  | +-----+ +-----+ +-----+
  | |userId|device|counts|
  | +-----+ +-----+ +-----+
  | |      1|iPhone|      6|
  | |      3|iPhone|      4|
  | +-----+ +-----+ +-----+
  """
.stripMargin
// It's really easy to put existing Java / Scala code here.
```

## Execution Plan - Dataset Typed APIs

```
ds.map { row => ErrorEvent(row.userId, row.device, row.counts + 1) }.filter { row =>
  row.device == "iPhone"
}.explain(true)
"""
|== Physical Plan ==
|*SerializeFromObject [
|  assertNotNull(input[0, com.apple.ErrorEvent, true]).userId AS userId#27L,
|  assertNotNull(input[0, com.apple.ErrorEvent, true]).device, true) AS device#28,
|  assertNotNull(input[0, com.apple.ErrorEvent, true]).counts AS counts#29L]
|+- *Filter <function1>.apply
|   +- *MapElements <function1>, obj#26: com.apple.ErrorEvent
|      +- *DeserializeToObject newInstance(class com.apple.ErrorEvent), obj#25:
|         com.apple.siri.ErrorEvent
|         +- *FileScan parquet [userId#3,device#4,counts#5]
|            Batched: true, Format: Parquet,
|            PartitionFilters: [], PushedFilters: [],
|            ReadSchema: struct<userId:int,device:string,counts:int>
|
|""".stripMargin
```



## Strongly Typed Pipeline

- Typed Dataset is used to guarantee the schema consistency
- Enables Java/Scala interoperability between systems
- Increases Data Scientist productivity

## Drawbacks of Strongly Typed Pipeline

- Dataset are slower than Dataframe <https://tinyurl.com/dataset-vs-dataframe>
- In Dataset, many POJO are created for each row resulting high GC pressure
- Data consumers typically query on subsets of data, but schema pruning and predicate pushdown are not working well in nested fields



## In Spark 2.3.1

```
case class FullName(first: String, middle: String, last: String)
```

```
case class Contact(id: Int,  
                   name: FullName,  
                   address: String)
```

```
sql("select name.first from contacts").where("name.first = 'Jane']").explain(true)
```

```
====  
|== Physical Plan ==  
|*(1) Project [name#10.first AS first#23]  
|+- *(1) Filter (isnotnull(name#10) && (name#10.first = Jane))  
|    +- *(1) FileScan parquet [name#10] Batched: false, Format: Parquet,  
PartitionFilters: [], PushedFilters: [IsNotNull(name)], ReadSchema:  
struct<name:struct<first:string,middle:string,last:string>>  
|  
====.stripMargin
```

## In Spark 2.4 with Schema Pruning

```
====  
|== Physical Plan ==  
|*(1) Project [name#10.first AS first#23]  
|+- *(1) Filter (isnotnull(name#10) && (name#10.first = Jane))  
|   +- *(1) FileScan parquet [name#10] Batched: false, Format: Parquet,  
PartitionFilters: [], PushedFilters: [IsNotNull(name)], ReadSchema:  
struct<name:struct<first:string>>  
  
|  
====".stripMargin
```

- [SPARK-4502], [SPARK-25363] Parquet nested column pruning

## In Spark 2.4 with Schema Pruning + Predicate Pushdown

```
====  
|== Physical Plan ==  
|*(1) Project [name#10.first AS first#23]  
|+- *(1) Filter (isnotnull(name#10) && (name#10.first = Jane))  
|   +- *(1) FileScan parquet [name#10] Batched: false, Format: Parquet,  
PartitionFilters: [], PushedFilters: [IsNotNull(name), EqualTo(name.first, Jane)],  
ReadSchema: struct<name:struct<first:string>>  
  
|  
====".stripMargin
```

- [SPARK-4502], [SPARK-25363] Parquet nested column pruning
- [SPARK-17636] Parquet nested Predicate Pushdown

# **Production Query - Finding a Needle in a Haystack**

# Spark 2.3.1

Spark 2.3.1 <span>Jobs</span> <b>Stages</b> <span>Storage</span> <span>Environment</span> <span>Executors</span> <span>SQL</span> <span>Spark shell application UI</span>									
<b>Stages for All Jobs</b>									
Completed Stages: 3									
Completed Stages (3)									
Stage Id ▾	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	<a href="#">parquet at &lt;console&gt;:26</a>	<a href="#">+details</a>	2018/09/20 20:38:21	1.2 h	<a href="#">72002/72002</a>	7.1 TB	31.8 MB		
1	<a href="#">parquet at &lt;console&gt;:23</a>	<a href="#">+details</a>	2018/09/20 20:37:41	1 s	<a href="#">1/1</a>				
0	Listing leaf files and directories for 6000 paths: <a href="#">hdfs://nameservice1/user/sirimeetrics_bot2/uberstream/20180802/part-00000-7264d04a-200e...</a> <a href="#">parquet at &lt;console&gt;:23</a>	<a href="#">+details</a>	2018/09/20 20:37:38	2 s	<a href="#">6000/6000</a>				





- 21x faster in wall clock time
- 8x less data being read
- Saving electric bills in many data centers



## Future Work

- Use Spark Streaming can be used to aggregate the request data in the edge first to reduce the data movement
- Enhance the Dataset performance by analyzing JVM bytecode and turn closures into Catalyst expressions
- Building a table format on top of Parquet using Spark's Datasource V2 API to tracks individual data files with richer metadata to manage versioning and enhance query performance



## Conclusions

With some work, engineering rigor and some optimizations  
Spark can run at very large scale in lightning speed

Thank you