

# RNNs for Recommendation & Personalization

Nick Pentreath  
Principal Engineer

*@MLnick*

**IBM**  
**CODE**

# About

@MLnick on Twitter & Github

Principal Engineer, IBM

CODAIT - Center for Open-Source Data & AI  
Technologies

Machine Learning & AI

Apache Spark committer & PMC

Author of *Machine Learning with Spark*

Various conferences & meetups



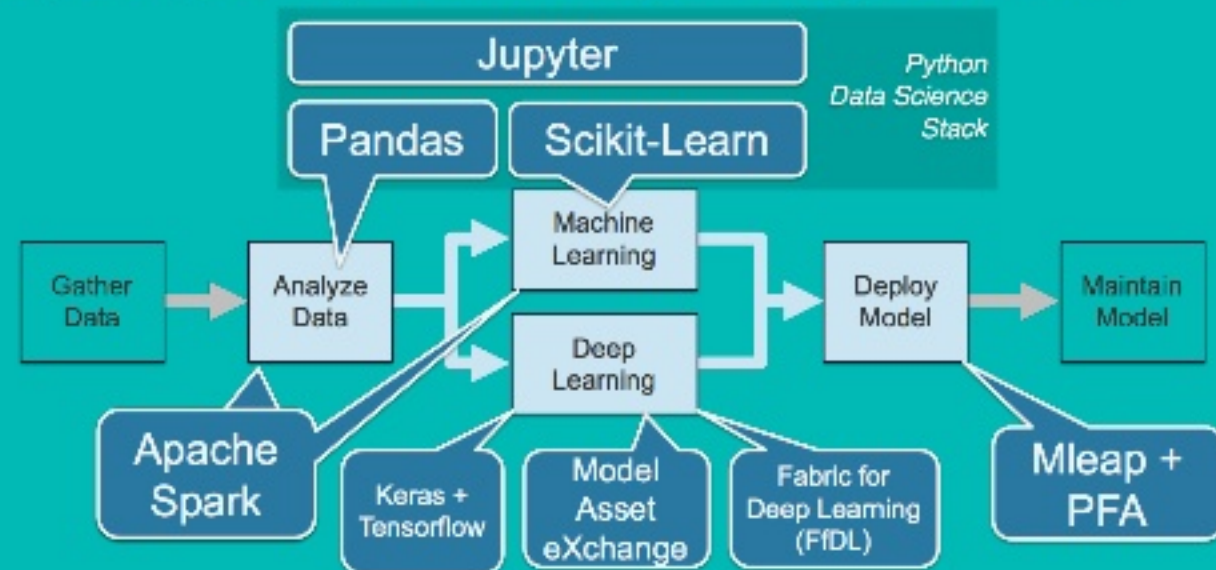


[codait.org](http://codait.org)

CODAIT aims to make AI solutions dramatically easier to create, deploy, and manage in the enterprise

Relaunch of the Spark Technology Center (STC) to reflect expanded mission

### Improving Enterprise AI Lifecycle in Open Source



# Agenda


Recommender systems overview

Deep learning and RNNs

RNNs for recommendations

Challenges and future directions

# Recommender Systems





# Users and Items



```
{  
  "user_id": "1",  
  "name": "Joe Bloggs",  
  "created_date": 1476884080,  
  "updated_date": 1476946916,  
  "last_active_date": 1476946962,  
  "age": 32,  
  "country": "GB",  
  "city": "London",  
  ...  
}
```



```
{  
  "item_id": "10",  
  "name": "LOL Cats",  
  "description": "catscatscats",  
  "category": ["Cat Videos", "Humour", "Animals"],  
  "tags": ["cat", "lol", "funny", "cats", "felines"],  
  "created_date": 1476884080,  
  "updated_date": 1476884080,  
  "last_played_date": 1476946962,  
  "likes": 100000,  
  "author_id": "321",  
  "author_name": "ilikecats",  
  "channel_id": "CatVideoCentral",  
  ...  
}
```

# Events

## Implicit preference data

- Online – page view, click, app interaction
- Commerce – cart, purchase, return
- Media – preview, watch, listen

## Explicit preference data


- Ratings, reviews

## Intent

- Search query

## Social

- Like, share, follow, unfollow, block



# Context




```
{  
  "user_id": "1",  
  "item_id": "10",  
  "event_type": "page_view",  
  "timestamp": 1476884080,  
  "referrer": "http://codait.org",  
  "ip": "123.12.12.12",  
  "device_type": "Smartphone",  
  "user_agent_os": "Android",  
  "user_agent_type": "Mobile Browser",  
  "user_agent_family": "Chrome Mobile",  
  "geo": "51.5085, 0.0298"  
  ...  
}
```



# Prediction

## Prediction is ranking

- Given a user and context, rank the available items in order of likelihood that the user will interact with them



## The *de facto* standard model

- Represent user ratings as a user-item matrix
- Find two smaller matrices (called the *factor* matrices) that approximate the full matrix
- Minimize the reconstruction error (i.e. rating prediction / completion)

- Efficient, scalable algorithms
  - Gradient Descent
  - Alternating Least Squares (ALS)
- Prediction is simple
- Can handle implicit data




# Cold Start

## New items

- No historical interaction data
- Typically use baselines (e.g. popularity) or item content

## New (or unknown) users

- Previously unseen or anonymous users have no user profile or historical interactions
- Have context data (but possibly very limited)
- Cannot directly use collaborative filtering models
  - Item-similarity for current item
  - Represent session as aggregation of items
  - Contextual models can incorporate short-term history





# Deep Learning and Recurrent Neural Networks




# Overview

Original theory from 1940s; computer models originated around 1960s; fell out of favor in 1980s/90s

Recent resurgence due to

- Bigger (and better) data; standard datasets (e.g. ImageNet)
- Better hardware (GPUs)
- Improvements to algorithms, architectures and optimization

Leading to new state-of-the-art results in computer vision (images and video); speech/text; language translation and more





# Modern Neural Networks

## Deep (multi-layer) networks

### Computer vision

- Convolution neural networks (CNNs)
- Image classification, object detection, segmentation

### Sequences and time-series



- Recurrent neural networks (RNNs)
- Machine translation, text generation
- LSTMs, GRUs

### Embeddings

- Text, categorical features

### Deep learning frameworks

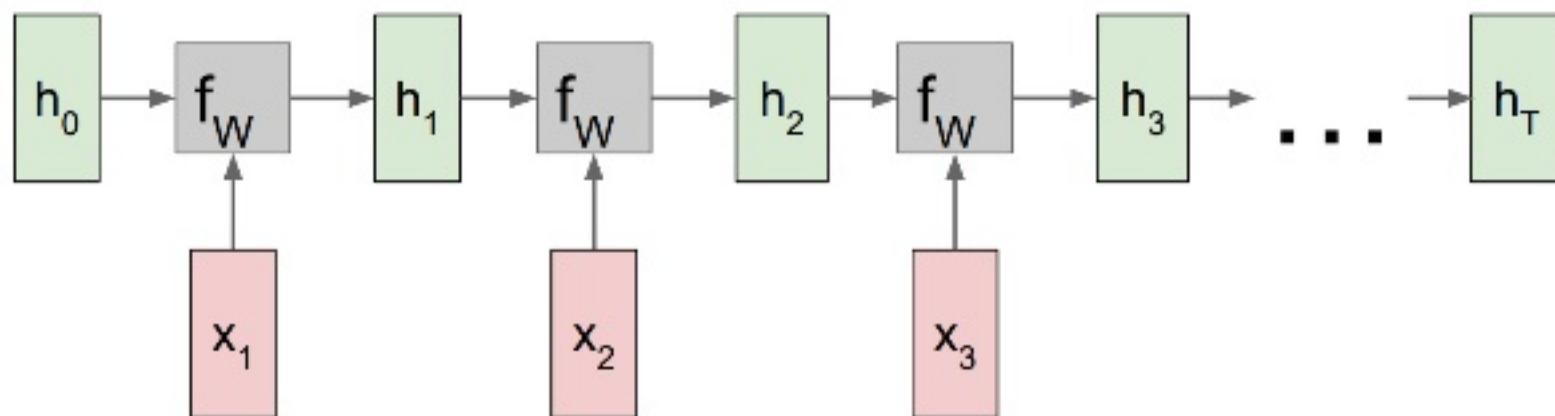

- Flexibility, computation graphs, auto-differentiation, GPUs



# Recurrent Neural Networks

## Neural Network on Sequences ...


- ... sequence of neural network (layers)
- Hidden layers (state) dependent on previous state as well as current input
- “memory” of what came before
- Share weights across all time steps
- Training using backpropagation through time (BPTT)




$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

# Recurrent Neural Networks


one to one




one to many




many to one



many to many



many to many




# Recurrent Neural Networks

## Issues

- Exploding gradients - clip / scale gradients
- Vanishing gradients

## Solutions

- Truncated BPTT
- Restrict sequence length
- Cannot encode very long term memory



# Recurrent Neural Networks

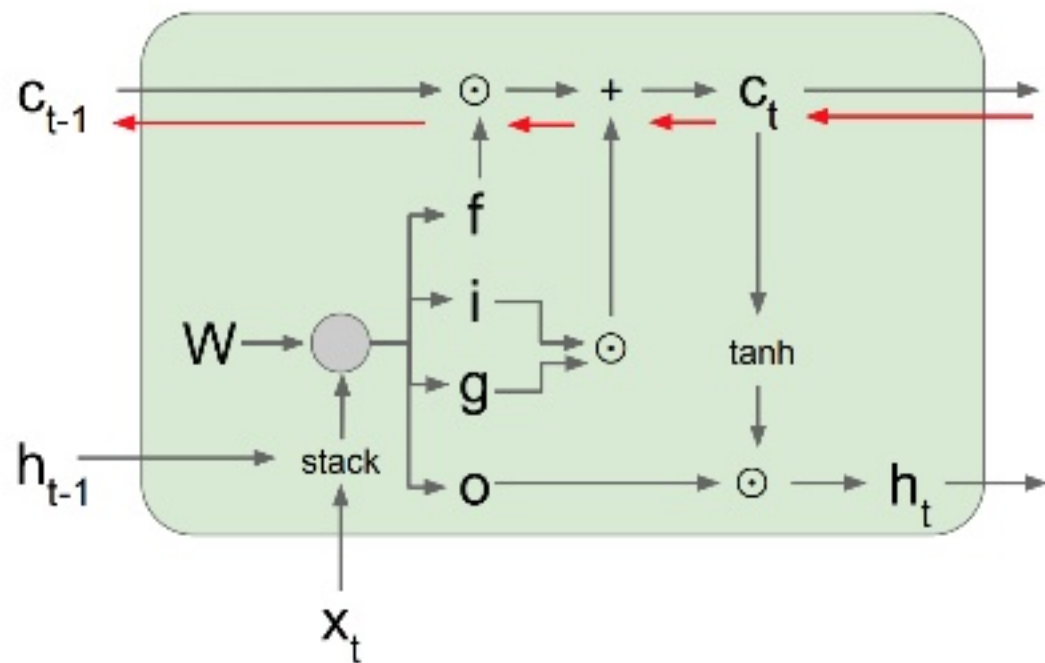
## Long Short Term Memory (LSTM)

- Replace simple RNN layer (activation) with a LSTM cell
- Cell has 3 gates - Input (i), Forget (f), Output (o)
- Activation (g)
- Backpropagation depends only on elementwise operations (no matrix operations over  $W$ )

## Gated Recurrent Unit (GRU)

- Effectively a simplified version of LSTM
- 2 gates instead of 3 - input and forget gate is combined into an update gate. No output gate

GRU has fewer parameters, LSTM may be more expressive



$$c_t = f \odot c_{t-1} + i \odot g$$



$$h_t = o \odot \tanh(c_t)$$





# Recurrent Neural Networks

## Variants

- Multi-layer (deep) RNNs
- Bi-directional
- Deep bi-directional
- Attention



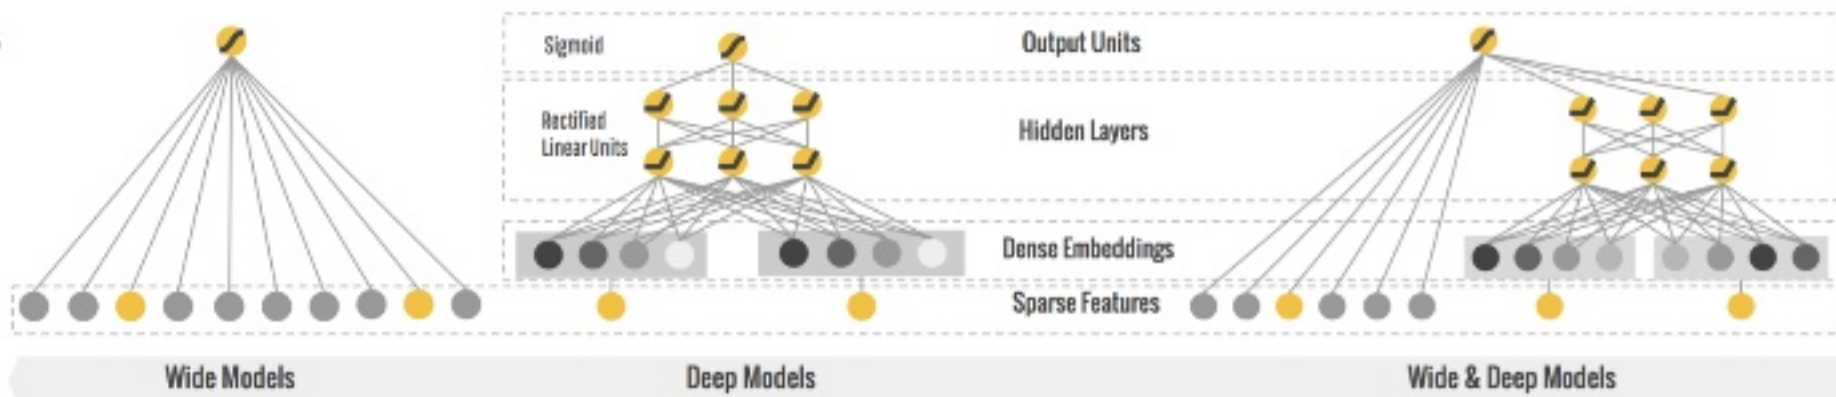

# RNNs for Recommendations



# Deep Learning for Recommenders Overview

Most approaches have focused on combining

- Performance of collaborative filtering models (especially matrix factorization)
  - Embeddings with appropriate loss = MF
- Power of deep learning for feature extraction
  - CNNs for image content, audio, etc.
  - Embeddings for categorical features
  - Linear models for interactions
  - RNNs for text

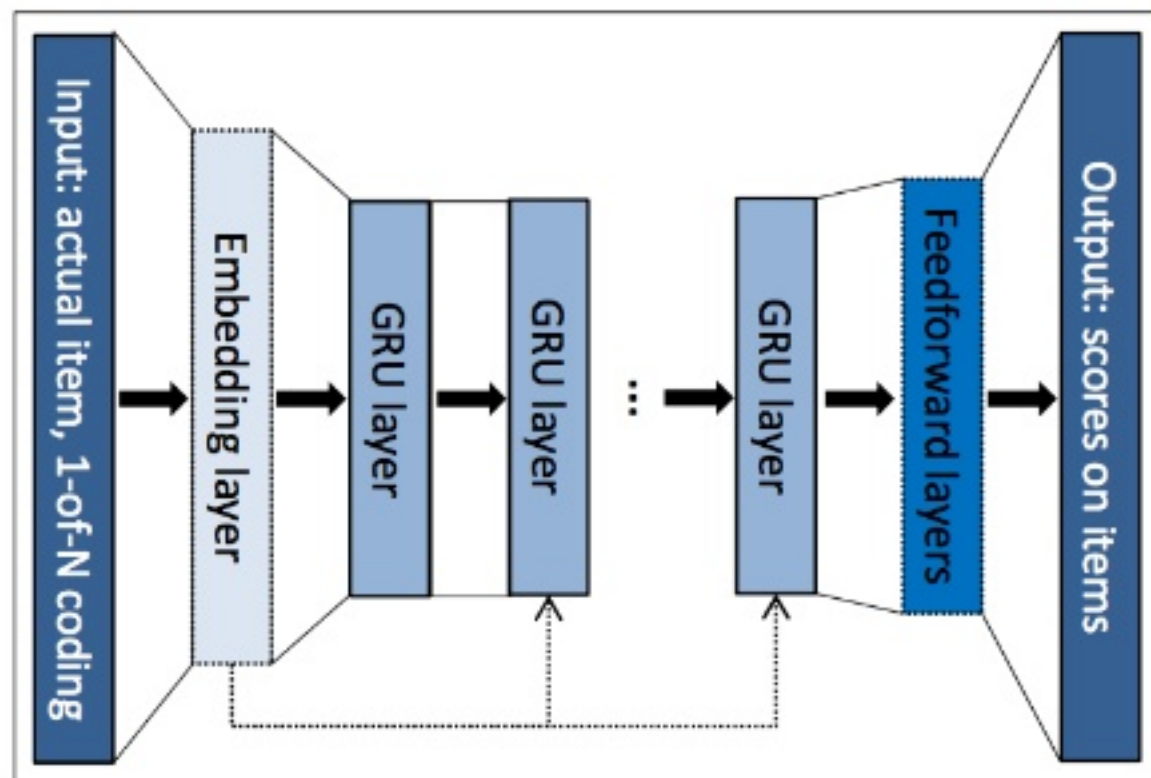




# Session-based recommendation

Apply the advances in sequence modeling from deep learning


- RNN architectures trained on the sequence of user events in a session (e.g. products viewed, purchased) to predict next item in session
- Adjustments for domain
  - Item encoding (1-of-N, weighted average)
  - Parallel mini-batch processing
  - Ranking losses – BPR , TOP1
  - Negative item sampling per mini-batch
- Report 20-30% accuracy gain over baselines



# Contextual Session-based models

Add contextual data to the RNN architecture

- Context included time, time since last event, event type
- Combine context data with input / output layer
- Also combine context with the RNN layers
- About 3-6% improvement (in Recall@10 metric) over simple RNN baseline
- Importantly, model is even better at predicting sales (vs view, add to cart events) and at predicting new / fresh items (vs items the user has already seen)

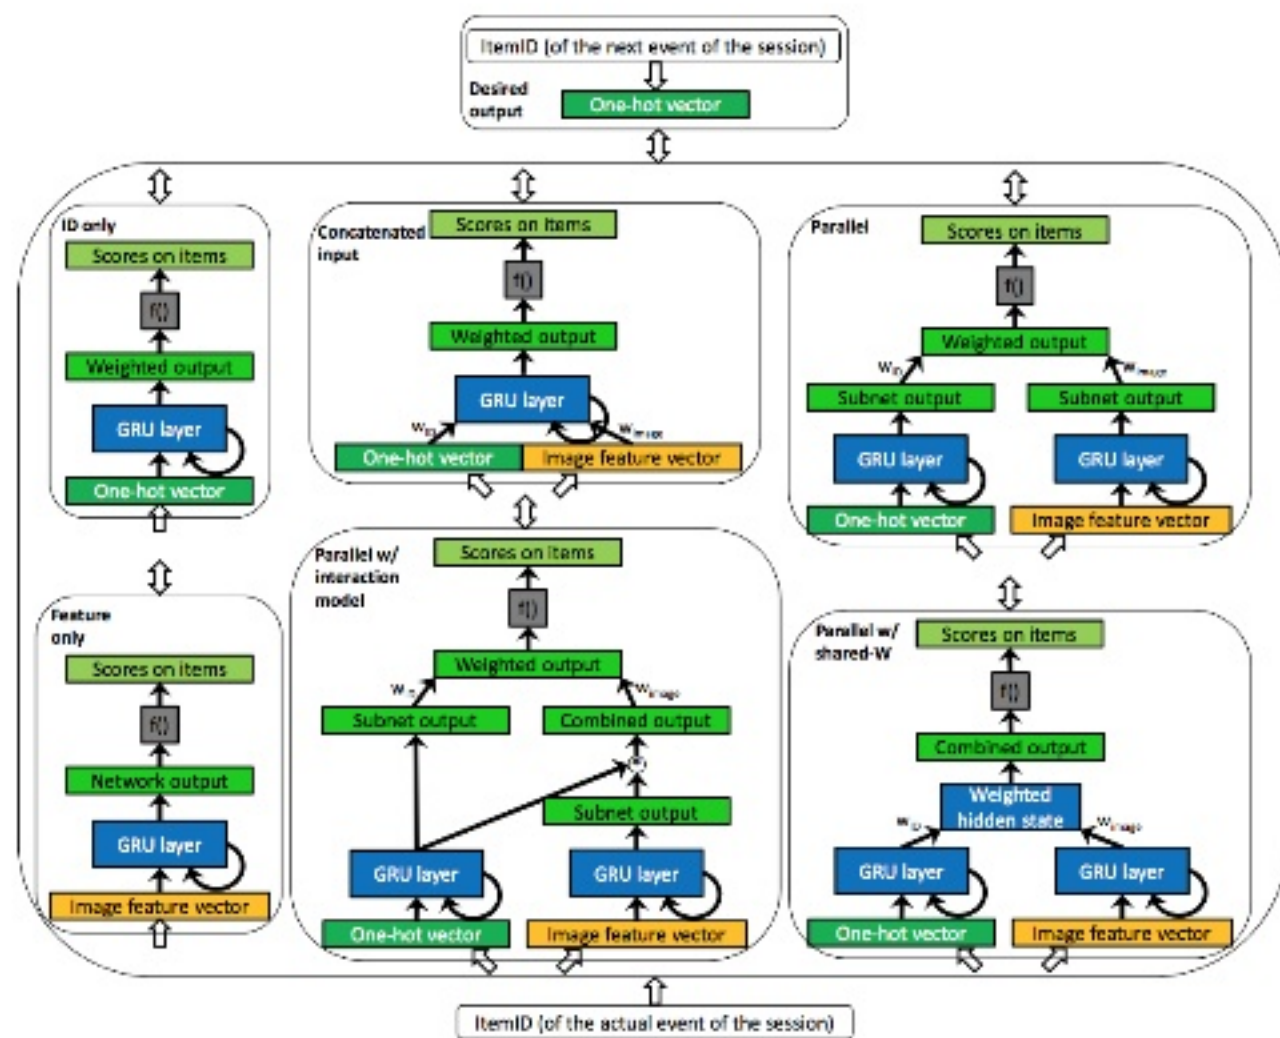




# Content and Session-based models

## Add content data to the RNN architecture

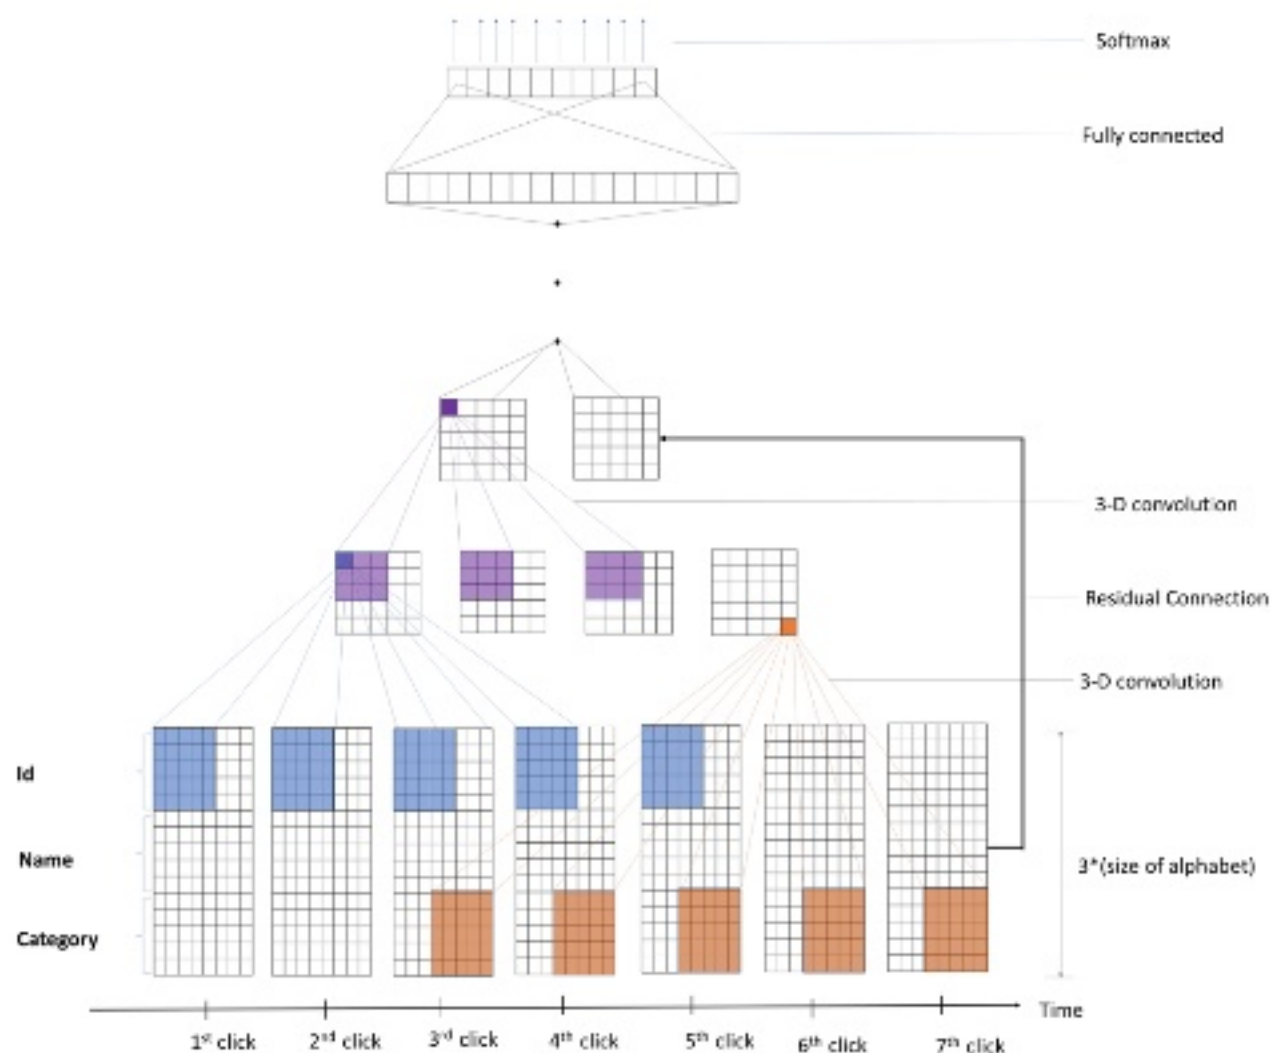
- Parallel RNN (p-RNN)
- Follows trend in combining DL architectures for content feature extraction with CF models for interaction data
  - CNN for image data
  - BOW for text (alternatives are Word2Vec-style models and RNN language models)
- Some training tricks
  - Alternating – keep one subnet fixed, train other
  - Residual – subnets trained on residual error
  - Interleaved – alternating training per mini-batch



# 3D CNNs for Session-based Recommendation

As we've seen in text / NLP, CNNs can also be effective in modeling sequences

- 3D convolutional models have been applied in video classification
- Potentially faster to train, easier to understand
- Use character-level encoding of IDs and item features (name, description, categories)
  - Compact representation
  - No embedding layer
- “ResNet” style architecture
- Show improvement over p-RNN



# Challenges

## Challenges particular to recommendation models

- Data size and dimensionality (input & output)
  - Sampling
- Extreme sparsity
  - Embeddings & compressed representations
- Wide variety of specialized settings
- Combining session, content, context and preference data
- Model serving is difficult – ranking, large number of items, computationally expensive
- Metrics – model accuracy and its relation to real-world outcomes and behaviors
- Need for standard, open, large-scale, datasets that have time and session data and are content- and context-rich
  - RecSys 15 Challenge – YouChoose dataset
- Evaluation – watch you baselines!
  - [When Recurrent Neural Networks meet the Neighborhood for Session-Based Recommendation](#)

# Future Directions

Most recent and future directions in research & industry

- Improved RNNs

- Cross-session models (e.g. Hierarchical RNN)
- Further research on contextual models, as well as content and metadata
- Attention models:
  - [Attentive Neural Architecture for Music Recommendation](#)
  - [Neural Attentive Session-based Recommendation](#)

- Combine sequence and historical models (long- and short-term user profiles)

- Domain-specific applications

- [Contextualized Location Sequence Recommender](#)

- [RecGAN](#) (yes, GANs **and** RNNS!)

- Applications at scale

- Dimensionality reduction, compressed encodings



# Summary

DL for recommendation is just getting started (again)

- Huge increase in interest, research papers. Already many new models and approaches
- DL approaches have generally yielded incremental % gains
  - But that can translate to significant \$\$\$
  - More pronounced in session-based
- Cold start scenarios benefit from multi-modal nature of DL models and explicit modeling of sequences
- Flexibility of DL frameworks helps a lot
- Benefits from advances in DL for images, video, NLP etc.
- Open-source libraries appearing (e.g. Spotlight)
- Check out DLRS workshops & tutorials @ RecSys 2016 / 2017, and upcoming in Oct, 2018
- RecSys challenges



Thank you!



[codait.org](https://codait.org)



[twitter.com/MLnick](https://twitter.com/MLnick)



[github.com/MLnick](https://github.com/MLnick)



[developer.ibm.com](https://developer.ibm.com)



FfDL

MAX



Sign up for IBM Cloud and try Watson Studio!

<https://ibm.biz/BdYbTY>

<https://datascience.ibm.com/>

# Links & References

[Wikipedia: Perceptron](#)

[Stanford CS231n Convolutional Neural Networks for Visual Recognition](#)

[Stanford CS231n – RNN Slides](#)

[Recurrent Neural Networks Tutorial](#)

[The Unreasonable Effectiveness of Recurrent Neural Networks](#)

[Understanding LSTM Networks](#)

[Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation](#)

[Long short-term memory](#)

[Attention and Augmented Recurrent Neural Networks](#)

IBM

**CODE**

DBG / Oct 3, 2018 / © 2018 IBM Corporation

# Links & References

[Deep Content-based Music Recommendation](#)

[Google's Wide and Deep Learning Model](#)

[Deep Learning for Recommender Systems Workshops @ RecSys](#)

[Deep Learning for Recommender Systems Tutorial @ RecSys 2017](#)

[Session-based Recommendations with Recurrent Neural Networks](#)

[Recurrent Neural Networks with Top-k Gains for Session-based Recommendations](#)

[Sequential User-based Recurrent Neural Network Recommendations](#)

IBM

**CODE**

DBG / Oct 3, 2018 / © 2018 IBM Corporation

# Links & References

[Personalizing Session-based Recommendations with Hierarchical Recurrent Neural Networks](#)

[Parallel Recurrent Neural Network Architectures for Feature-rich Session-based Recommendations](#)

[Contextual Sequence Modeling for Recommendation with Recurrent Neural Networks](#)

[When Recurrent Neural Networks meet the Neighborhood for Session-Based Recommendation](#)

[3D Convolutional Networks for Session-based Recommendation with Content Features](#)

[Spotlight: Recommendation models in PyTorch](#)

[RecSys 2015 Challenge – YouChoose Dataset](#)

IBM

**CODE**

DBG / Oct 3, 2018 / © 2018 IBM Corporation



