

# Spark SQL Catalyst Code Optimization Using Function Outlining



Kavana N Bhat, Senior Systems Developer &  
Madhusudanan Kandasamy, STSM, PowerAI  
Power Systems, IBM

#SAISDD15

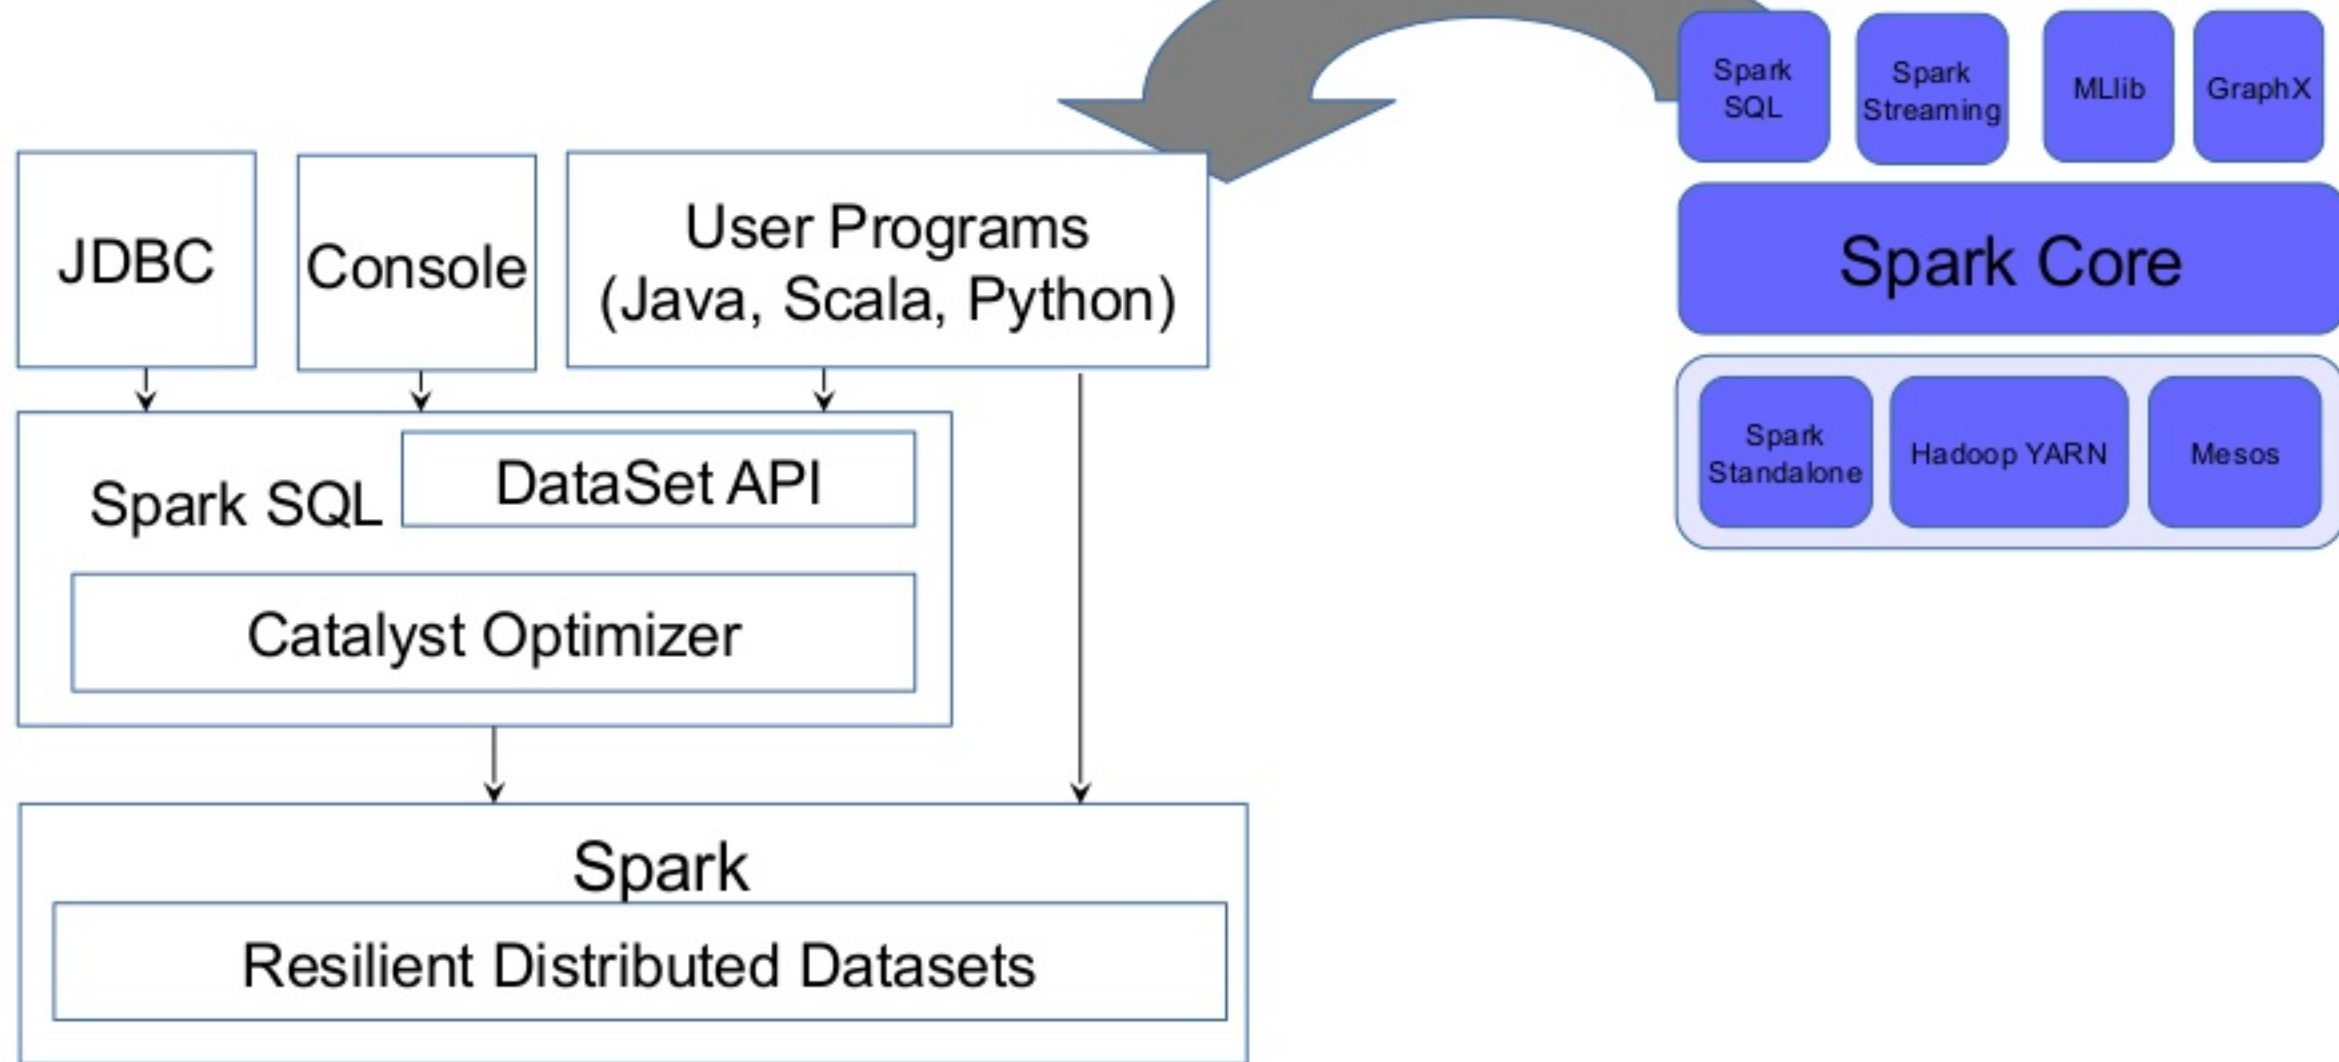
# About Presenters:

- **Kavana N Bhat**
  - Senior Systems Developer at IBM
  - Working for IBM Power Systems over 14 years
    - ➔ AIX OS Development
    - ➔ Apache Spark & Alluxio Optimization for Power Systems
    - ➔ ML/DL Frameworks for Enterprise Systems
  - Has 5 Patents, 5 Disclosure Publications
  - Email: [kavana.bhat@in.ibm.com](mailto:kavana.bhat@in.ibm.com)
- **Madhusudanan Kandasamy**
  - PowerAI STSM at IBM Systems
  - Working for IBM Power Systems over 15 years
    - ➔ AIX & Linux OS Development
    - ➔ Apache Spark Optimization for Power Systems
    - ➔ Distributed ML/DL Framework with GPU & NVLink
  - IBM Master Inventor (20+ Patents, 18 Disclosure Publications)
  - Github: <https://github.com/kmadhugit> Email: [madhusudanan@in.ibm.com](mailto:madhusudanan@in.ibm.com)

# Outline

- Spark SQL Overview
- Catalyst Optimizer in Depth
- Java JIT Compiler Overview
- Spark Performance Profile of TPC-DS Benchmark Query
- Function Outlining
- Catalyst Code Generation Optimizations
- Performance Statistics
- Conclusion
- References

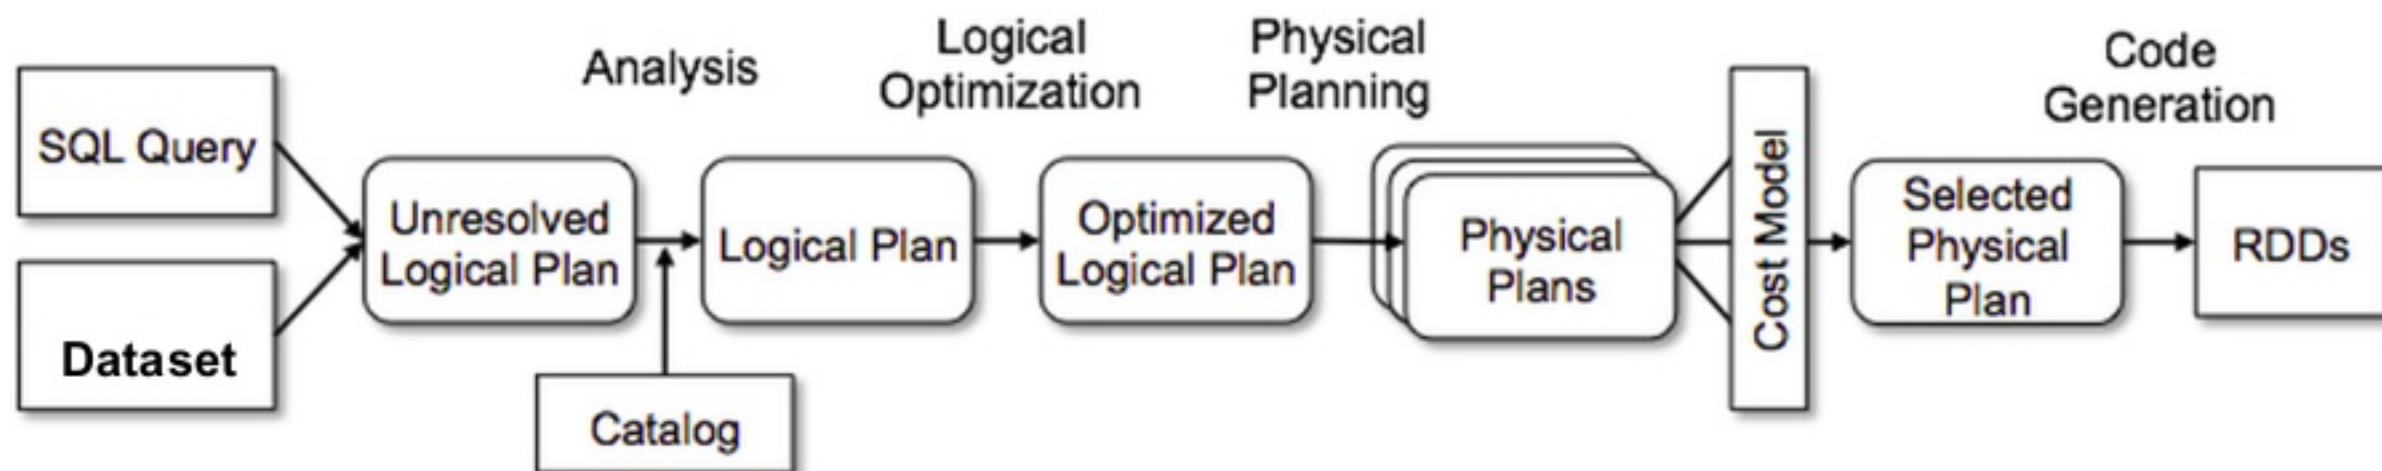
# Spark SQL Interaction with Spark



# Spark SQL Architecture

## Terminology:

- Logical and Physical plans are trees representing query evaluation
- Logical plan is higher-level and algebraic
- Physical plan is lower-level and operational






# Catalyst Optimizer Phases

- Analysis:
  - resolve unresolved attribute references or relations
- Logical Optimization:
  - standard sql query optimizations like constant folding, predicate pushdown, project pruning are applied.
- Physical Planning:
  - one or more physical plans are formed from the optimized logical plan, using physical operator matching the Spark execution engine
- Code Generation:
  - fuses multiple physical operators together into a single optimized function
  - to avoid large amount of branches and virtual function calls, generates Java code to run on each machine
  - generated code is compiled using Janino compiler

# Logical Plan Optimization Example

- An example query


```
SELECT count(*)  
FROM (  
    SELECT id, name  
    FROM People ) p  
WHERE p.id = 1
```



# Logical Plan Optimization Example (Contd..)

## ➤ Optimization Rules example

1. Find Filters on top of projections.
2. Check that the filter can be evaluated without the result of the project.
3. If so, switch the operators.

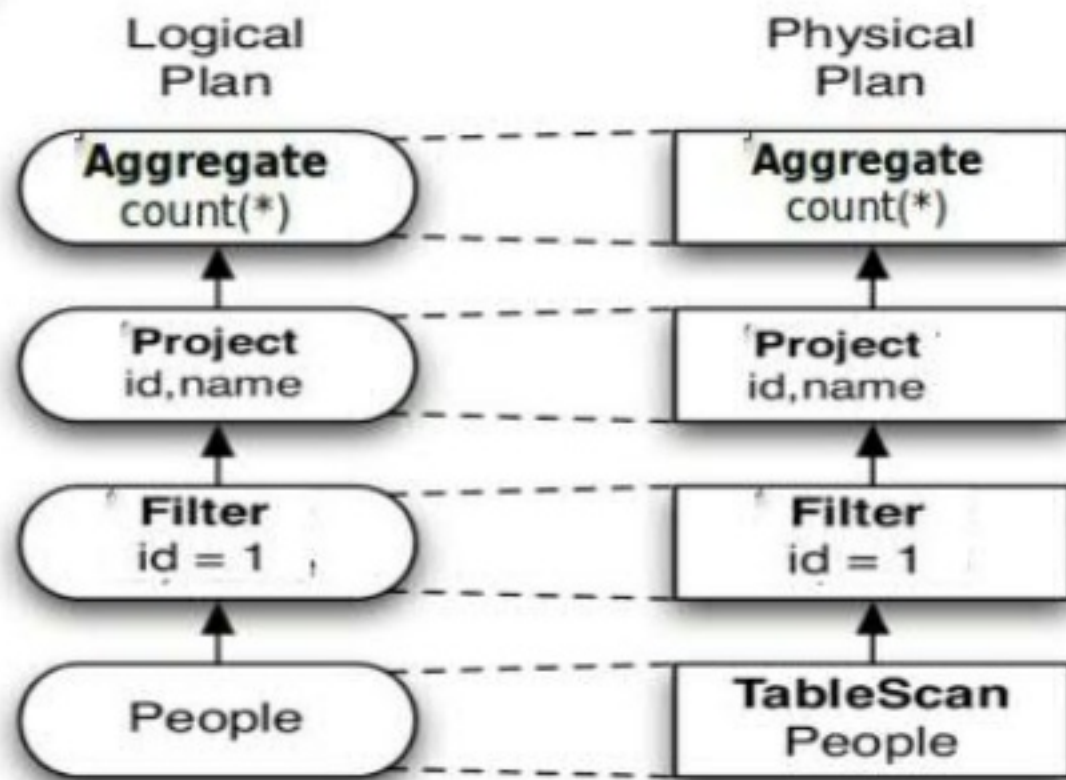




# Physical Planning

- Native query planning


```
SELECT count(*)  
FROM (  
    SELECT id, name  
    FROM People ) p  
WHERE p.id = 1
```



# Catalyst Code Generation

- Without Code Generation, each of the operator in the tree has to be interpreted for each row of data by walking down the tree
  - large amounts of branches and virtual function calls
  - unnecessary memory allocations/de-allocations from creating new rows
- With Code Generation, this can be rewritten as one iterator which takes in a row
  - creates one row object per output row, and doesn't use any extraneous method calls
  - allows the compiler/processor optimizations to kick in


```
int count = 0;
for (Row r : rows) {
    if (<filter>) {
        count++;
    }
}
return new Row(count)
```



# Traditional Java Compilation and Execution

- A Java Compiler compiles high level Java source code to Java bytecode readable by Java Virtual Machine (JVM)
- JVM interprets bytecode to machine instructions at runtime
- Advantages
  - platform independence (JVM present on most machines)
- Drawbacks
  - needs memory
  - not as fast as running pre-compiled machine instructions

# Java JIT Compiler Overview:



- A just-in-time (JIT) compiler is a compiler that compiles code into machine instructions during program execution, rather than ahead of time.
- Combines speed of compiled code w/ flexibility of interpretation
- JVM interprets the application code initially
- Collects execution statistics for each method
- Invokes JIT compiler for identified hot methods
- JIT code generally offers far better performance than interpreted code



# Spark Performance Profile of TPC-DS Benchmark Queries

- TPC-DS is the de-facto industry standard benchmark for measuring the performance of SQL queries
- Performance profile of TPC-DS query #72 showed one of the catalyst generated routine to be hot.
- Logs from JVM PrintCompilation flag for the workload were analyzed
  - This flag shows basic information on when Hotspot compiles methods
- The identified hot method was reported as 'COMPILE SKIPPED'
  - Indicates the method was identified by JVM as hot, but JIT failed to compile.
  - One reason could be the code cache is full
- The job-stage-task which took the longest time was identified
- The catalyst code modified to log the dynamically generated java code with file names tagged with job/stage/task Id.
- Detailed look at the catalyst generated code for the hot method indicated that it was a huge routine (>300 java LOC)

## Stages for All Jobs

Completed Stages: 29

Completed Stages (29)

Stage Id	Description	Submitted		Duration	Tasks: Succeeded/Total
		00	15:25:09		
28	select i_item_desc ,w_v processCmd at CliDrive	00	15:24:56	13 s	200/200
27	select i_item_desc ,w_v processCmd at CliDrive	00	14:59:24	26 min	200/200
26	select i_item_desc ,w_warehouse processCmd at CliDriver.java:376	00	14:58:00	1.7 min	343/343

```

207669 12478 % 3 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
207681 12478 % 3 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
207697 12481 % 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
207720 12482 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext (1605 bytes)
207758 12483 % 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
207949 12486 % 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
208282 12497 % 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
208886 12513 % 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
209538 12516 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext (1605 bytes)
1485428 12926 4 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext (1605 bytes)
197821 12391 % 3 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
197842 12391 % 3 org.apache.spark.sql.catalyst.expressions.GeneratedClass$GeneratedIterator::processNext @ 59 (1605 bytes)
    
```

COMPILE SKIPPED: CodeBuffer overflow (retry at different tier)

COMPILE SKIPPED: CodeBuffer overflow (retry at different tier)



# Function Outlining

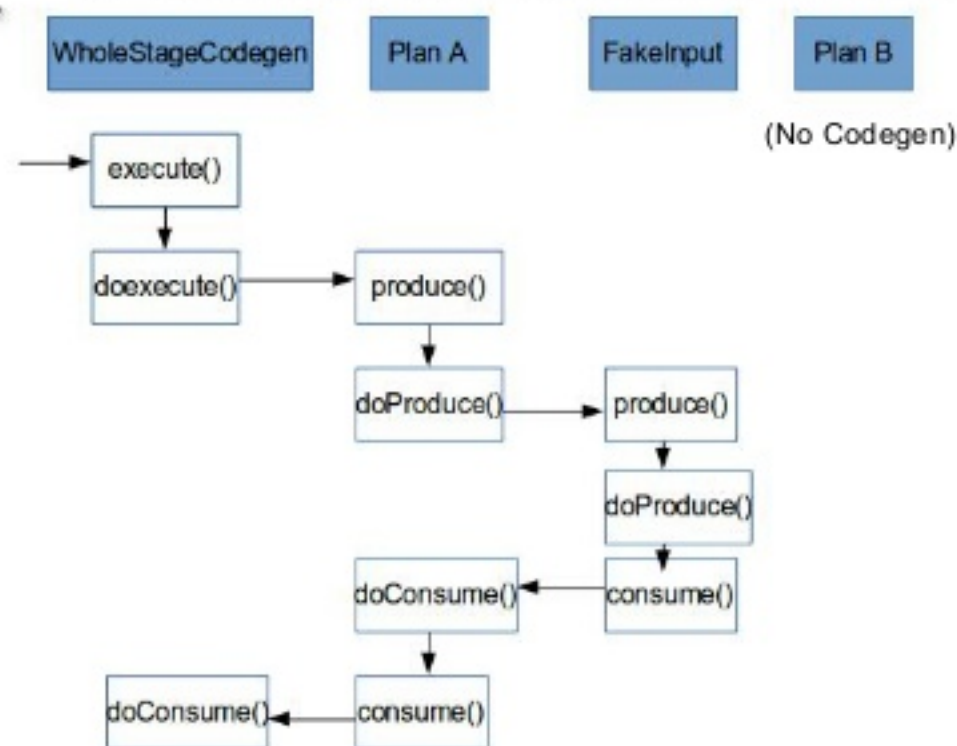
- It is a technique that splits a code region into a new, independent function and replaces it with a function call to new function.
- Benefits :
  - JVM loves small methods– can improve performance as they might enable more inlining
  - Improve code locality.
- Drawbacks :
  - Extra function calls are introduced to transfer control between the outlined region and the other parts of the program unit.

# Catalyst Code Generation Optimizations

- Query #72 used Broadcast Hash Join/Sort Merge Join
- The Spark catalyst code generation for the identified hot routine was modified to generate the same logic using multiple small routines.
- Some of the logic pulled into smaller routines referenced some locals of the initial routine.
- The issue was resolved by making such local variables to be class private variables.
- The performance was analyzed with TieredCompilation enabled IBM JDK

# Catalyst Code Generation Optimizations – High Level Code Modifications

- WholeStageCodegen compiles a subtree of plans that support codegen together into single java function
- Three code generation paths:
  - Non-whole-stage-codegen path
  - Whole-stage-codegen "produce" path - Java code that reads the rows from the input RDDs
  - Whole-stage-codegen "consume" path - Java source code to consume the generated columns or a row from a physical operator.



Call Graph to generate Java Code

*sql/core/src/main/scala/org/apache/spark/sql/execution/joins/SortMergeJoin  
Exec.scala*

```
override def doProduce(ctx: CodegenContext): String = {
  .....
  val forLoop_sub = ctx.freshName("forLoop")
  ctx.addNewFunction(forLoop_sub,
    s"""
      |private void $forLoop_sub(scala.collection.Iterator<UnsafeRow> iterator)
      |throws java.io.IOException {
      |  while ($iterator.hasNext()) {
      |    InternalRow $rightRow = (InternalRow) $iterator.next();
      |    ${condCheck.trim}
      |    $numOutput.add(1);
      |    ${consume(ctx, leftVars ++ rightVars)}
      |  }
      |}""".stripMargin)
  s"""
    |while (findNextInnerJoinRows($leftInput, $rightInput)) {
    |  ${beforeLoop.trim}
    |  scala.collection.Iterator<UnsafeRow> $iterator =
    |  $matches.generateIterator();
    |  $forLoop_sub($iterator);
    |  if (shouldStop()) return;
    |}
    |""".stripMargin
```



# Catalyst Code Generation Optimizations – A peek into catalyst generated code

Original Catalyst Generated Code	Modified Catalyst Generated Code
<pre> /* 001 */ public Object generate(Object[] references) { /* 002 */   return new GeneratedIterator(references); /* 003 */ } /* 004 */ /* 005 */ final class GeneratedIterator extends org.apache.spark.sql.execution.BufferedRowIterator { ..... /* 317 */ protected void processNext() throws java.io.IOException { /* 318 */   if (scan_batch == null) { /* 319 */     scan_nextBatch(); /* 320 */   } /* 321 */   while (scan_batch != null) { /* 322 */     int numRows = scan_batch.numRows(); /* 323 */     while (scan_batchIdx &lt; numRows) { ..... /* 360 */       scala.collection.Iterator bhj_matches = bhj_isNull ? null : (scala.collection.Iterator)bhj_relation.get(bhj_value); /* 361 */       if (bhj_matches == null) continue; /* 362 */       while (bhj_matches.hasNext()) { /* 363 */         UnsafeRow bhj_matched = (UnsafeRow) bhj_matches.next(); ..... /* 633 */       } /* 634 */       if (shouldStop()) return; /* 635 */     } /* 636 */     scan_batch = null; /* 637 */     scan_nextBatch(); /* 638 */   } /* 639 */   scan_scanTime.add(scan_scanTime1 / (1000 * 1000)); </pre>	<pre> /* 001 */ public Object generate(Object[] references) { /* 002 */   return new GeneratedIterator(references); /* 003 */ } /* 004 */ /* 005 */ final class GeneratedIterator extends org.apache.spark.sql.execution.BufferedRowIterator { ..... /* 273 */ private void scan_processNextSub() throws java.io.IOException { /* 274 */   int numRows = scan_batch.numRows(); ..... /* 305 */   bhj_genJK(); /* 306 */   if (bhj_matches == null) continue; /* 307 */   bhj_wloop(); /* 308 */   if (shouldStop()) return; /* 309 */ } /* 310 */ } .... /* 873 */ protected void processNext() throws java.io.IOException { /* 874 */   if (scan_batch == null) { /* 875 */     scan_nextBatch(); /* 876 */   } /* 877 */   while (scan_batch != null) { /* 878 */     scan_processNextSub(); /* 879 */     scan_batch = null; /* 880 */     scan_nextBatch(); /* 881 */   } /* 882 */   scan_scanTime.add(scan_scanTime1 / (1000 * 1000)); /* 883 */   scan_scanTime1 = 0; /* 884 */ } </pre>

## Performance on Power8 cluster with IBM JDK

IBM JDK version	Time Taken in min - Default Spark	Time Taken in min – Catalyst AutoSplit Code	Improvement in %
Version SR3	56	45	19
Version SR5	42	38	9.5



# Conclusion

- Function outlining of huge routines proves to be beneficial when
  - It is very heavily used
  - JVM is unable to accommodate the compiled code into code cache.
- Automatic split of huge routines at source may/may not lead to improvements in all cases due to additional branches
- JIT compilers need to evaluate automatic use of function outlining as needed for hot functions.
- Monitoring and tuning JVM code cache and JIT compiler warm-up period can help tune Spark application performance
  - Use of Tiered Compilation. Option : `-XX:+TieredCompilation`
  - Maximum code cache size (in bytes) for JIT-compiled code. Option : `-XX:ReservedCodeCacheSize=<size>`
  - Number of interpreted method invocations before compilation.
    - ➔ In Server JVM, default is 10,000 and Client JVM, its 1,500. Option: `-XX:CompileThreshold=<invocations>`
  - Compiler threads to use for compilation
    - ➔ By default, is set based on the number of cores. Option: `-XX:CICompilerCount=threads`
  - Useful Logging Options - Option: `-XX:+LogCompilation` or `-XX:+PrintCompilation`

# References

- Good Reads for Spark SQL's Catalyst Optimizer
  - [Spark SQL: Relational Data Processing in Spark](#) by Michael Armbrust, Reynold S. Xin, Cheng Lian, Yin Huai, Davies Liu, Joseph K. Bradley, Xiangrui Meng, Tomer Kaftan, Michael J. Franklin, Ali Ghodsi, Matei Zahari
  - Deep Dive into Spark SQL's Catalyst Optimizer by Michael Armbrust, Yin Huai, Cheng Liang, Reynold Xin and Matei Zaharia
  - [Deep dive into the new Tungsten execution engine](#) by Sameer Agarwal, Davies Liu and Reynold Xin
  - [Deep Dive Into Catalyst: Apache Spark's Optimizer](#) by Yin Huai
  - [Cost-Based Optimizer in Apache Spark 2.2](#) By Ron Hu, Zhenhua Wang, Sameer Agarwal, Wenchen Fan
- Github link to log all dynamically generated catalyst code:  
[https://github.com/kmadhugit/spark/tree/catalyst\\_debug\\_v2.1.0-rc5](https://github.com/kmadhugit/spark/tree/catalyst_debug_v2.1.0-rc5)
- Github link to Automatic Catalyst Code Split : <https://github.com/kavanabhat/spark/tree/catalyst-autosplit>

# Questions?

# Thank You