

Using Deep Learning in Production Pipelines to Predict Consumers' Interest

Mathieu DESPRIEE, Tinyclues

#SAISML2

Who am I ?

Mathieu DESPRIEE

ML Engineering Manager at **tinyclues**
In charge of ML Scalability & Operability

Contributor to Spark, MxNET

TINYCLUES IN A FEW WORDS



**AI-FIRST
SAAS SOLUTION**



**DEEP LEARNING FOR CAMPAIGN
TARGETING & PLANNING**



**ACROSS ALL
CHANNELS**



**DRIVE MORE REVENUE
AND ENGAGEMENT**



**DESIGNED FOR
MARKETERS**



**USING ANONYMIZED
FIRST PARTY DATA**



**SEAMLESSLY INTEGRATED WITH
YOUR MARKETING STACK**



**FIRST CAMPAIGNS IN
2 WEEKS**

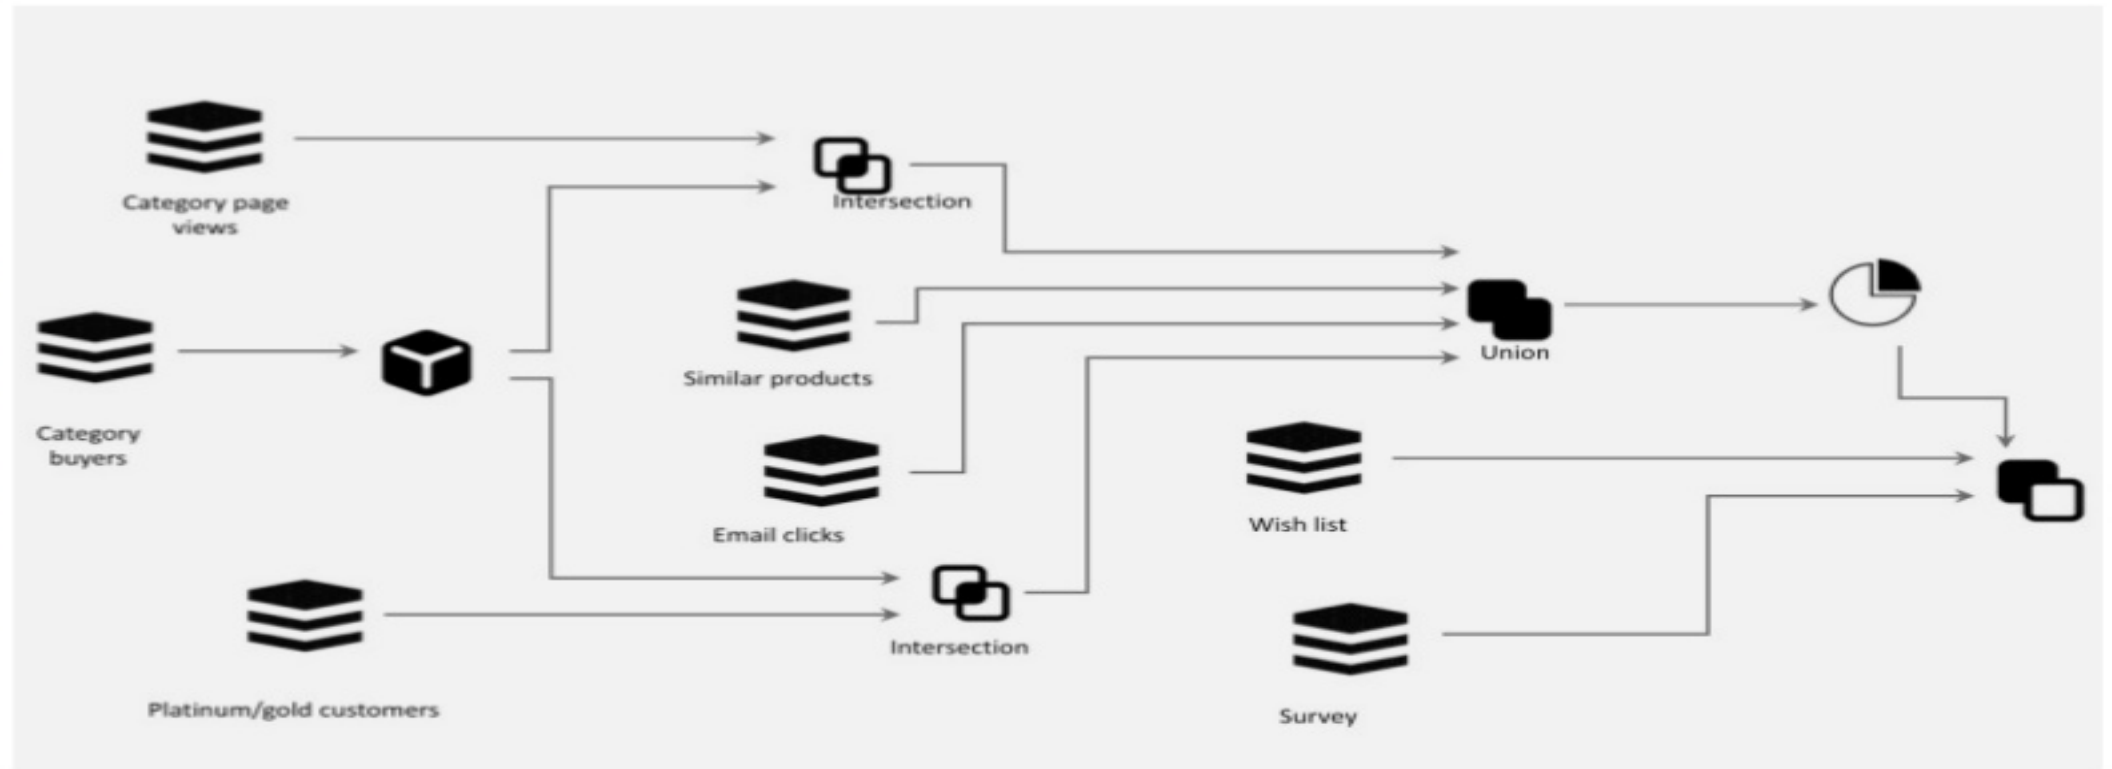
Instead of...

INTUITION

RFM

DEMOGRAPHICS

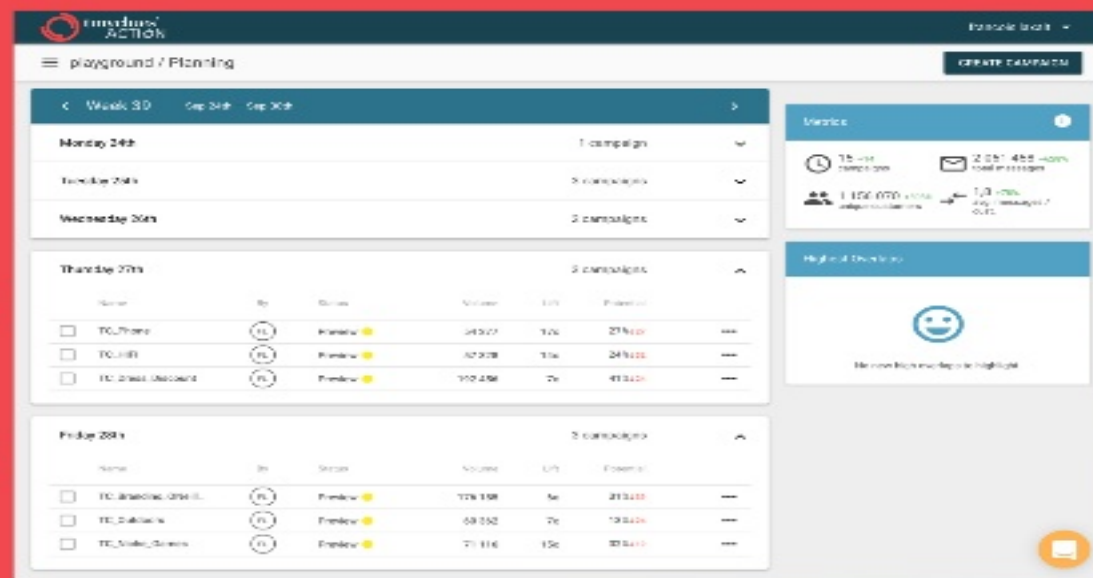
LAST PURCHASE



AI-FIRST: INTELLIGENT CAMPAIGNS IN JUST A FEW CLICKS

SPOT ON

TARGET THE FUTURE BUYERS FOR
ANY PRODUCT, IN THE DAYS
FOLLOWING A CAMPAIGN



START FROM YOUR
CAMPAIGN IDEAS
AND BUSINESS
GOALS

STRATEGIC
PRIORITIZE BASED ON VALUE
& MARKETING FATIGUE

EASY

BUILD SEGMENTS IN MINUTES –
NO GUESSWORK &
NO DATA-SCIENCE NEEDED

QUANTIFIED
ANTICIPATE & MAXIMIZE
CAMPAIGN PERFORMANCE

SUCCESS STORIES

+115%

CAMPAIGN REVENUE



TRAVEL

10M+
customers

\$1B+

+30%

CAMPAIGN REVENUE



RETAIL

10M+
customers

\$1B+

+151%

REVENUE PER EMAIL



RETAIL / FASHION

\$1B+

+178%

CAMPAIGN REVENUE



HOSPITALITY

5M+
customers

\$1B+

+30%

CAMPAIGN REVENUE



HOSPITALITY

\$1B+

+20%

CAMPAIGN REVENUE
ONLINE & IN-STORE



RETAIL / FASHION

\$1B+

+80%

CAMPAIGN REVENUE



E-COMMERCE

\$1B+

10M+
customers

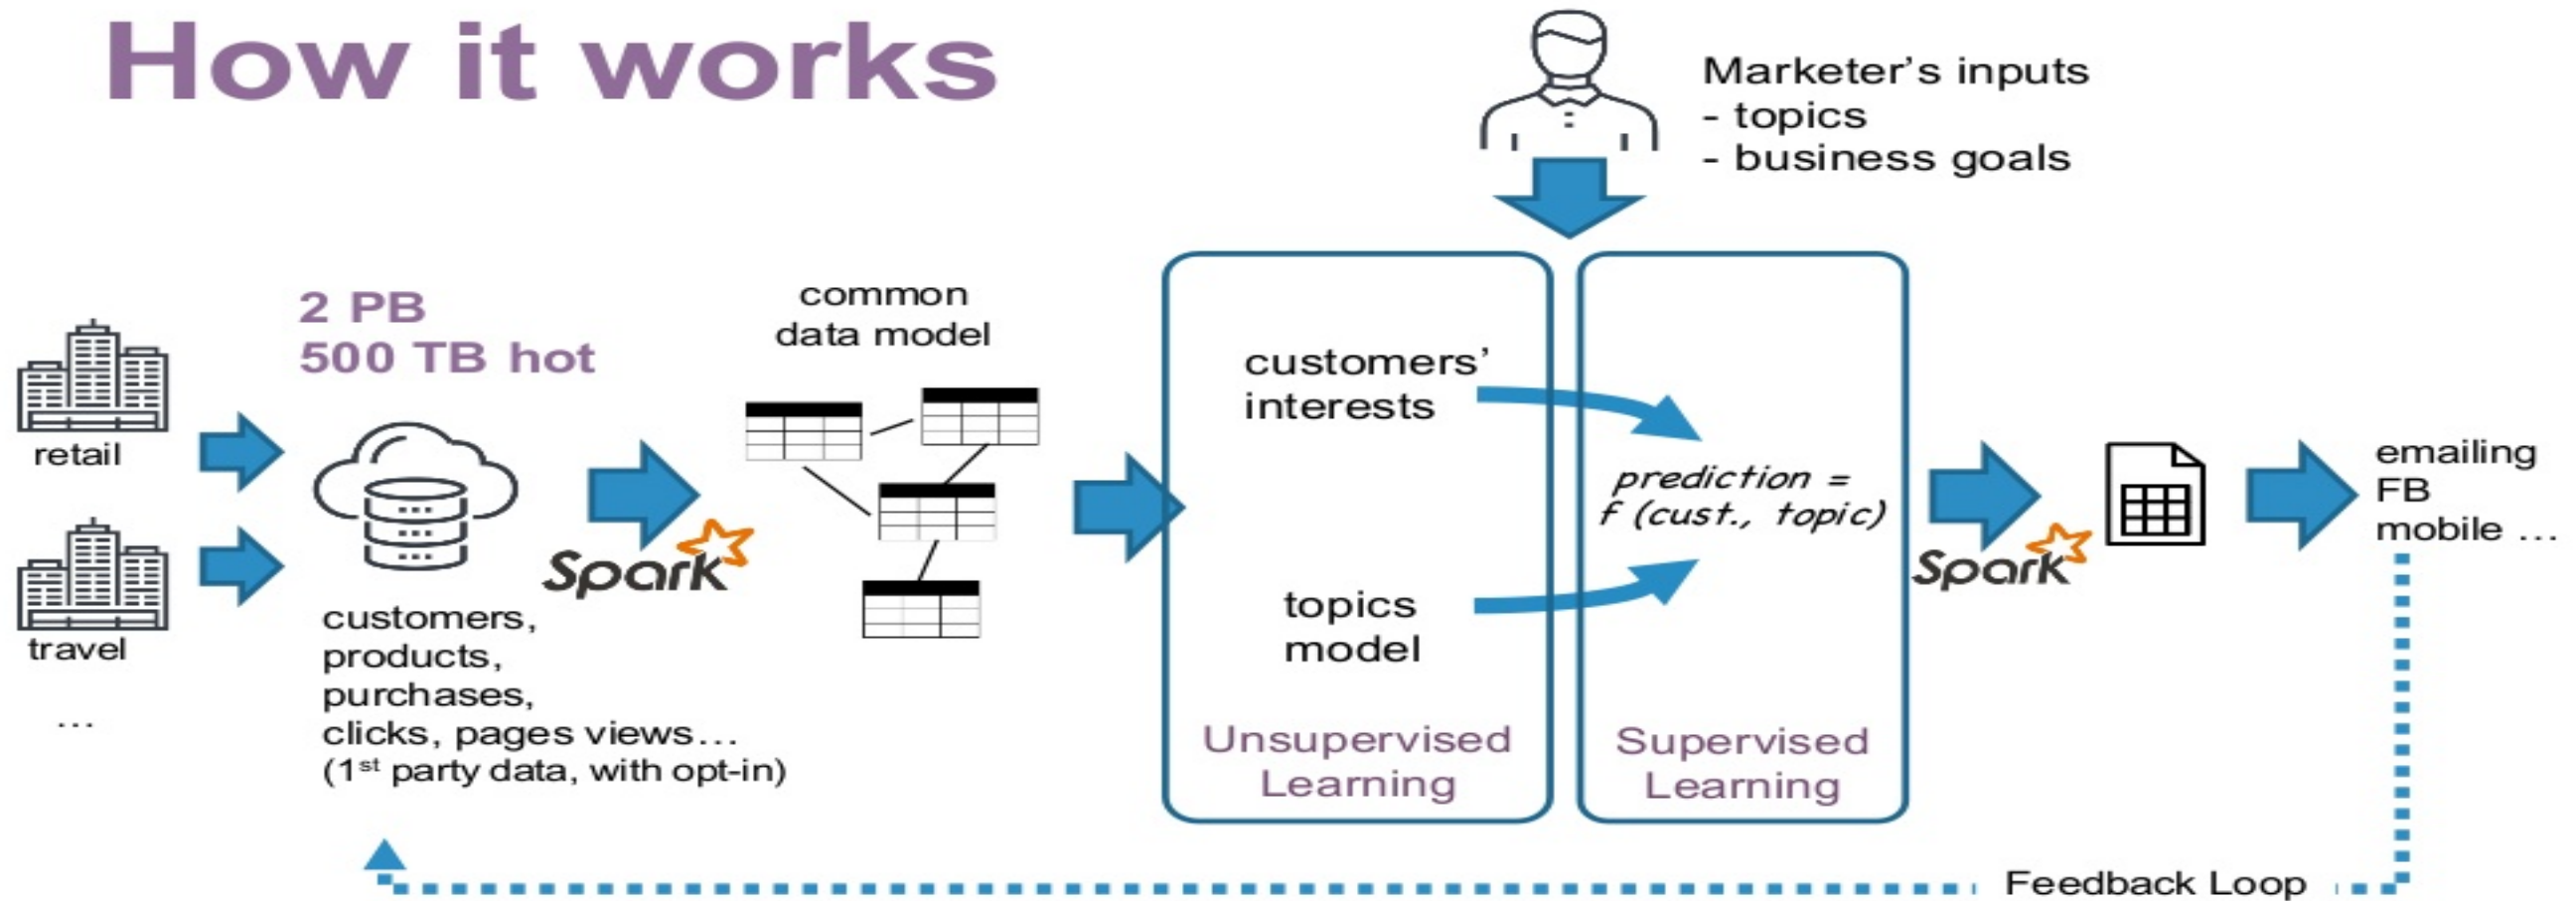
+60%

REVENUE PER EMAIL

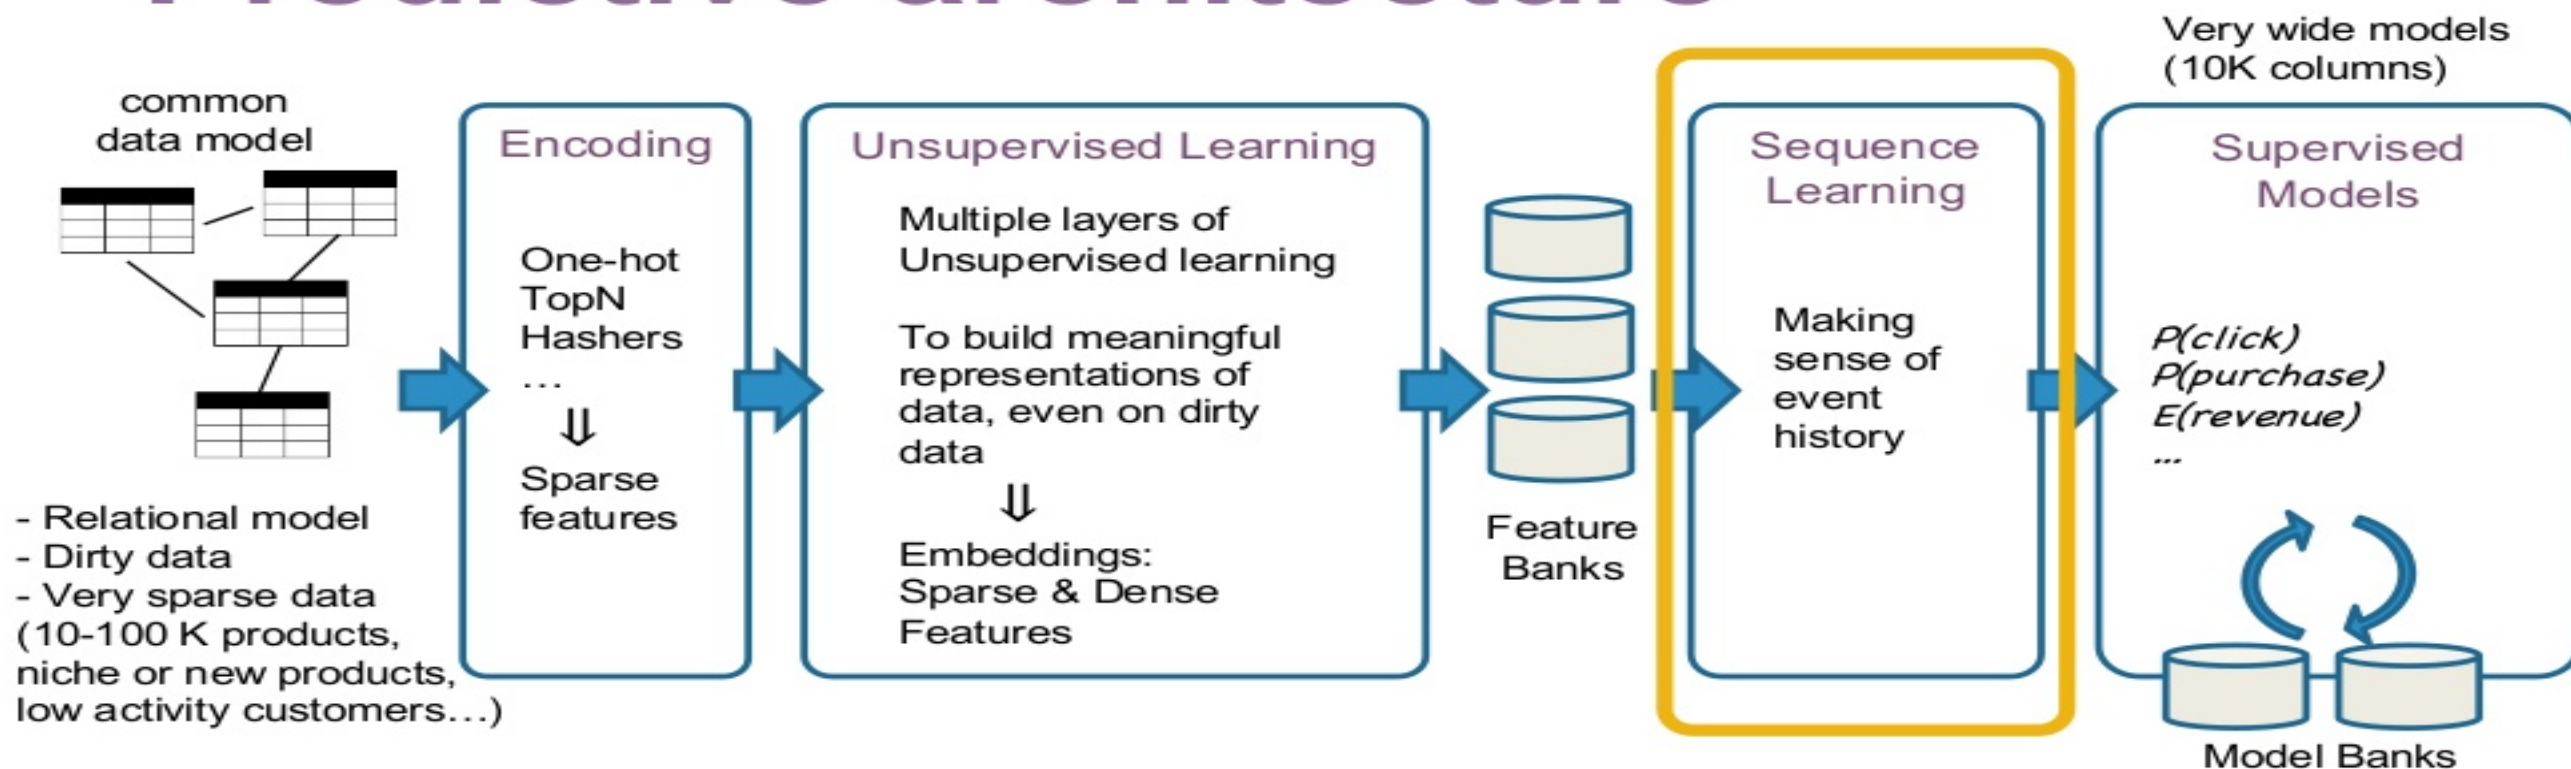


RETAIL / FASHION

How it works



Predictive architecture




Working with events

- Events (page views, clicks...) are not always meaningful individually, but sequences are
- Events are ordered, and are not independent from each other
- Sequences are not all of the same length
- Can we find a model architecture leveraging on these characteristics ?

RNN: Recurrent Neural Networks

- Networks with loops in them, allowing information to persist.




- They also allow input sequences of variable length
 - you just need to pad data with zeroes $x_0 \dots x_k$

Illustrations by Christopher Olah - Check out his blog! colah.github.io

LSTM: Long Short-Term Memory

- Classical RNN can't learn long-term dependencies with gradient descent because of vanishing or exploding gradients: they tend to forget in the long-term
- LSTMs are explicitly designed to avoid this problem
 - Invented by Hochreiter & Schmidhuber (1997)
 - Many variants exist
- LSTM have great results in speech recognition, translation, and are building block of assistants.

LSTM

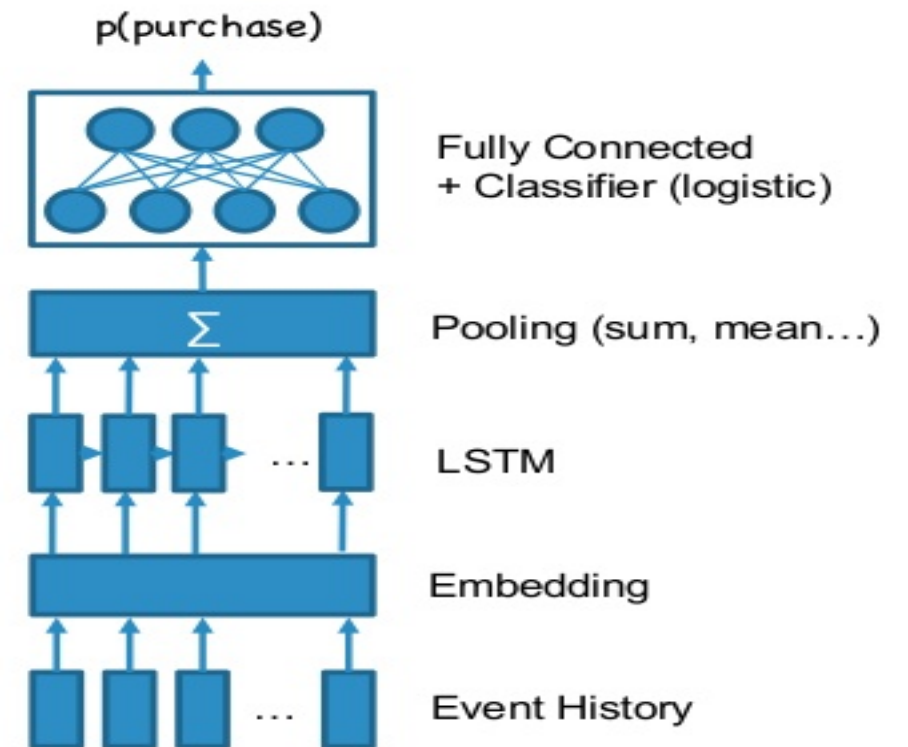


- Cell-state is transmitted from cell to cell (top line) without going through a non-linear activation function
- The rest of the cell controls:
 - how the input and previous outputs modifies the cell state,
 - how the state and input are combined together to form output

Example of LSTM for purchase prediction

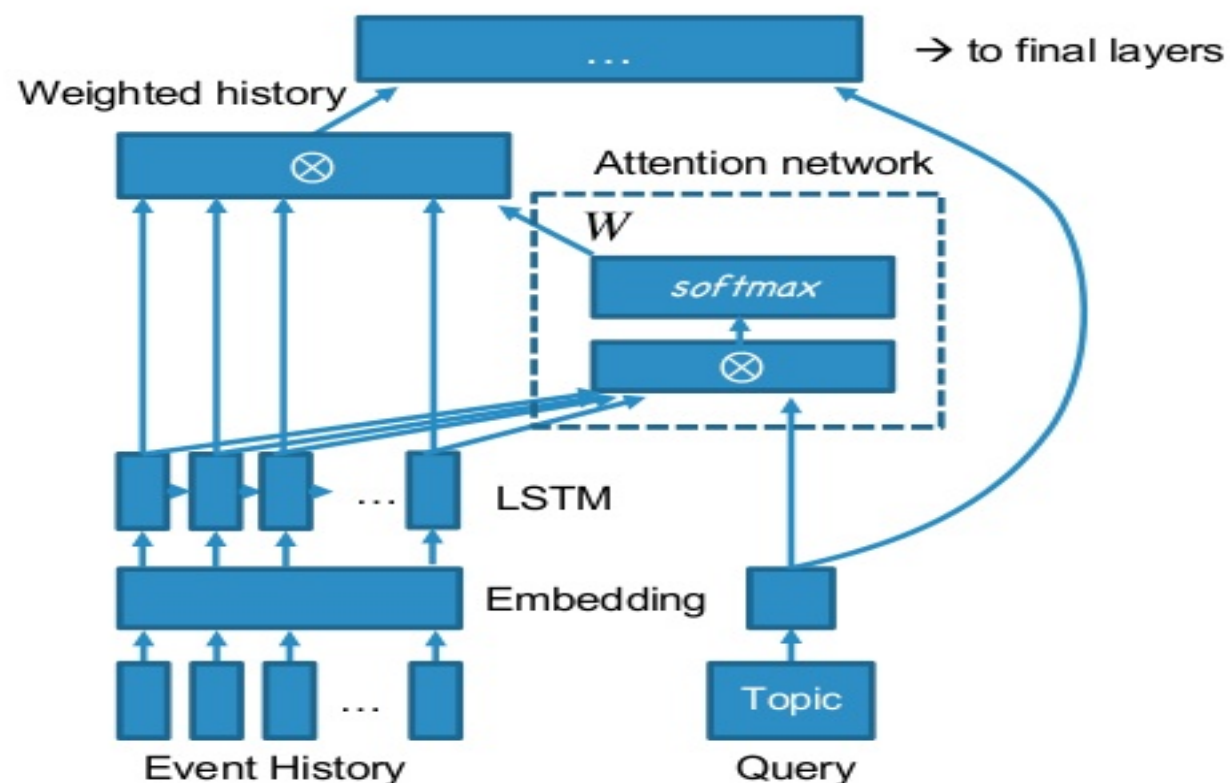
- Inputs are events (page views, clicks), padded to a max sequence length, and passed through an embedding layer
- LSTM output is aggregated with a pooling layer, and some densely connected layers
- Through gradient descent, the network will learn what event sequences are more likely to lead to the target (purchase)

"embedding": representing an information by a vector in a space, capturing its essential meaning



Bringing some focus: Attention Model

- Attention models are a way to make the network focus on elements of interest
- Introduce the notion of “Query” = what the network should focus on (topic)
- The attention network outputs weights W that modulate the importance of each LSTM output, given the topic
- This weighted and interpreted history is then used in subsequent layers for classification



Implementation & Execution Engine

- Desired technology for production:
 - Easily distributable with Spark
 - Callable from Scala
 - Interoperability with Python (used in our Lab)
 - High performance out-of-the-box
 - should benefit from native libs
 - ability to use GPU for training phases
 - Good support of Sparse tensors

Apache MxNET

- Deep Learning framework, incubated at Apache
- Back-end in C++
- Multiple front-end languages: scala, python, Julia, R, ...
- Integration with linalg native libs, including Intel MKL
- Benchmarks with GPUs are very good
- Aims at interoperability with others frameworks (ONNX, MxNET as Keras backend, etc.)

LSTM w/ Attention in Python

```
class NetLSTM(mx.gluon.nn.HybridSequential):
    def __init__(self, embedding_shape, **kwargs):
        super(NetLSTM, self).__init__(**kwargs)
        with self.name_scope():
            self.embedding_shape = embedding_shape
            self.embeddedding = self.params.get("embeddedding", shape=embedding_shape, grad_req='null')
            self.lstm = mx.gluon.rnn.LSTM(hidden_size=16, layout="NTC")
            self.lstm_pool = mx.gluon.nn.Dense(units=1, flatten=False)
            self.deep1 = mx.gluon.nn.Dense(units=10, activation="relu", flatten=False)
            self.deep2 = mx.gluon.nn.Dense(units=1, activation="sigmoid", flatten=False)

    def hybrid_forward(self, F, hist, query, *args, **kwargs):
        embeddedding = kwargs["embeddedding"]
        query = F.reshape(query, shape=(0, 1))

        history = F.Embedding(data=hist, weight=embeddedding, \
                               input_dim=self.embedding_shape[0], output_dim=self.embedding_shape[1])
        query = F.Embedding(data=query, weight=embeddedding, \
                             input_dim=self.embedding_shape[0], output_dim=self.embedding_shape[1])

        simh = F.broadcast_mul(history, query)
        lstm = self.lstm(simh)
        lstm_pool = self.lstm_pool(lstm)
        softmax = F.exp(F.sum(lstm_pool, axis=2))
        softmax = F.broadcast_div(softmax, F.sum(softmax, axis=1, keepdims=True))
        attention_weights = F.reshape(softmax, shape=(0, 1, -1))
        aggregate_history = F.batch_dot(F.L2Normalization(attention_weights, mode='instance'), \
                                         F.L2Normalization(history, mode='instance'))

        interaction = query * aggregate_history
        concat = F.concat(aggregate_history, interaction, query, dim=2)
        deep = self.deep1(concat)
        deep = F.concat(concat, deep, dim=2)
        deep2 = self.deep2(deep)
        out = F.reshape(deep2, shape=(0,))
        return out
```

- Written using (the very young) Gluon API
- Similar to Keras
- Can be compiled into a “Symbol” model (symbolic graph of execution)

Using in Scala a model trained in Python

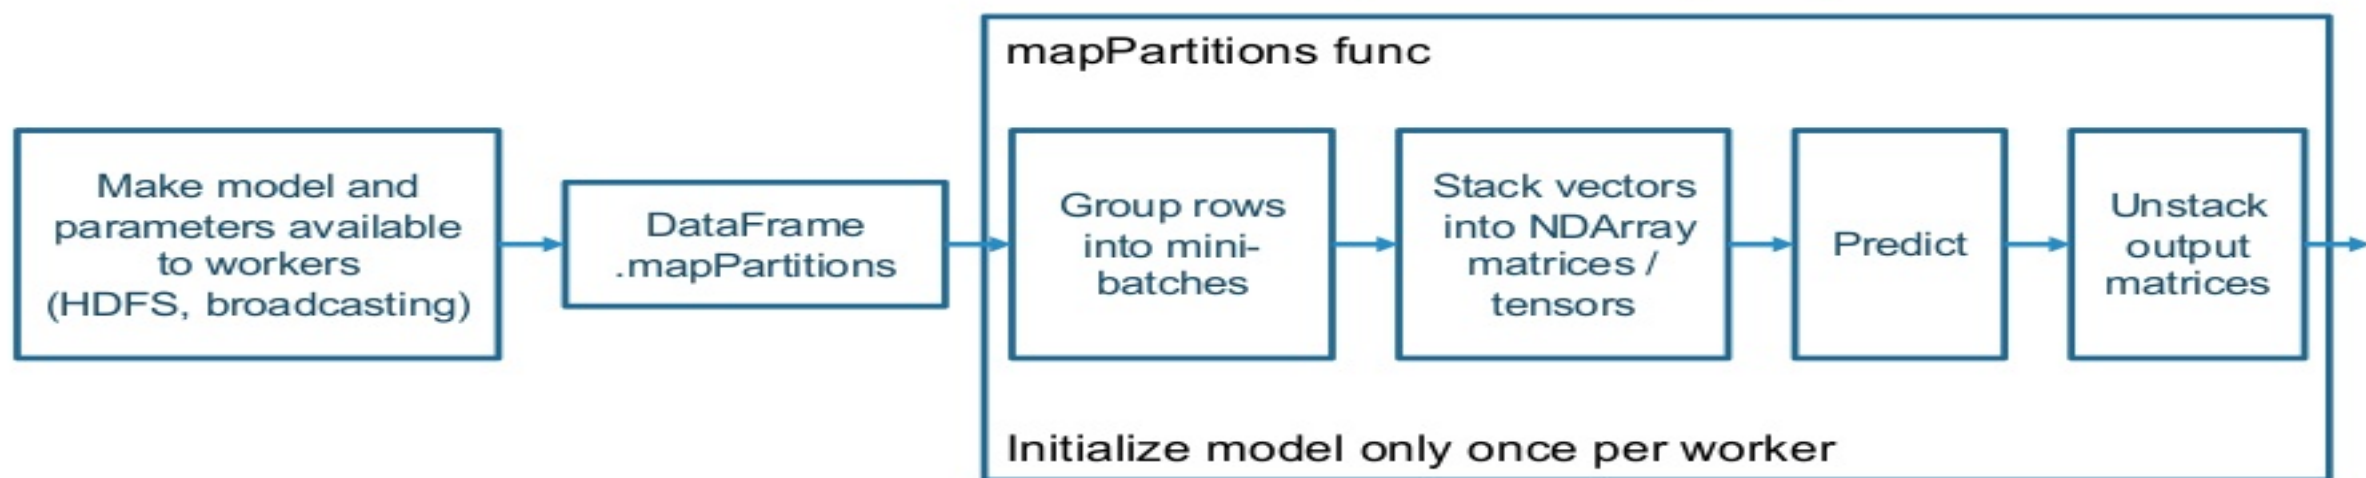
```
In [29]: model.export(epoch=40, path="s3://test-mde-us/mxnet/model")
```

```
val symbol = Symbol.load(symbol_path)
val params = NDArray.load2Map(params_path)

/* ... */
val mod = new Module(symbol, dataNames = inFields, labelNames = IndexedSeq(),
    fixedParamNames = Some(params.keySet))
mod.bind(shapes, labelShapes = None, forTraining = false)
mod.initParams(new ZeroInit(), argParams = params)

/* ... */
while (batchIter.hasNext) {
    val batch = batchIter.next()
    mod.forward(batch, Some(false))
}
```


Inference from a Spark pipeline



- Pay attention to resource usage on workers
 - Ensure backend processing will have available CPU cores and plenty of RAM
- Pay attention to data copy operation that will happen between Spark Vectors (from java memory) and MxNET native backend

Learnings about MxNET in Scala

- MxNET is very young
 - Mixed API styles
 - Scala API is lagging Python's
 - No sparse tensors yet
 - Unhelpful error messages, especially when they come from the C++ backend
 - In Lab context (notebook, fast iteration), it feels heavy
 - Some non-trivial boilerplate code to write
 - ONNX interop has limitations
- Yet, it's very promising, and fast out-of-the-box
 - Many binaries available with hw optimization
 - Usable in Scala/Spark pipelines

Take aways

- To find the “tiny clues” within your data, you need powerful learning algorithms to discover their latent meaning
- LSTM with Attention Models are especially good when it comes to sequences of events
- Using inference models from Scala/Spark code is easy and performant with MxNET, even if this framework has still a lot of limitations

Thank you !


Questions ?

Reach me on twitter @mdesprie

Backup


Classical RNN have limitations

- Vanilla RNN can't learn long-term dependencies with gradient descent because of vanishing or exploding gradients




Illustrations by Christopher Olah (check out his blog!) <http://colah.github.io/>

RNN




Input from upstream layers:
Each x is a value at a point in
time, represented by a Feature
Vector


*"embedding": a compact
mathematical representation of an
information into a vector, capturing
its semantics*



MKL impact




CPU times: user 4min 12s, sys: 10.4 s, total: 4min 22s Wall time: 2min 2s



CPU times: user 15min 41s, sys: 40.6 s, total: 16min 21s Wall time: 59.6 s

Performance of training



Disclaimer: These measurements are indicative, and vary a lot depending on many parameters.